

Automatically Scheduling Halide Image Processing Pipelines

Ravi Teja Mullapudi (CMU)

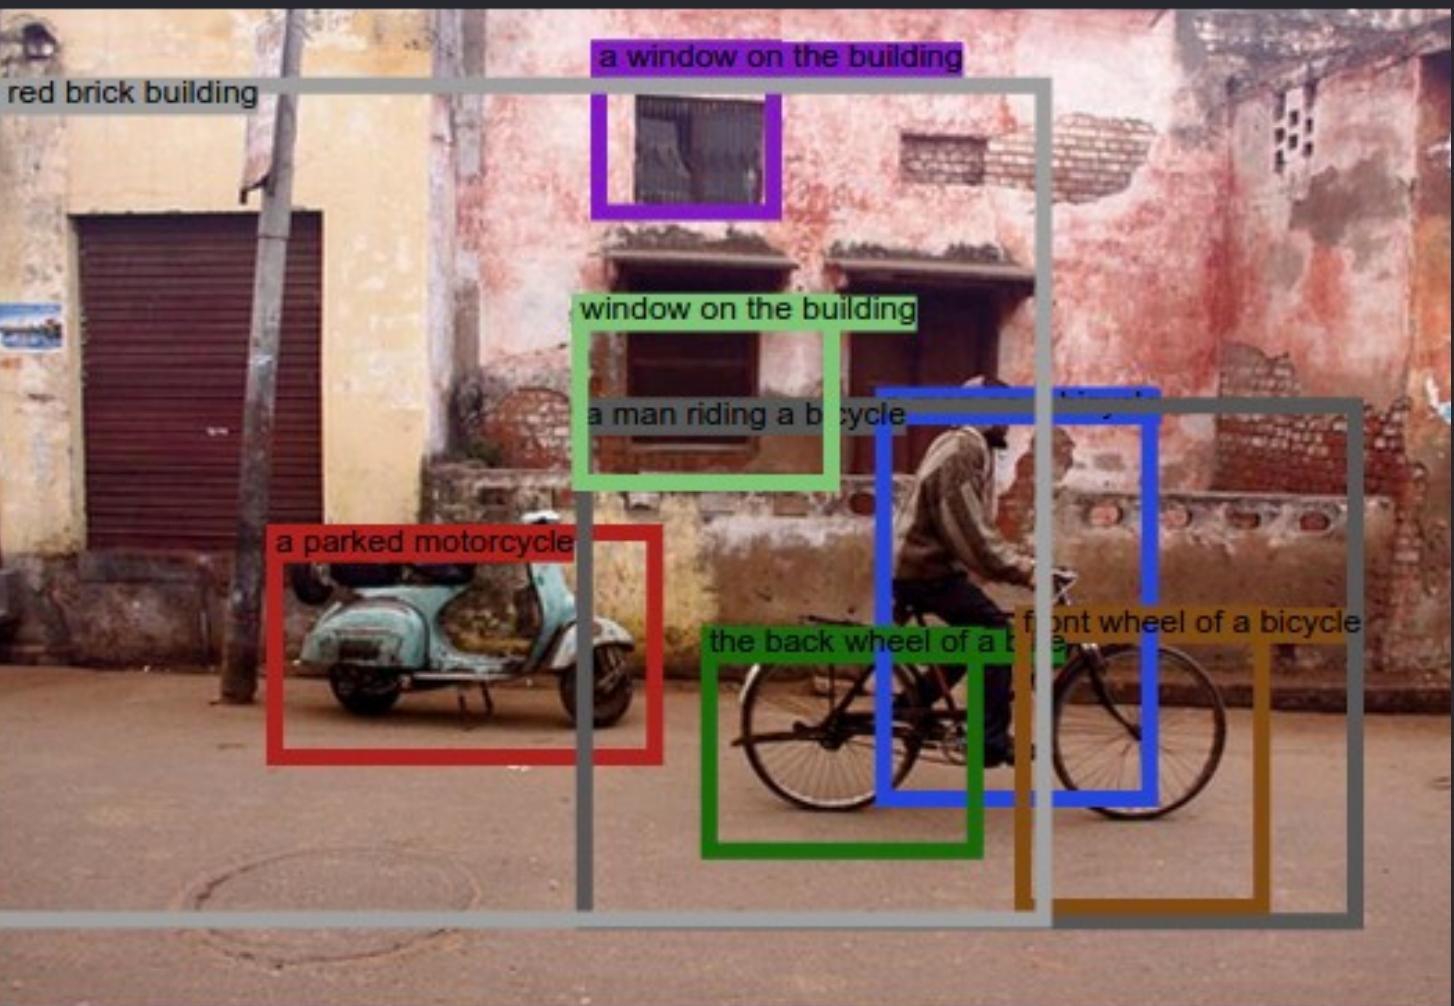
Andrew Adams (Google)

Dillon Sharlet (Google)

Jonathan Ragan-Kelley (Stanford)

Kayvon Fatahalian (CMU)

High demand for efficient image processing

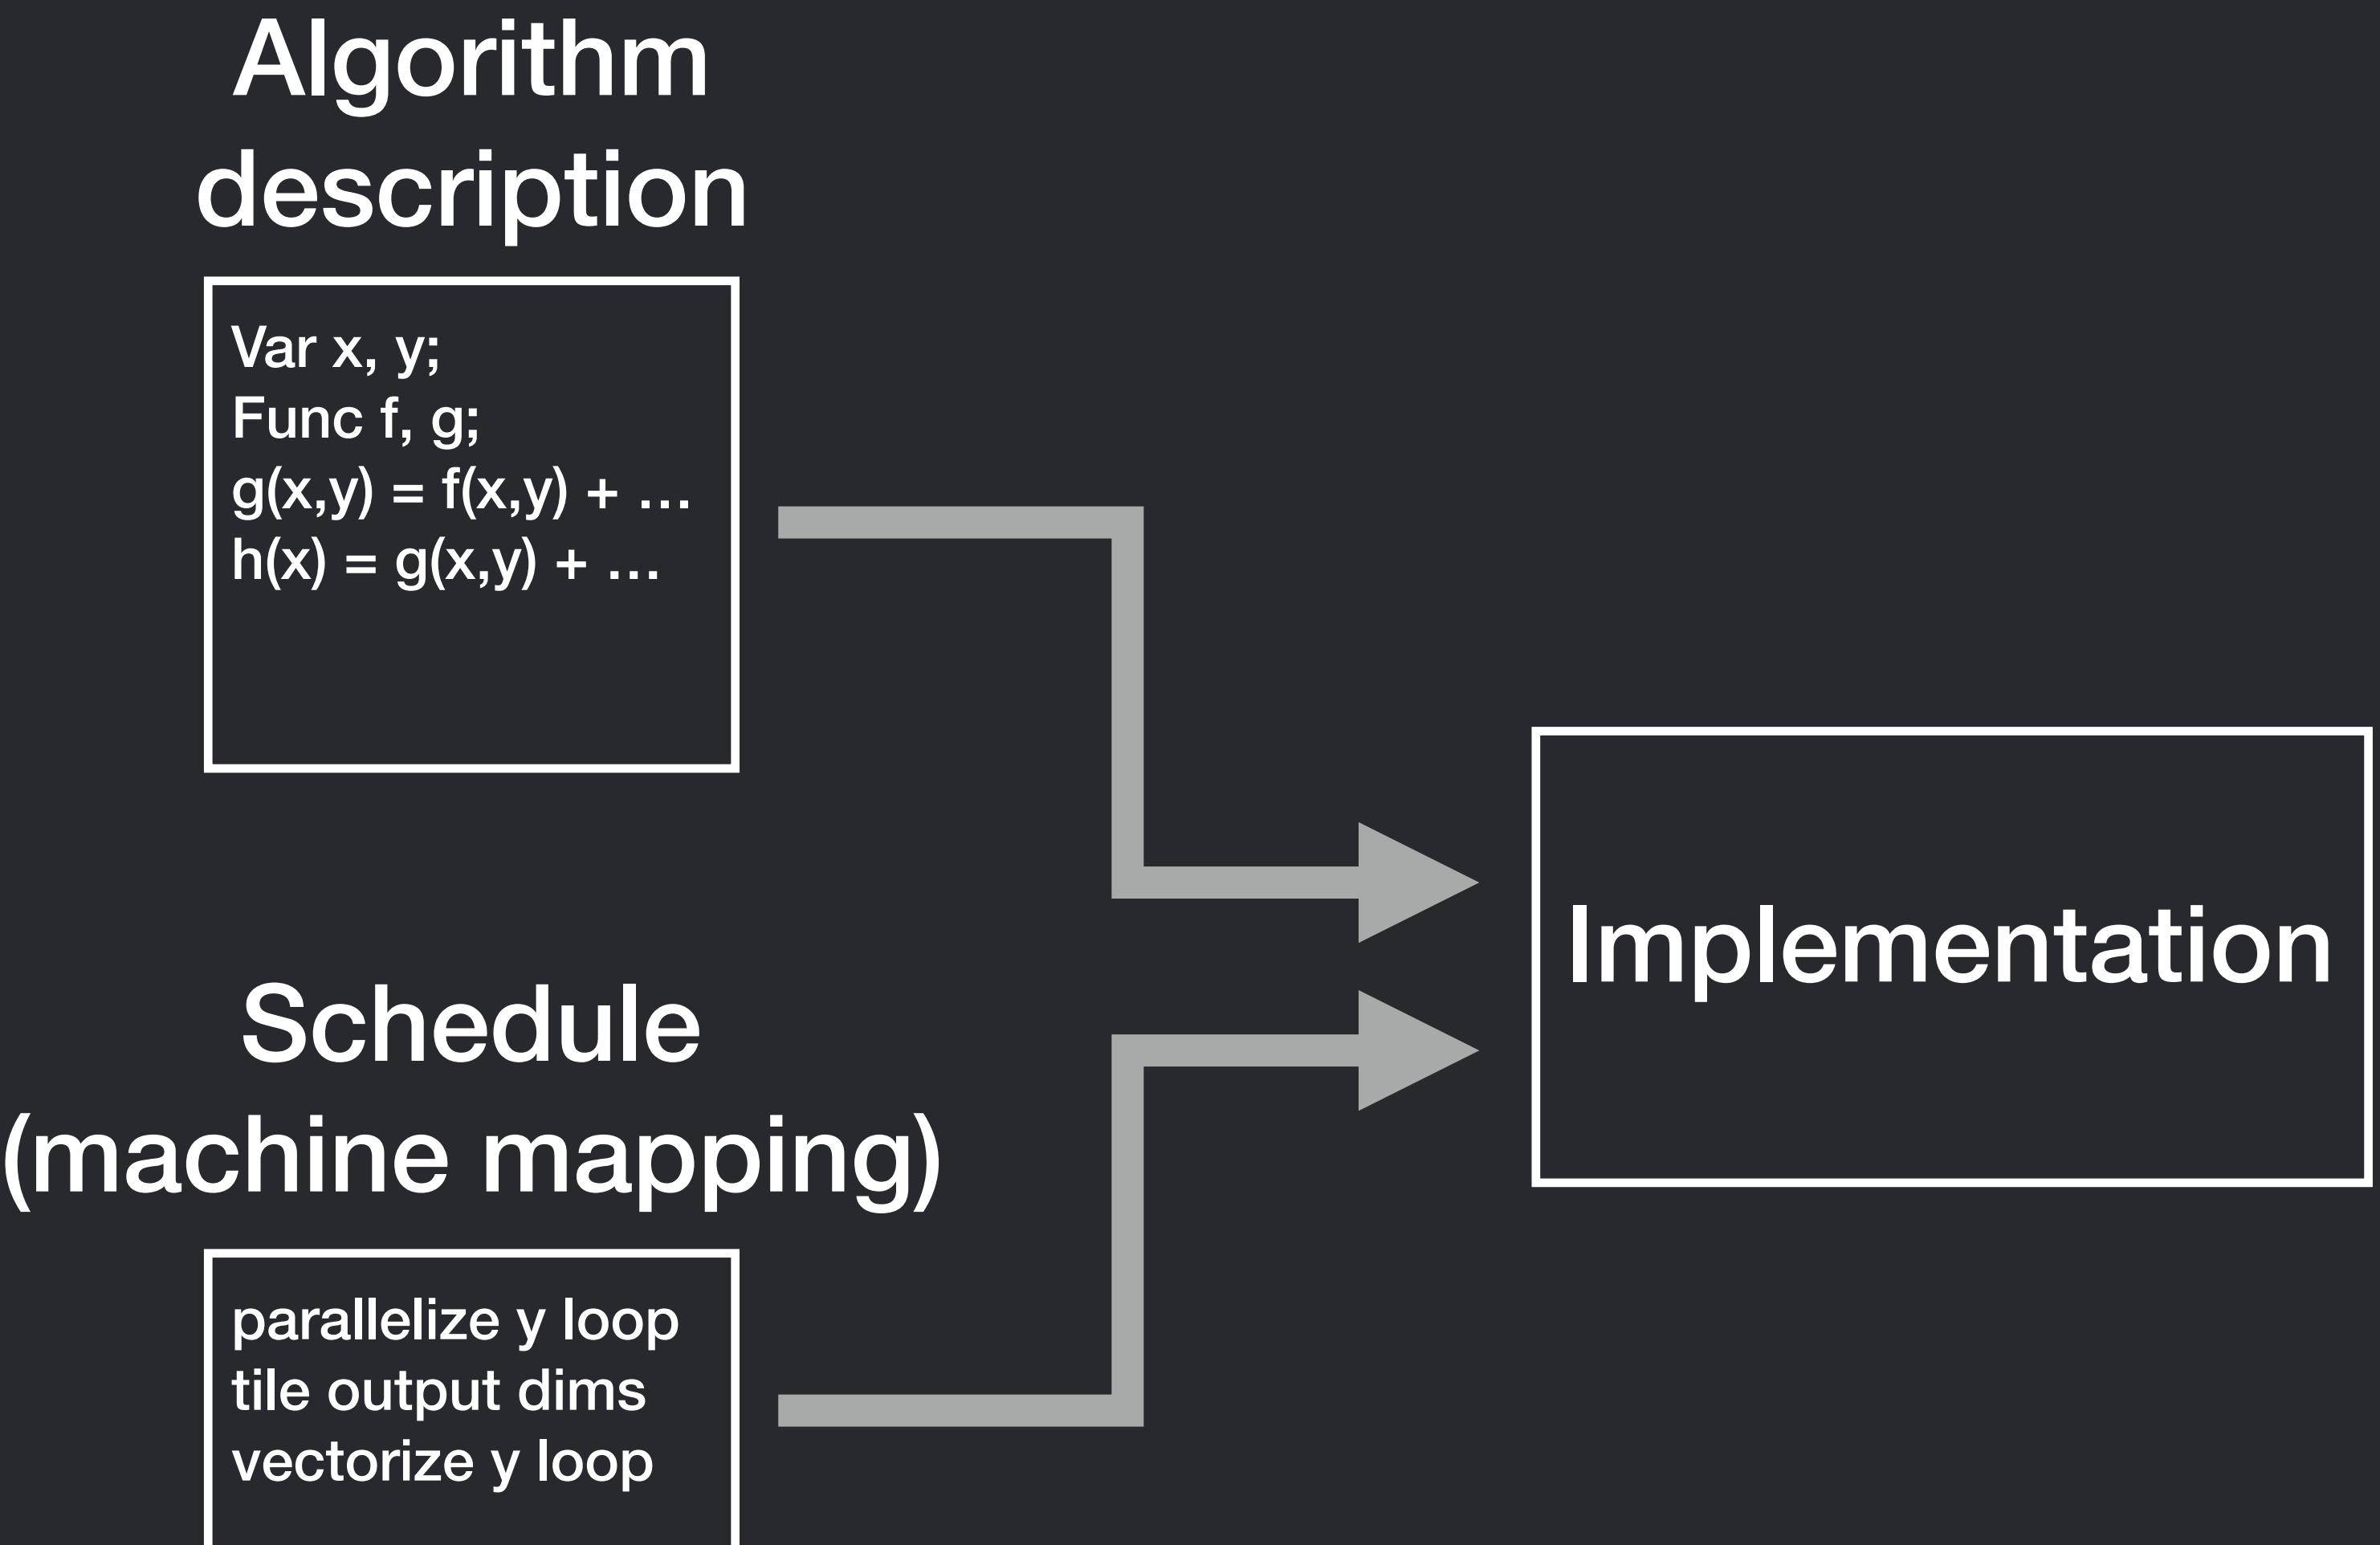


Scheduling image processing algorithms

Algorithm description

```
Var x, y;  
Func f, g;  
g(x,y) = f(x,y) + ...  
h(x) = g(x,y) + ...
```

Scheduling image processing algorithms



Scheduling image processing algorithms

Image processing
algorithm developers



Algorithm
description

```
Var x, y;  
Func f, g;  
g(x,y) = f(x,y) + ...  
h(x) = g(x,y) + ...
```

Schedule
(machine mapping)

```
parallelize y loop  
tile output dims  
vectorize y loop
```

Implementation

Few developers have the skill set to author highly optimized schedules

Image processing algorithm developers



Algorithm description

```
Var x, y;  
Func f, g;  
g(x,y) = f(x,y) + ...  
h(x) = g(x,y) + ...
```

Schedule
(machine mapping)

parallelize y loop
tile output dims
vectorize y loop

> 10x Faster Implementation

Contribution: automatic scheduling of image processing pipelines

Image processing algorithm developers



Algorithm description

```
Var x, y;  
Func f, g;  
g(x,y) = f(x,y) + ...  
h(x) = g(x,y) + ...
```

Generates expert-quality schedules in seconds

Scheduling Algorithm

> 10x Faster Implementation

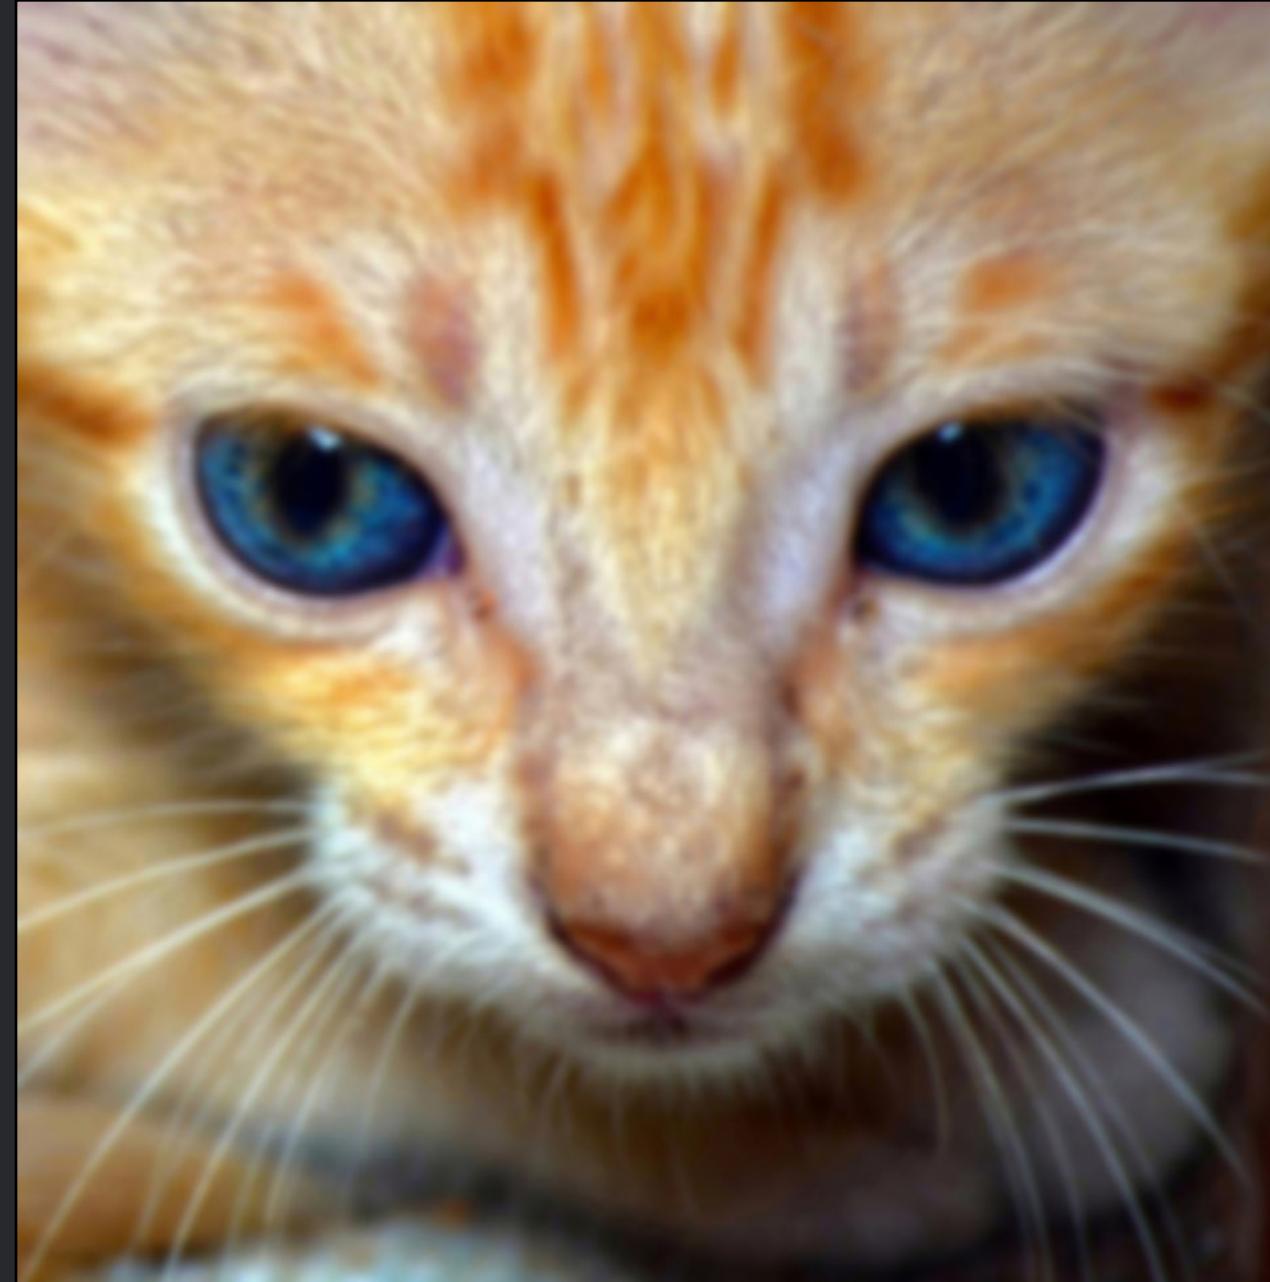
**Why is it challenging to schedule
image processing pipelines?**

Algorithm: 3x3 box blur



in

Algorithm: 3x3 box blur

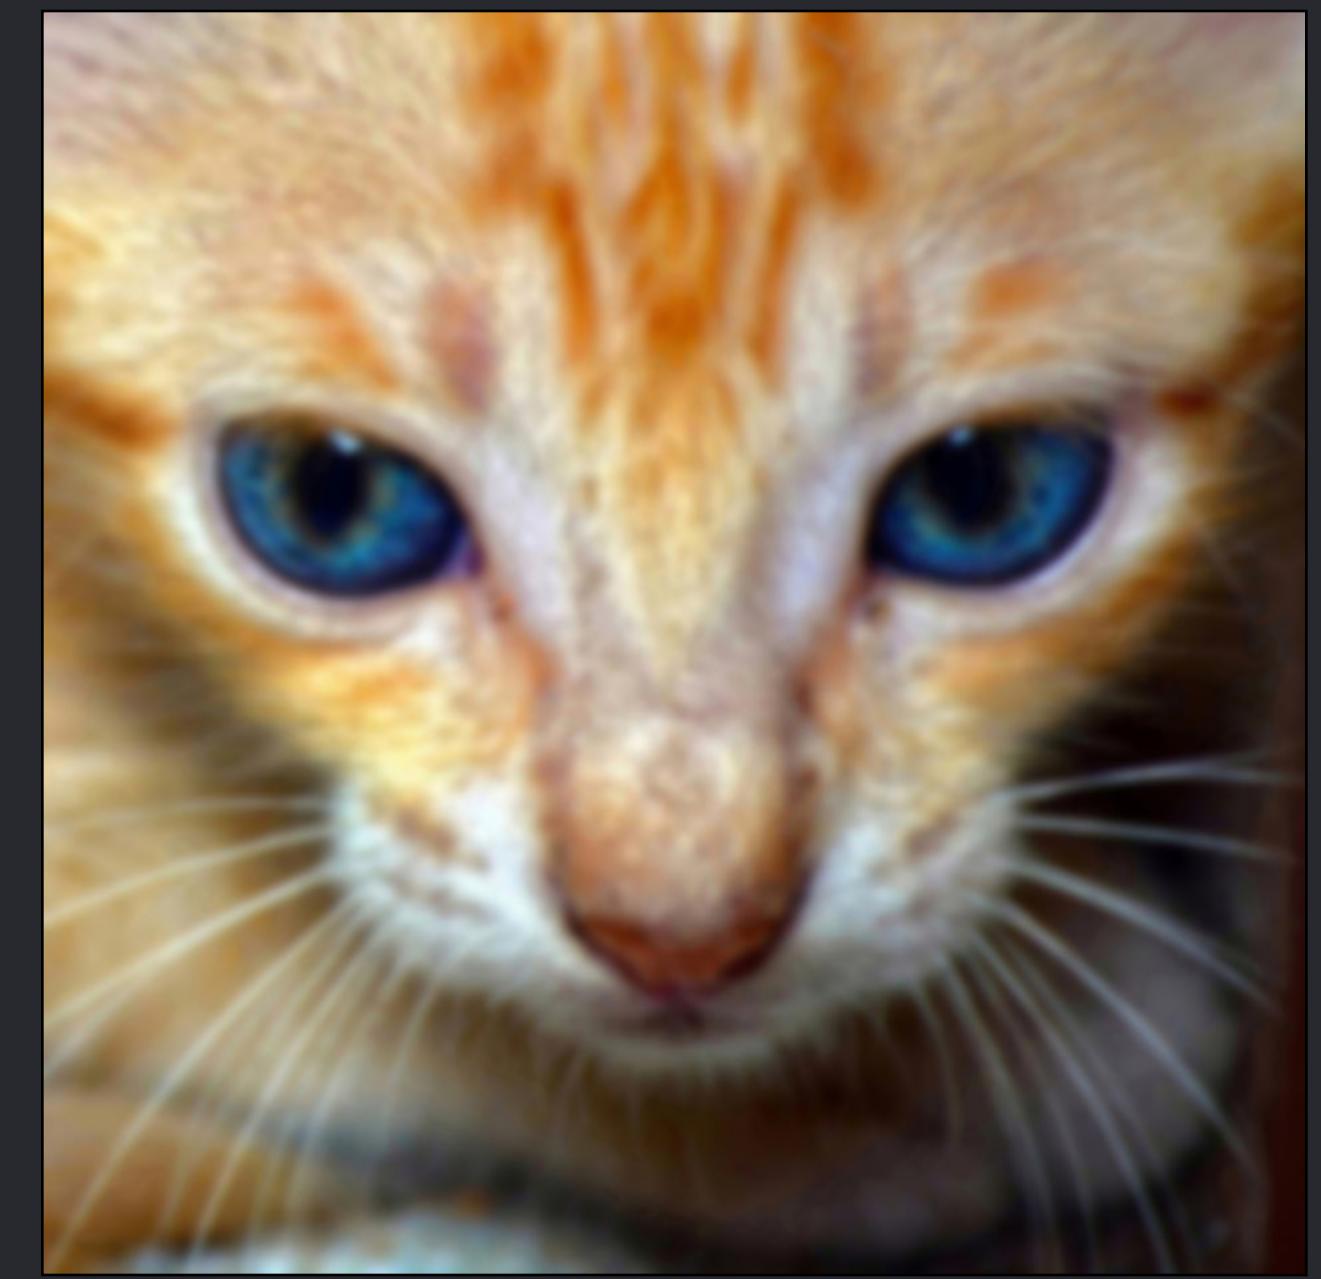
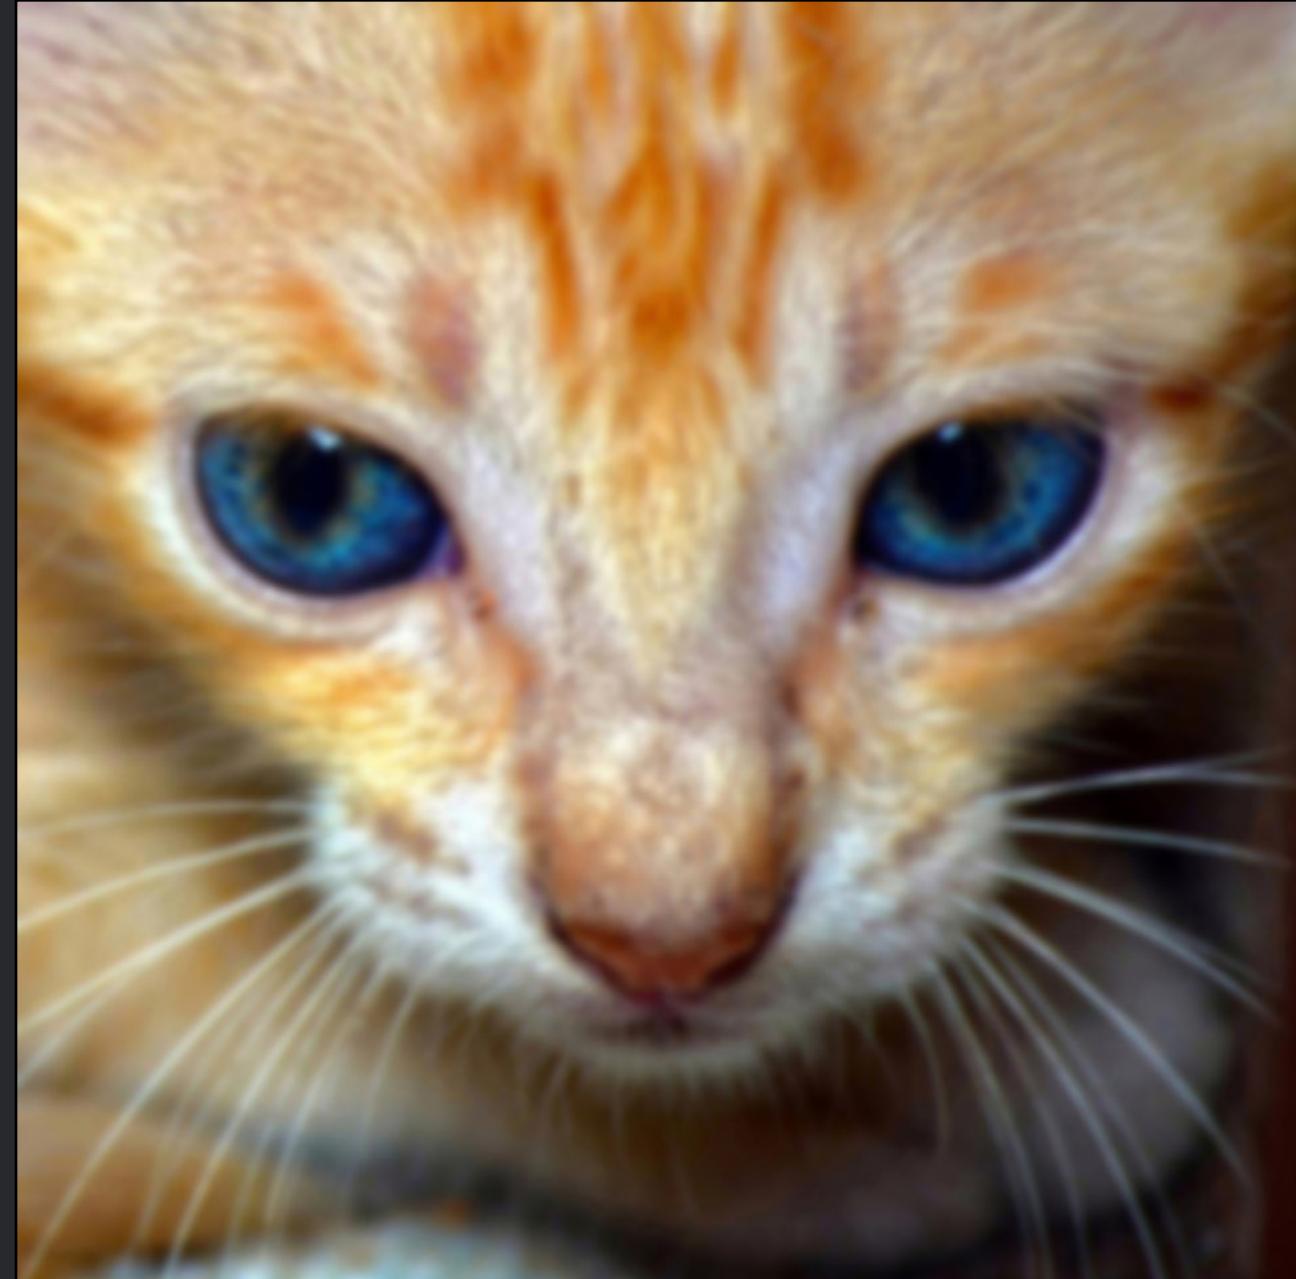


in

bx

$$bx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y)) / 3$$

Algorithm: 3x3 box blur



in

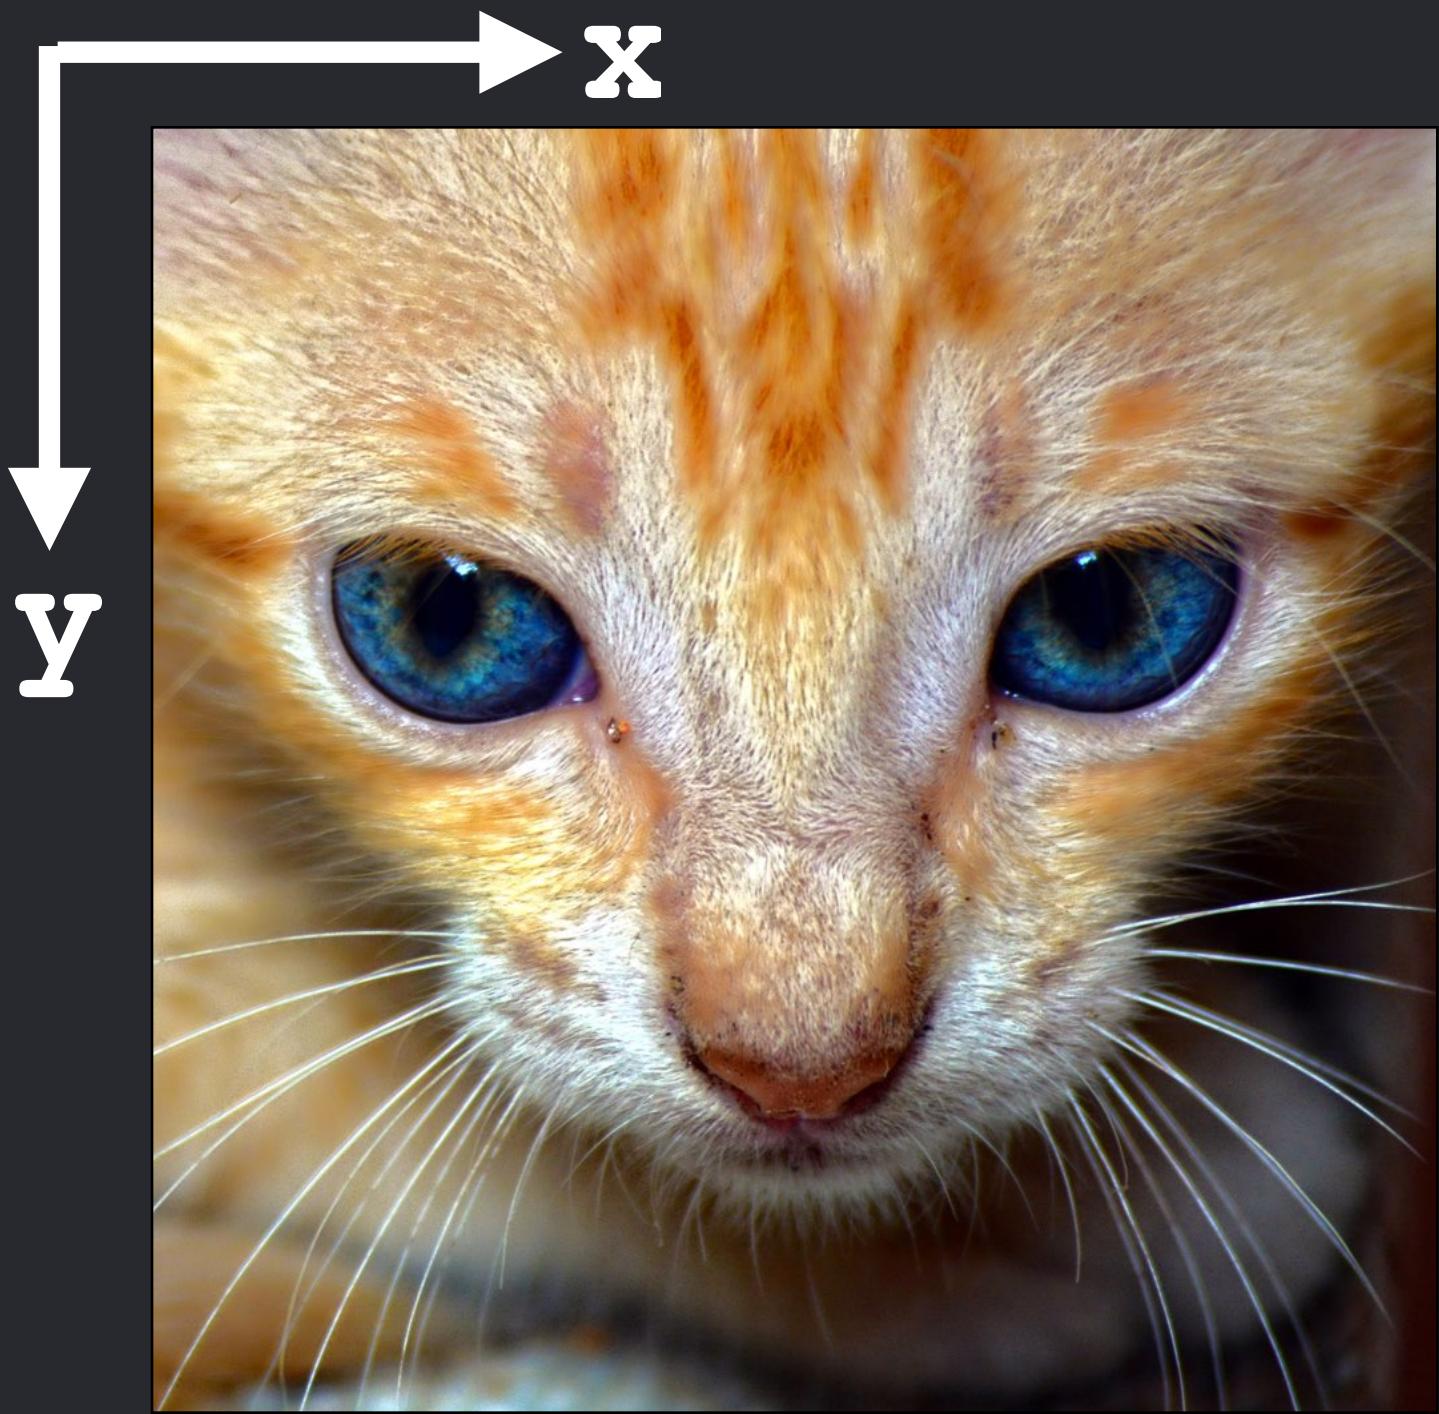
bx

out

$$bx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y)) / 3$$

$$out(x, y) = (bx(x, y-1) + bx(x, y) + bx(x, y+1)) / 3$$

A basic (slow) schedule

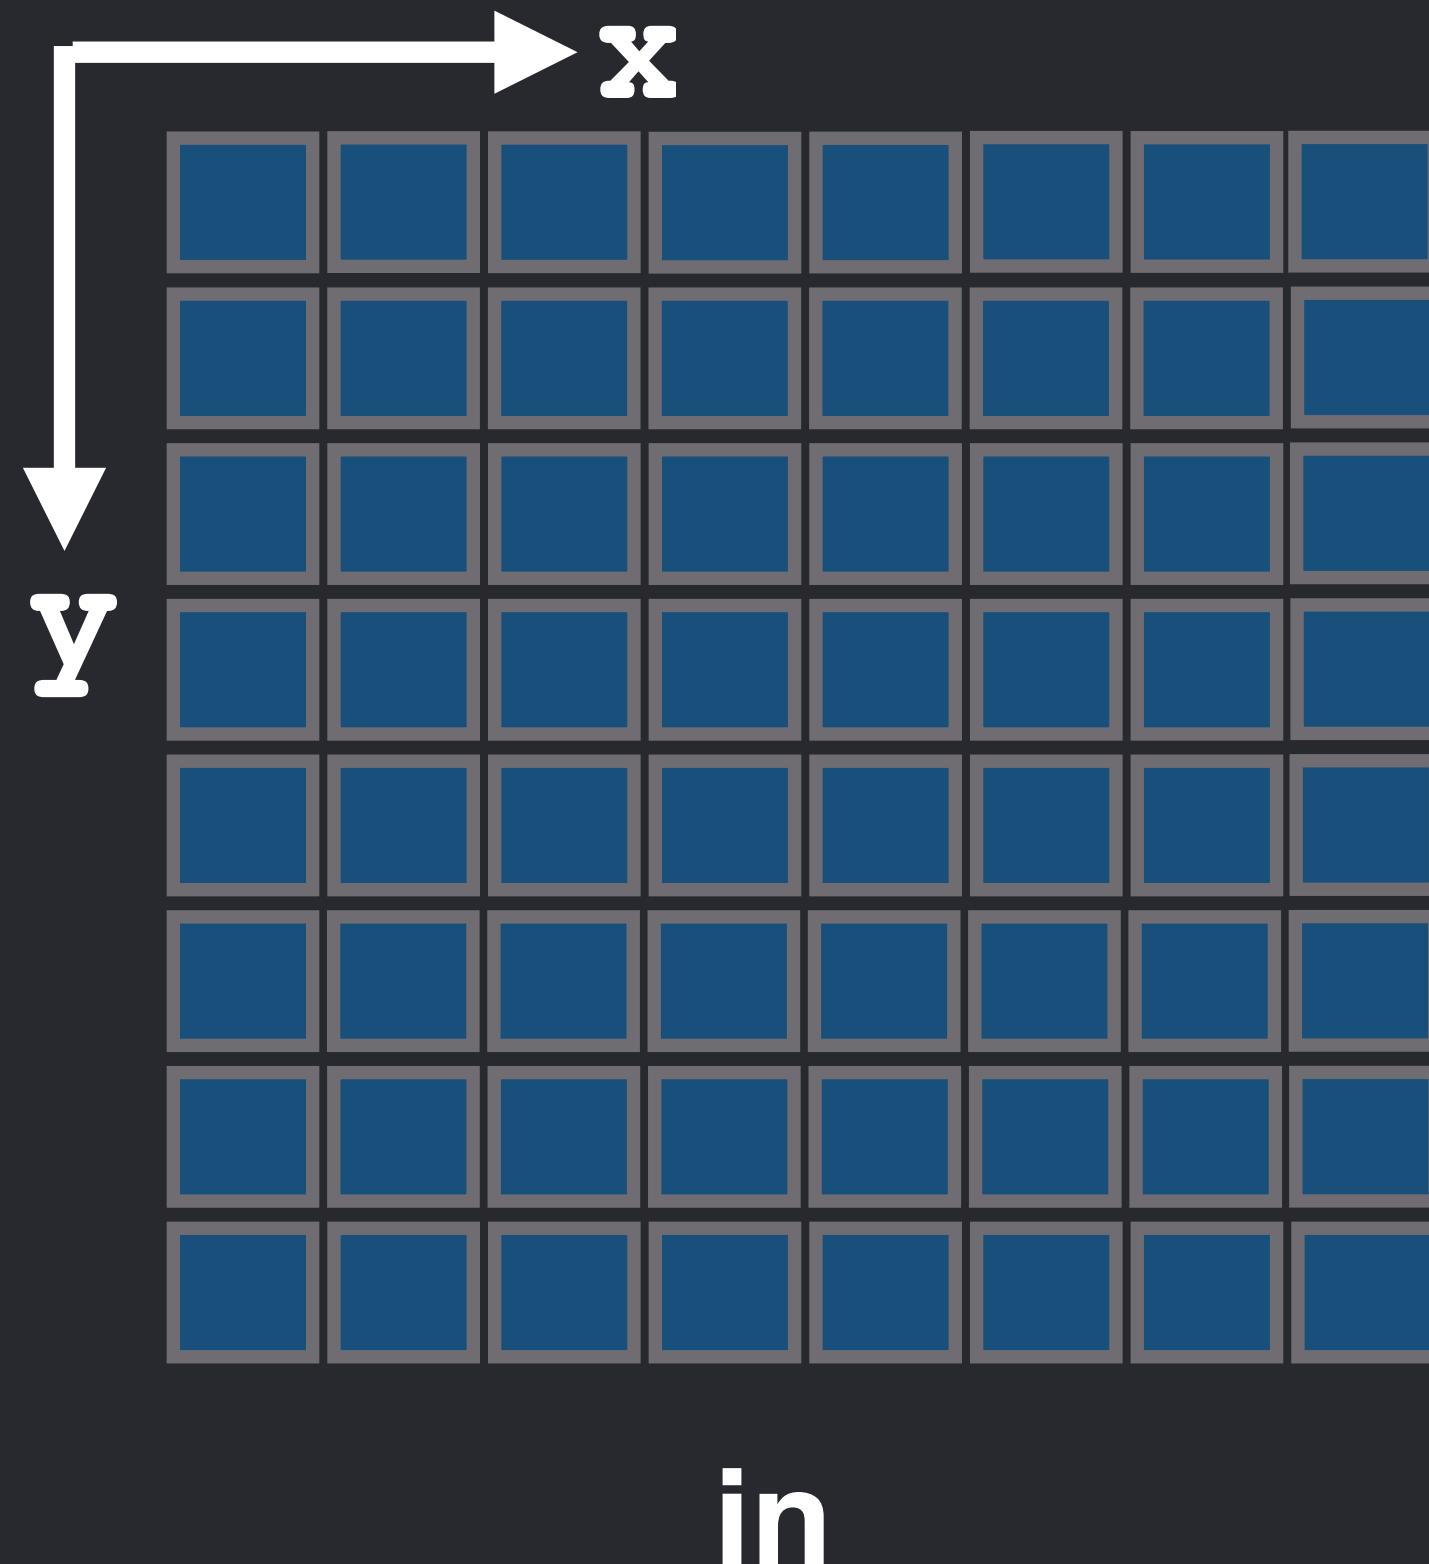


in

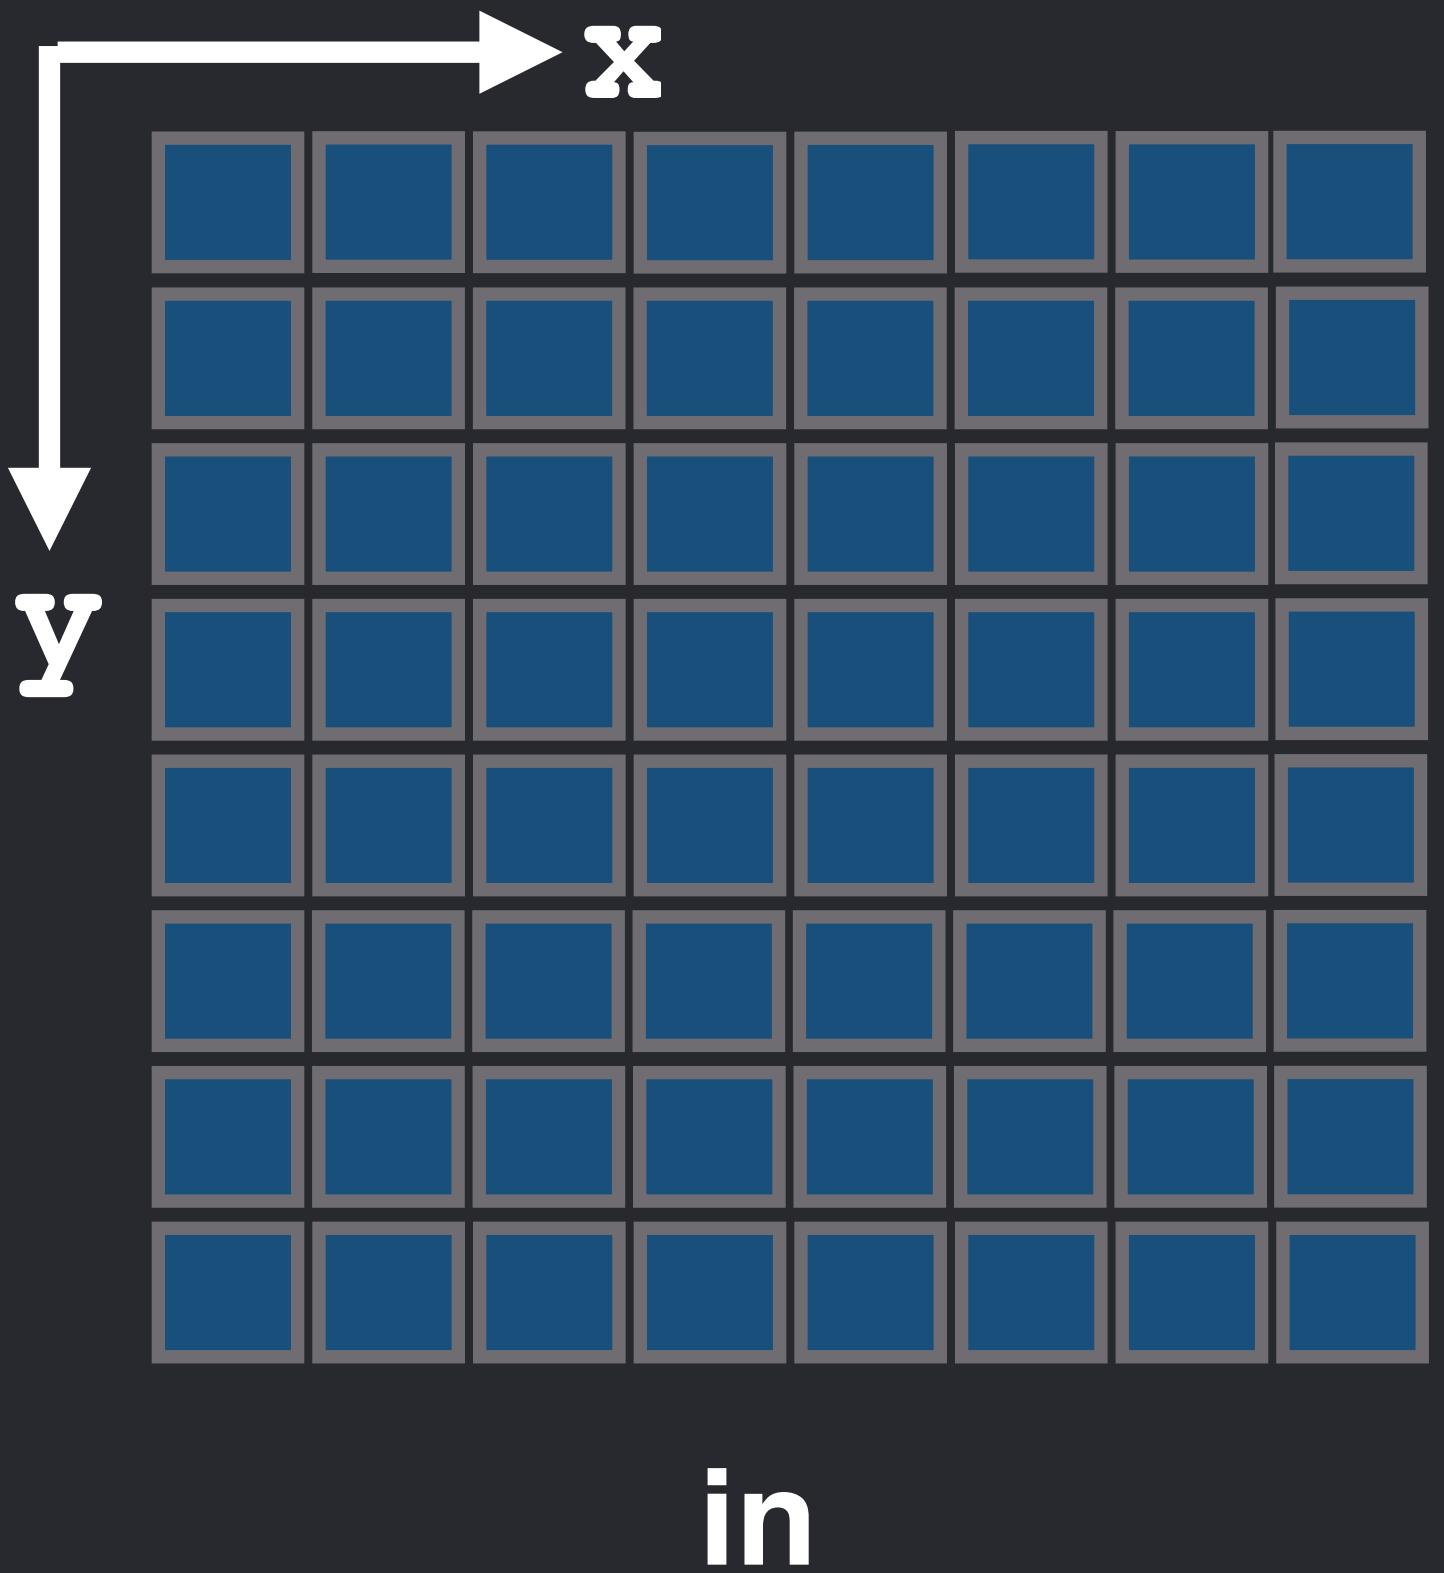
compute all pixels of bx , in parallel
compute all pixels of by , in parallel

A basic (slow) schedule

compute all pixels of bx , in parallel
compute all pixels of by , in parallel

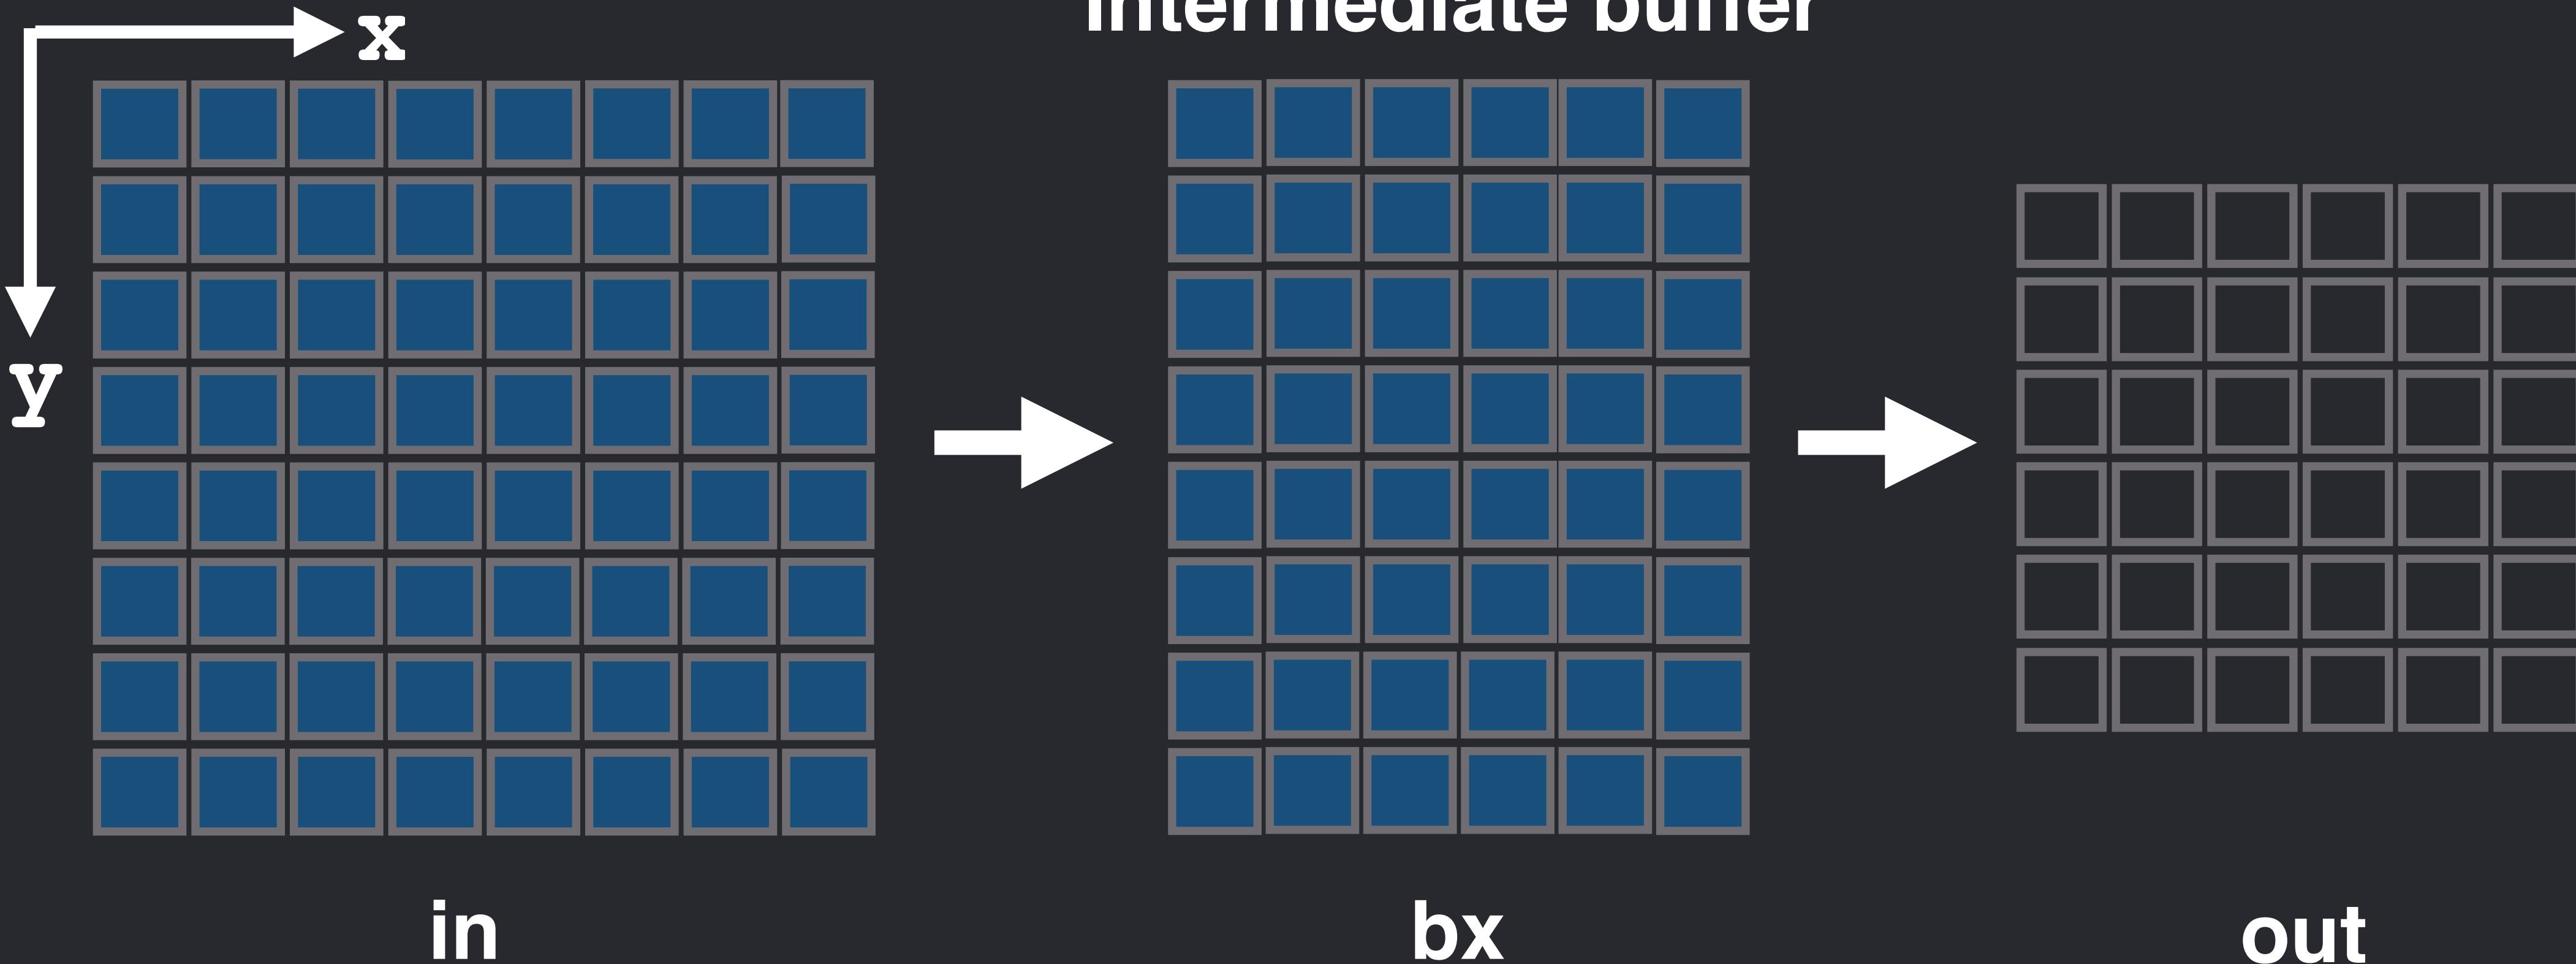


A basic (slow) schedule

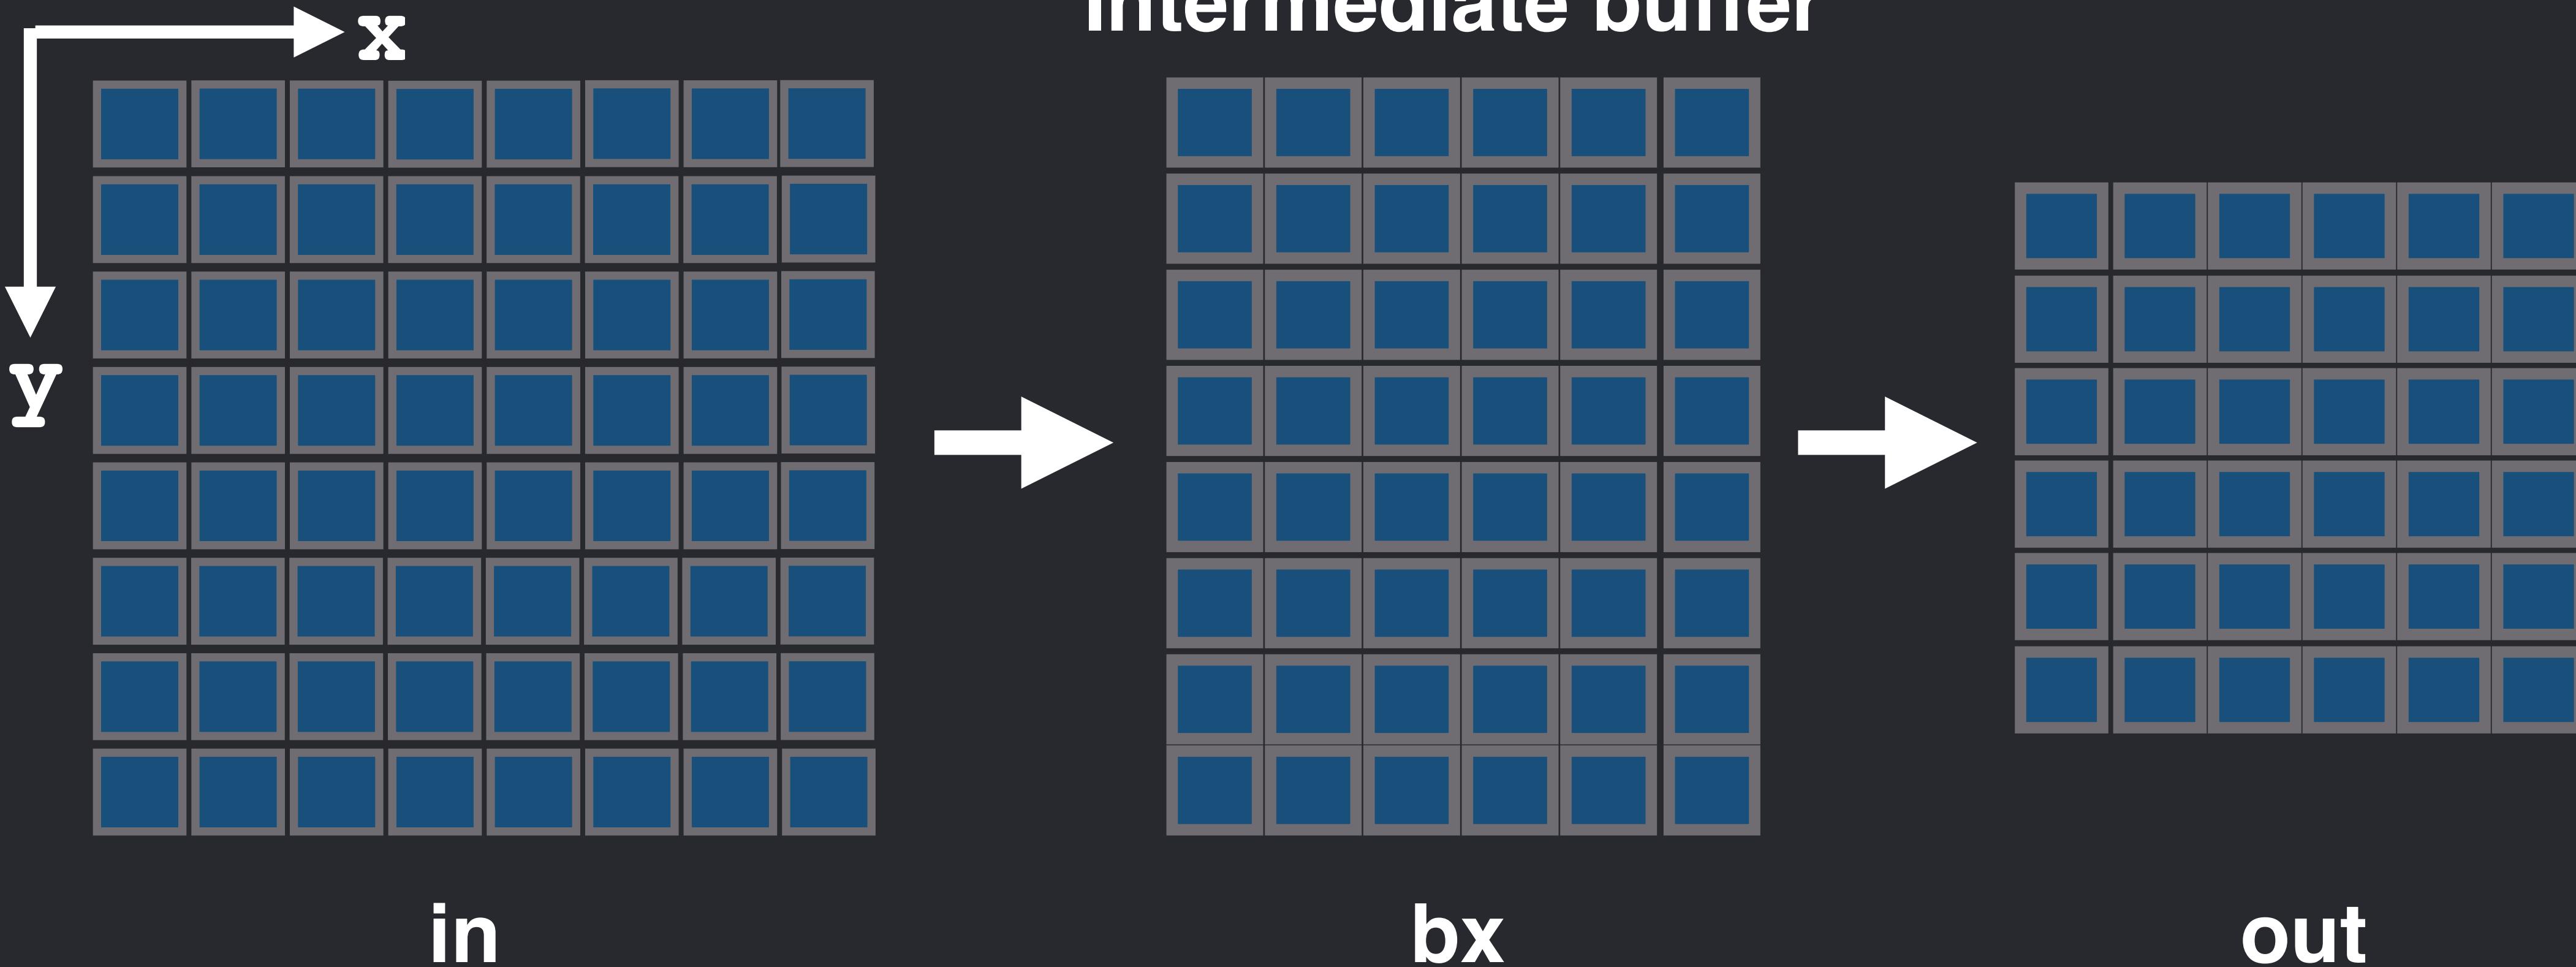


compute all pixels of bx , in parallel
compute all pixels of by , in parallel

A basic (slow) schedule

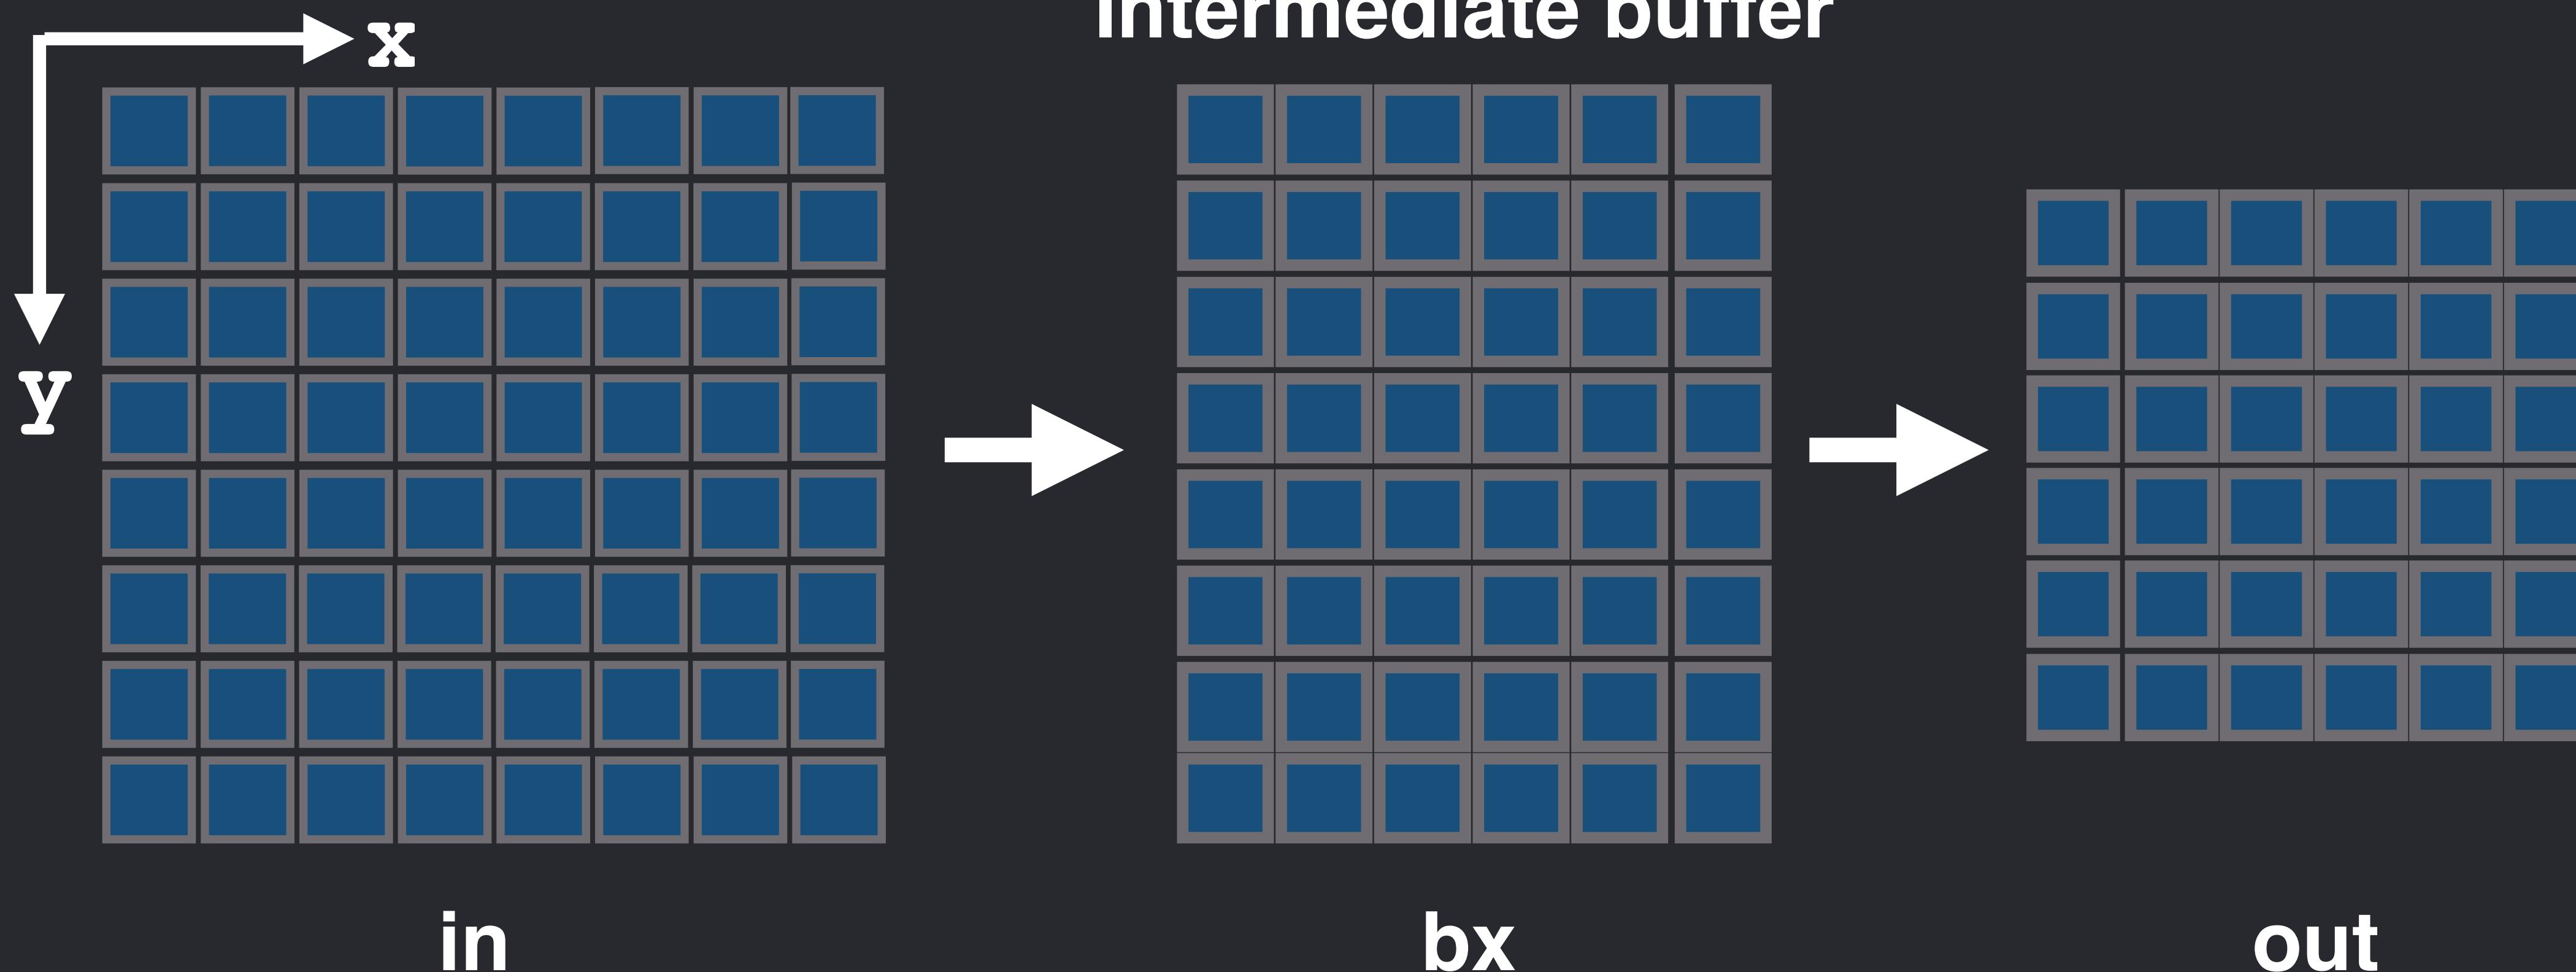


A basic (slow) schedule

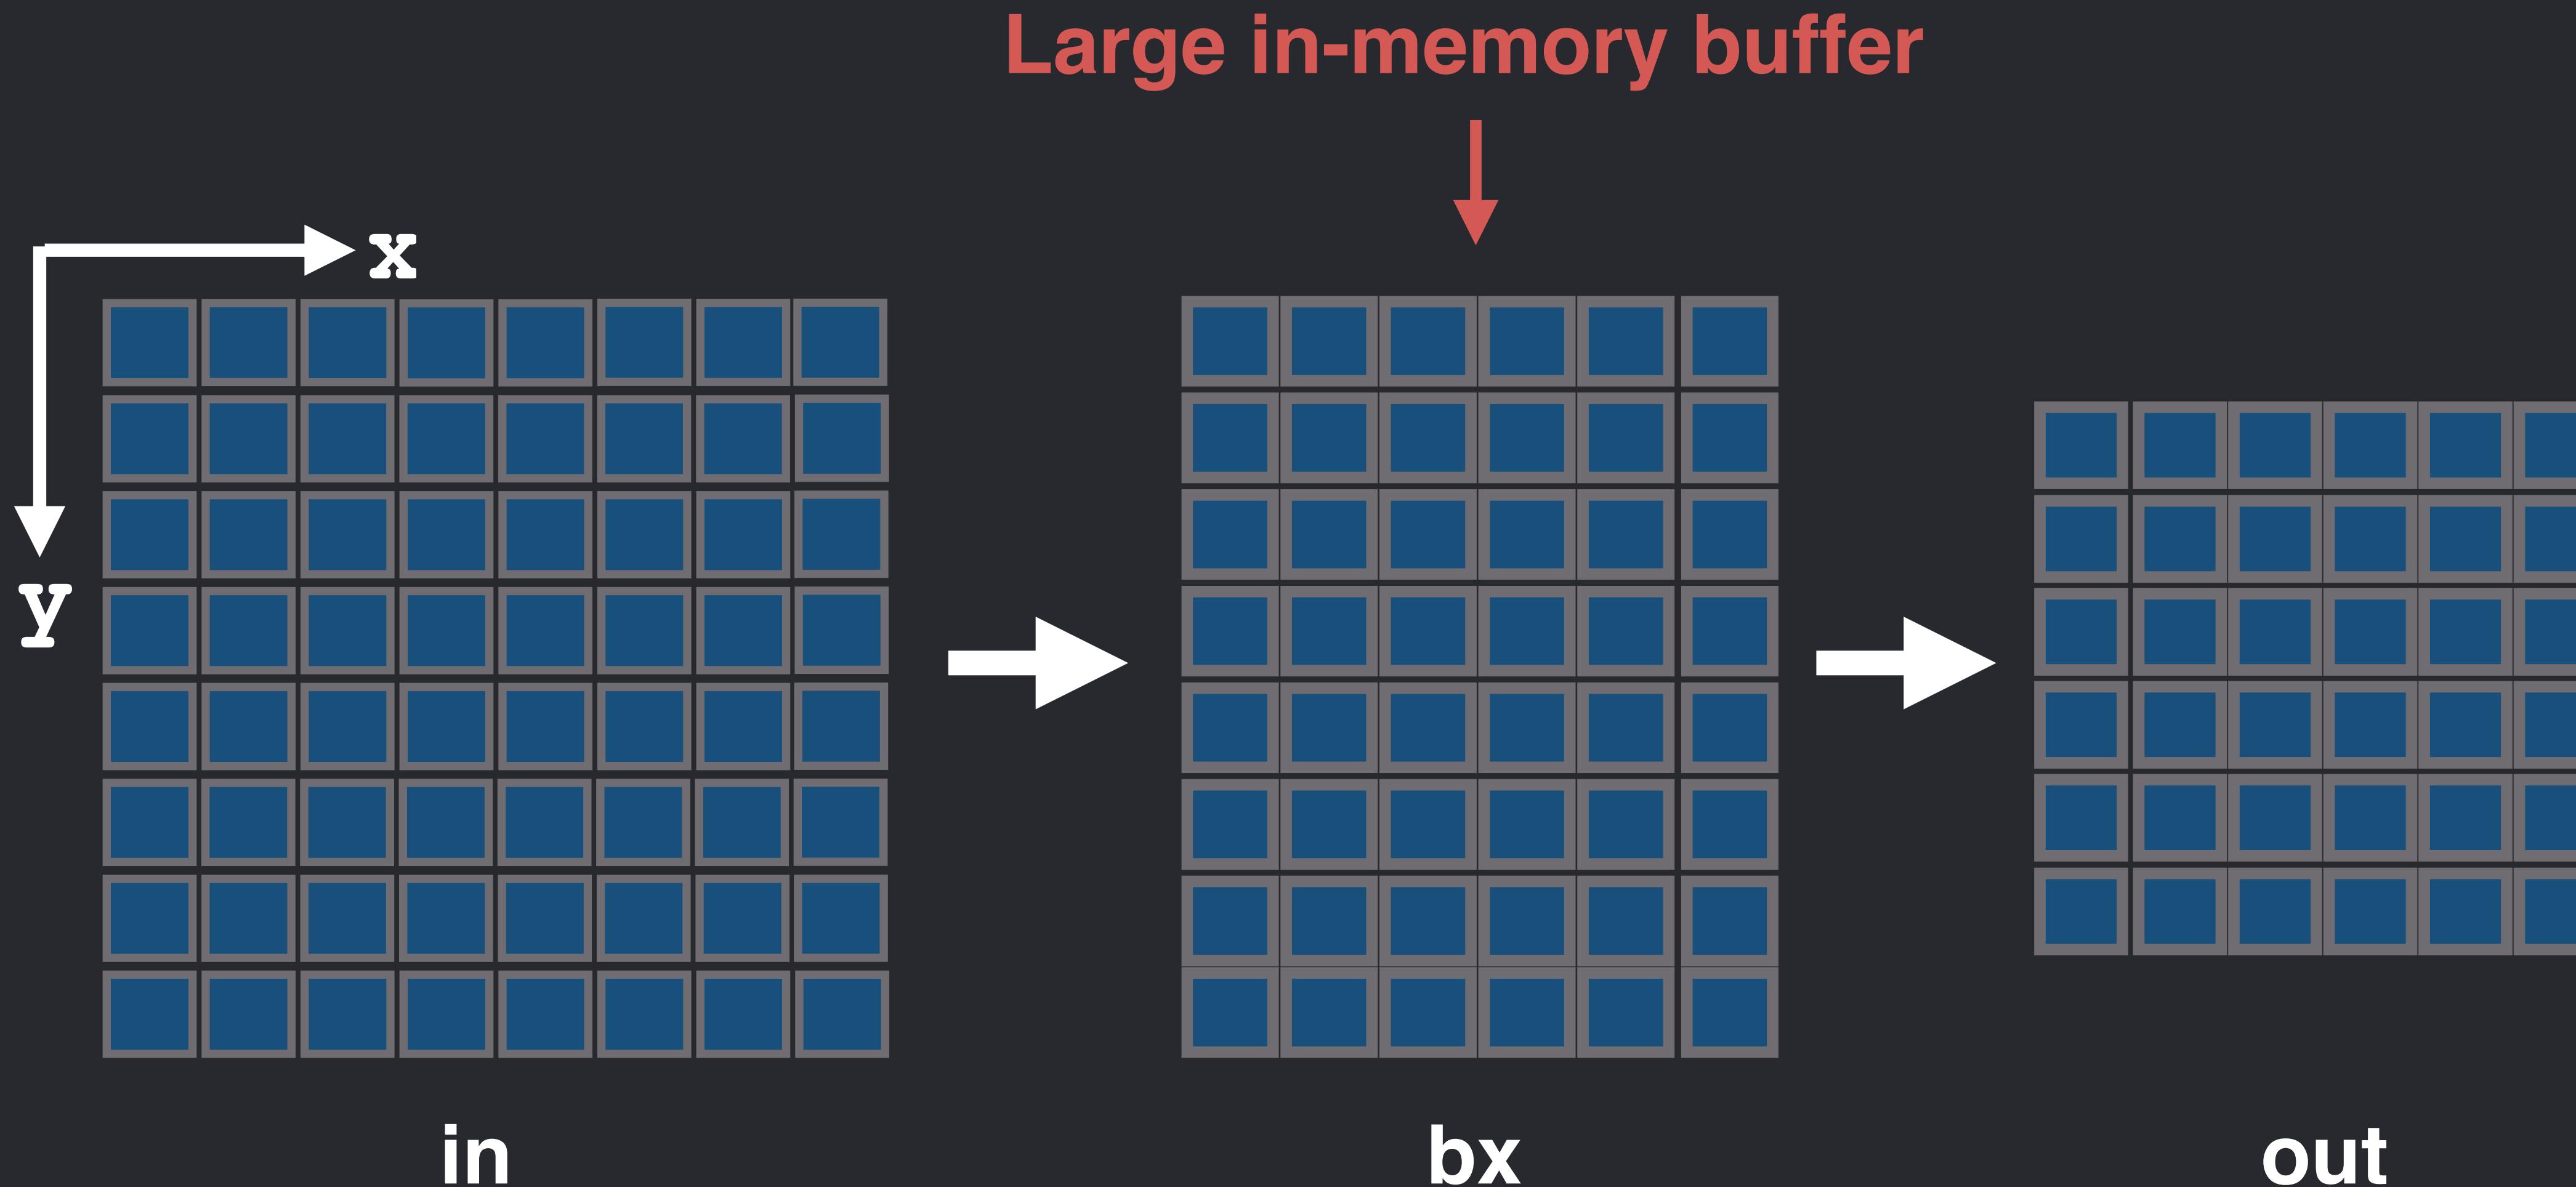


A basic (slow) schedule

compute all pixels of bx , in parallel
compute all pixels of by , in parallel

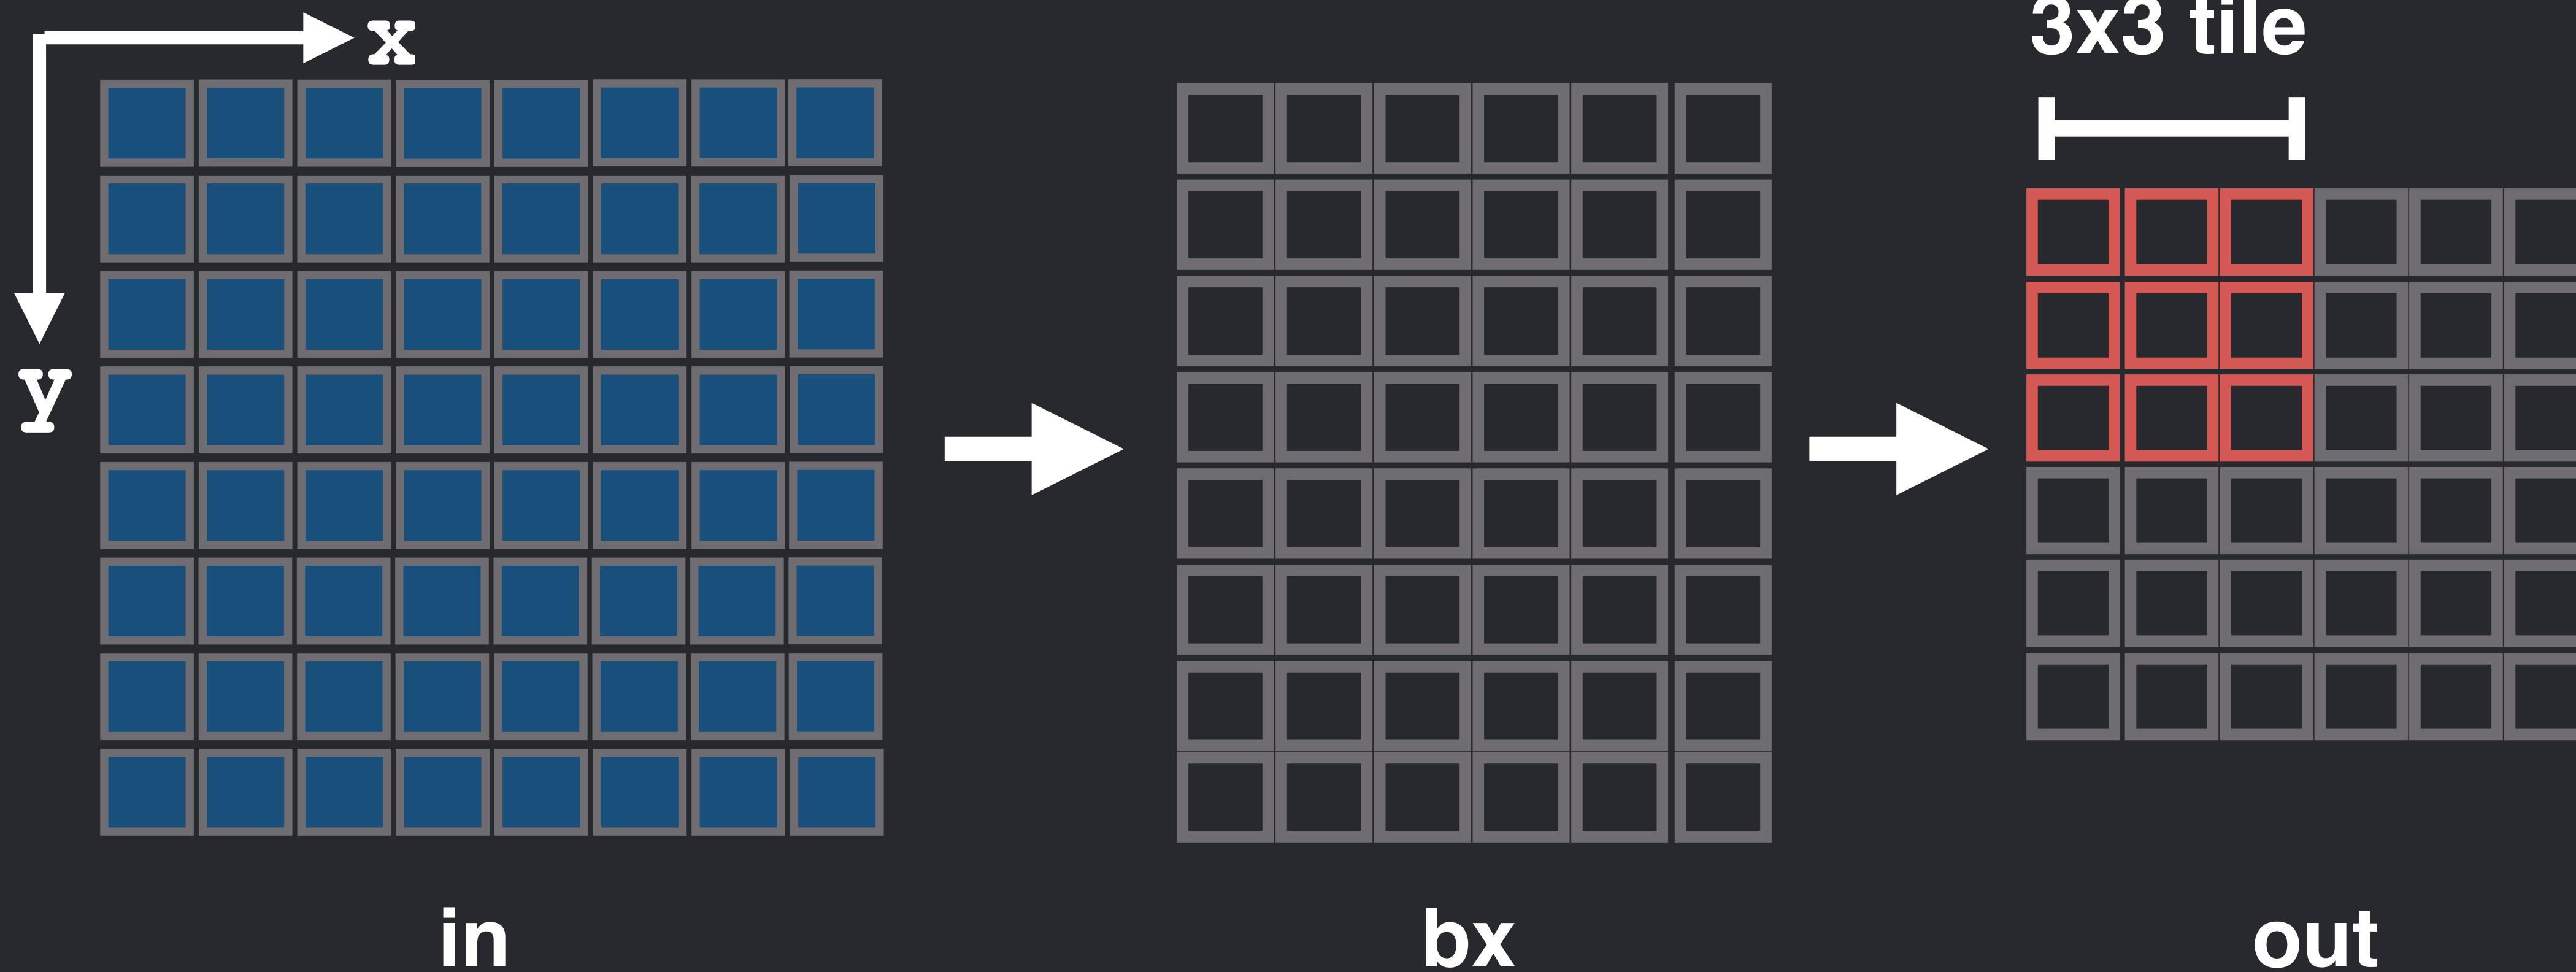


Low performance: bandwidth bound



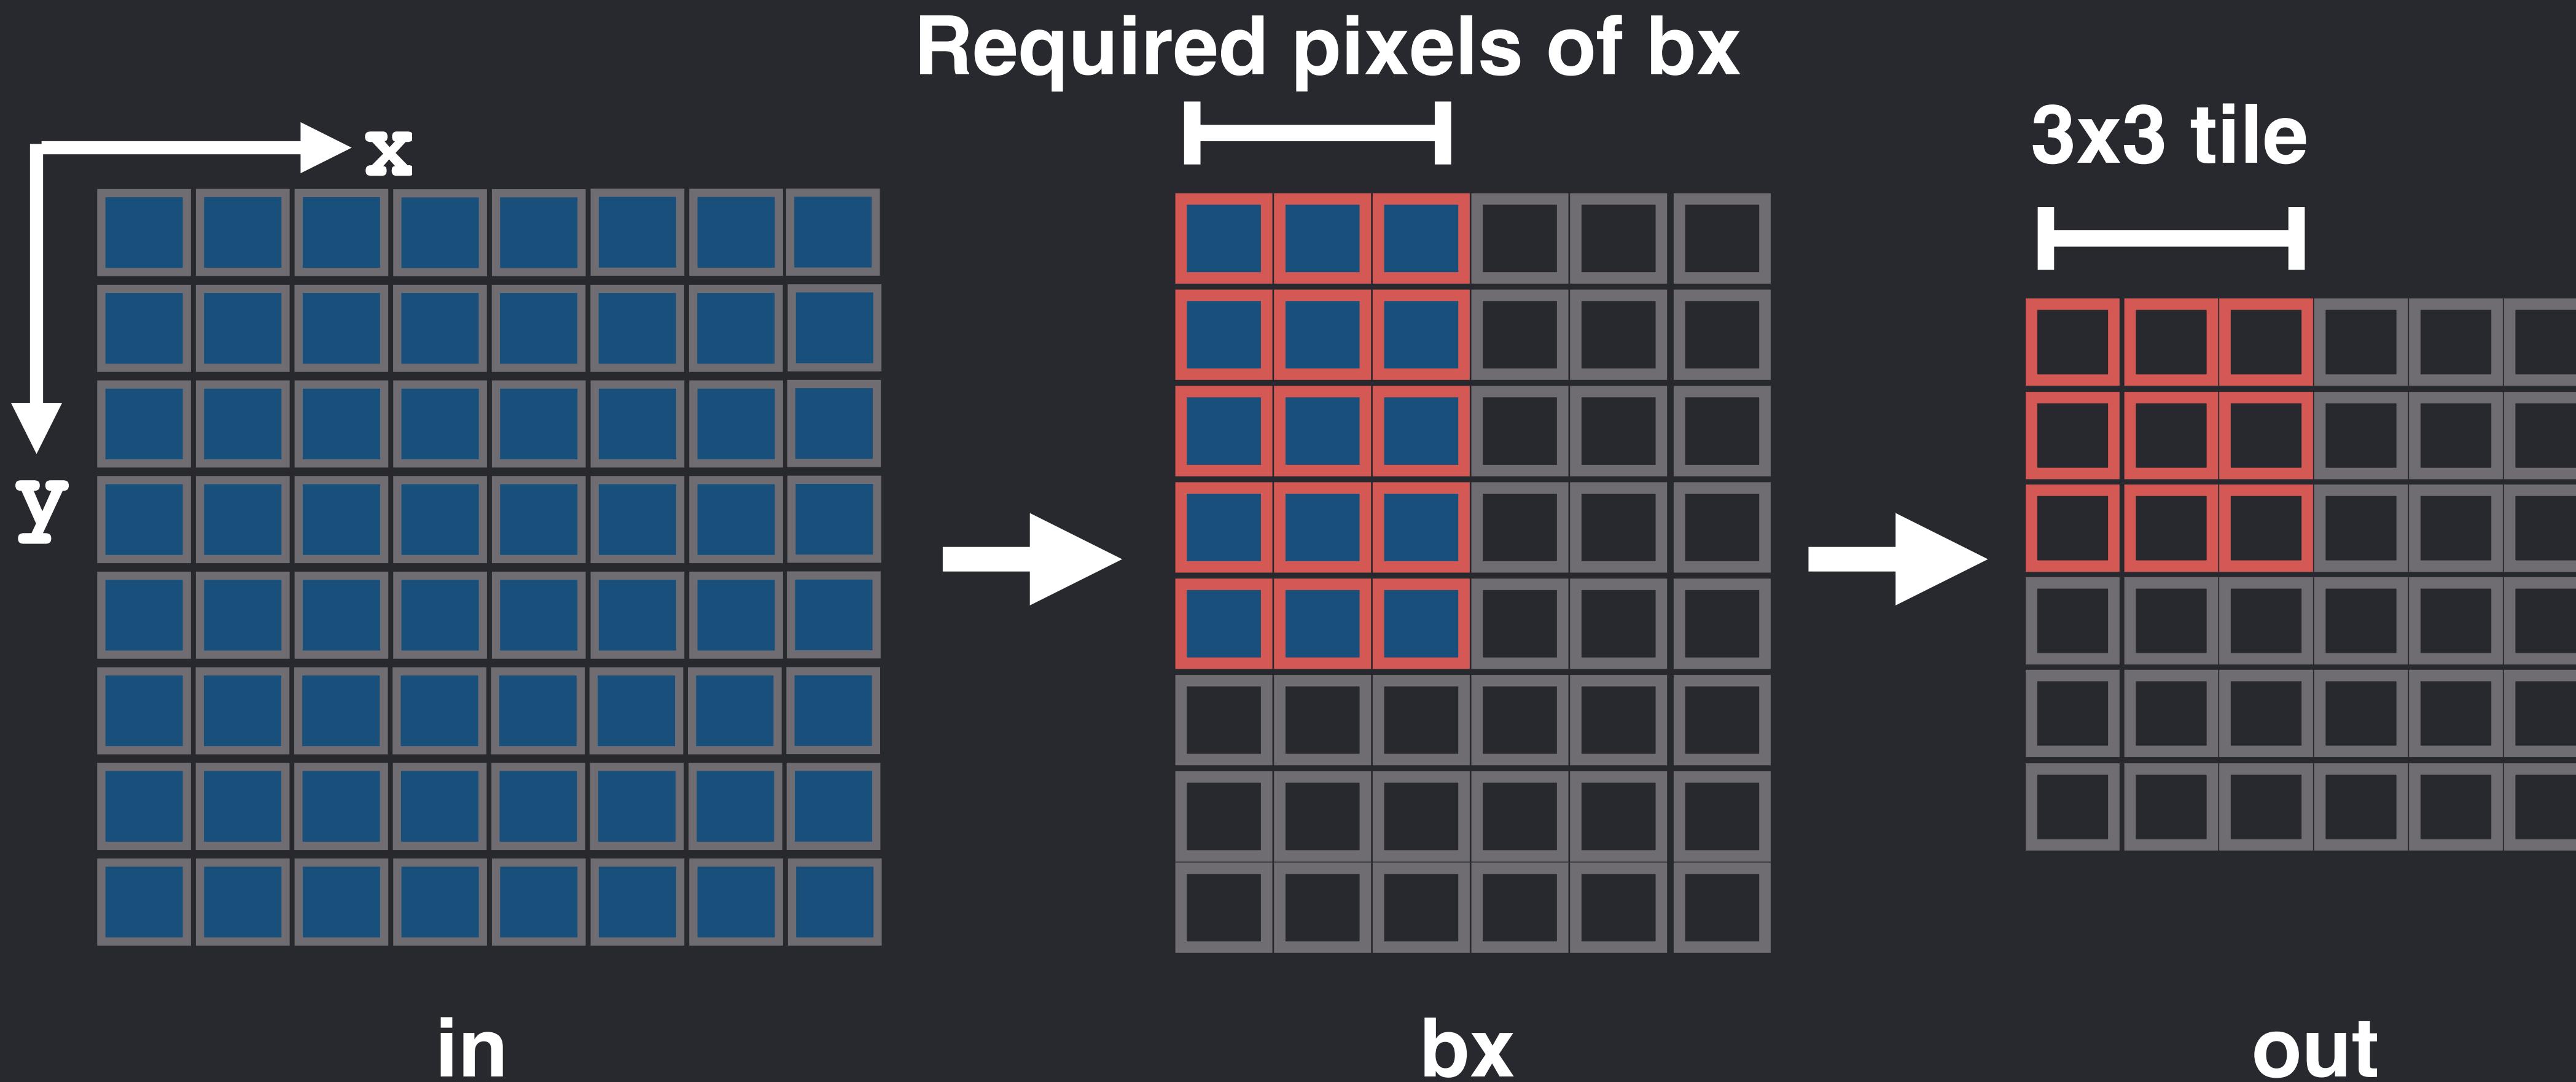
Tiling to improve data locality

for each 3x3 tile, in parallel
compute required pixels of bx
compute pixels of out in tile



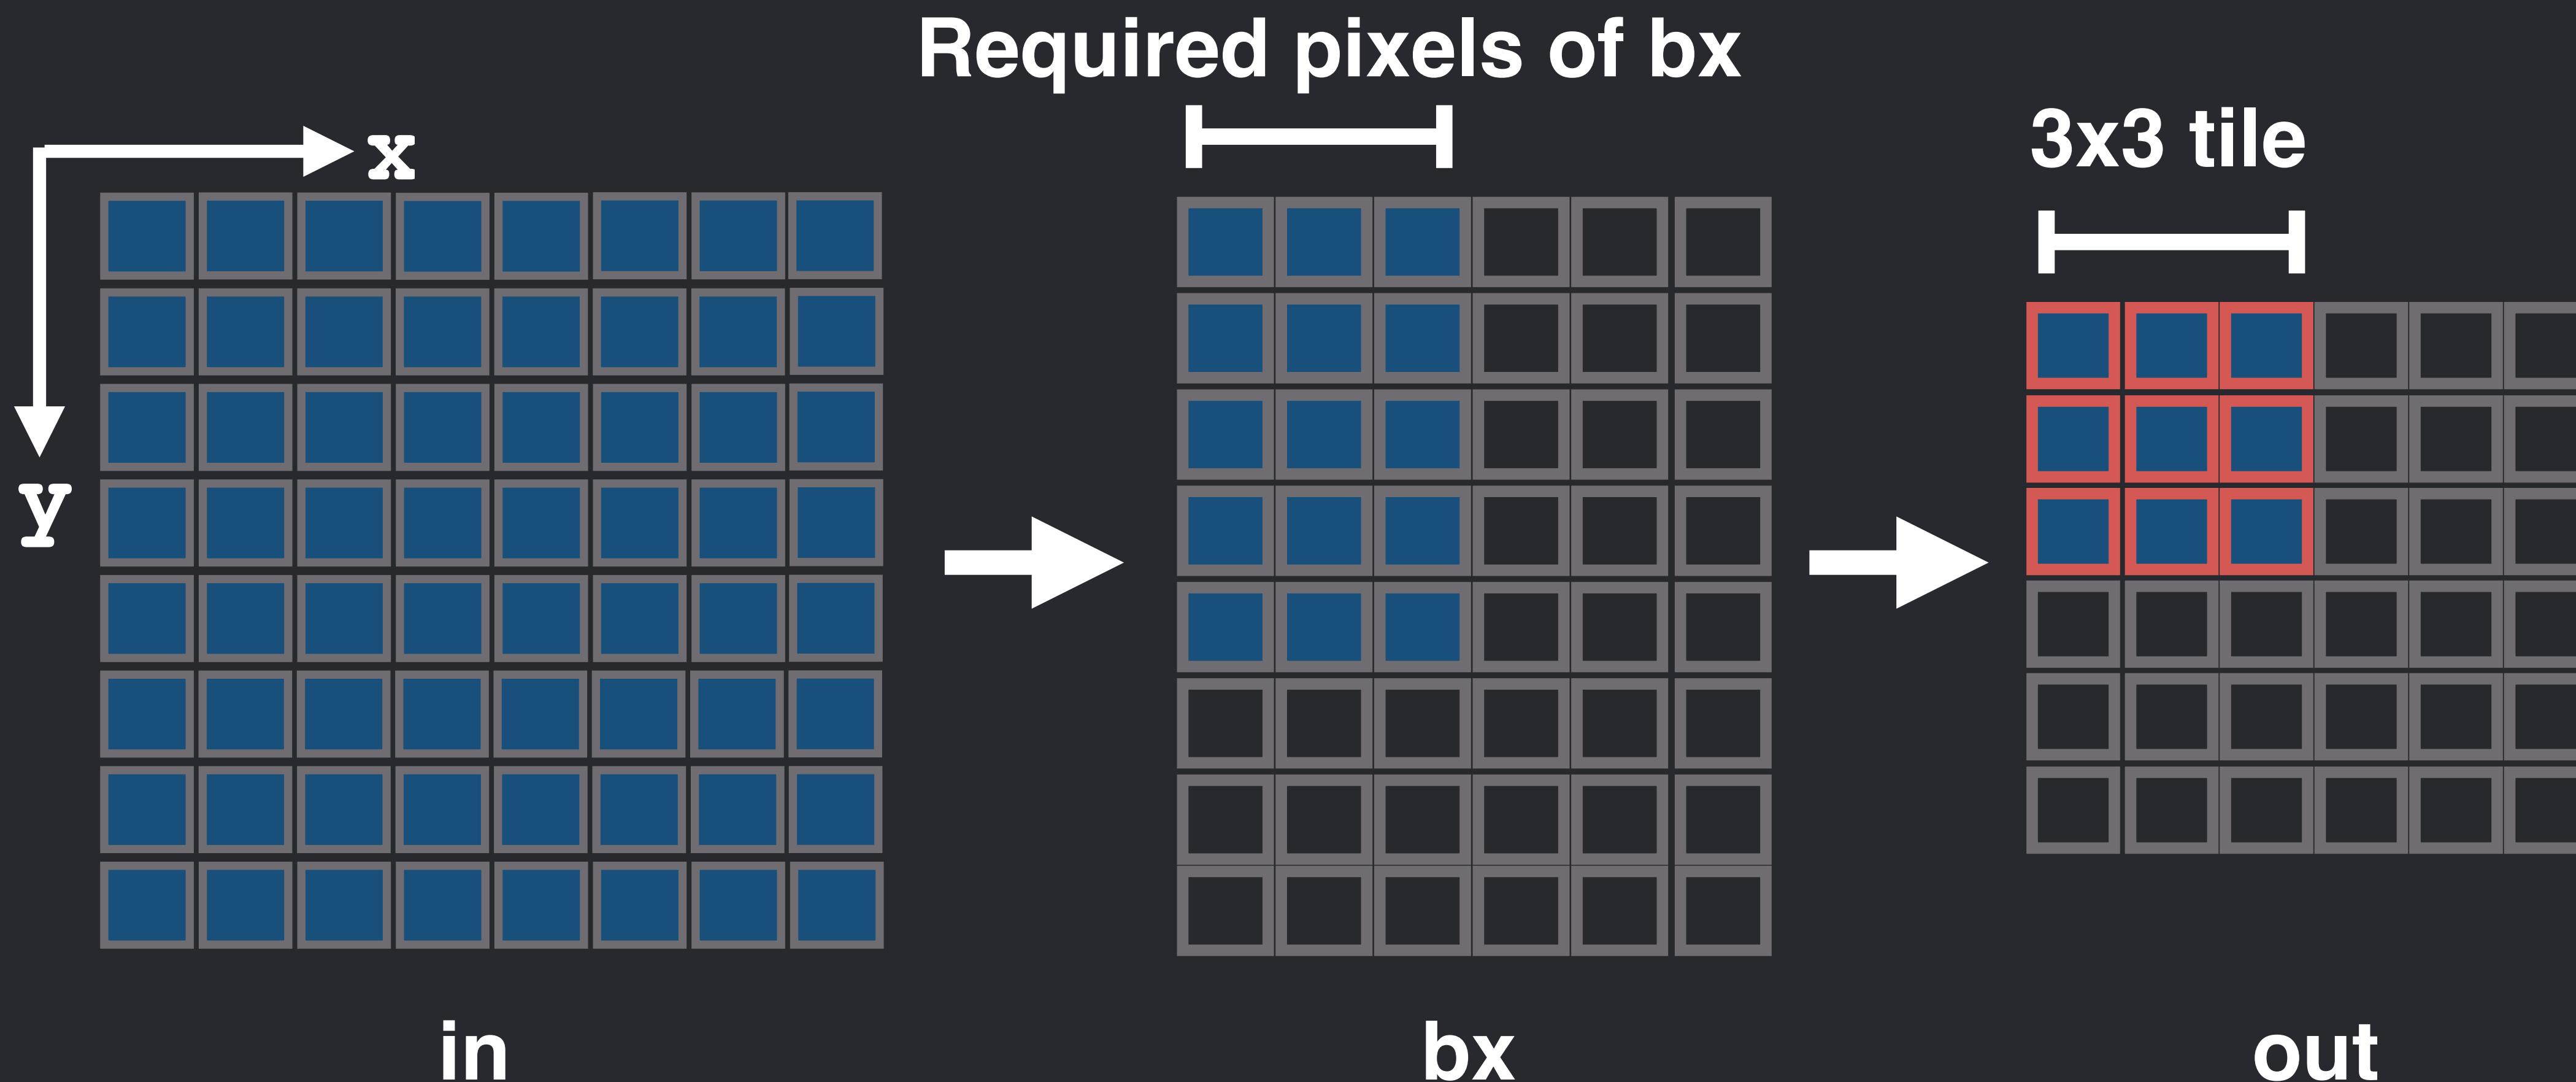
Tiling to improve data locality

for each 3x3 tile, in parallel
compute required pixels of bx
compute pixels of out in tile



Tiling to improve data locality

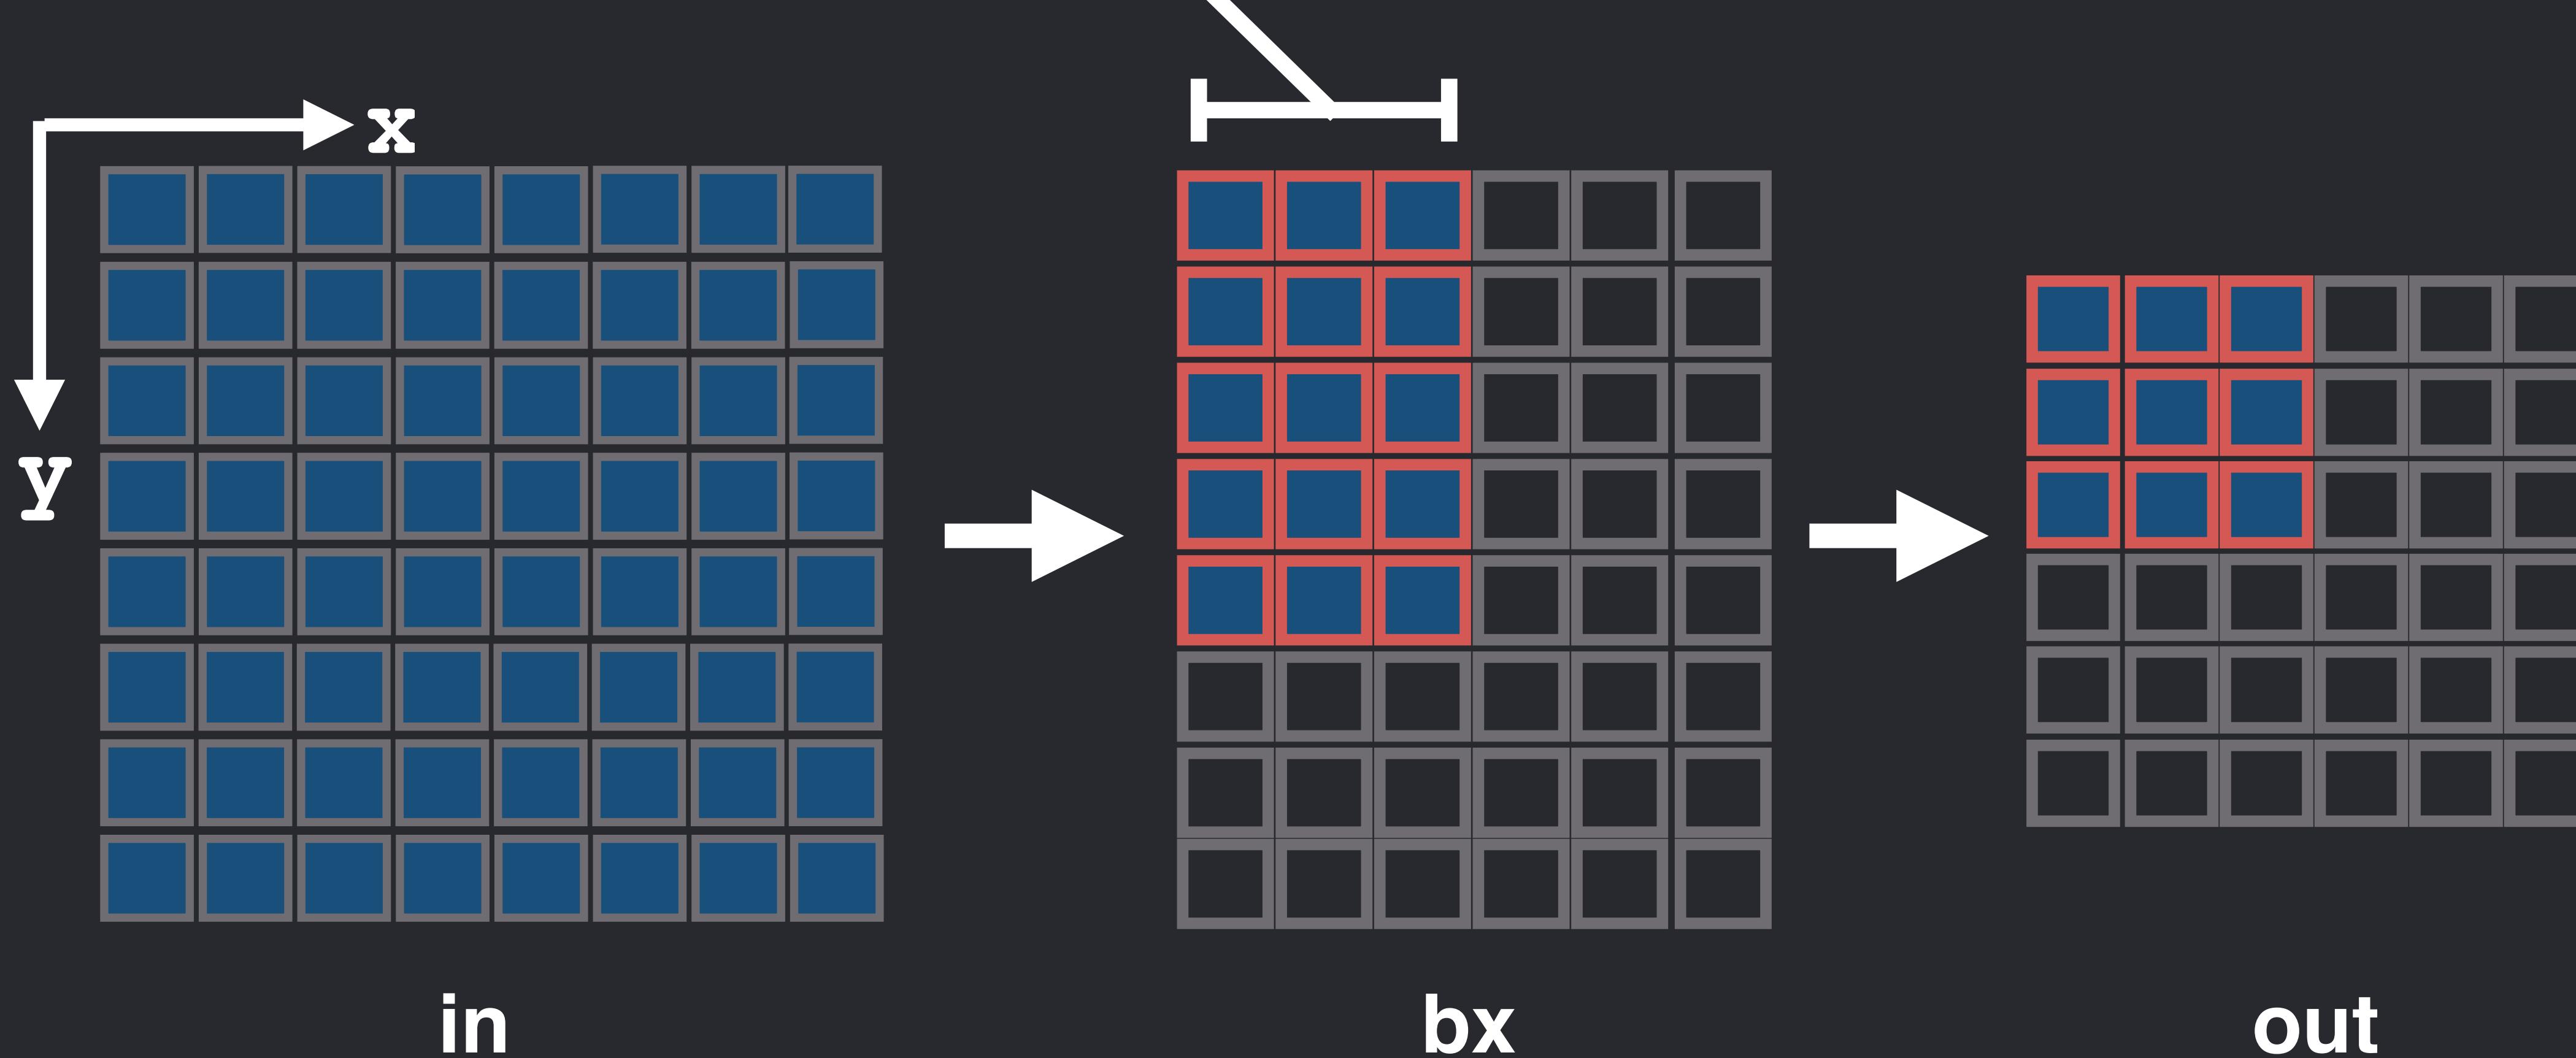
for each 3x3 tile, in parallel
compute required pixels of bx
compute pixels of out in tile



Tiling to improve data locality

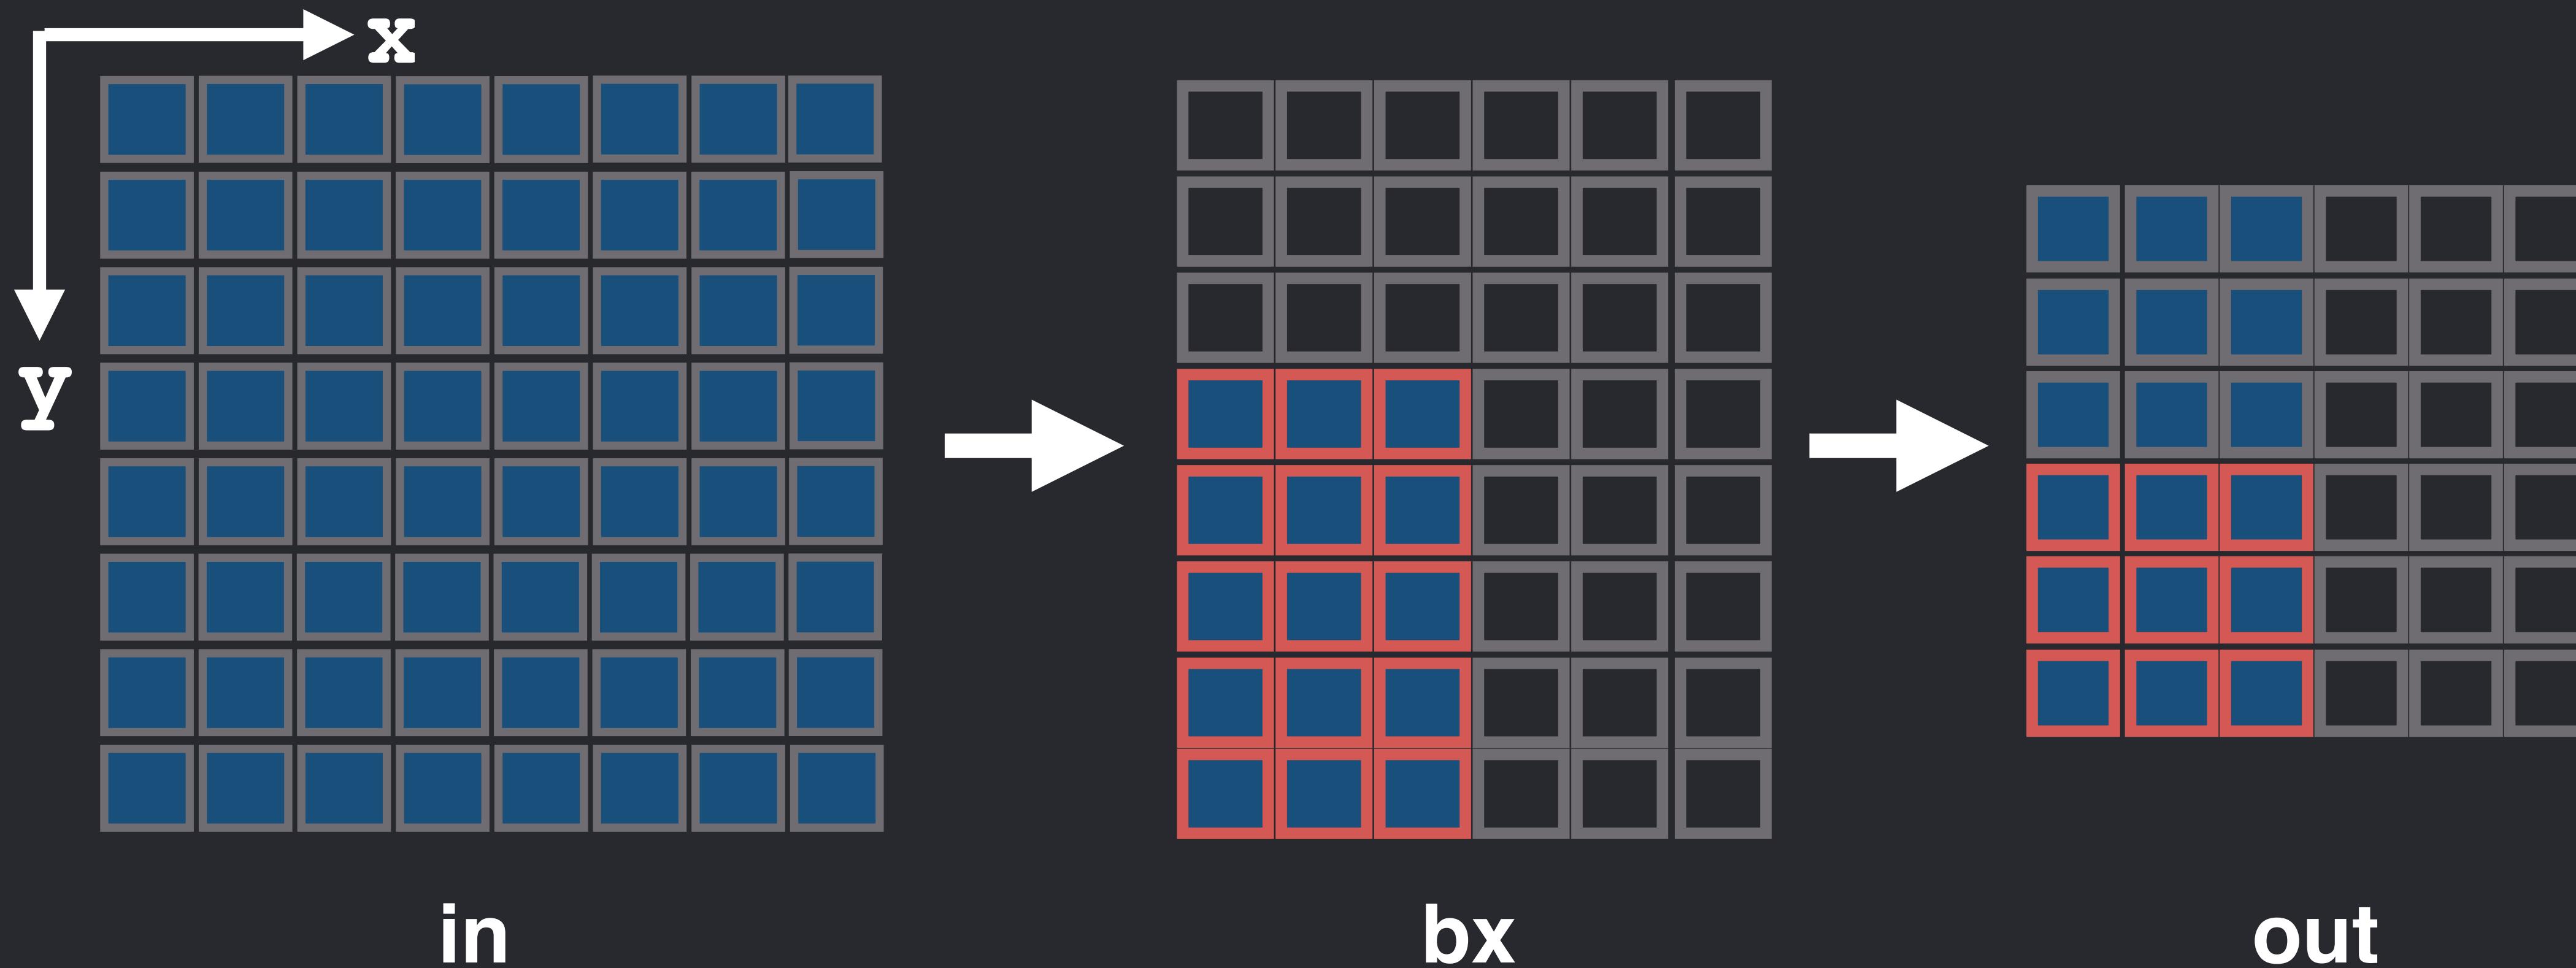
Intermediate buffer:
fits in fast on-chip storage

for each 3x3 tile, in parallel
compute required pixels of bx
compute pixels of out in tile



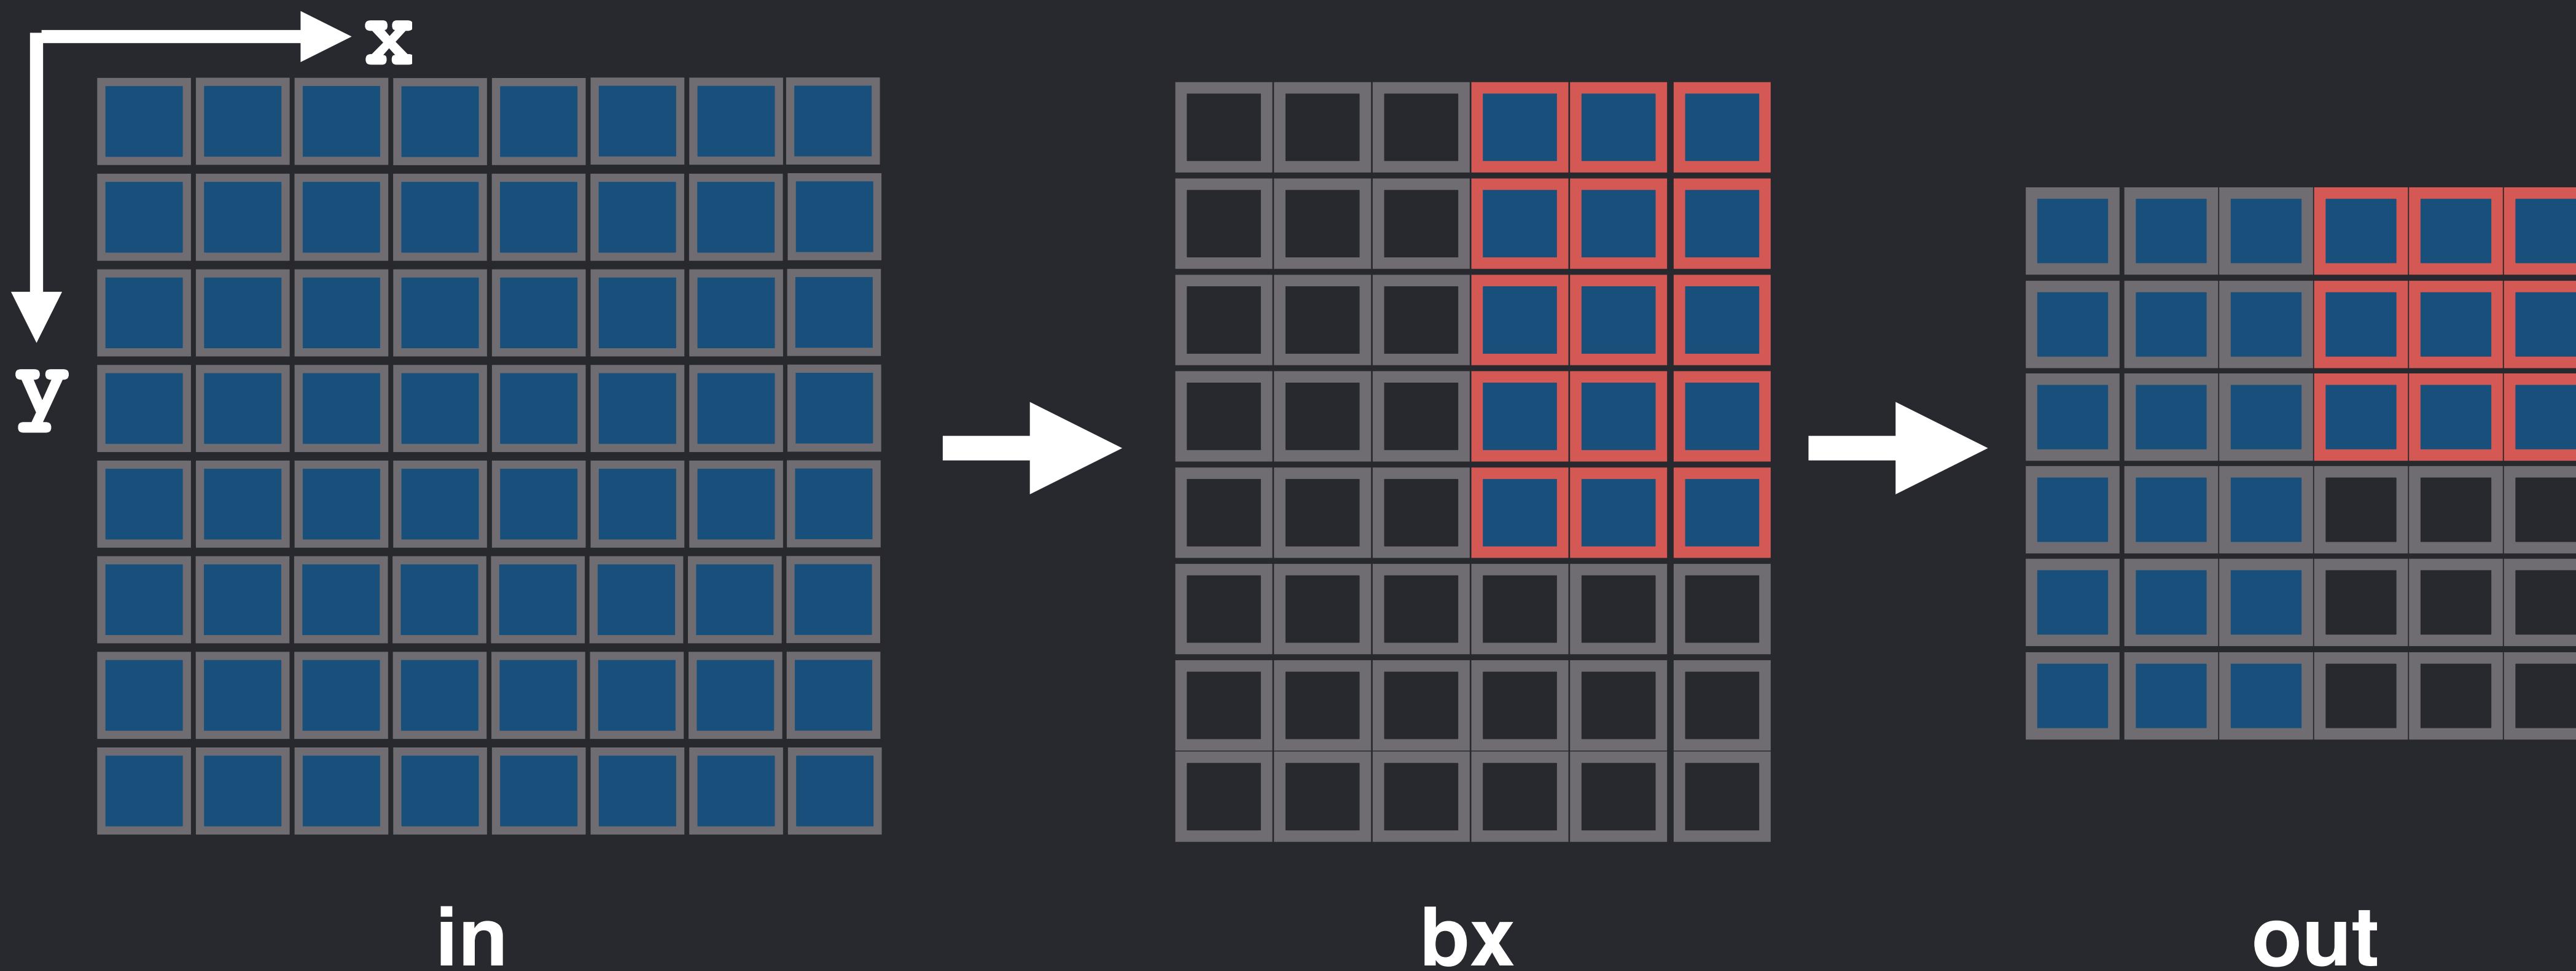
Tiling to improve data locality

for each 3x3 tile, in parallel
compute required pixels of bx
compute pixels of out in tile



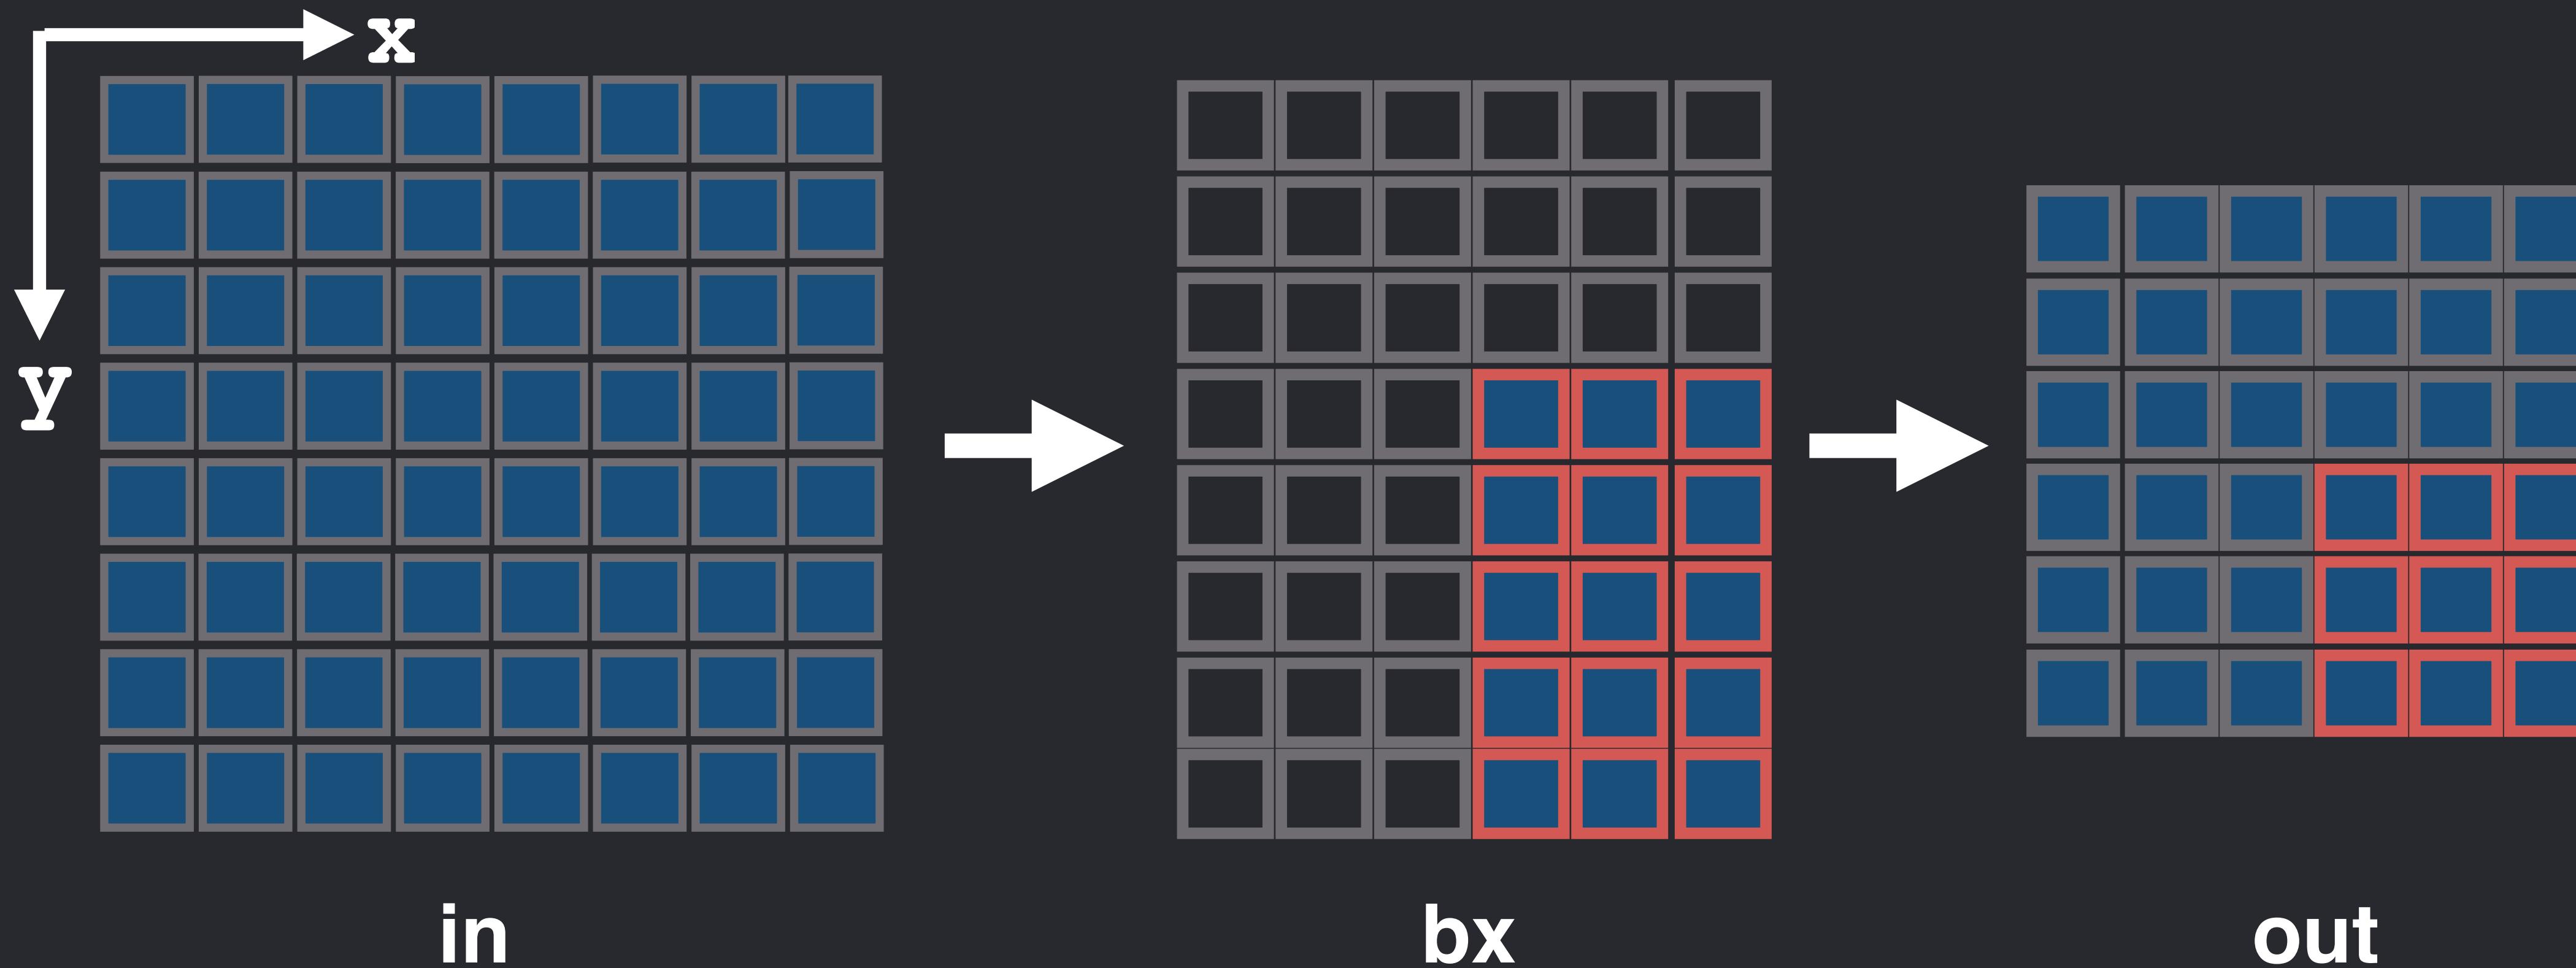
Tiling to improve data locality

for each 3x3 tile, in parallel
compute required pixels of bx
compute pixels of out in tile



Tiling to improve data locality

for each 3x3 tile, in parallel
compute required pixels of bx
compute pixels of out in tile

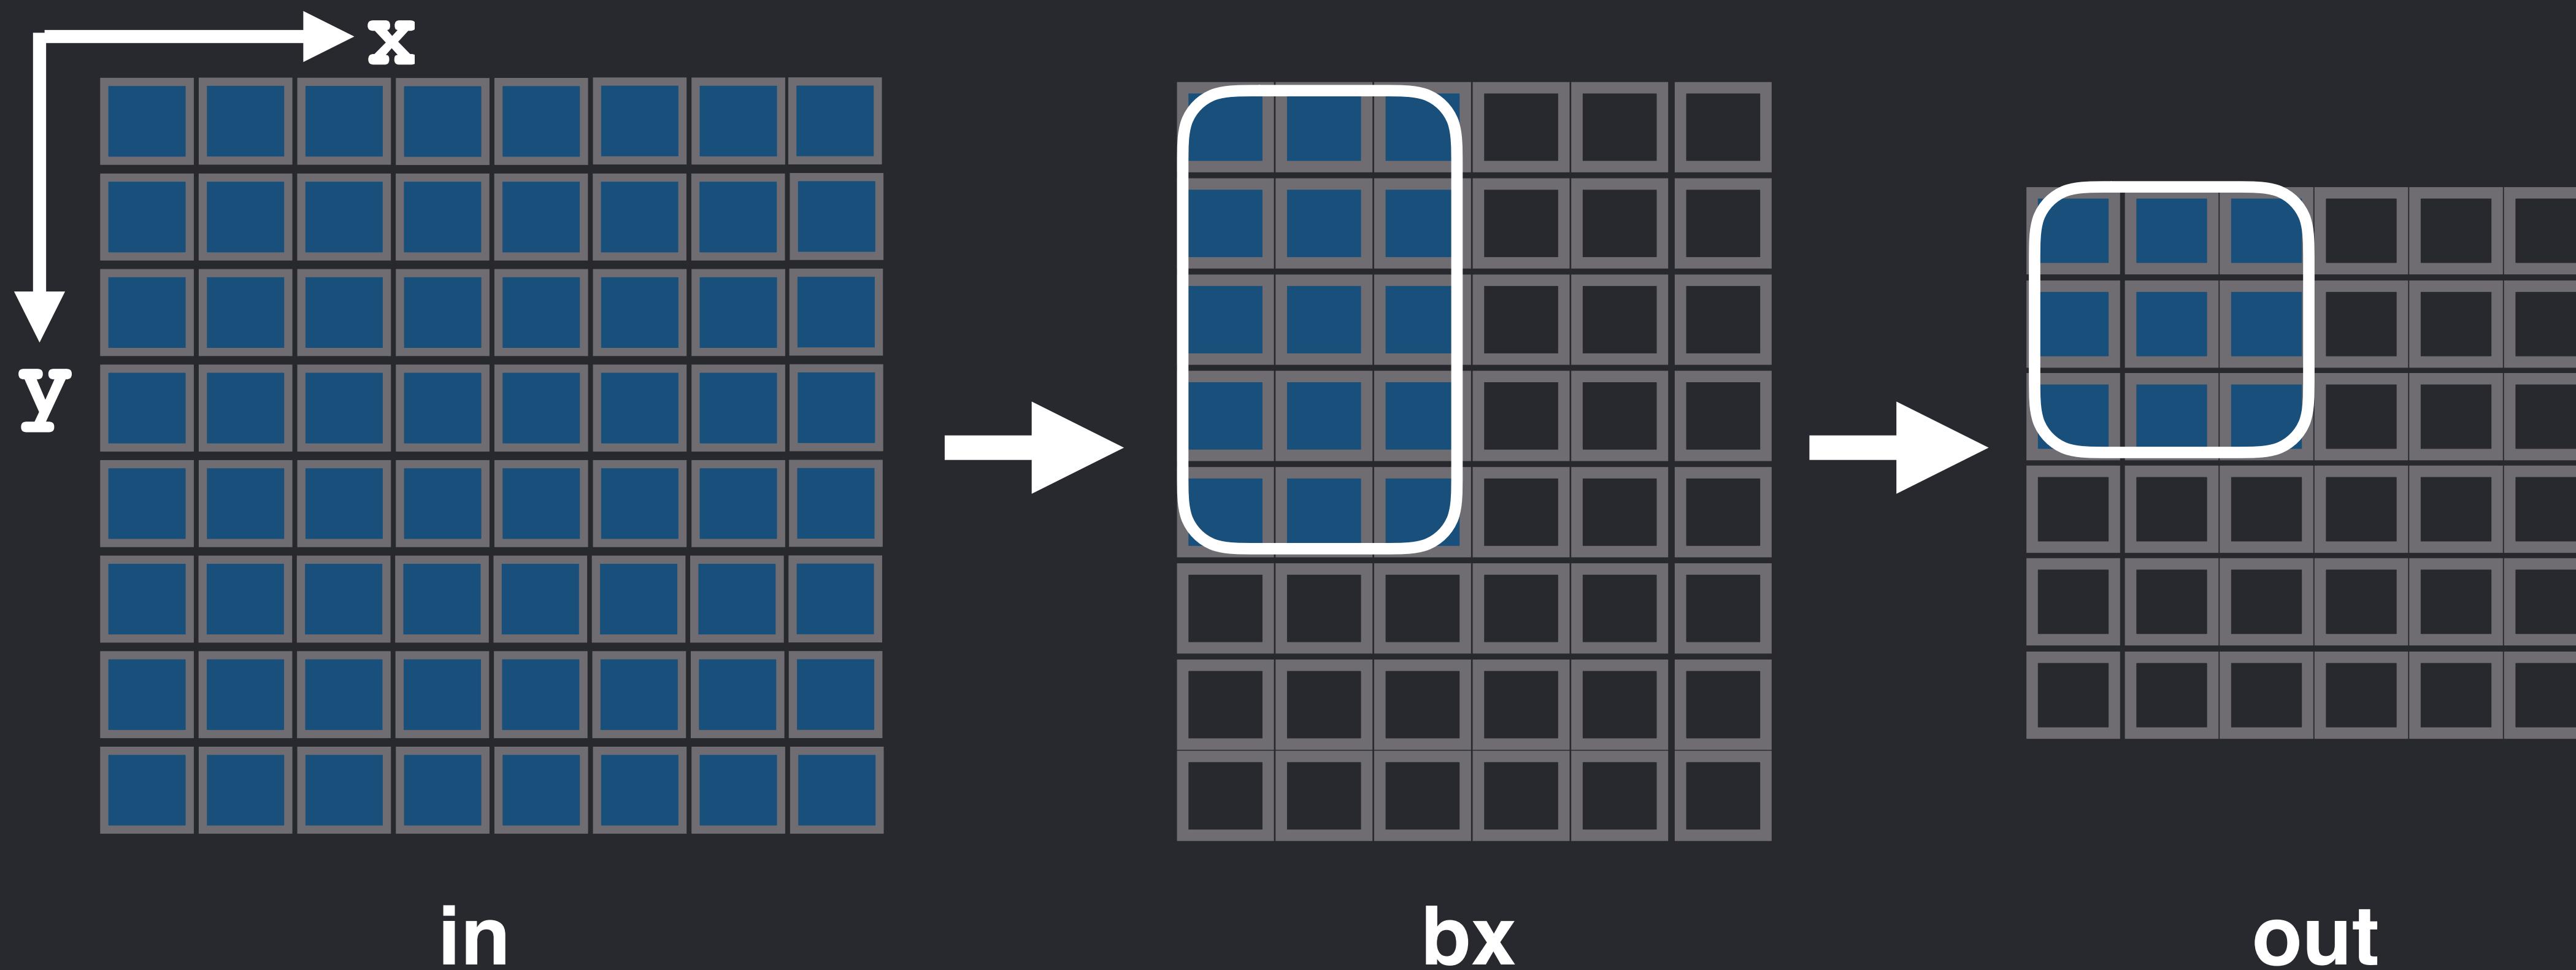


Tiling to improve data locality

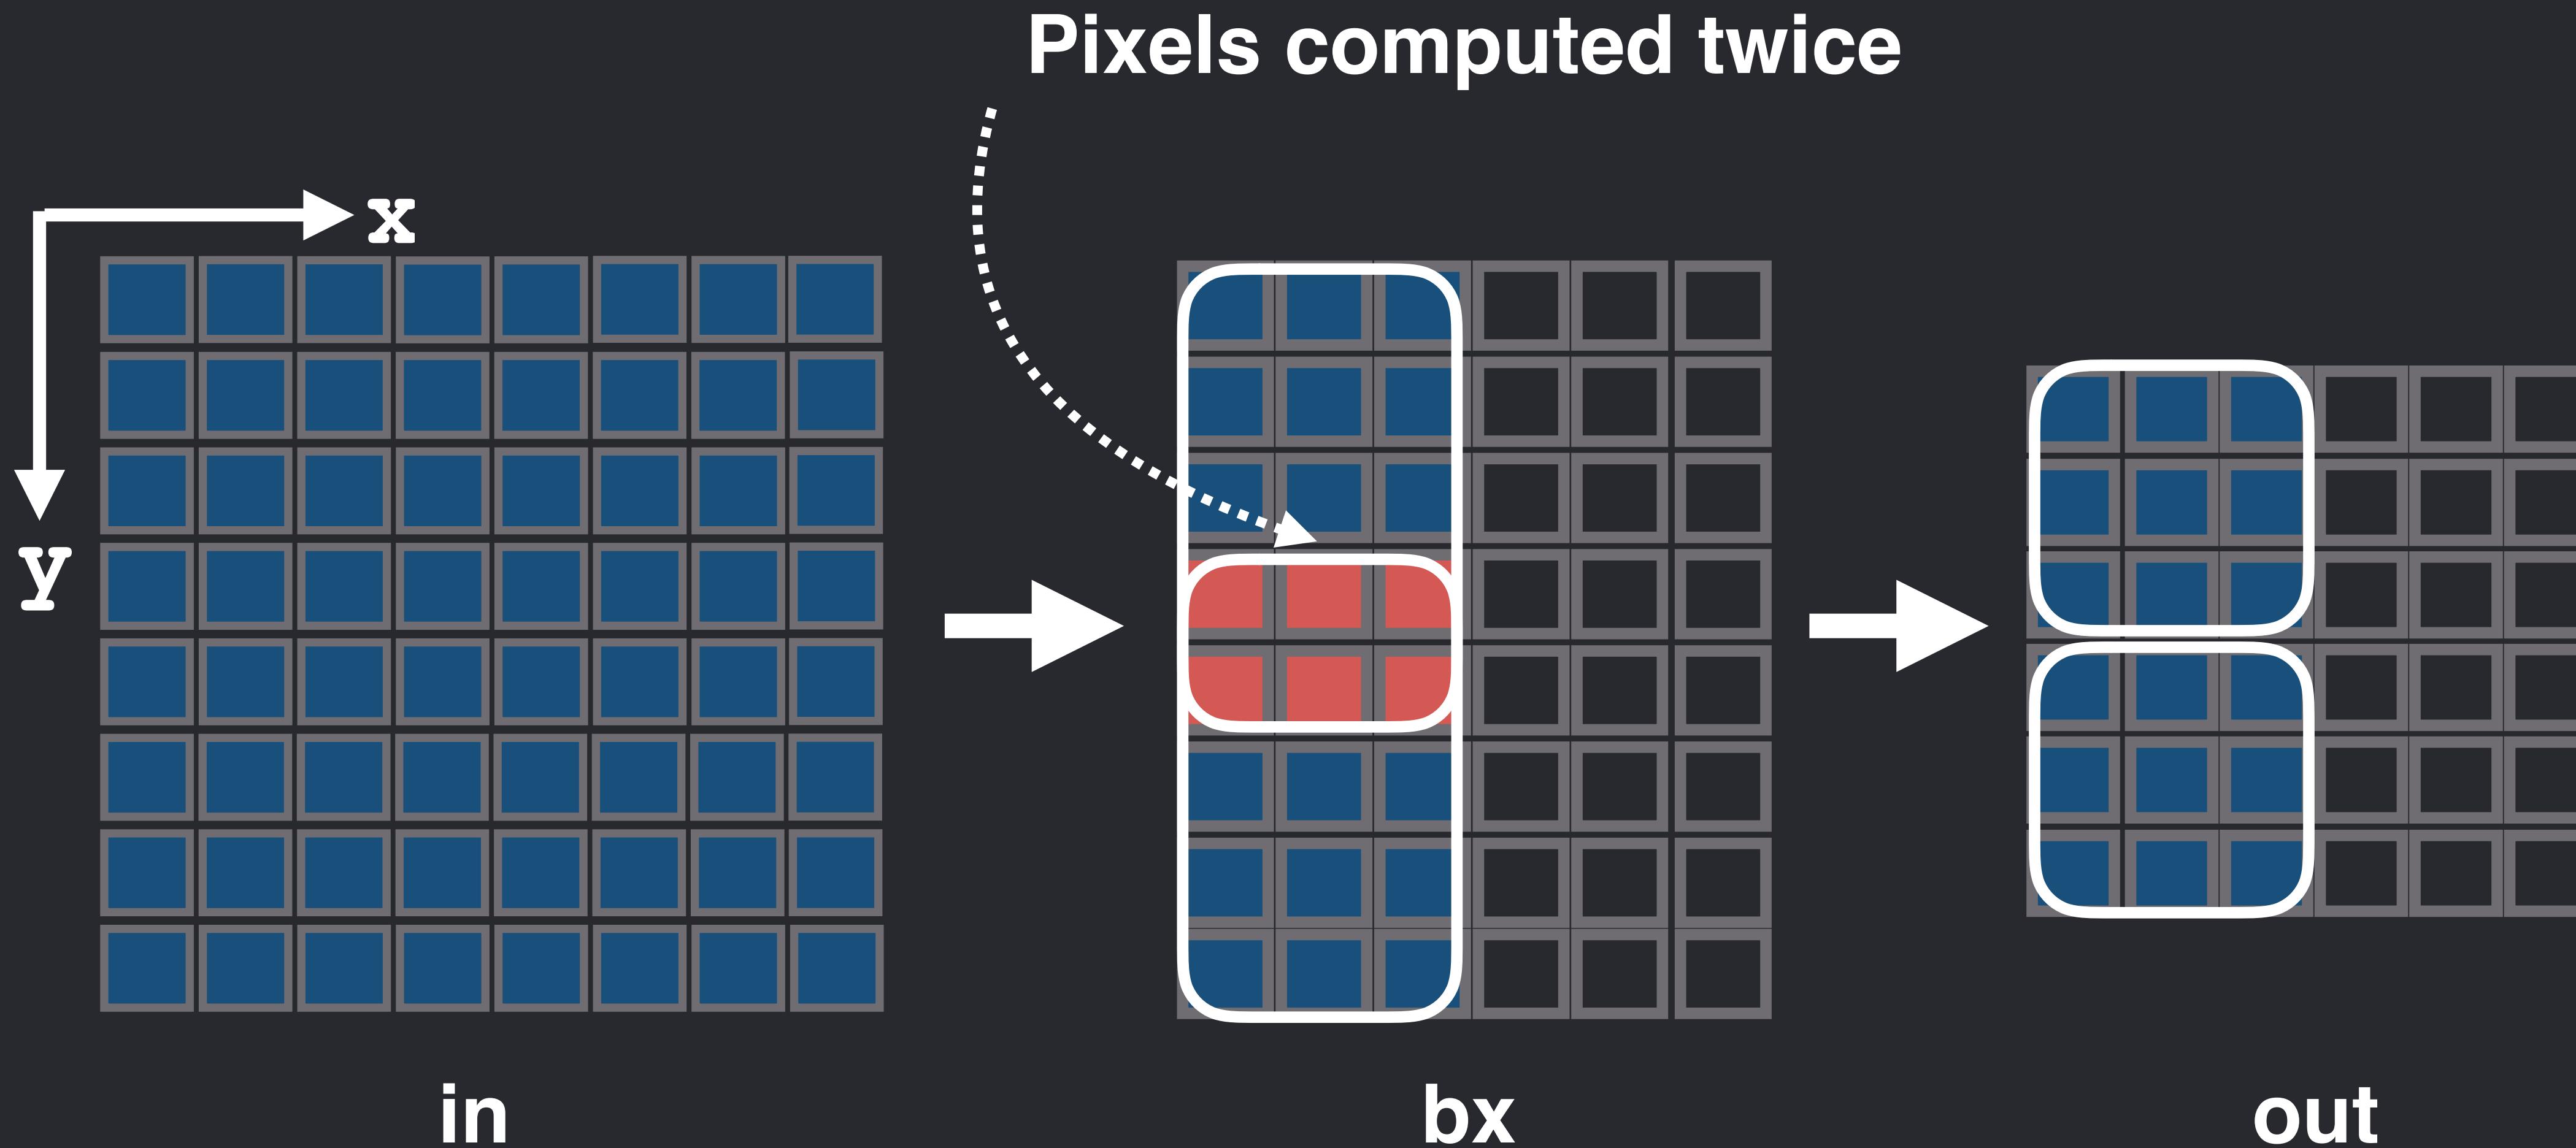
for each 3x3 tile, in parallel
compute required pixels of bx
compute pixels of out in tile



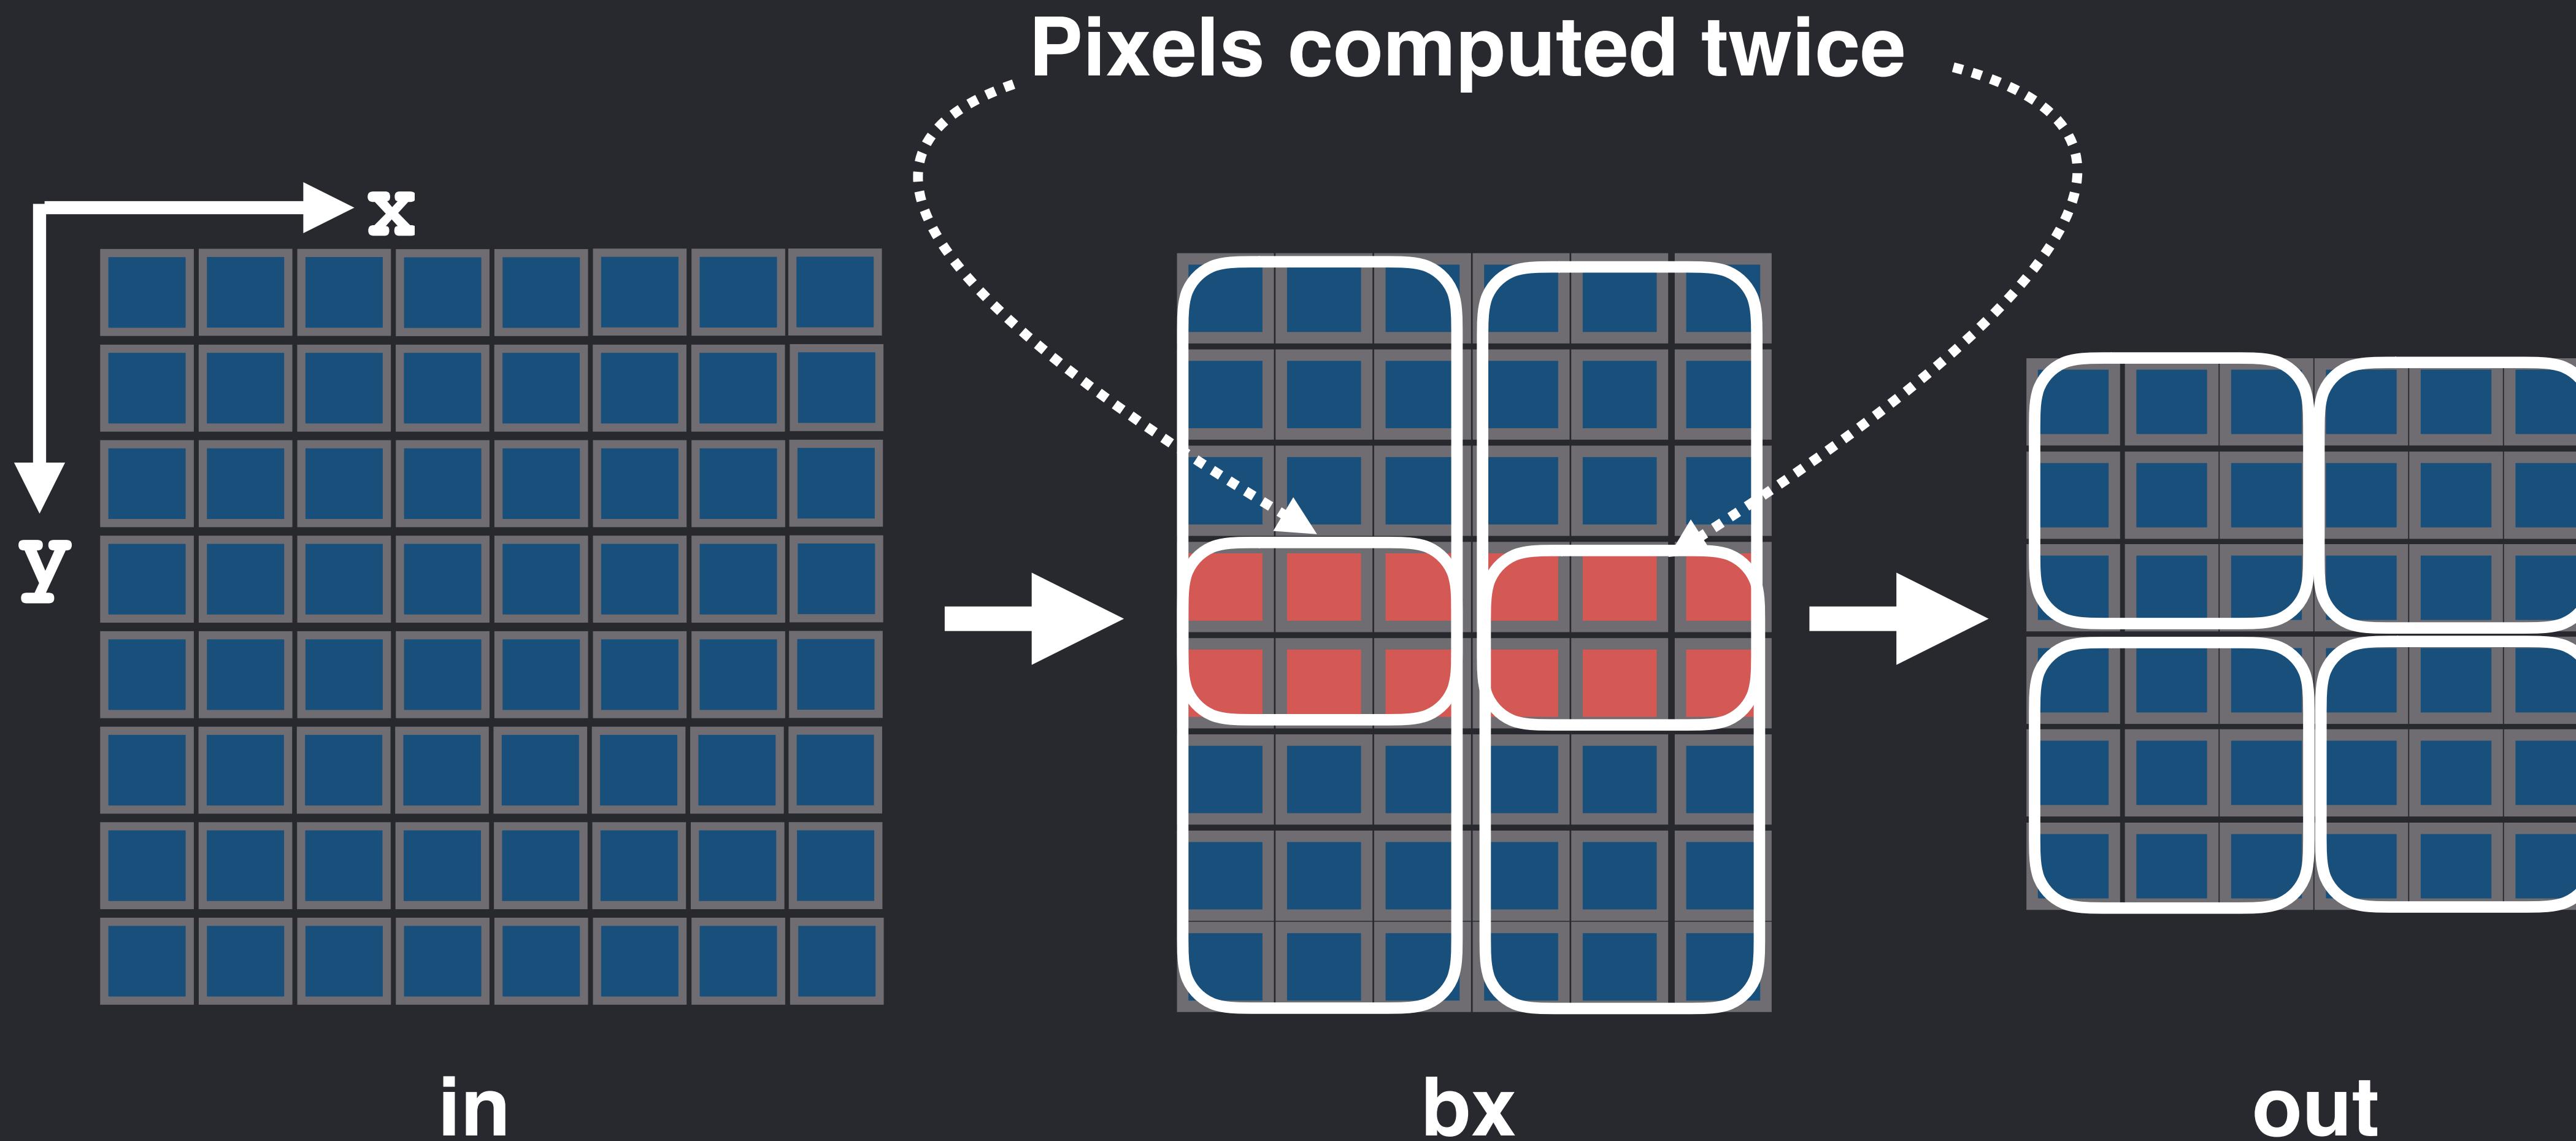
Tiling introduces redundant work



Tiling introduces redundant work

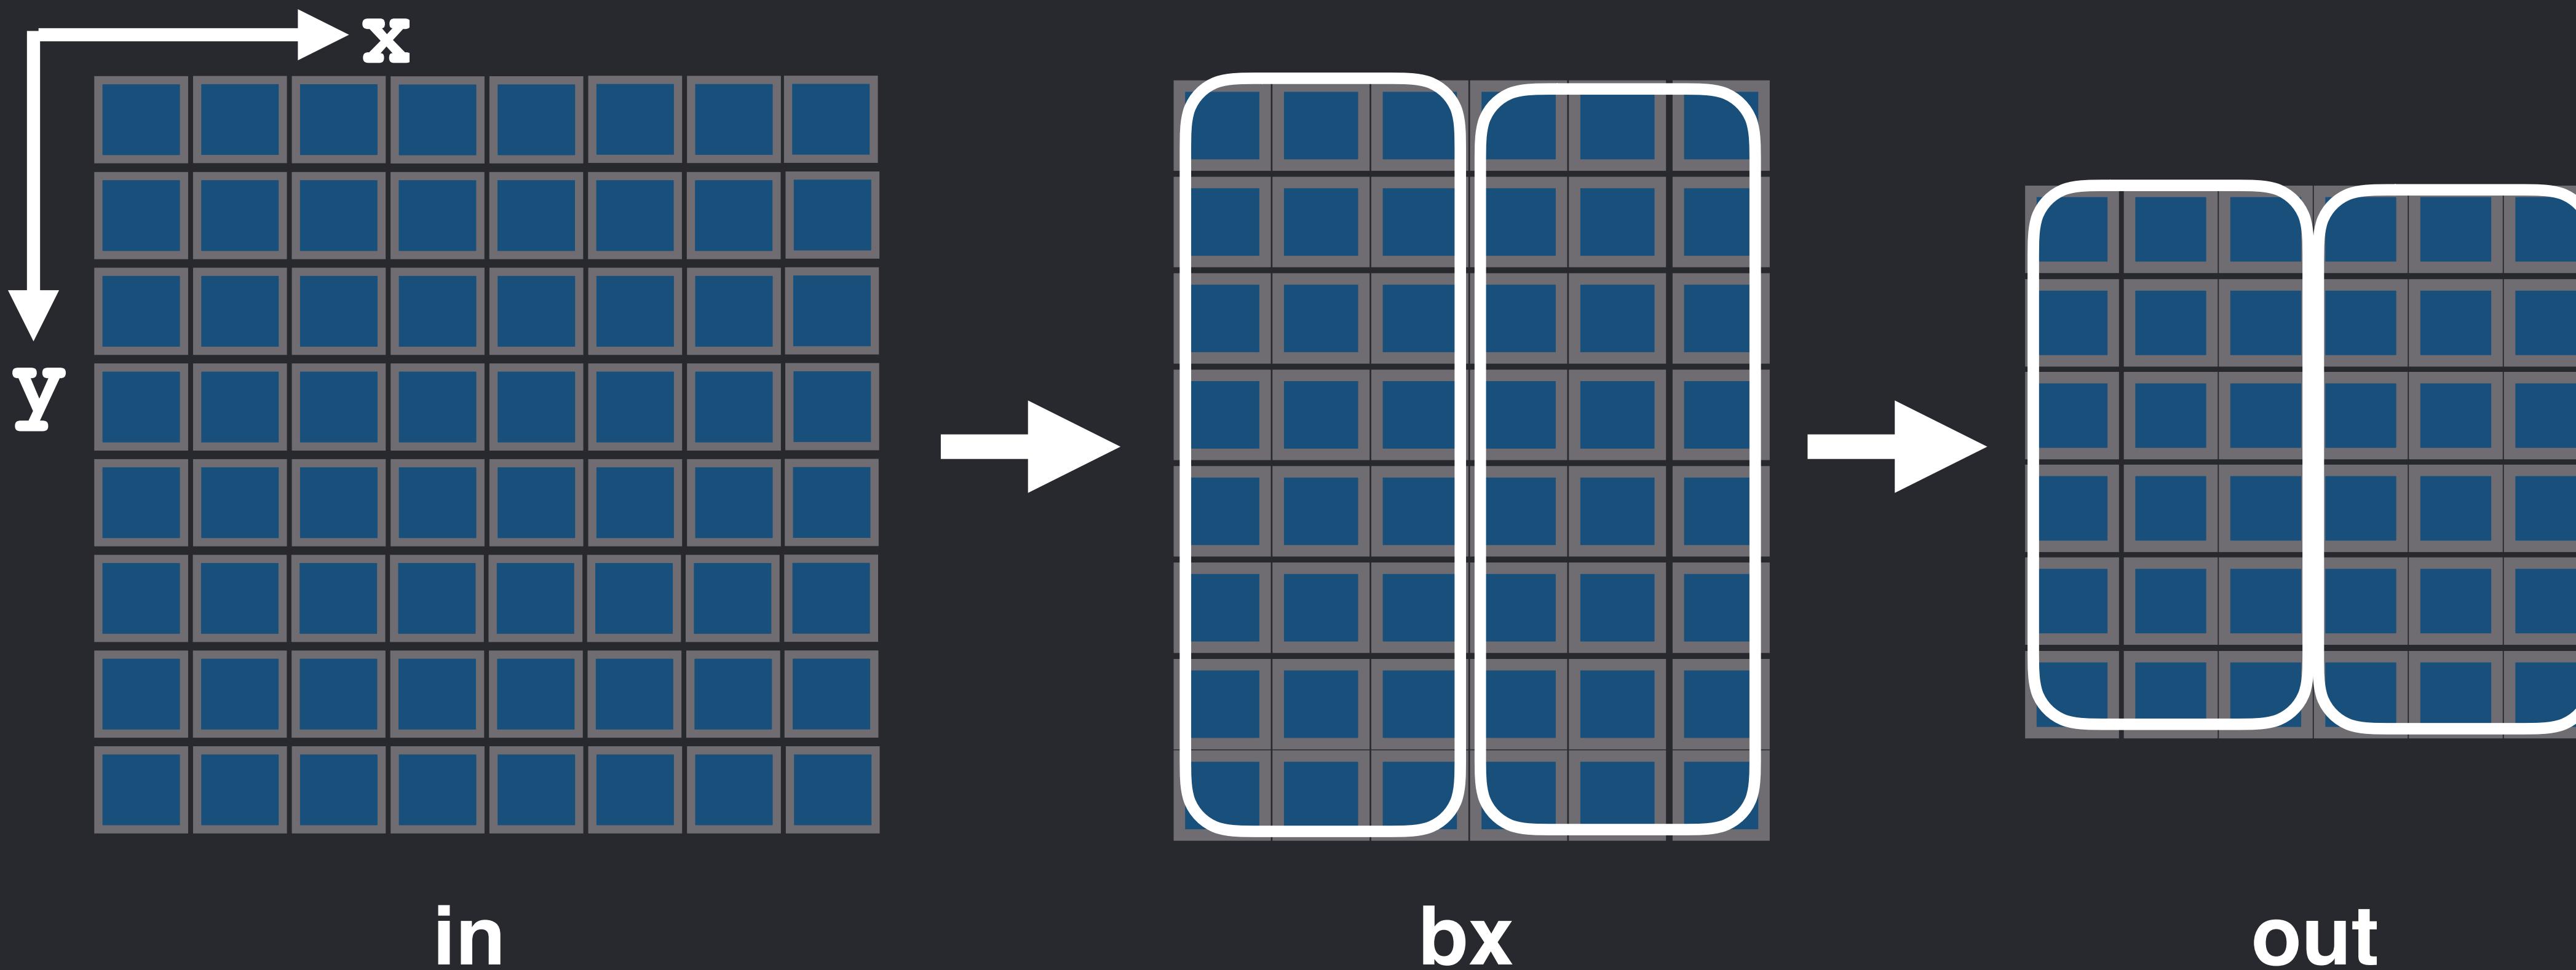


Tiling introduces redundant work



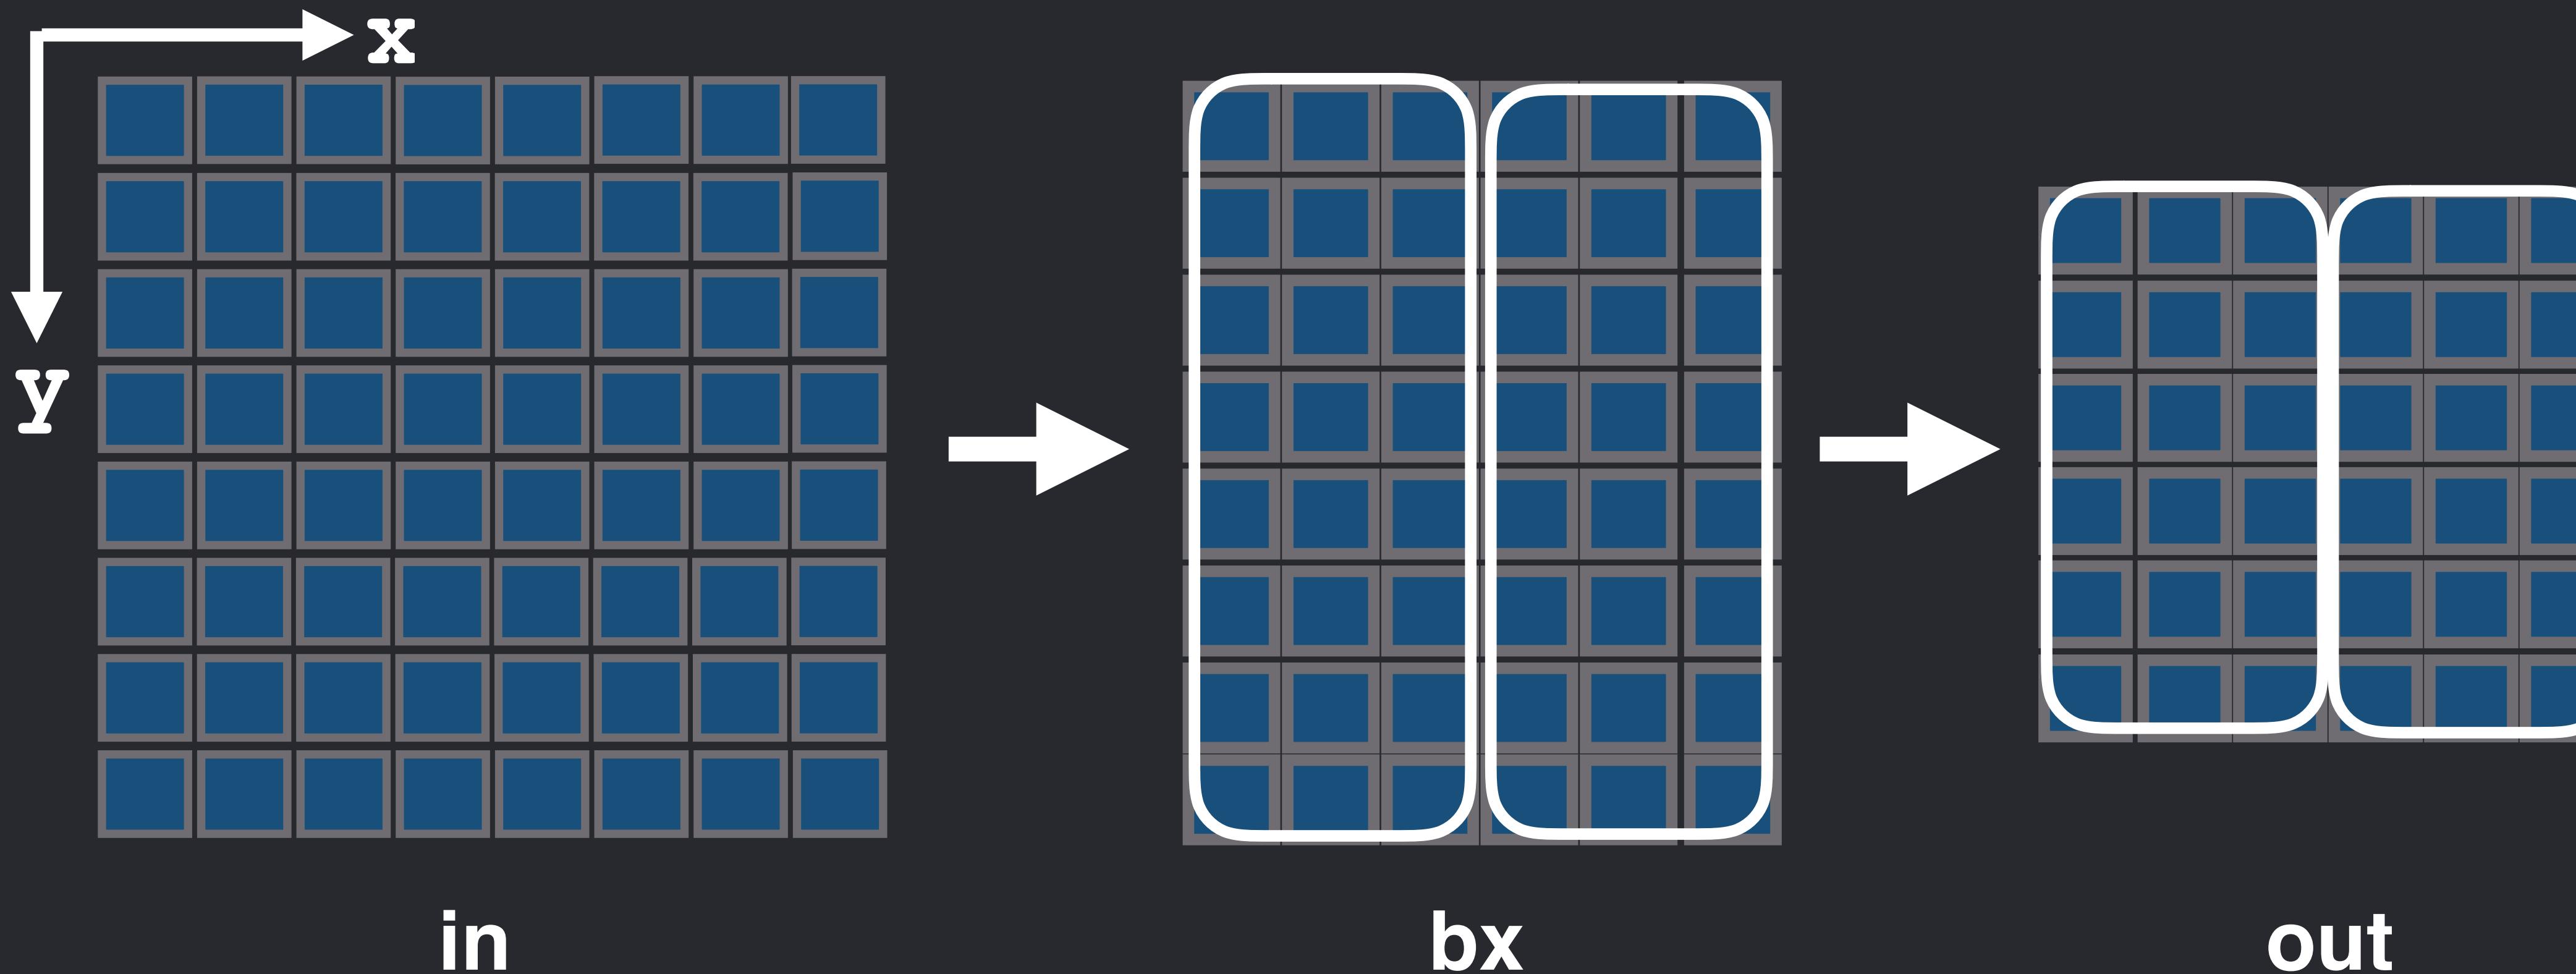
Larger tiles reduce redundant work

for each **3x6 tile**, in parallel
compute required pixels of bx
compute pixels in tile of out



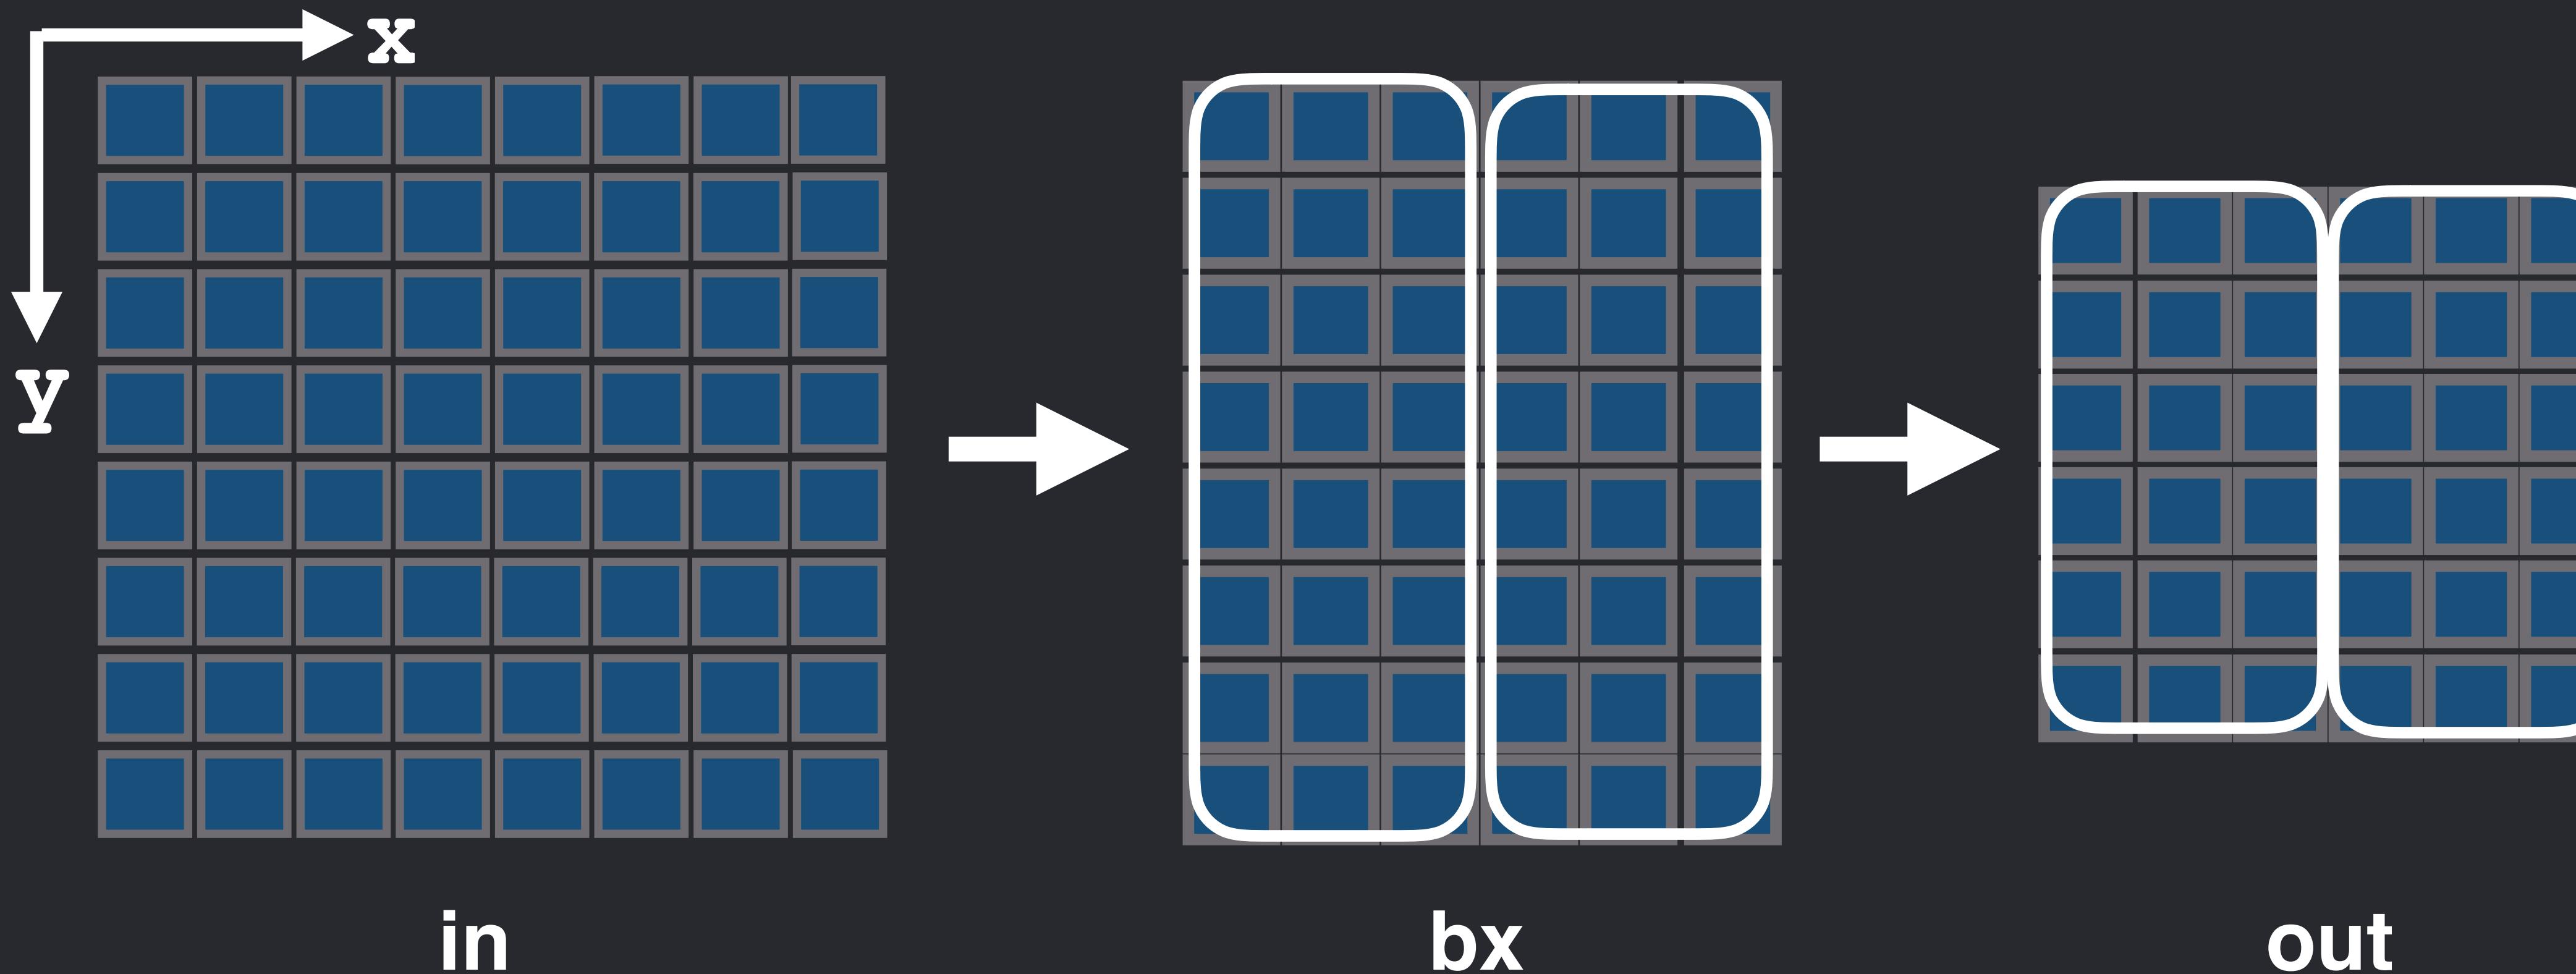
Goal: balance parallelism, locality, work

for each **3x6 tile**, in parallel
compute required pixels of bx
compute pixels in tile of out



Goal: balance parallelism, locality, work

for each **3x6 tile**, in parallel
compute required pixels of **bx**
compute pixels in tile of **out**

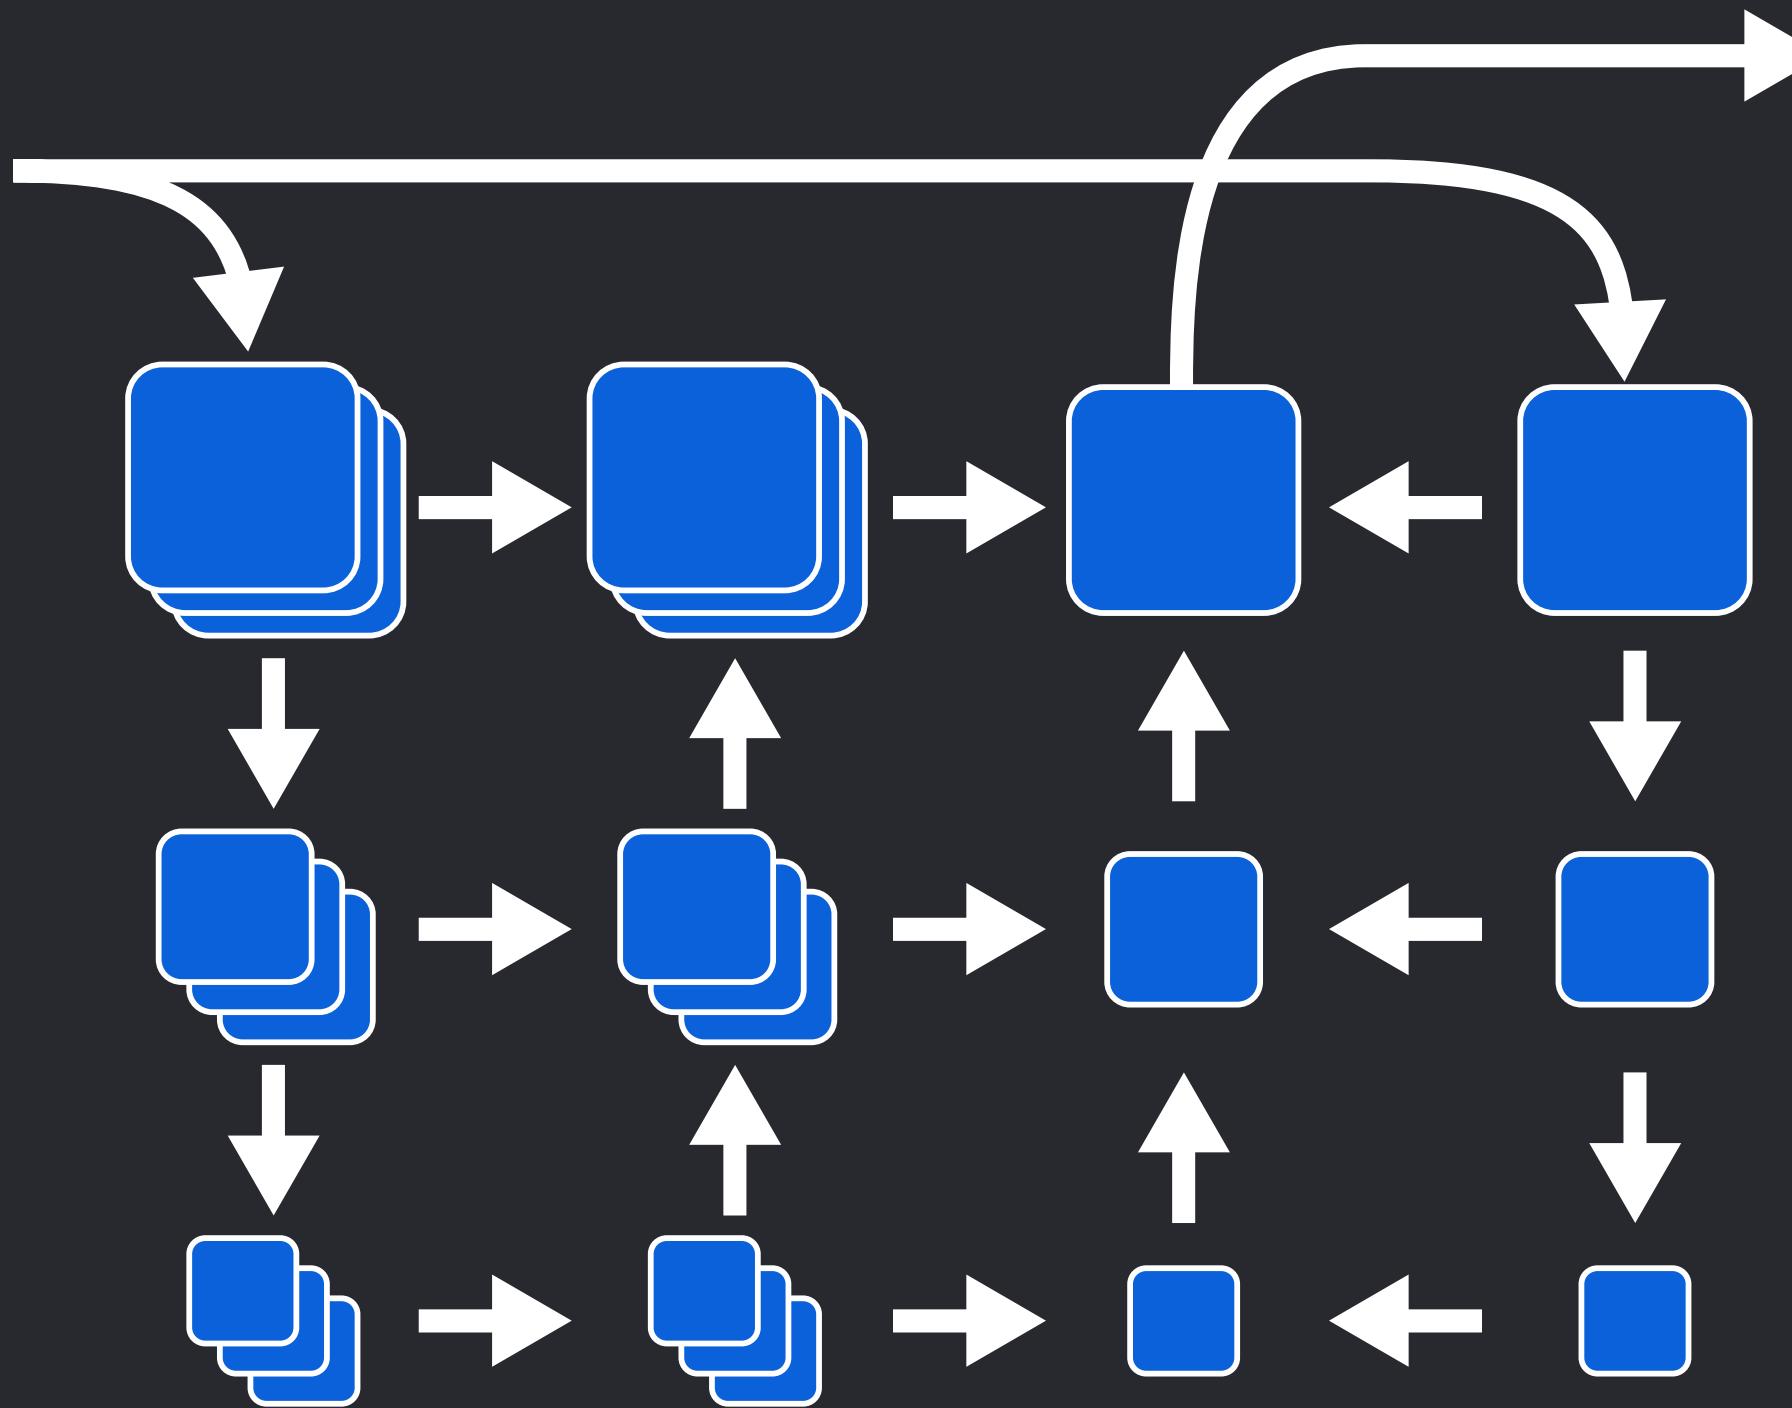


Represent image processing pipelines as graphs



DAG representation of the two-stage blur pipeline

Real world pipelines are complex graphs



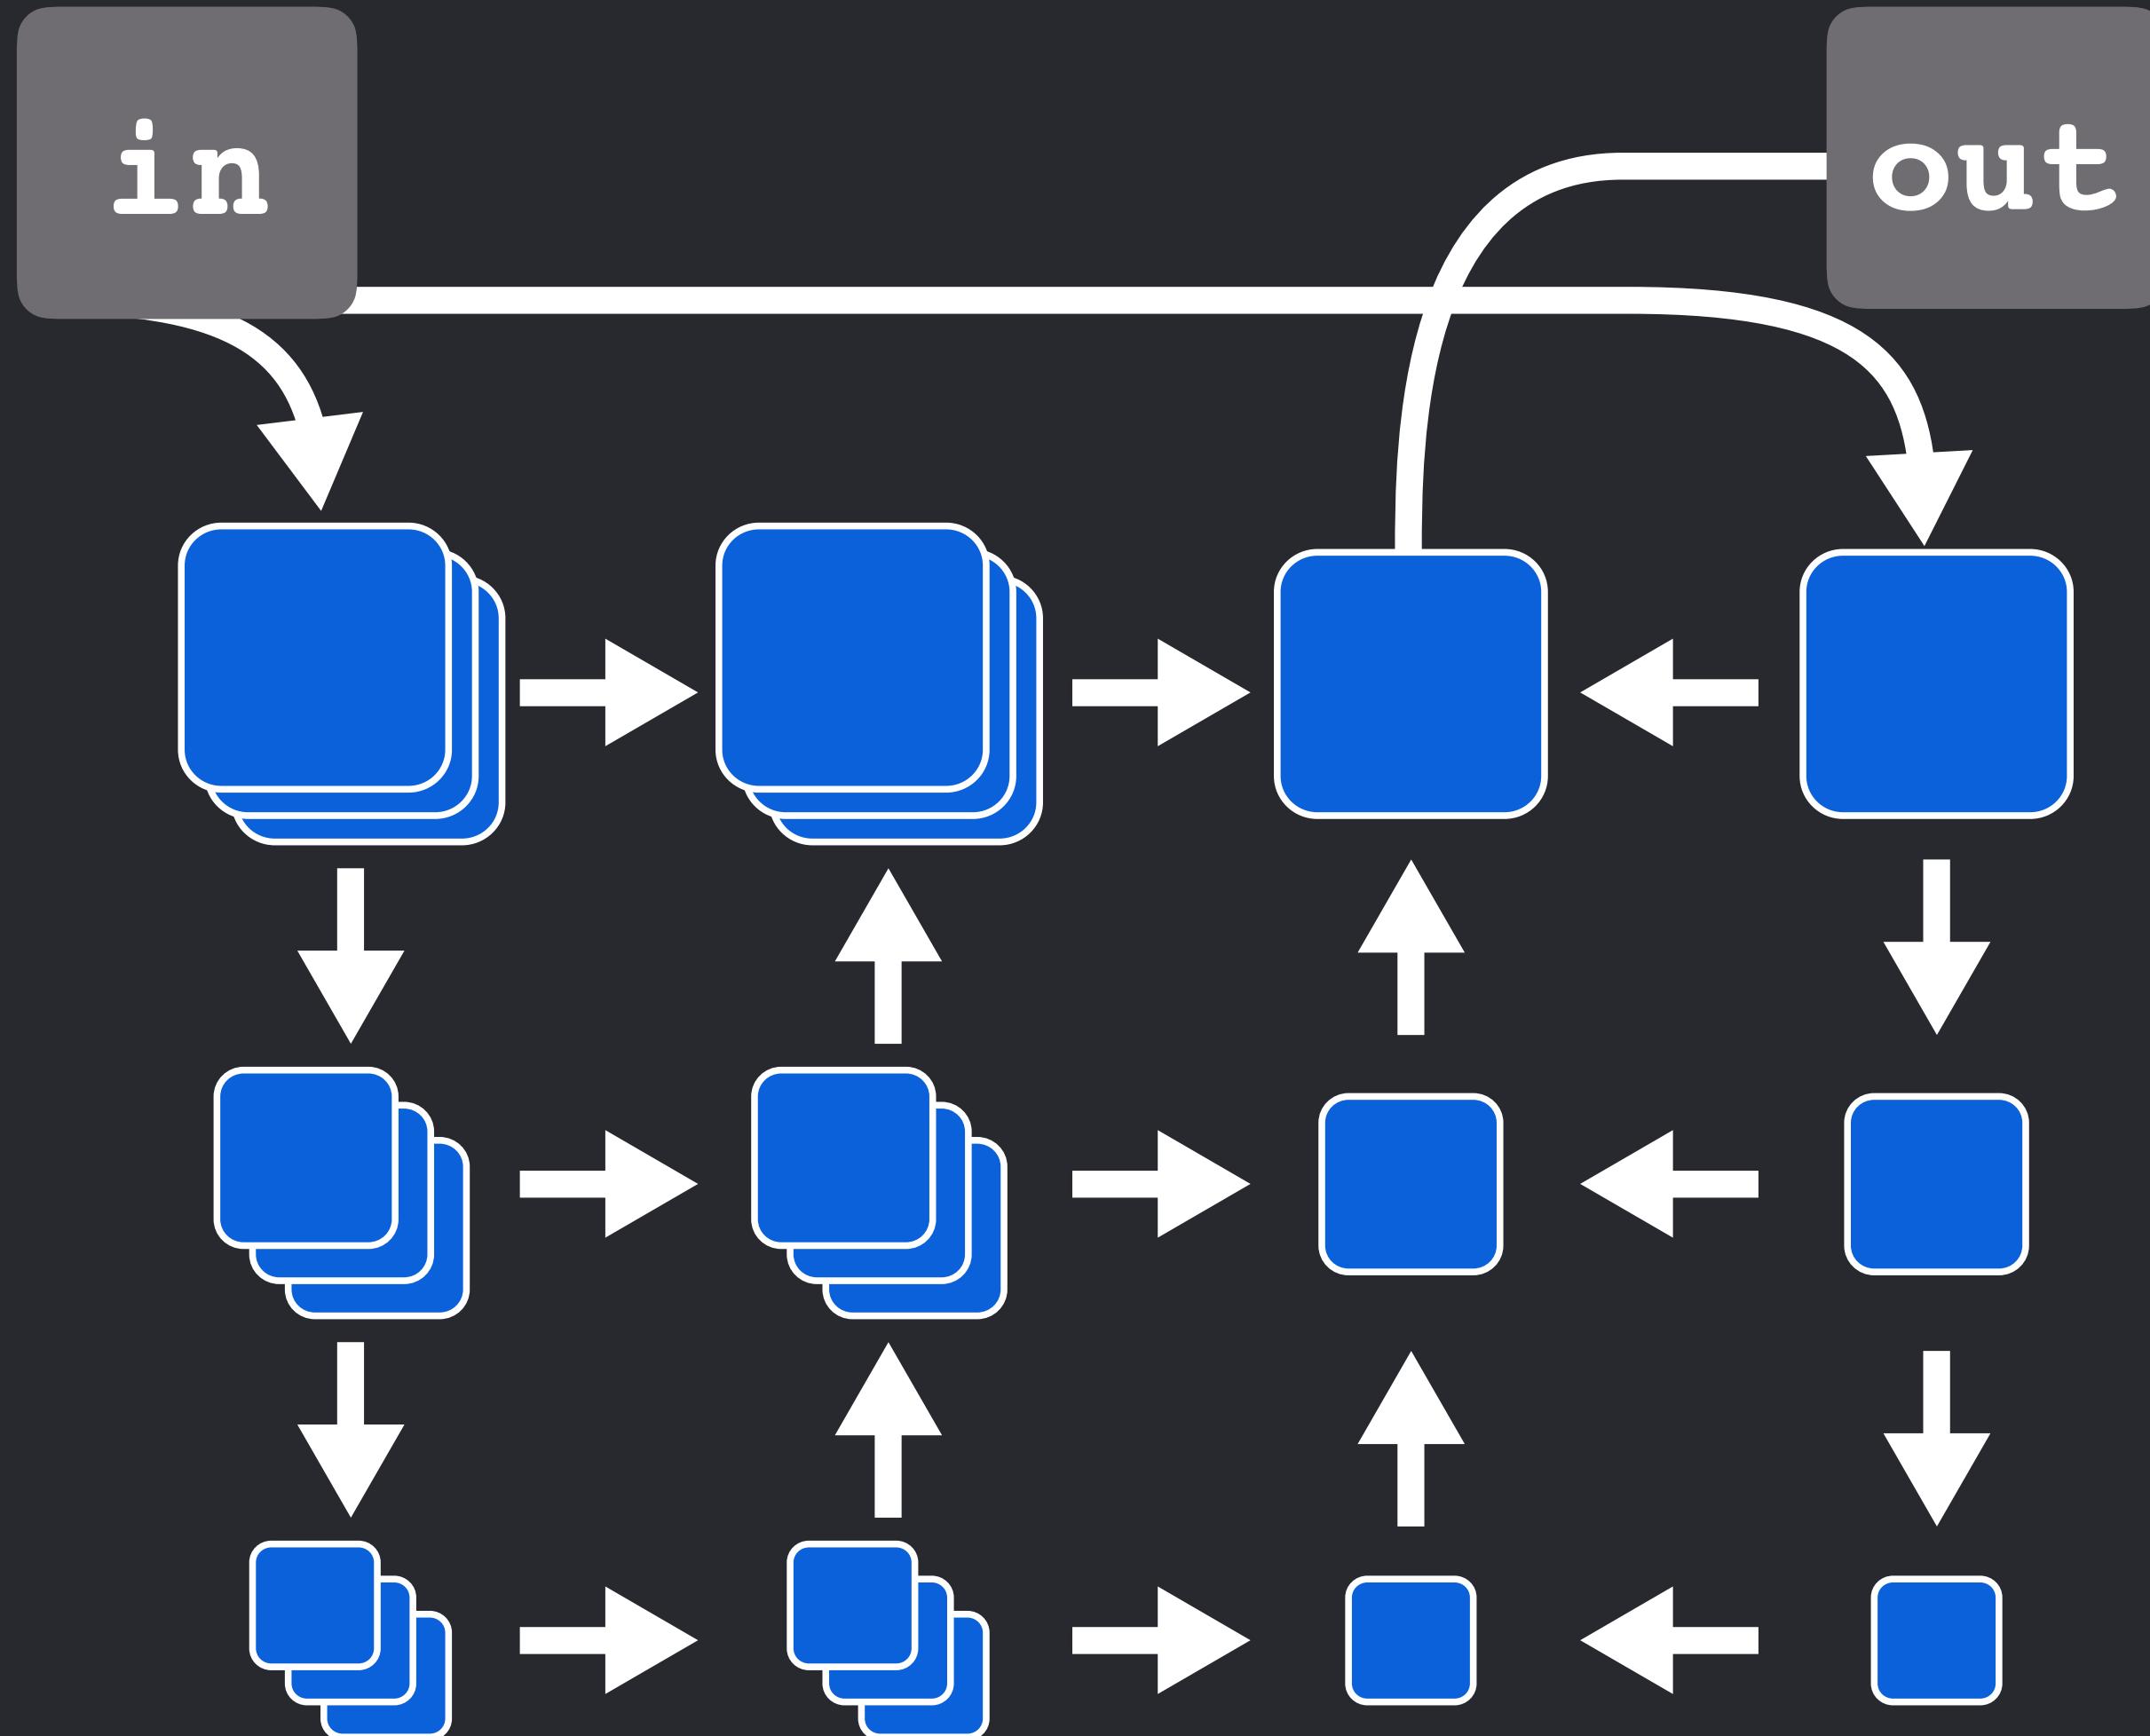
Local Laplacian filters

[Paris et al. 2010, Aubry et al. 2011]

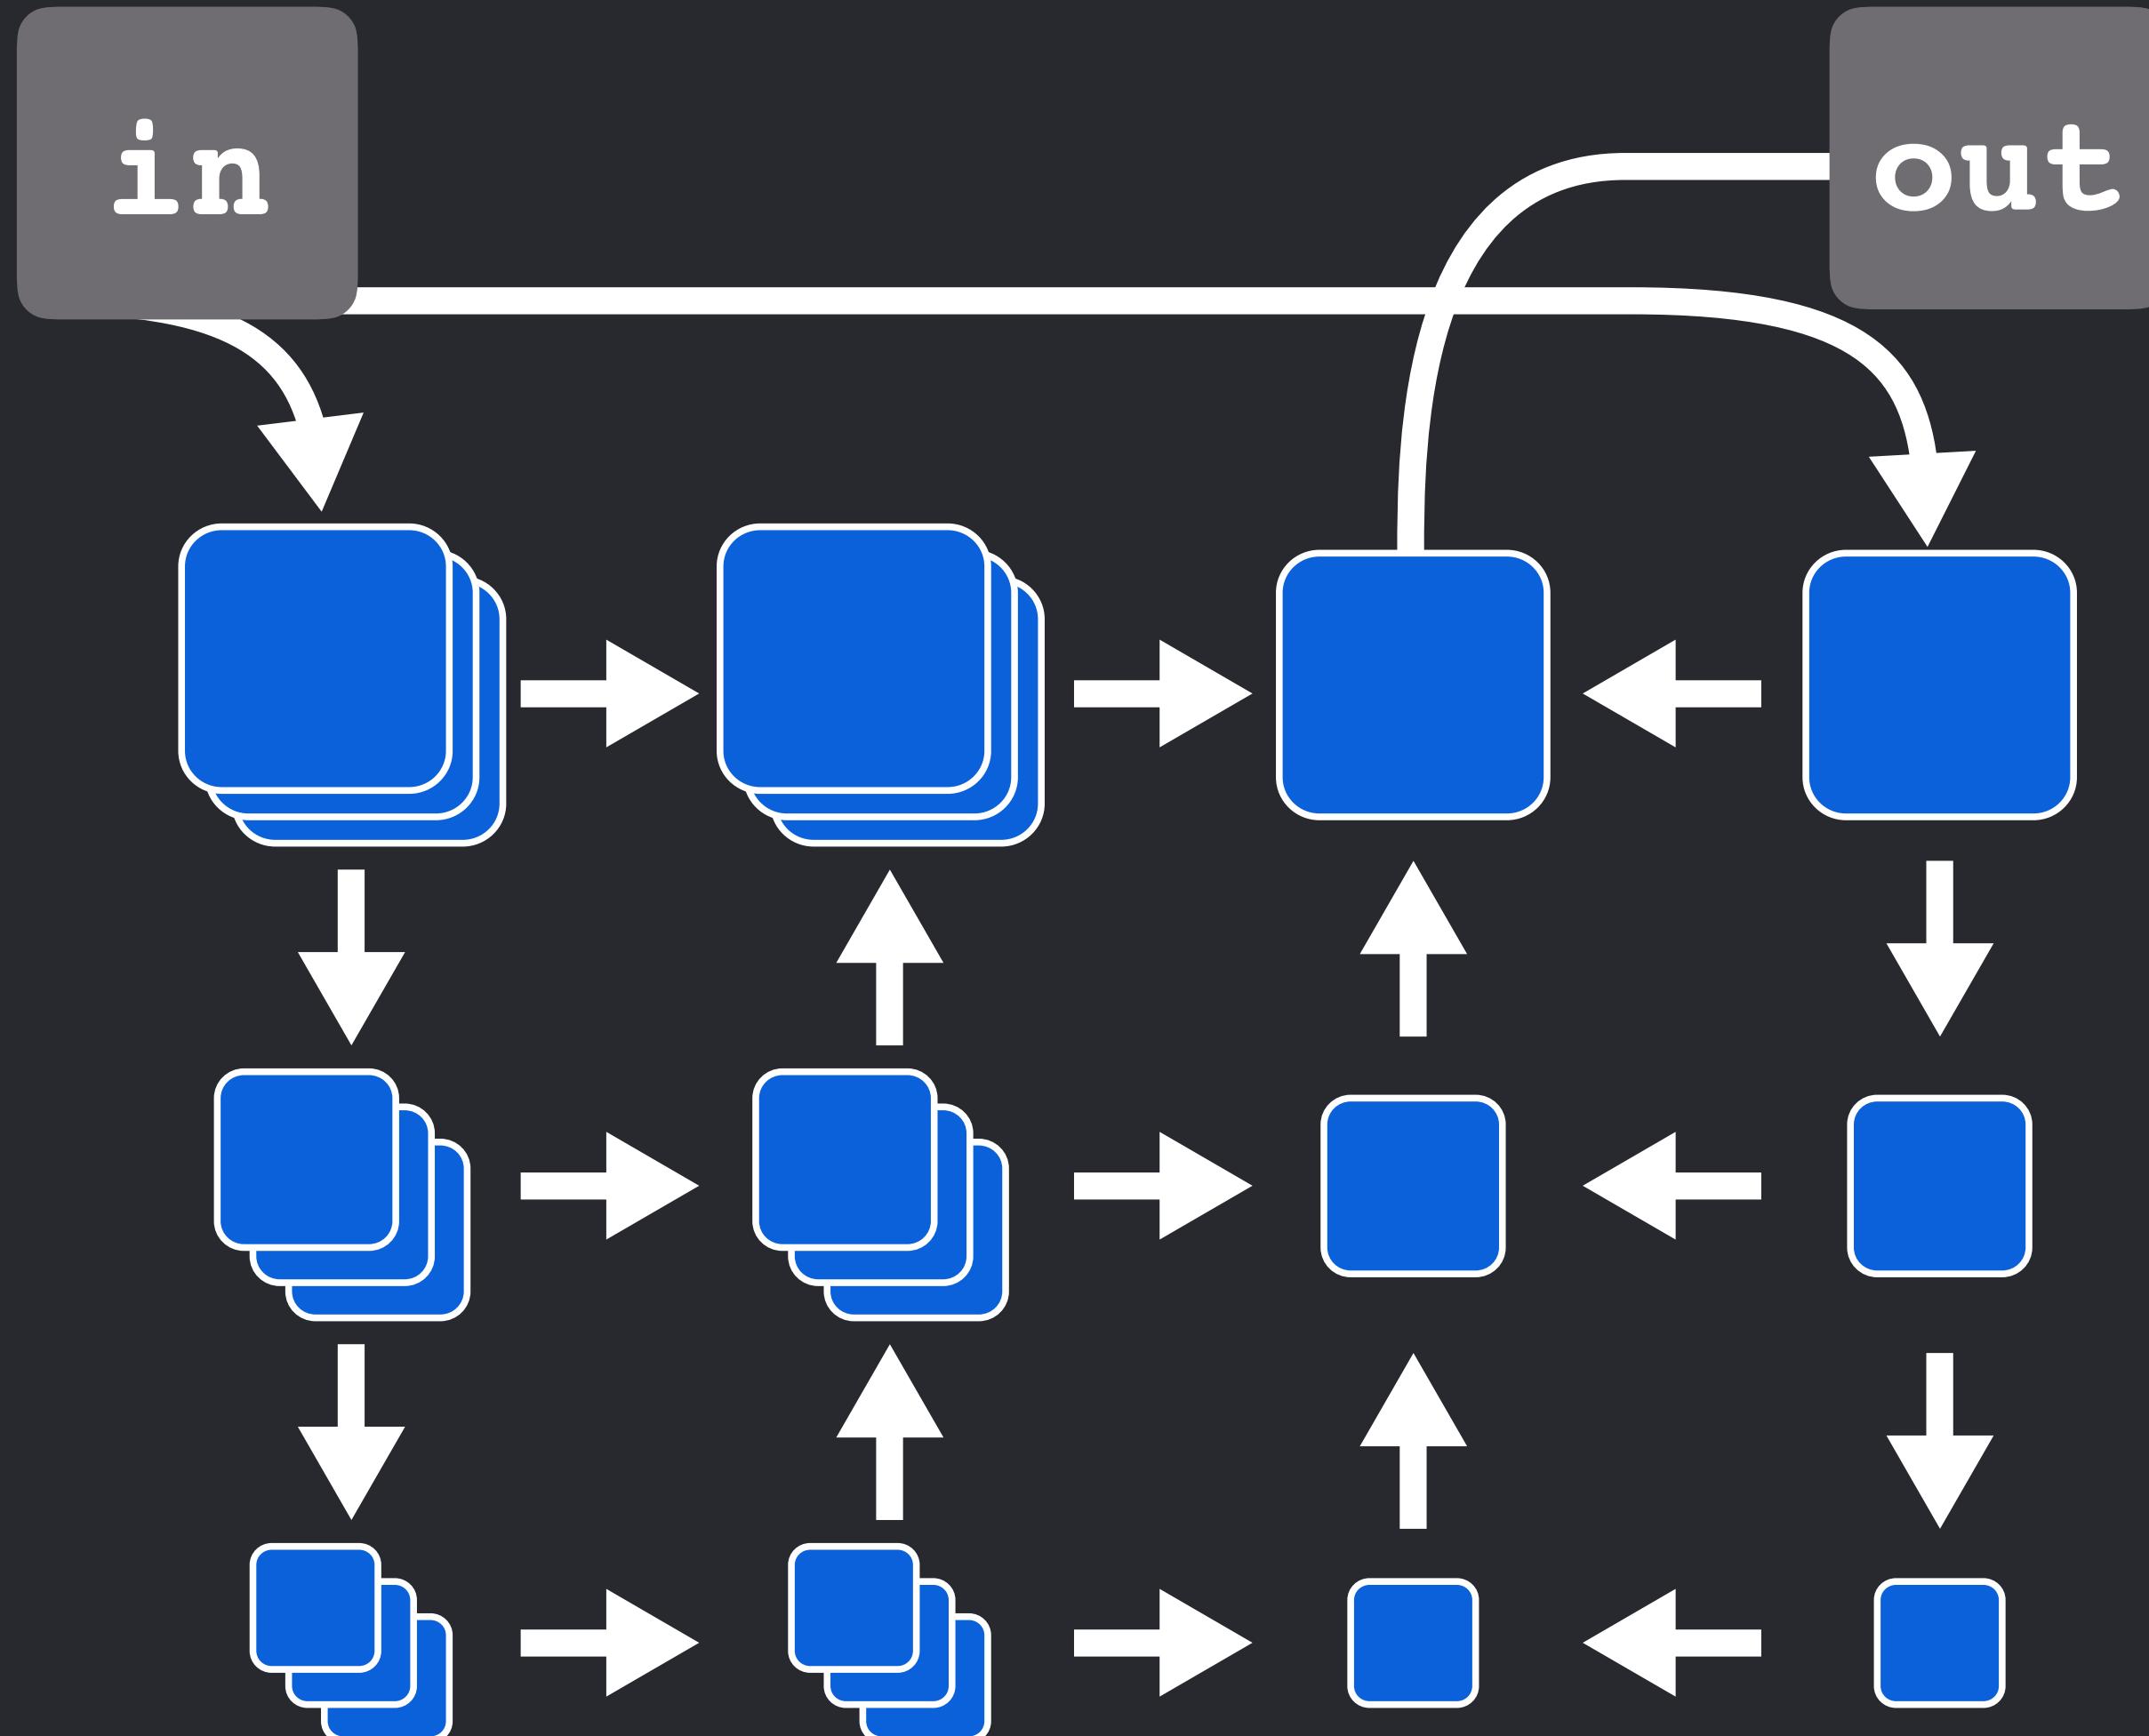
100 stages

Google Nexus HDR+ mode: over 2000 stages!

Key aspects of scheduling

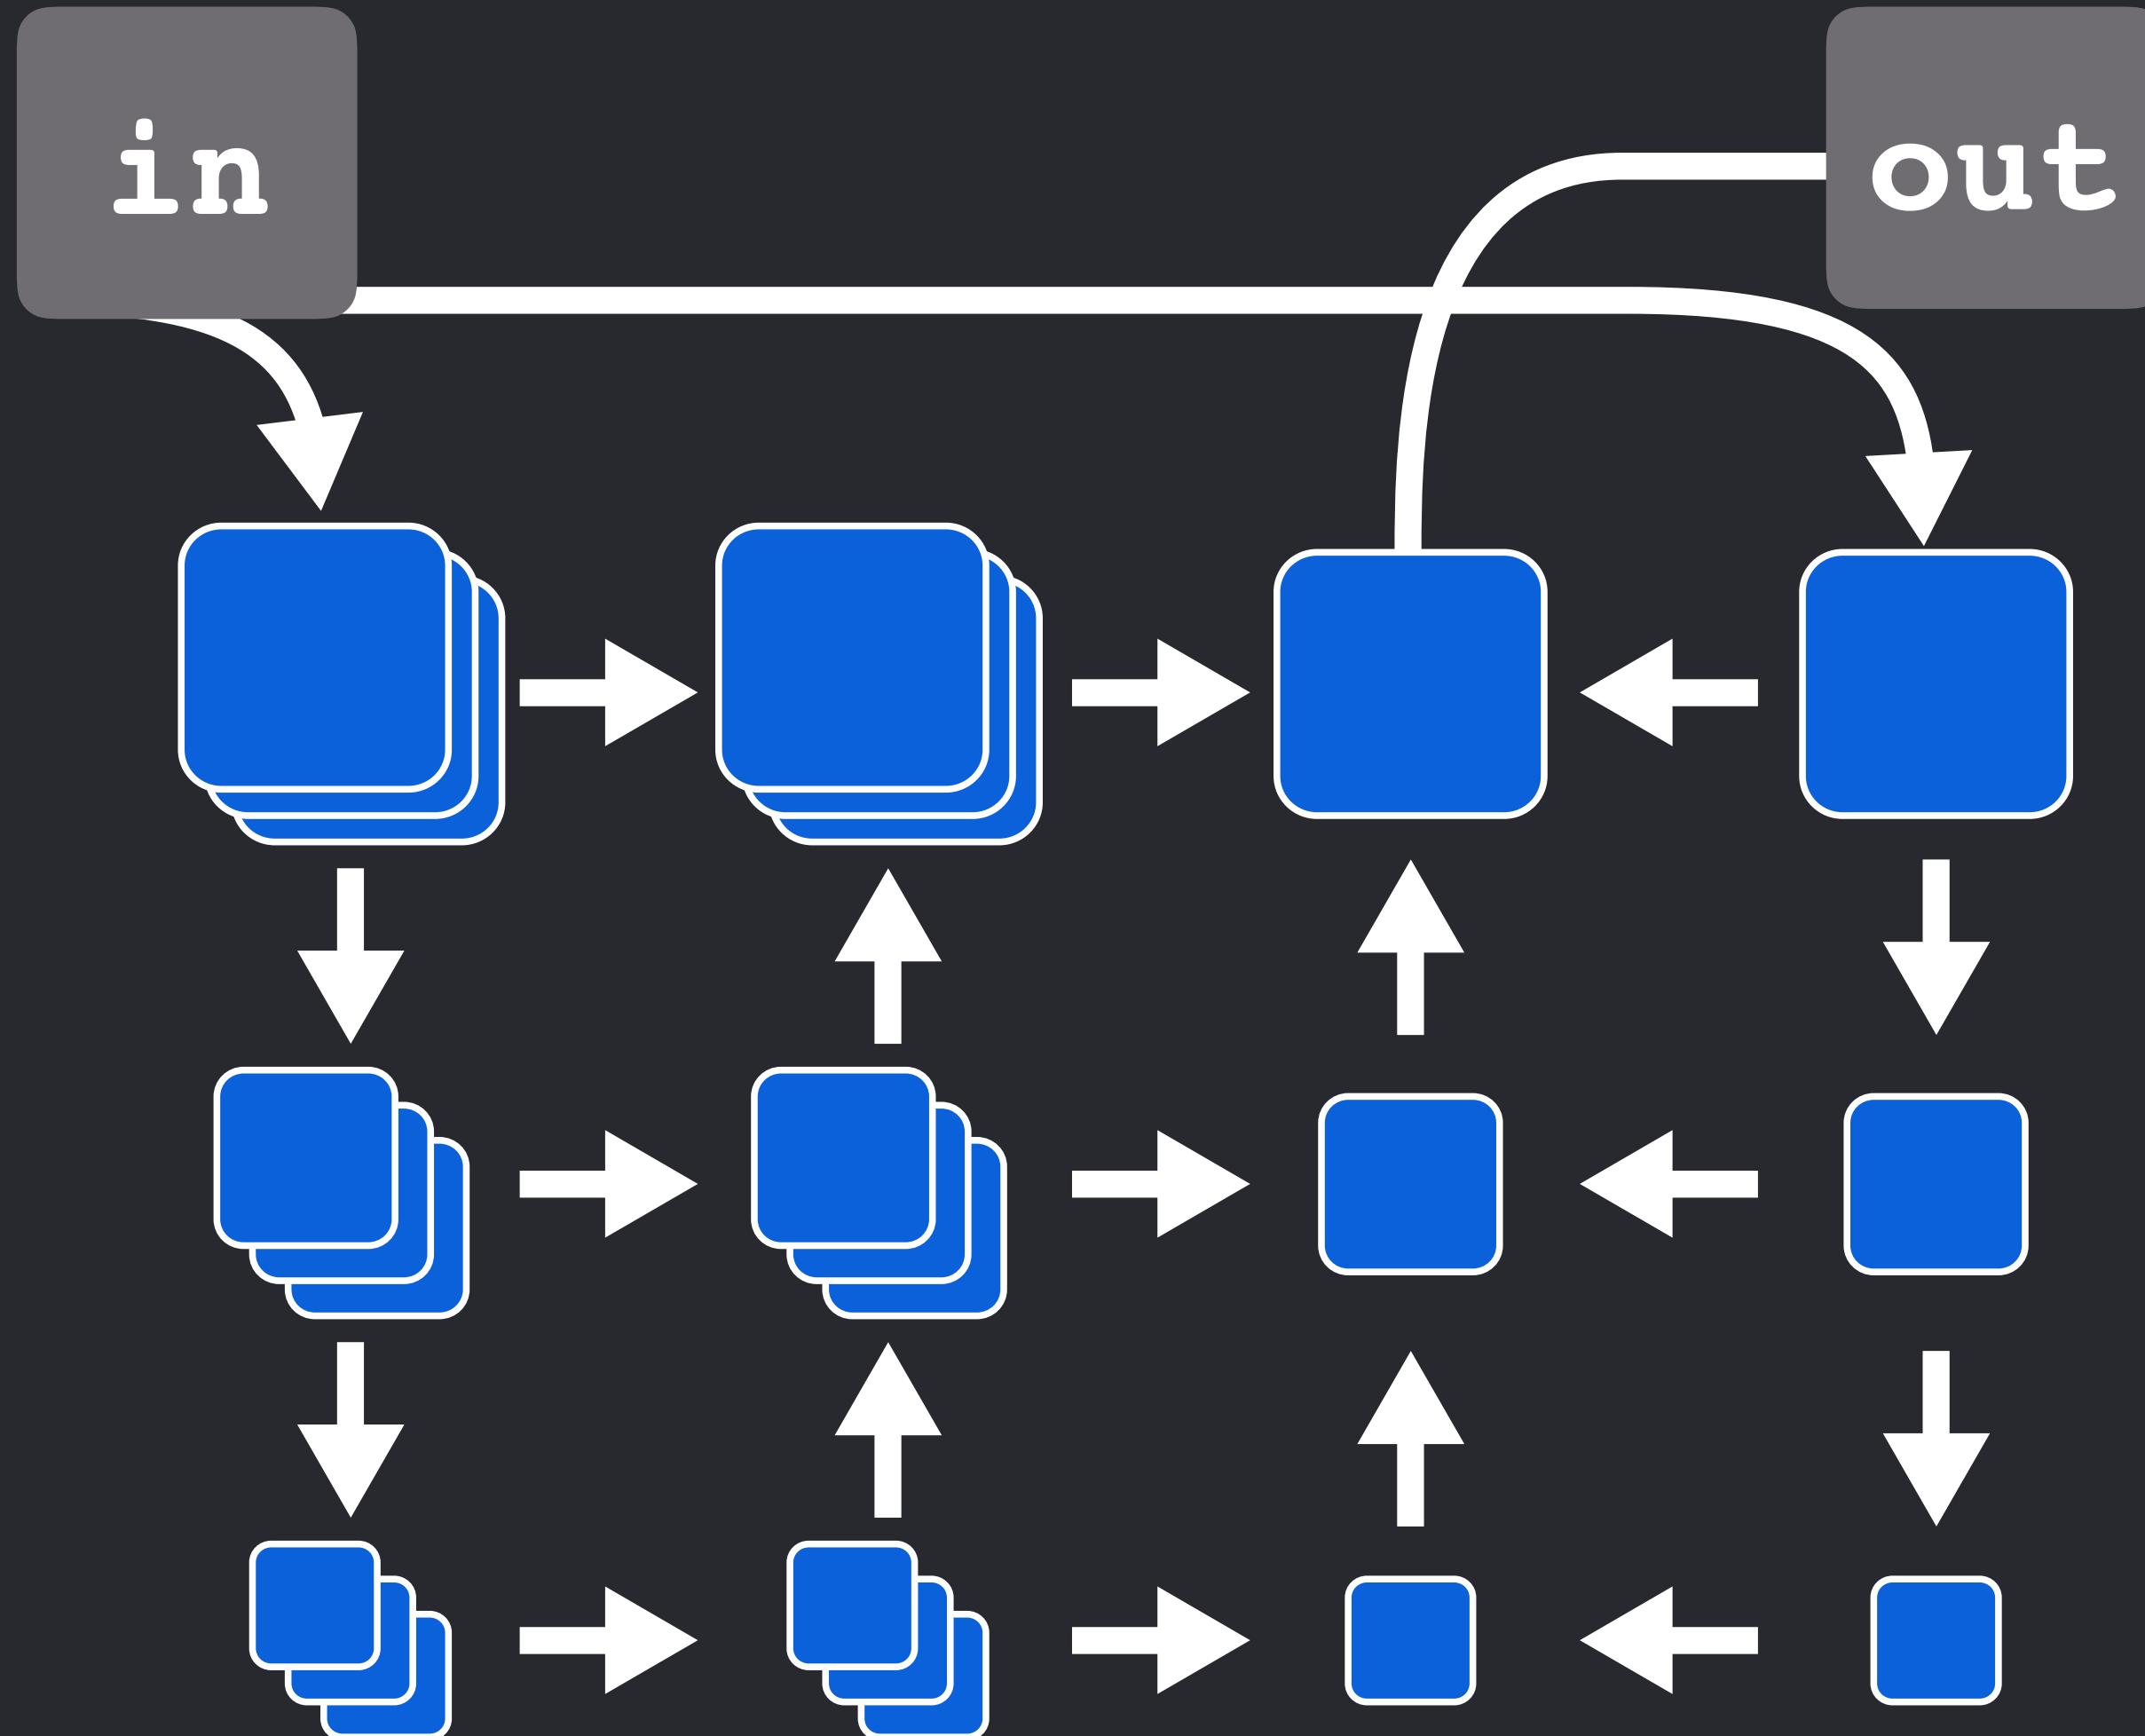


Key aspects of scheduling



Deciding which stages to
interleave for better data
locality

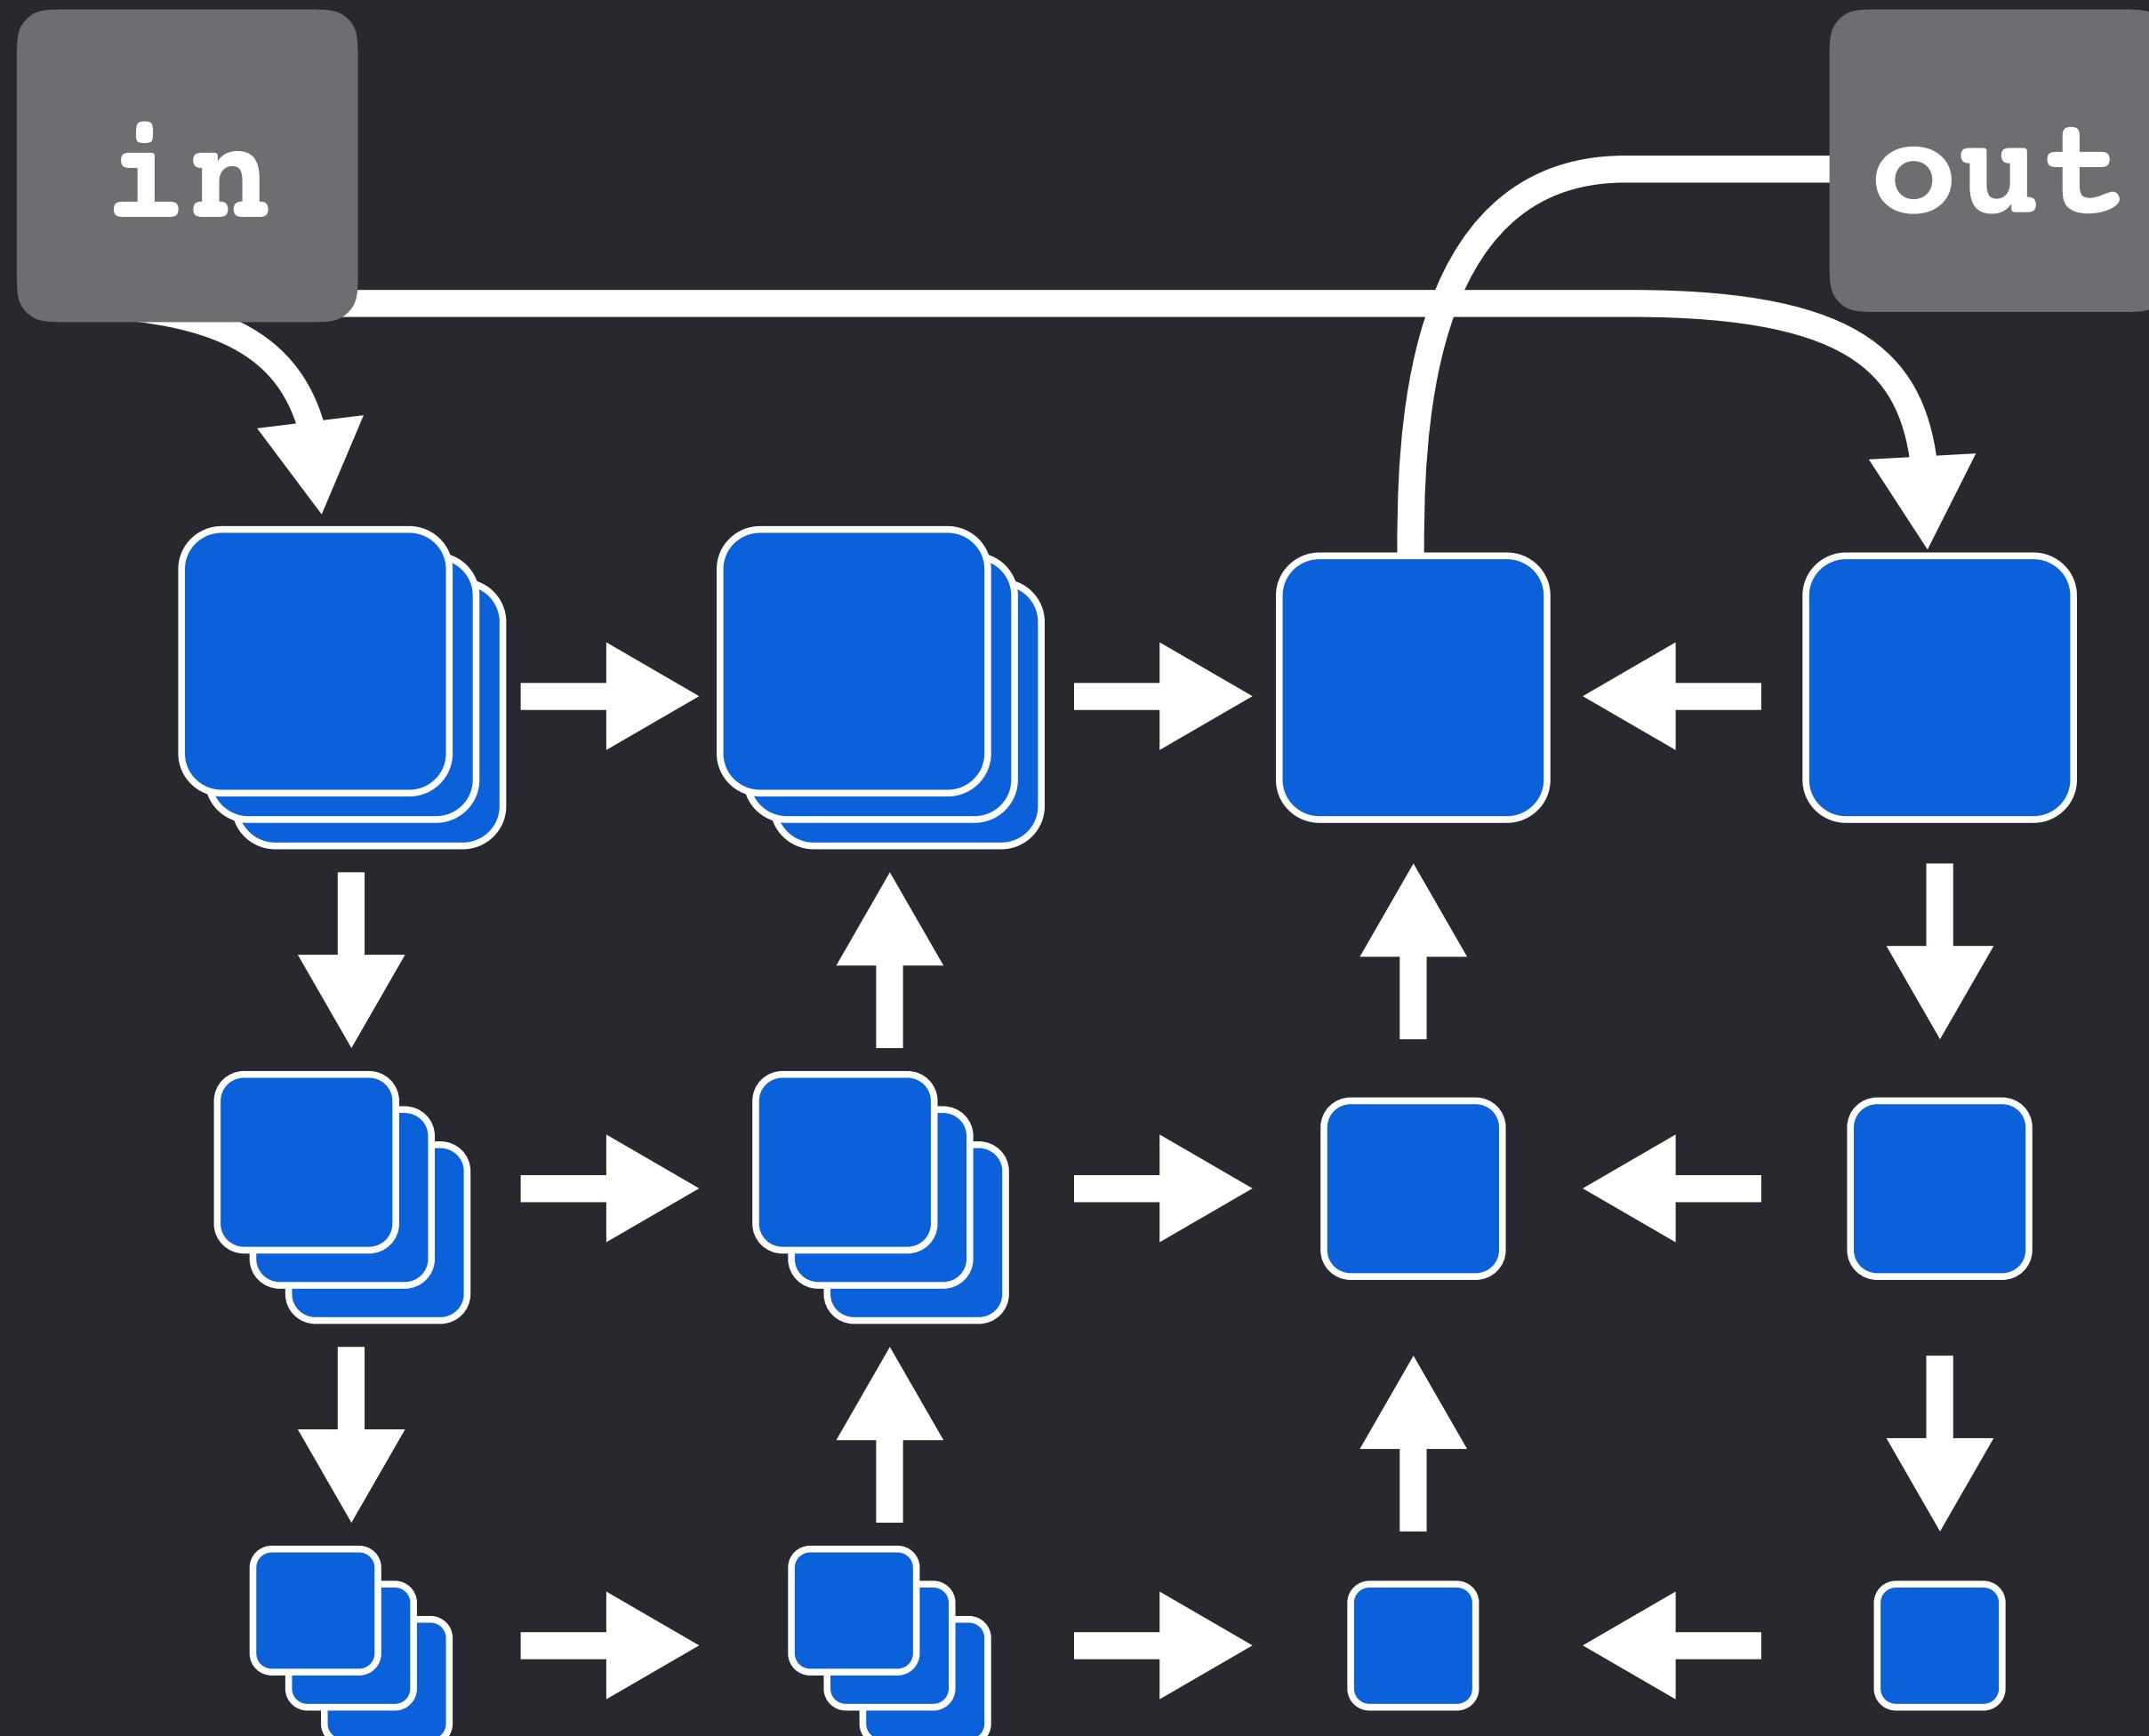
Key aspects of scheduling



Deciding which stages to interleave for better data locality

Picking tiles sizes to trade-off locality and re-computation

Key aspects of scheduling



Deciding which stages to
interleave for better data
locality

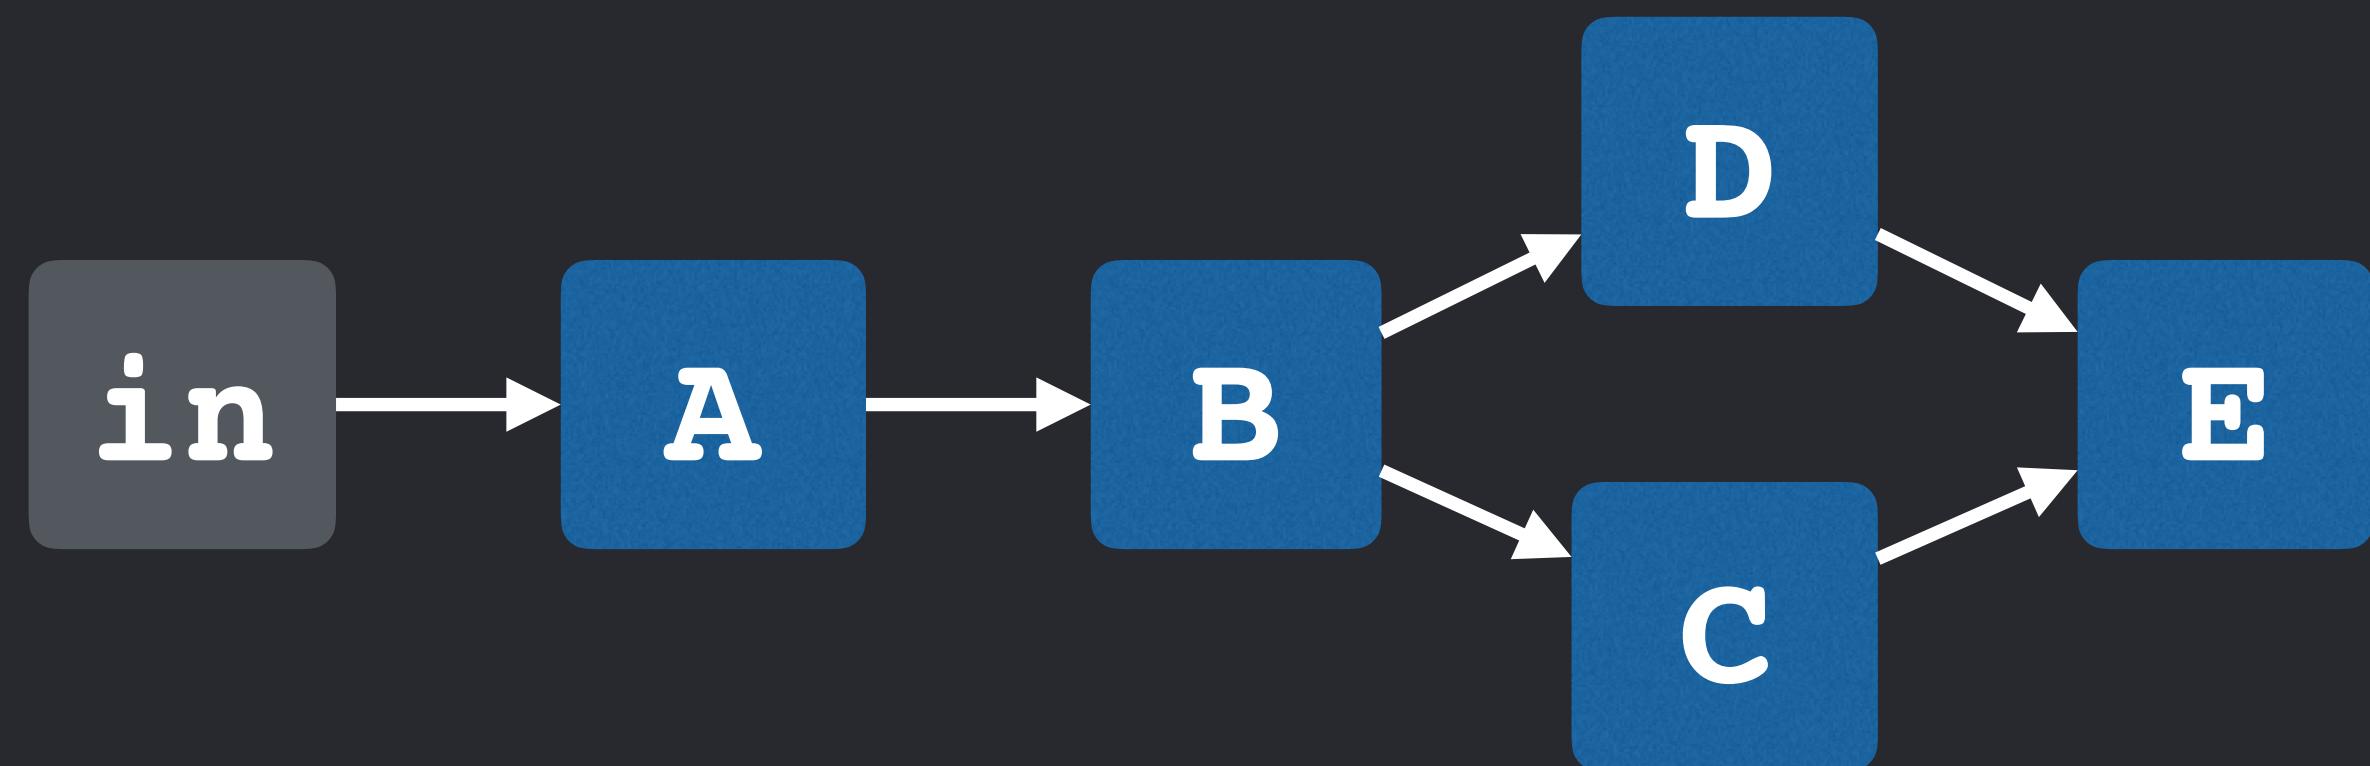
Picking tiles sizes to trade-off
locality and re-computation

Maintain ability to execute in
parallel

An Algorithm for Scheduling Image Processing Pipelines

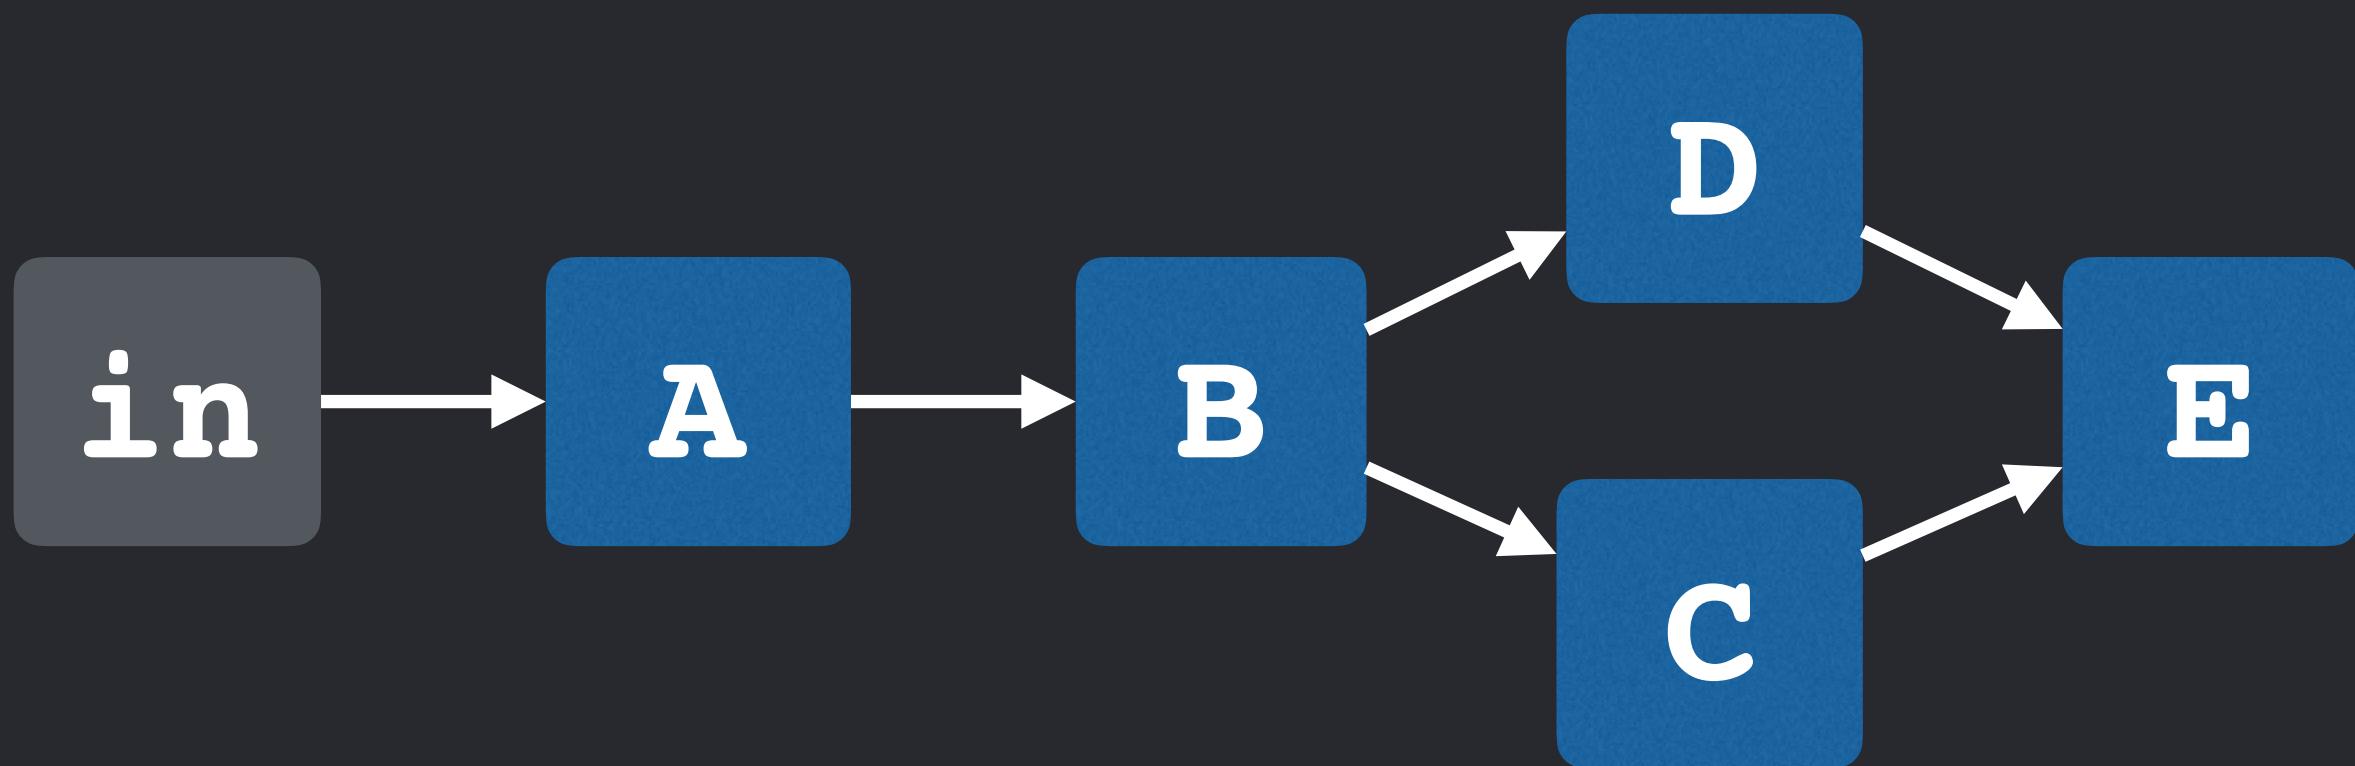
Algorithm

Input: DAG of pipeline stages



Algorithm

Input: DAG of pipeline stages



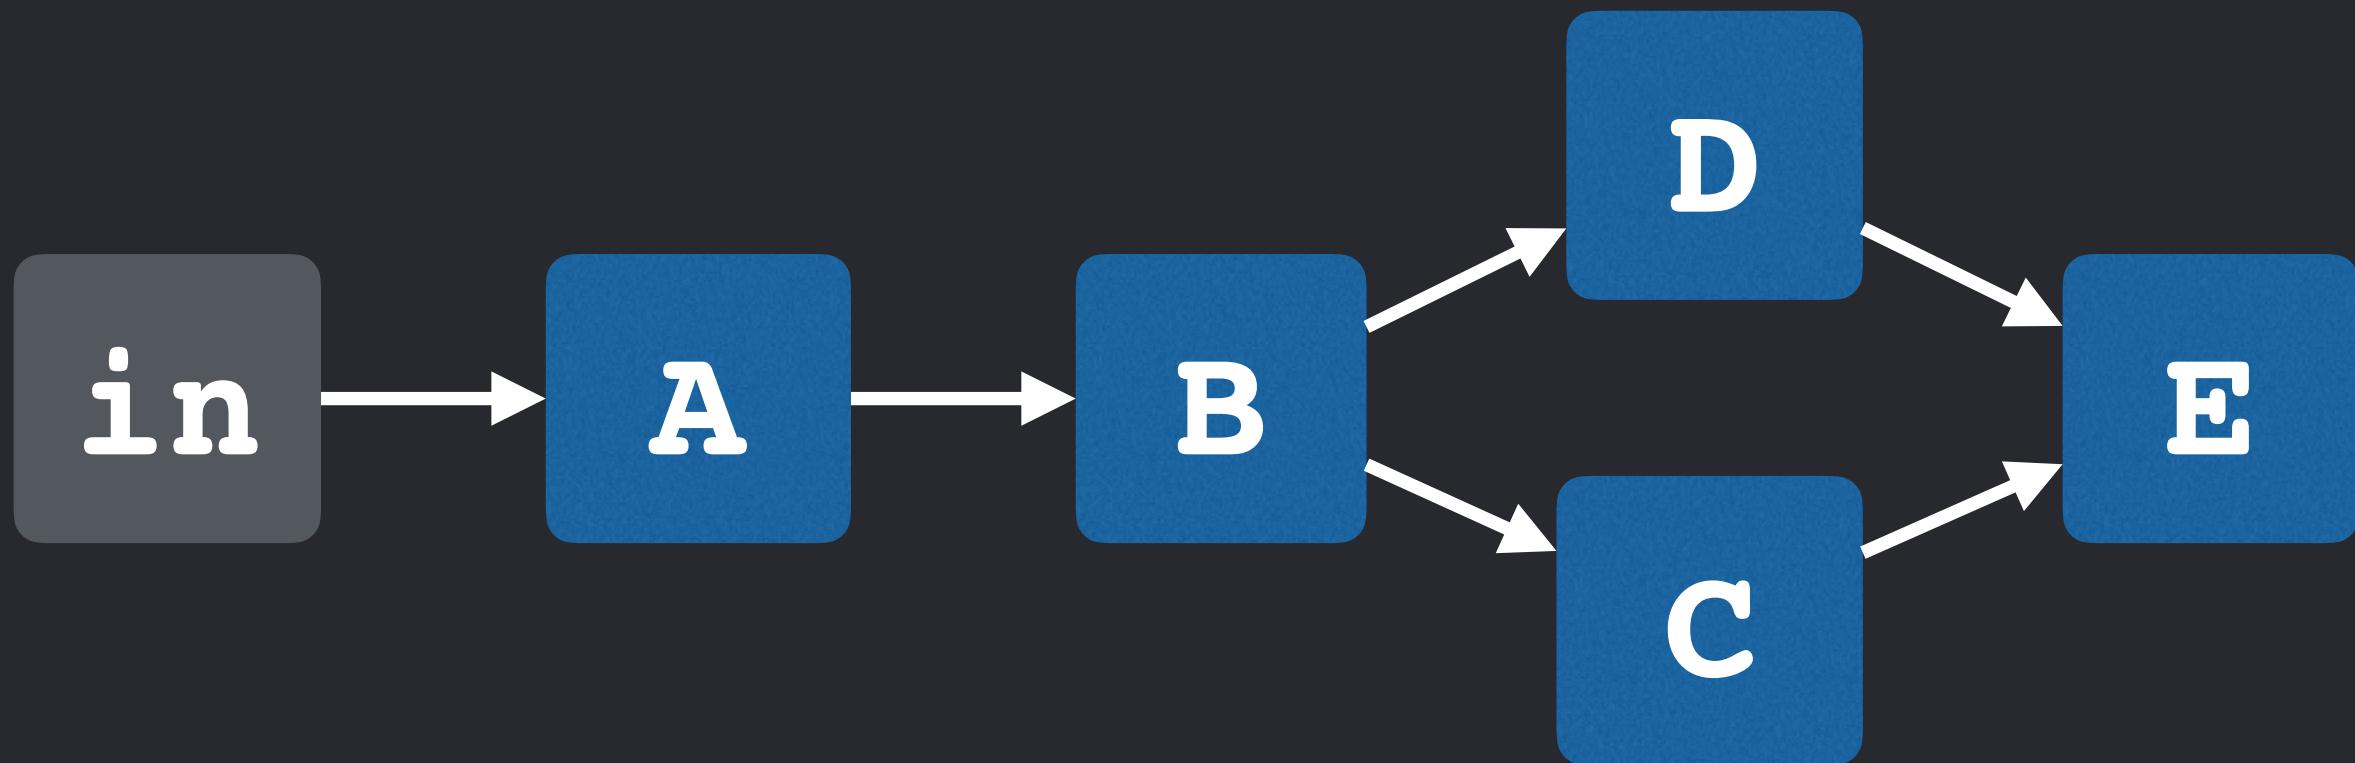
Output: Optimized schedule

```
for each 8x128 tile in parallel  
  compute required pixels of A  
  compute pixels in tile of B
```

```
for each 8x8 tile in parallel  
  compute required pixels of C  
  compute required pixels of D  
  compute pixels in tile of E
```

Algorithm

Input: DAG of pipeline stages



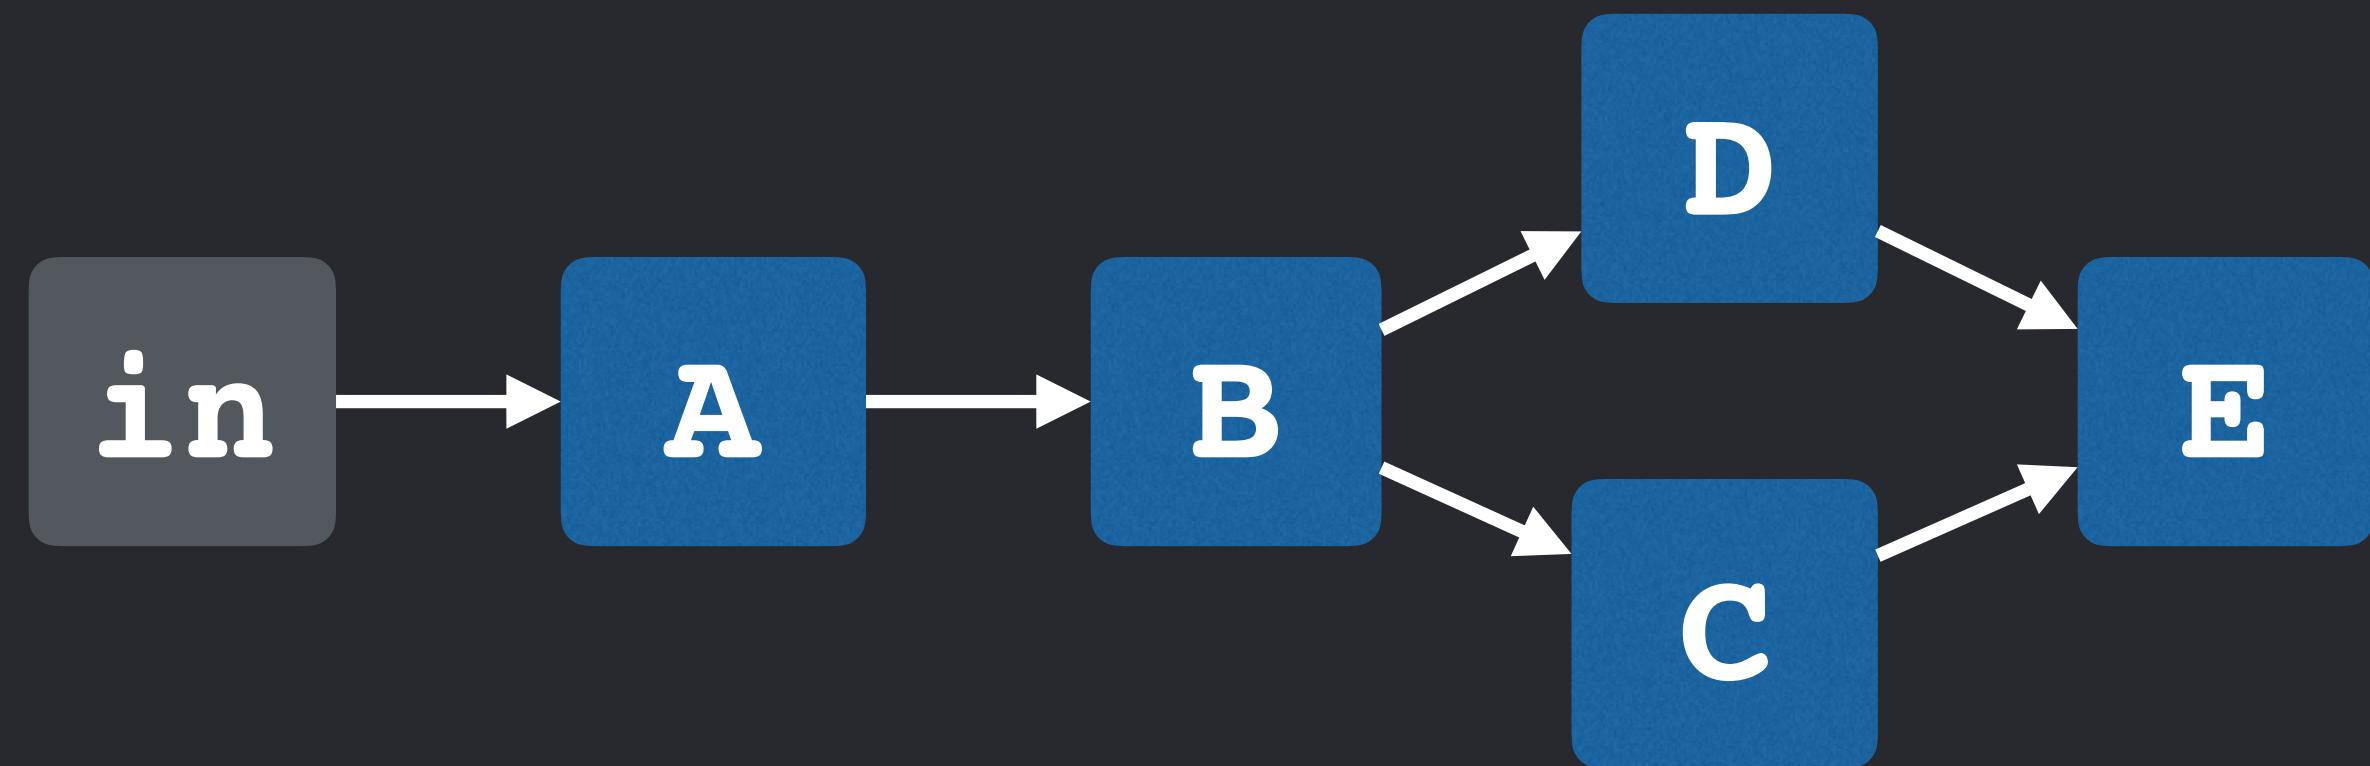
Output: Optimized schedule

for each 8x128 tile in parallel
 compute required pixels of A
 compute pixels in tile of B

for each 8x8 tile in parallel
 compute required pixels of C
 compute required pixels of D
 compute pixels in tile of E

Algorithm

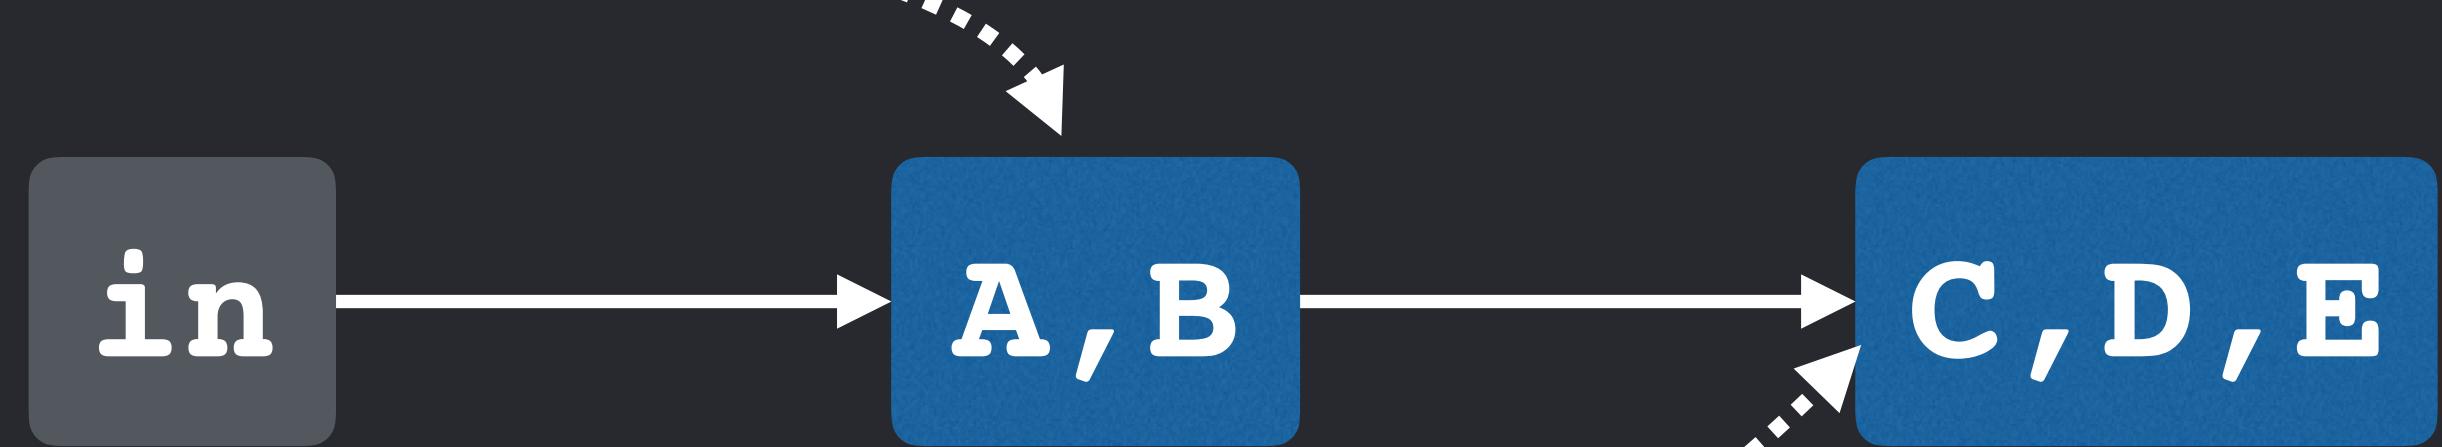
Input: DAG of pipeline stages



Output: Optimized schedule

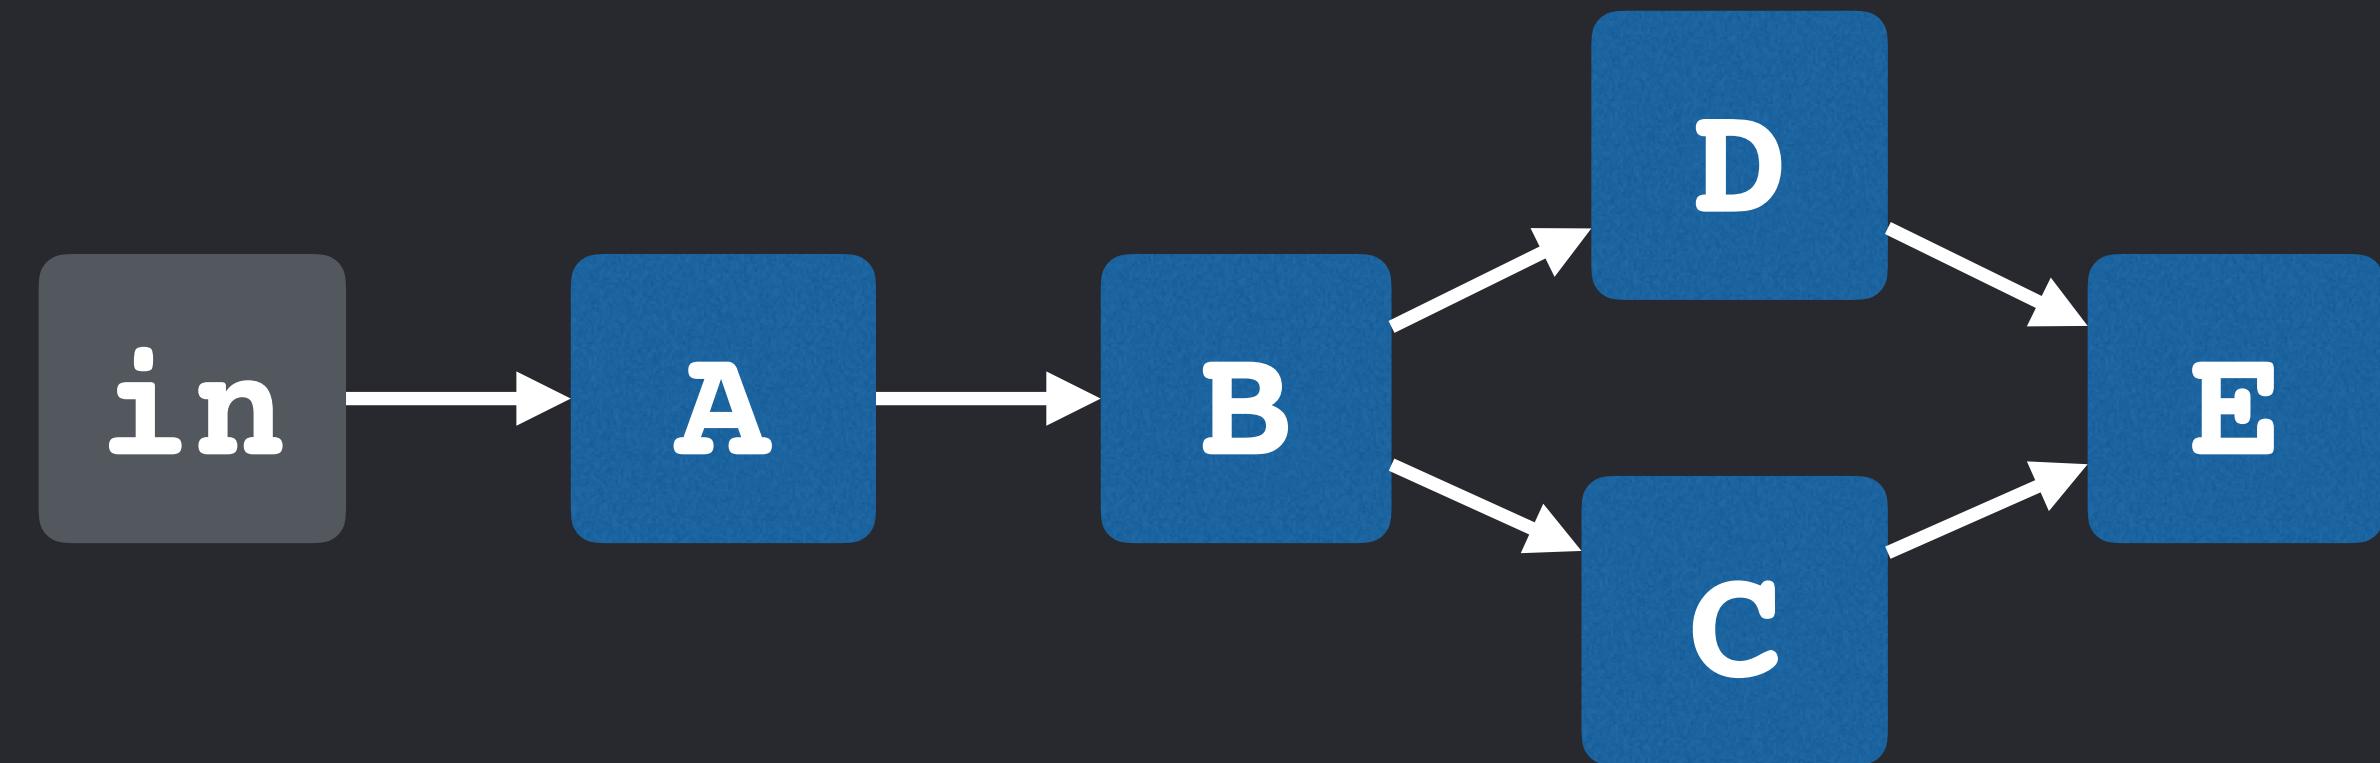
for each 8x128 tile in parallel
 compute required pixels of A
 compute pixels in tile of B

for each 8x8 tile in parallel
 compute required pixels of C
 compute required pixels of D
 compute pixels in tile of E



Algorithm

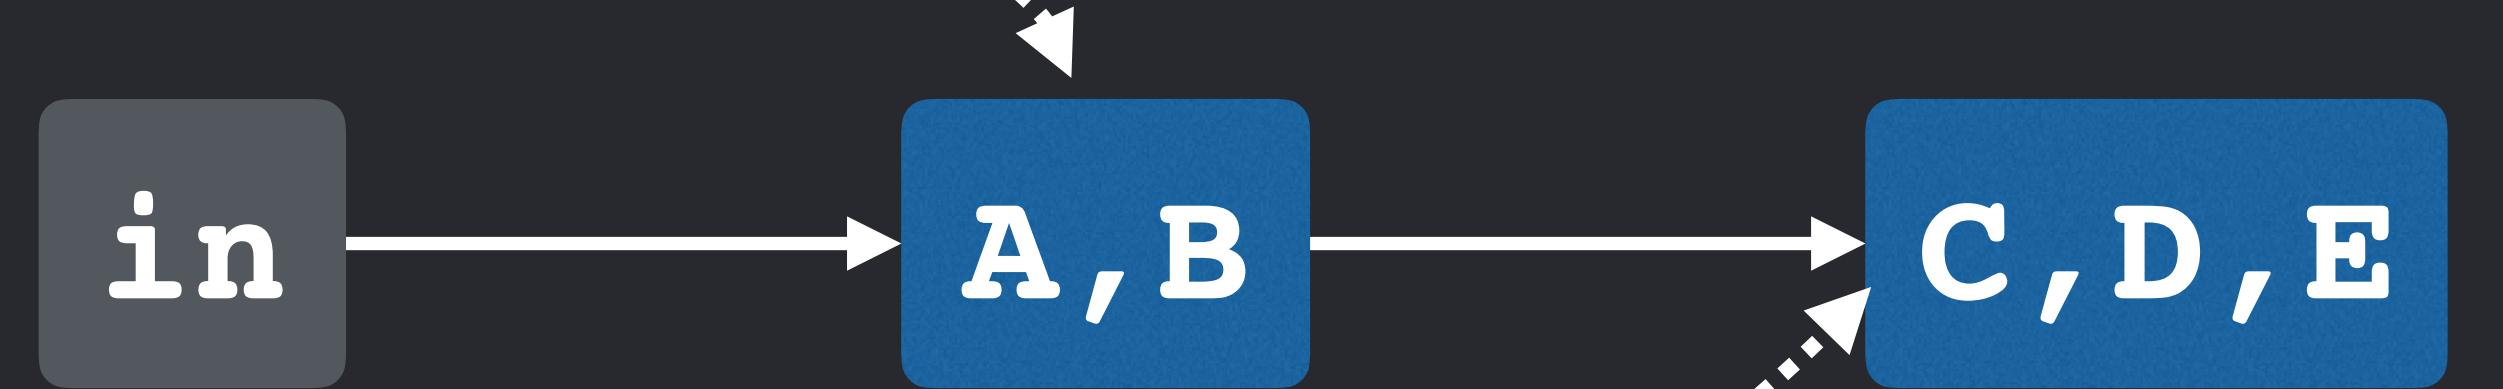
Input: DAG of pipeline stages



Output: Optimized schedule

for each 8x128 tile in parallel
 compute required pixels of A
 compute pixels in tile of B

for each 8x8 tile in parallel
 compute required pixels of C
 compute required pixels of D
 compute pixels in tile of E



Tile size: 8 x 128

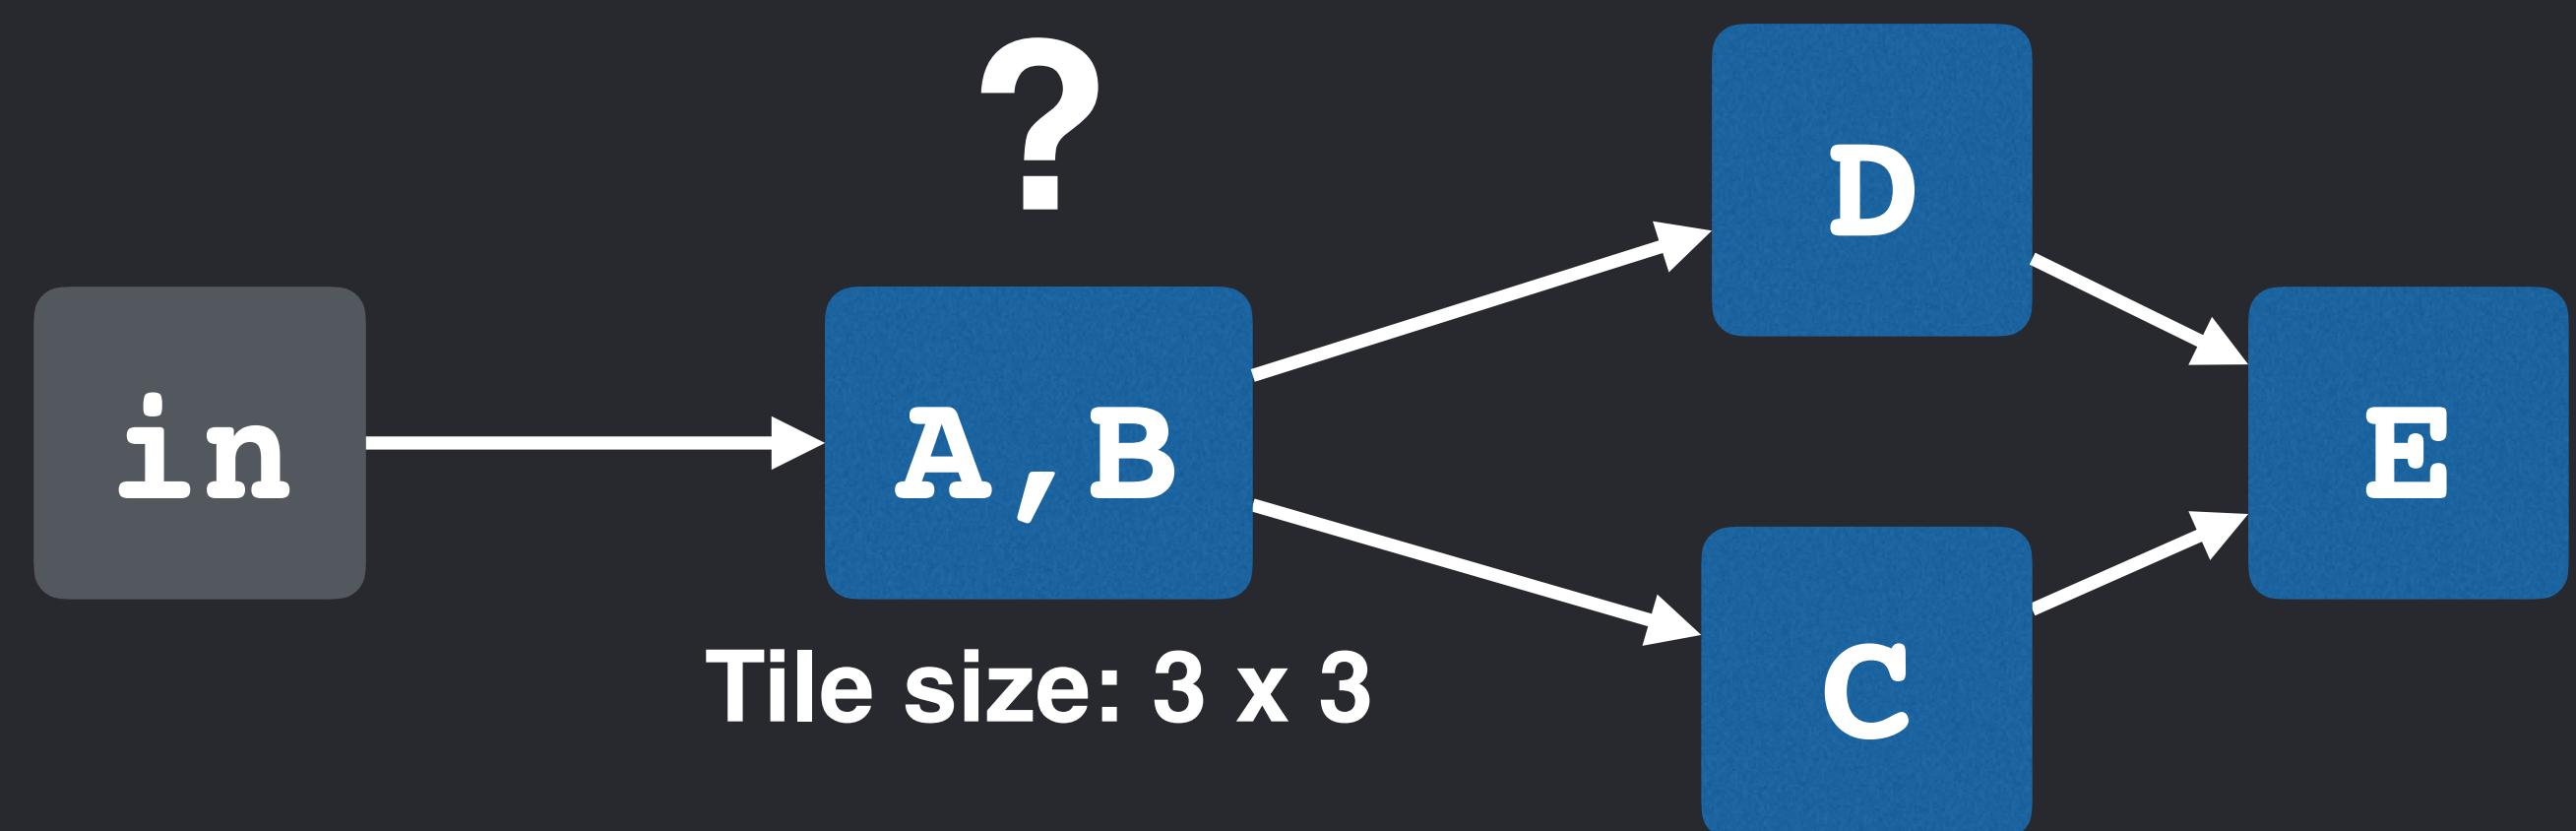
Tile size: 8 x 8

Scheduling the DAG for better locality

Determine which stages to group together?

How to tile stages in each group?

When to group stages?

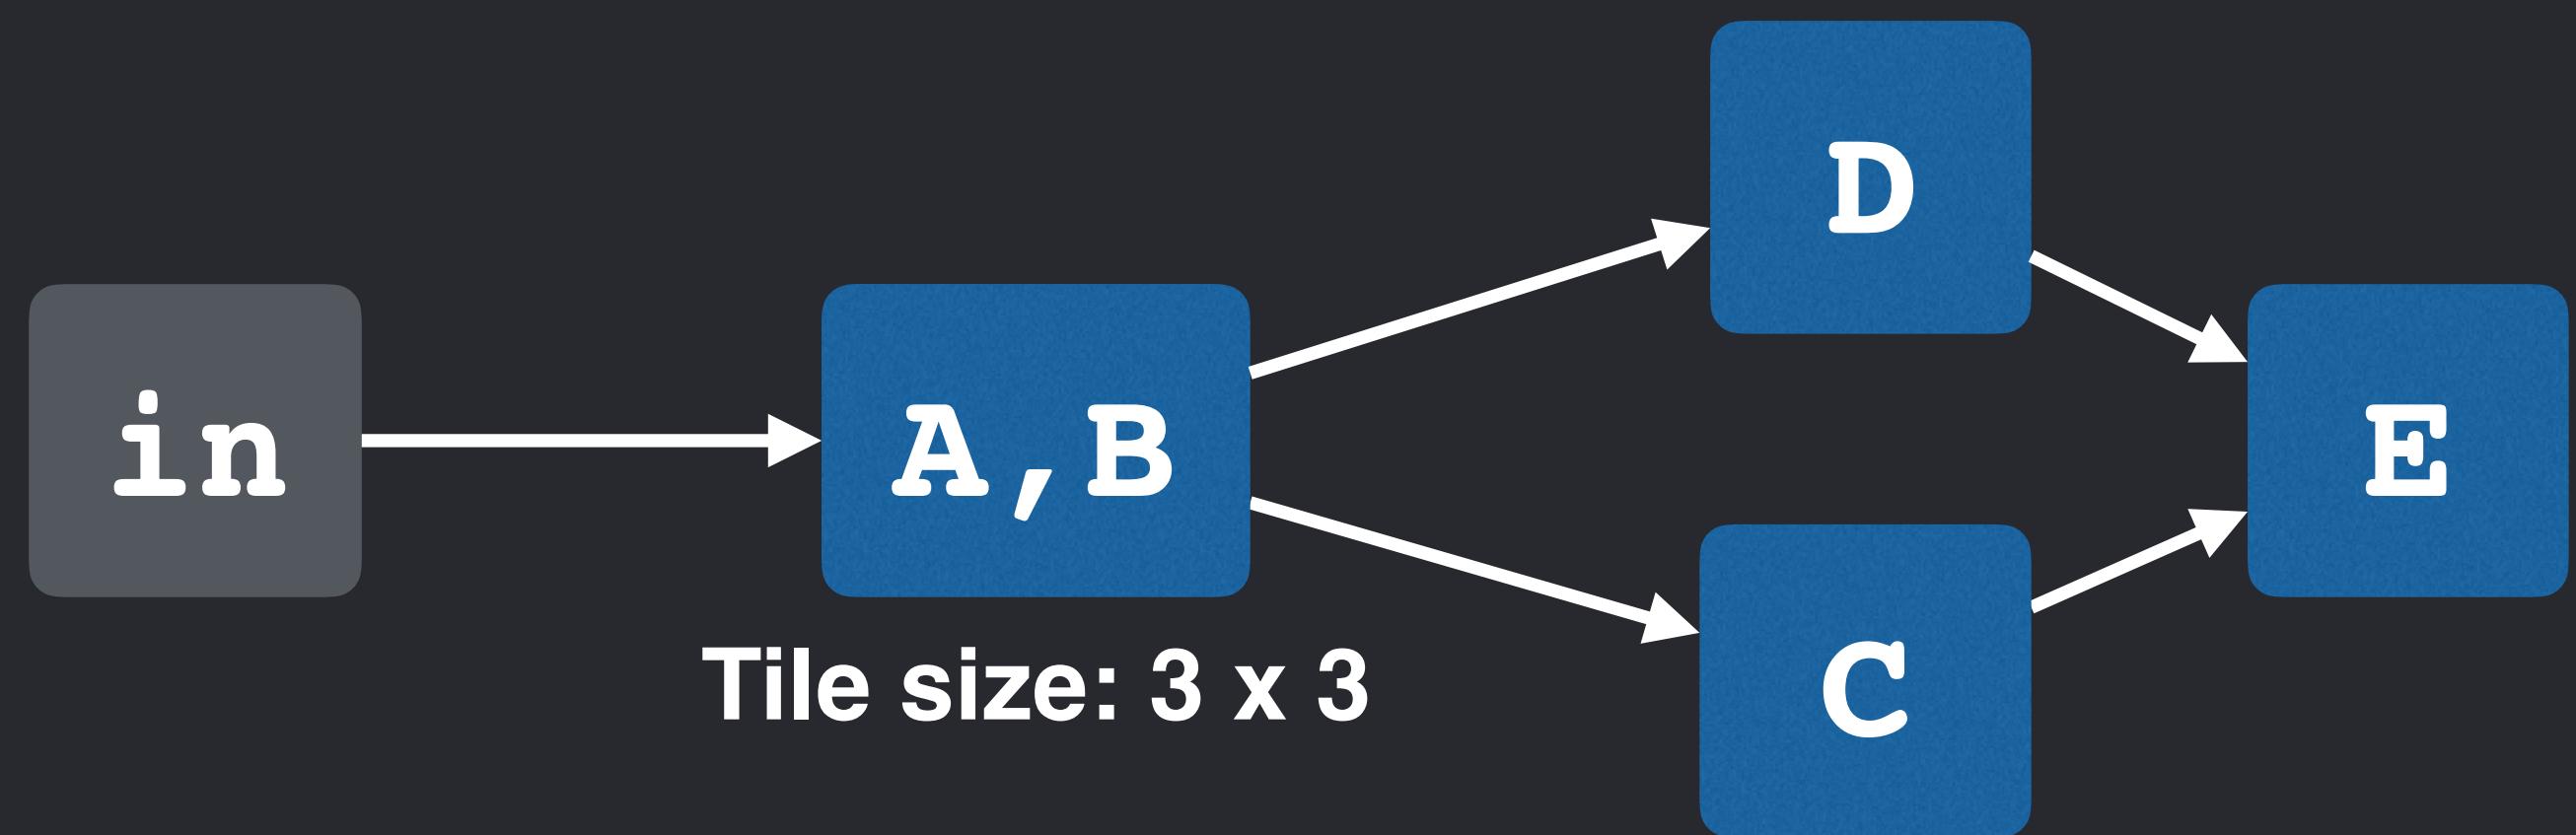


for each 3x3 tile in parallel
compute required pixels of A
compute pixels in tile of B

compute all pixels of C, in parallel
compute all pixels of D, in parallel
compute all pixels of E, in parallel

Grouping A and B together can either improve or degrade performance

Quantifying the cost of a group

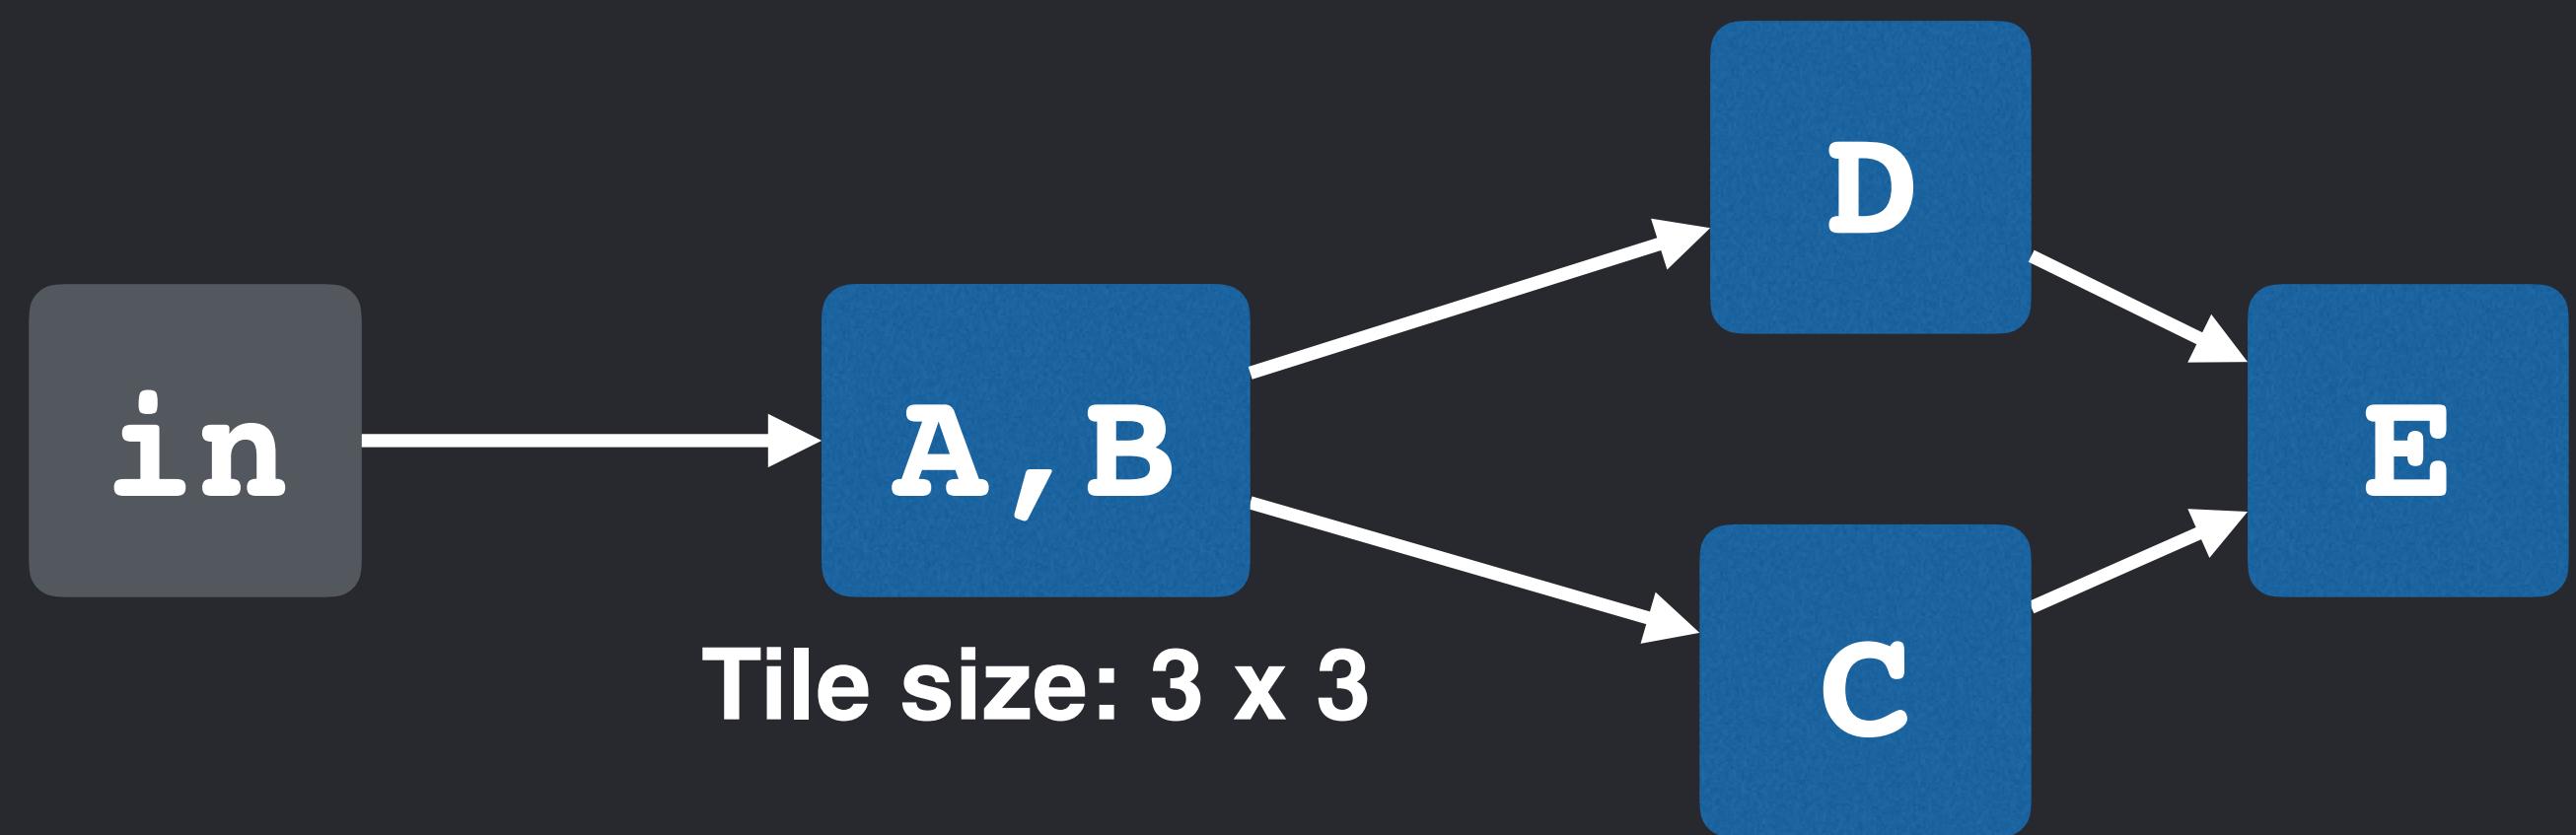


for each 3x3 tile in parallel
compute required pixels of A
compute pixels in tile of B

compute all pixels of C, in parallel
compute all pixels of D, in parallel
compute all pixels of E, in parallel

Cost = Cost of arithmetic + Cost of memory

Quantifying the cost of a group

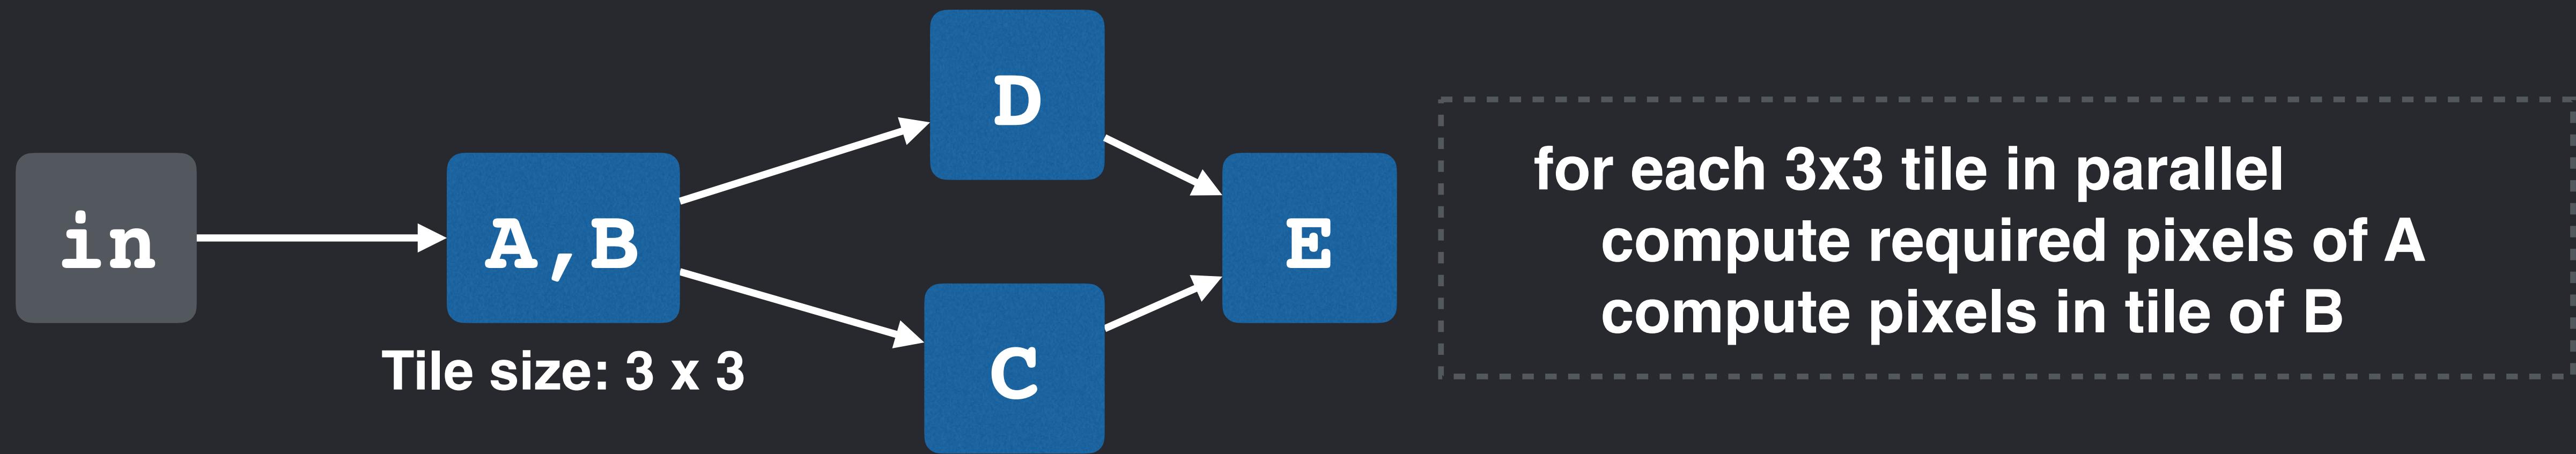


for each 3x3 tile in parallel
compute required pixels of A
compute pixels in tile of B

compute all pixels of C, in parallel
compute all pixels of D, in parallel
compute all pixels of E, in parallel

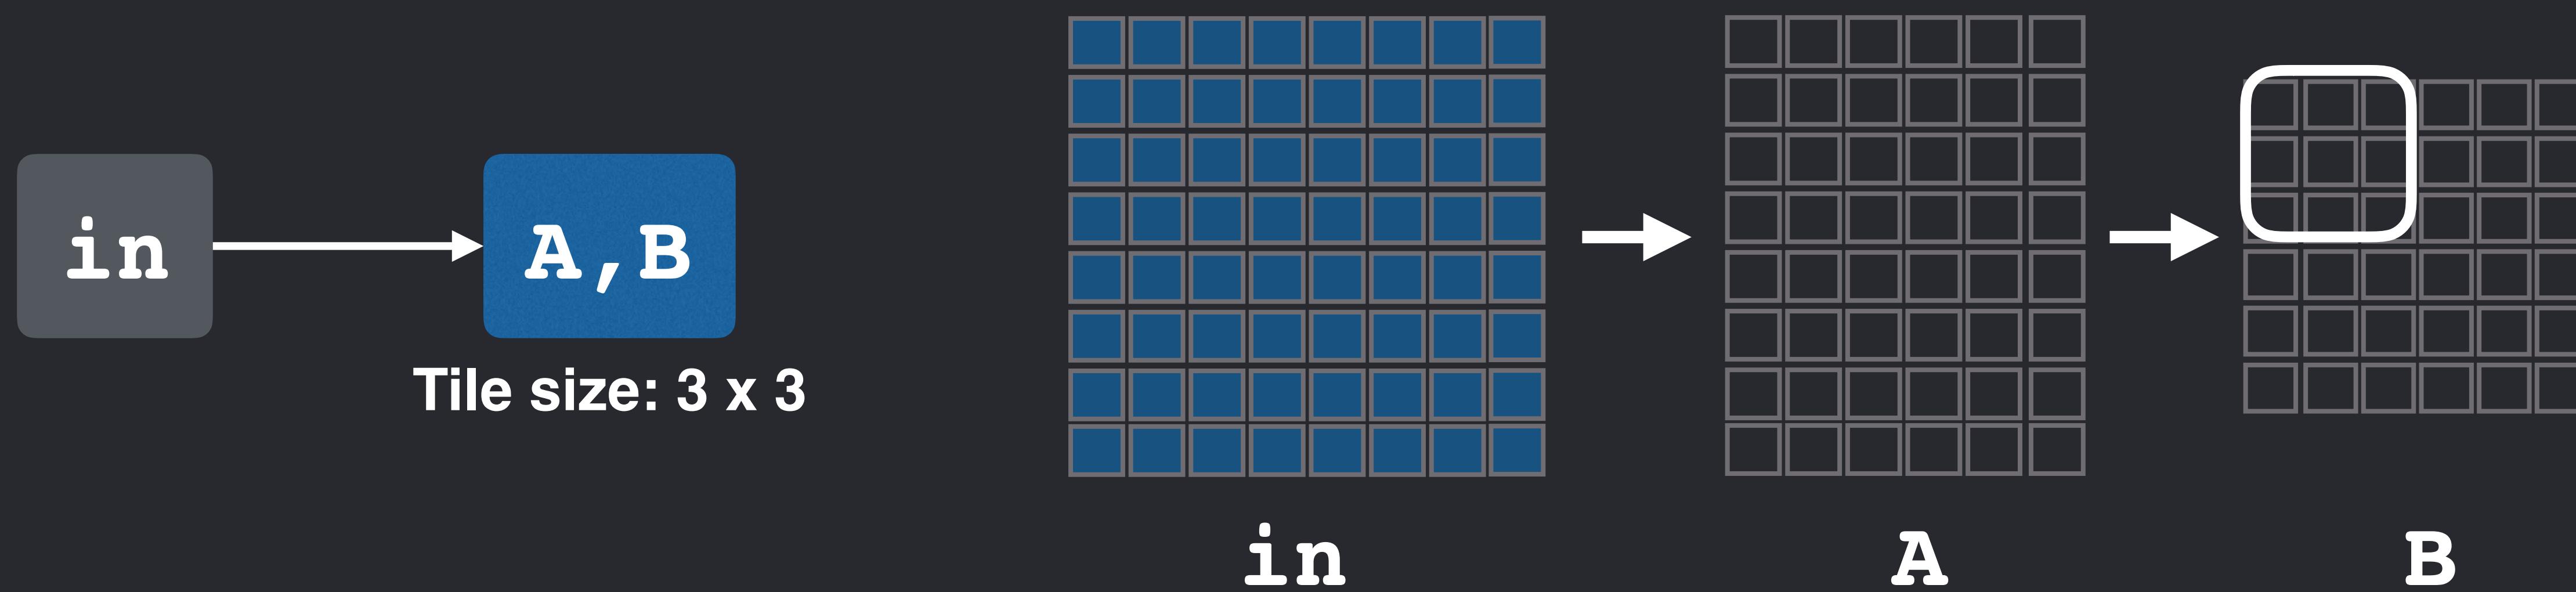
$$\text{Cost} = (\text{Number of arithmetic operations}) + (\text{Number of memory accesses}) \times (\text{LOAD COST})$$

Quantifying the cost of a group



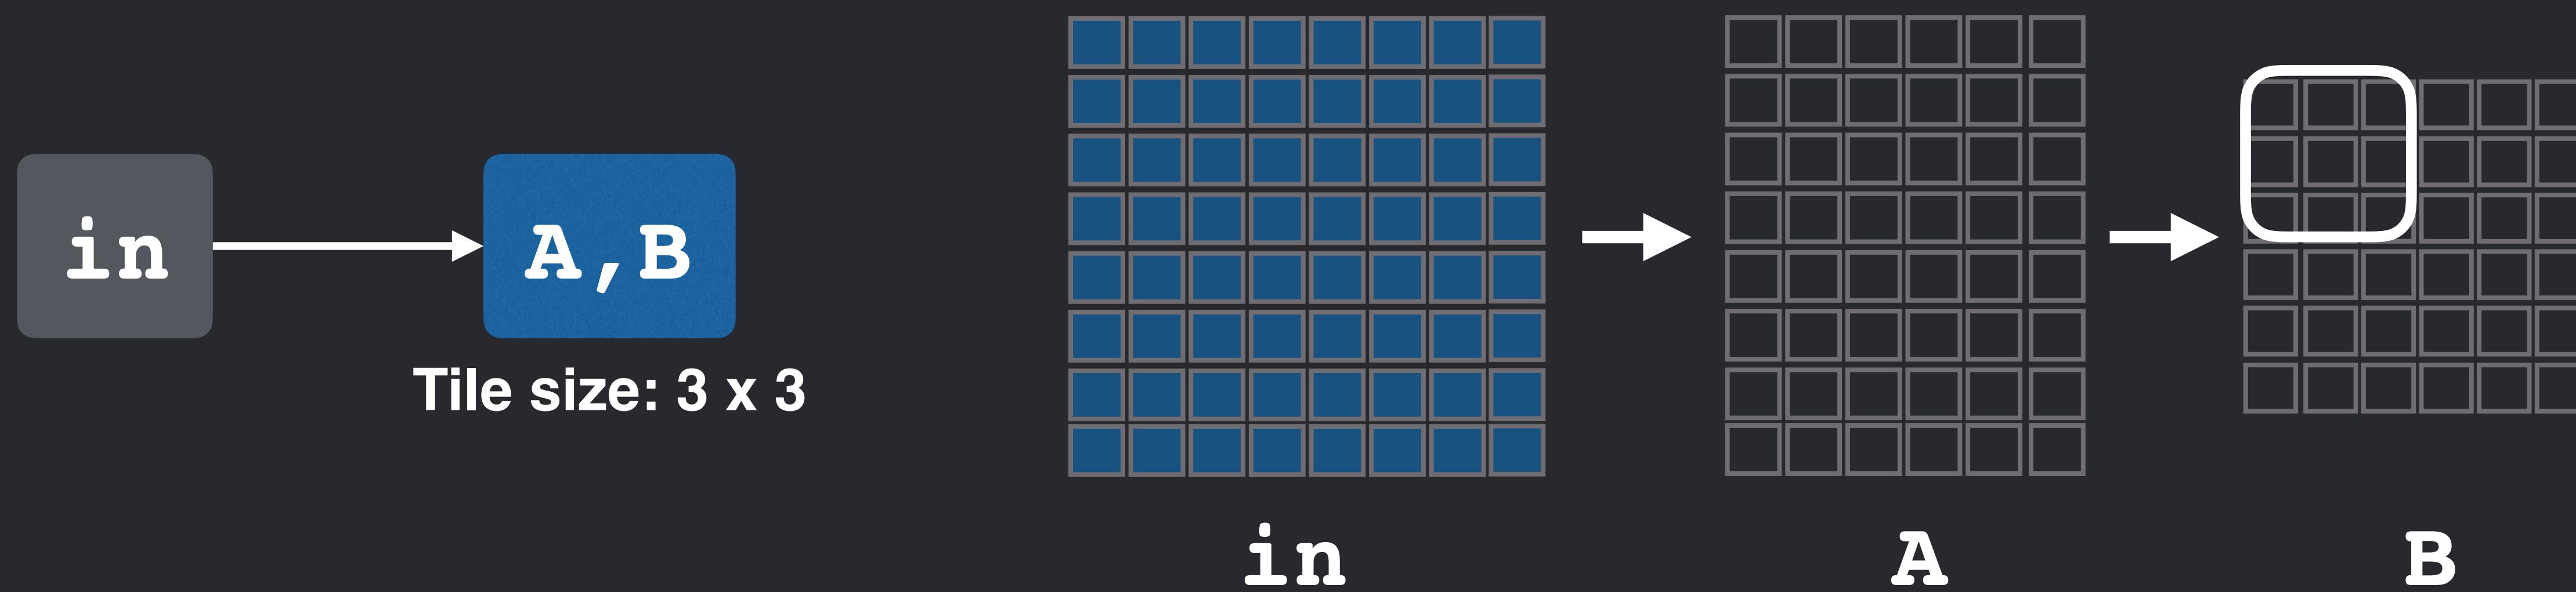
Cost = (Number of arithmetic operations) +
(Number of memory accesses) x (LOAD COST)

Estimating cost using interval analysis



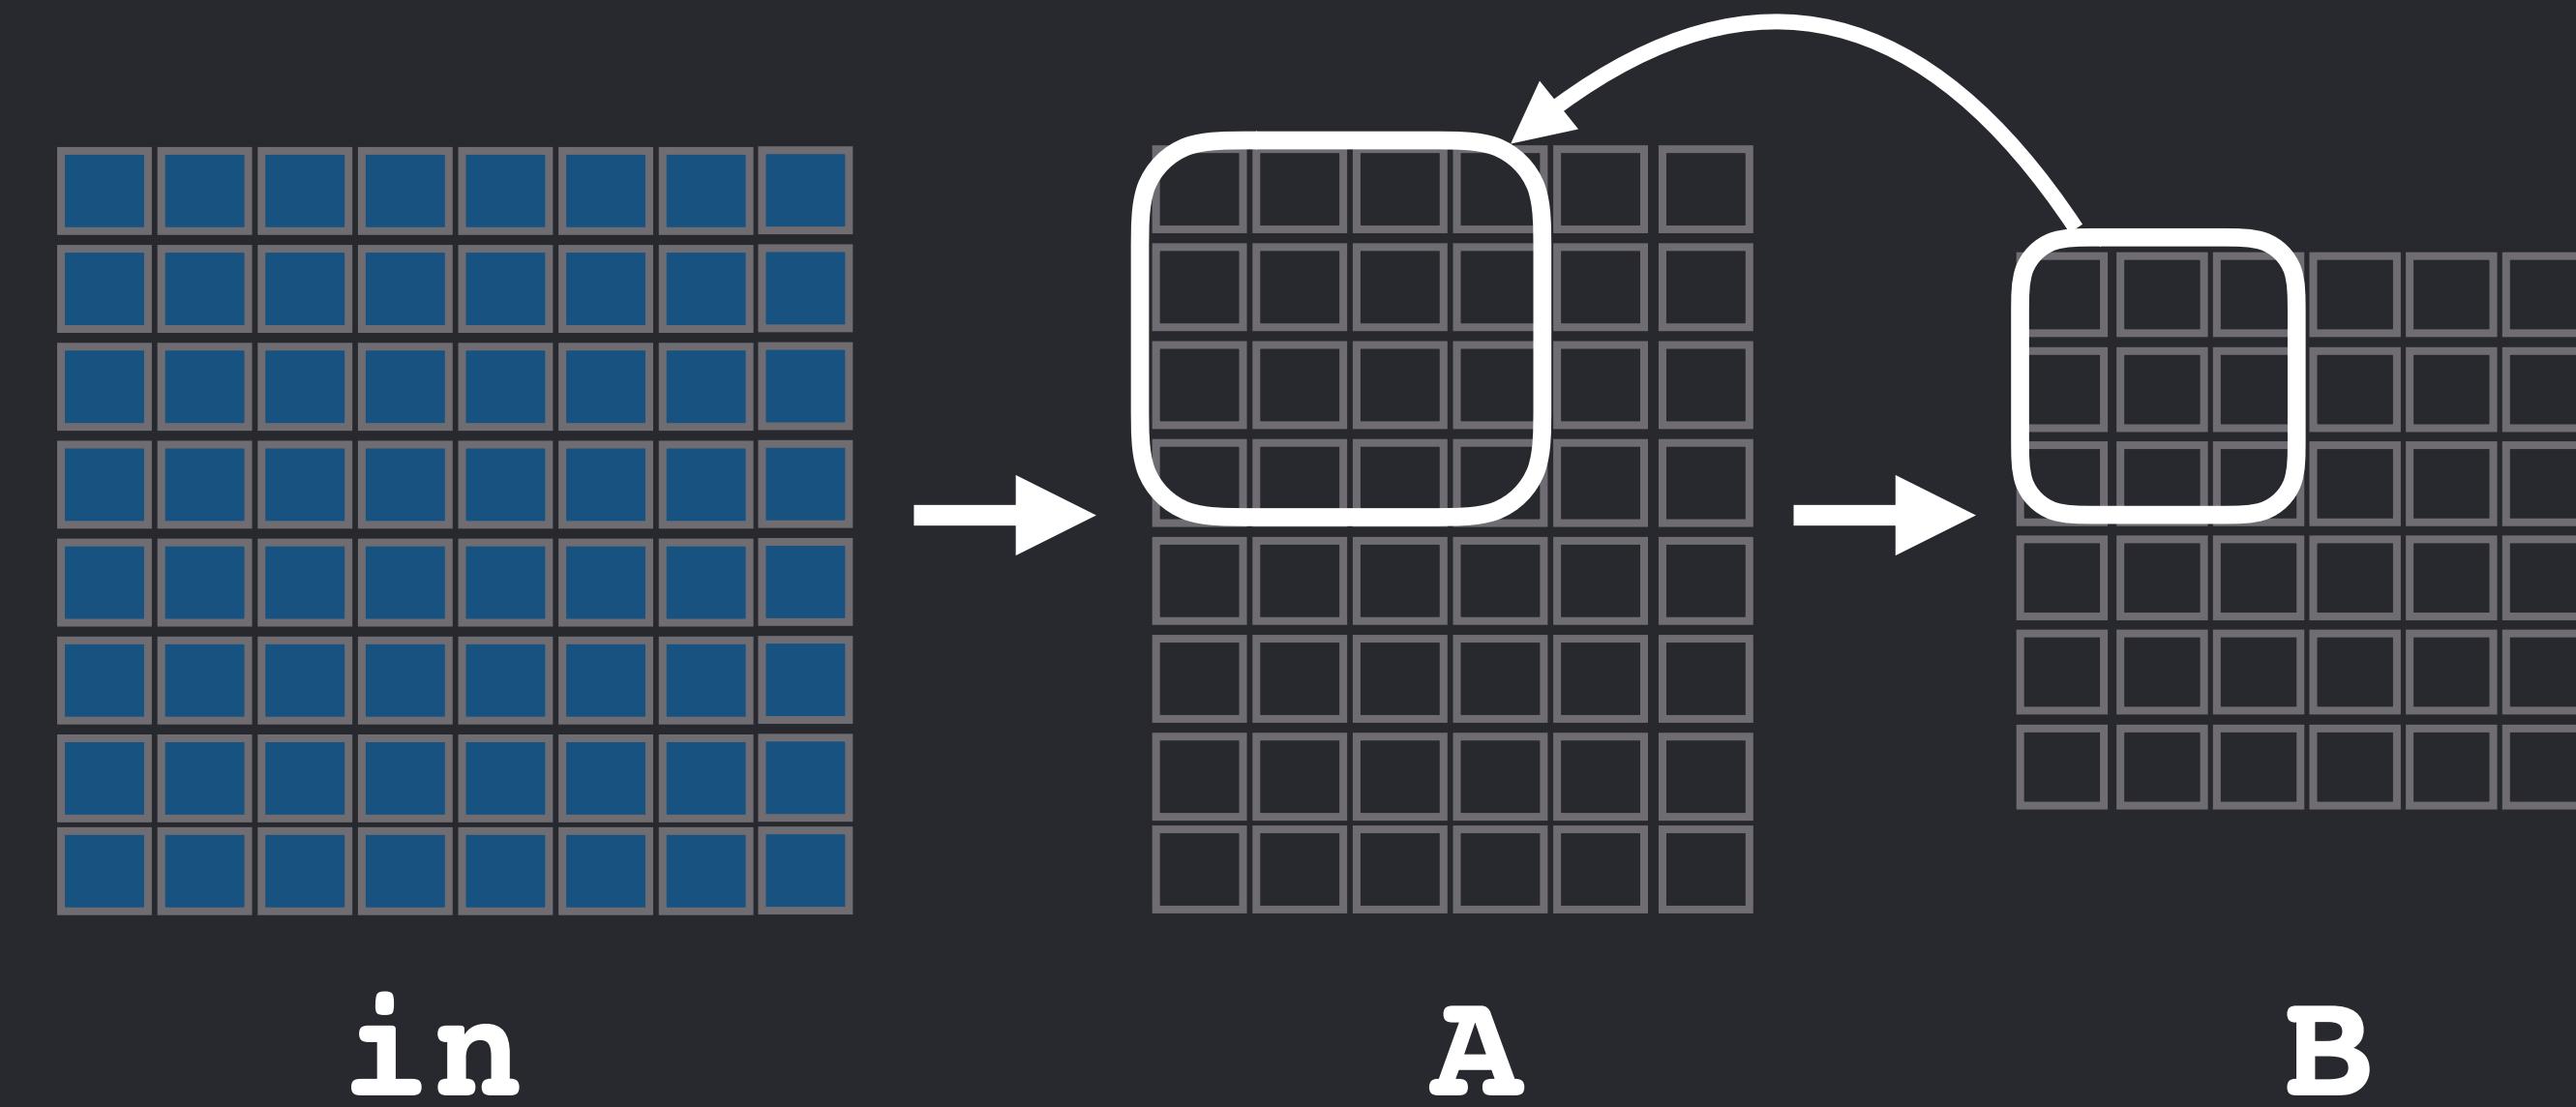
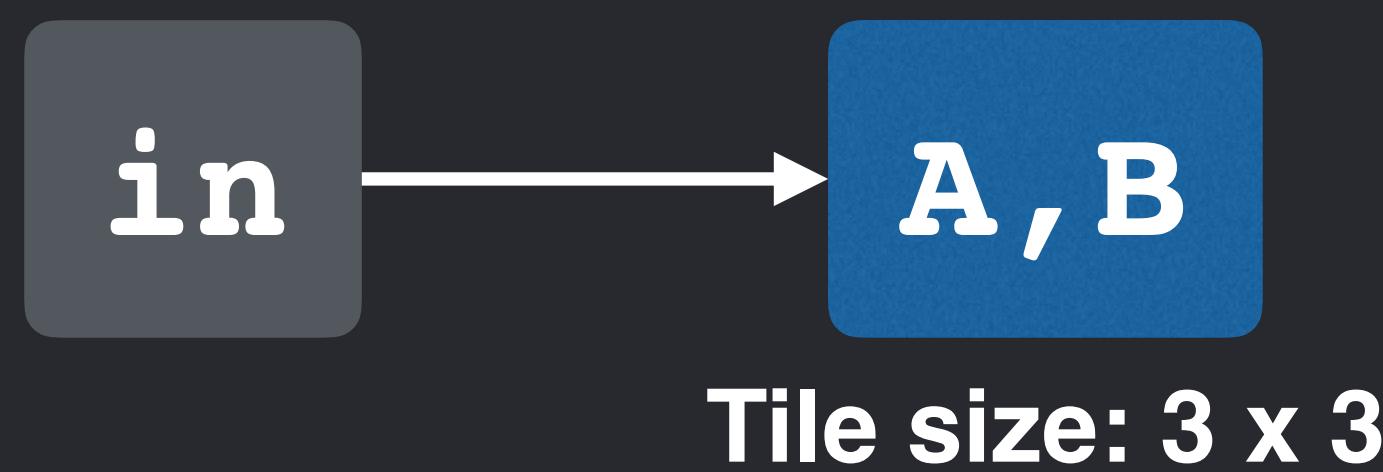
Cost = (Number of arithmetic operations) +
(Number of memory accesses) x (LOAD COST)

Estimating cost using interval analysis



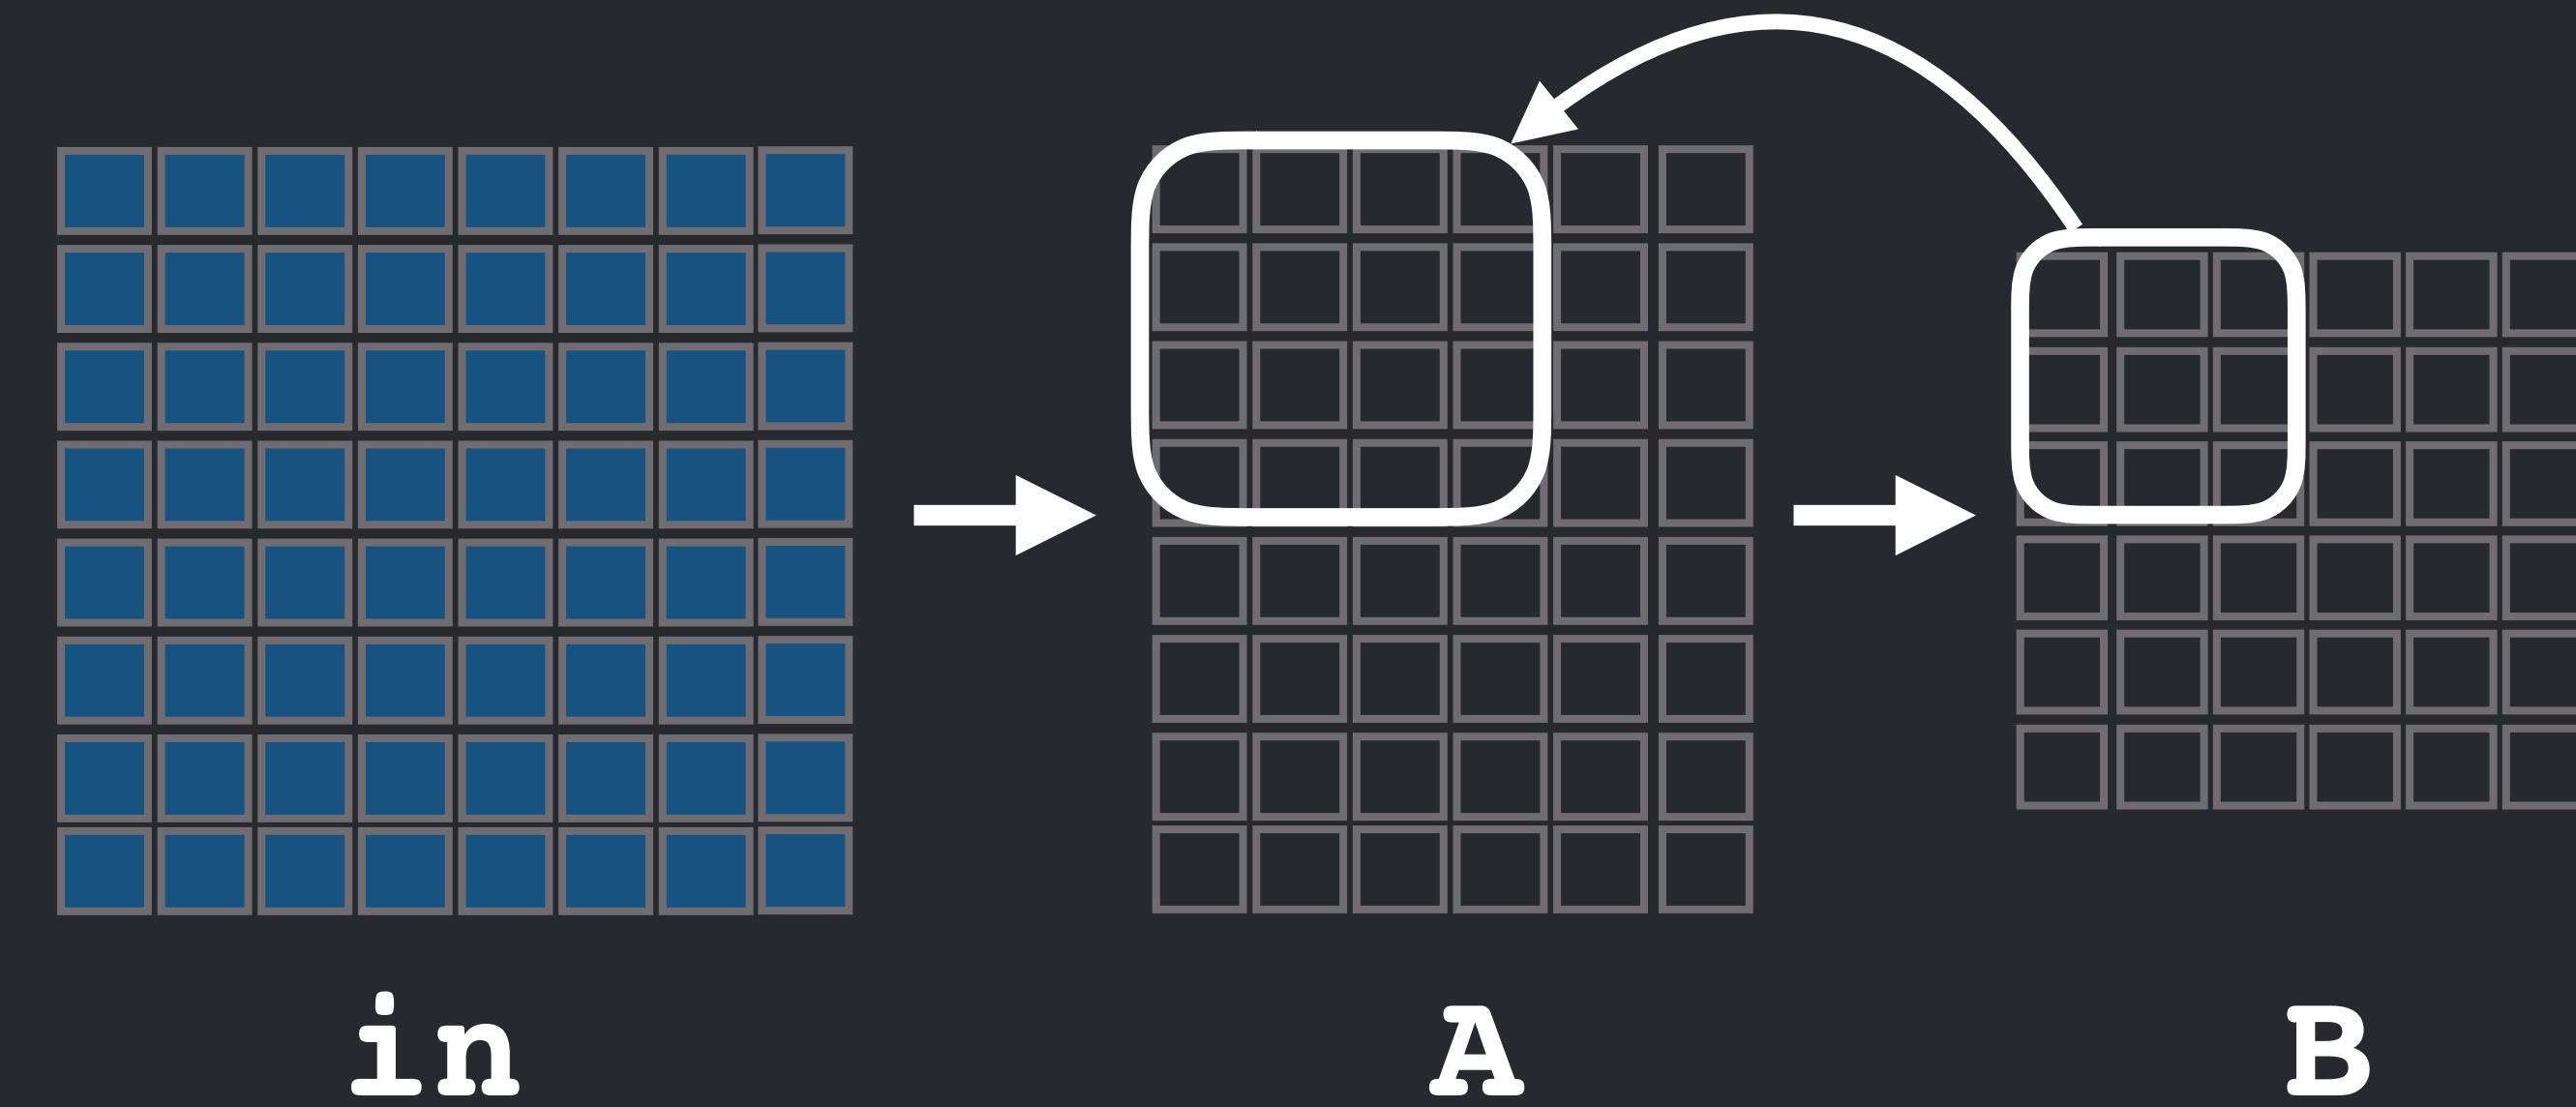
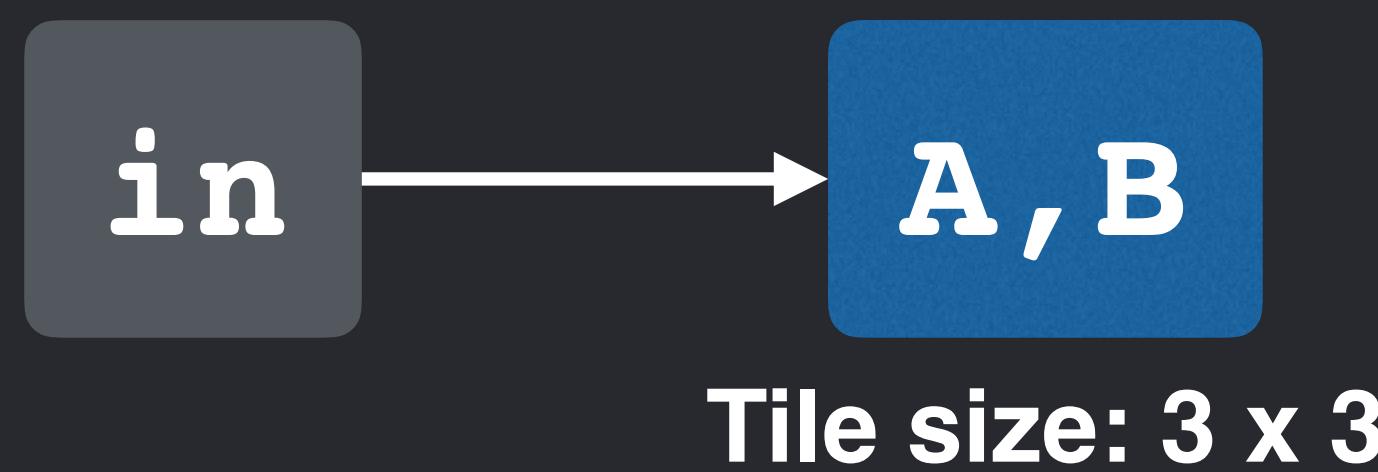
Cost = **Number of arithmetic operations** +
(Number of memory accesses) x (LOAD COST)

Estimating cost using interval analysis



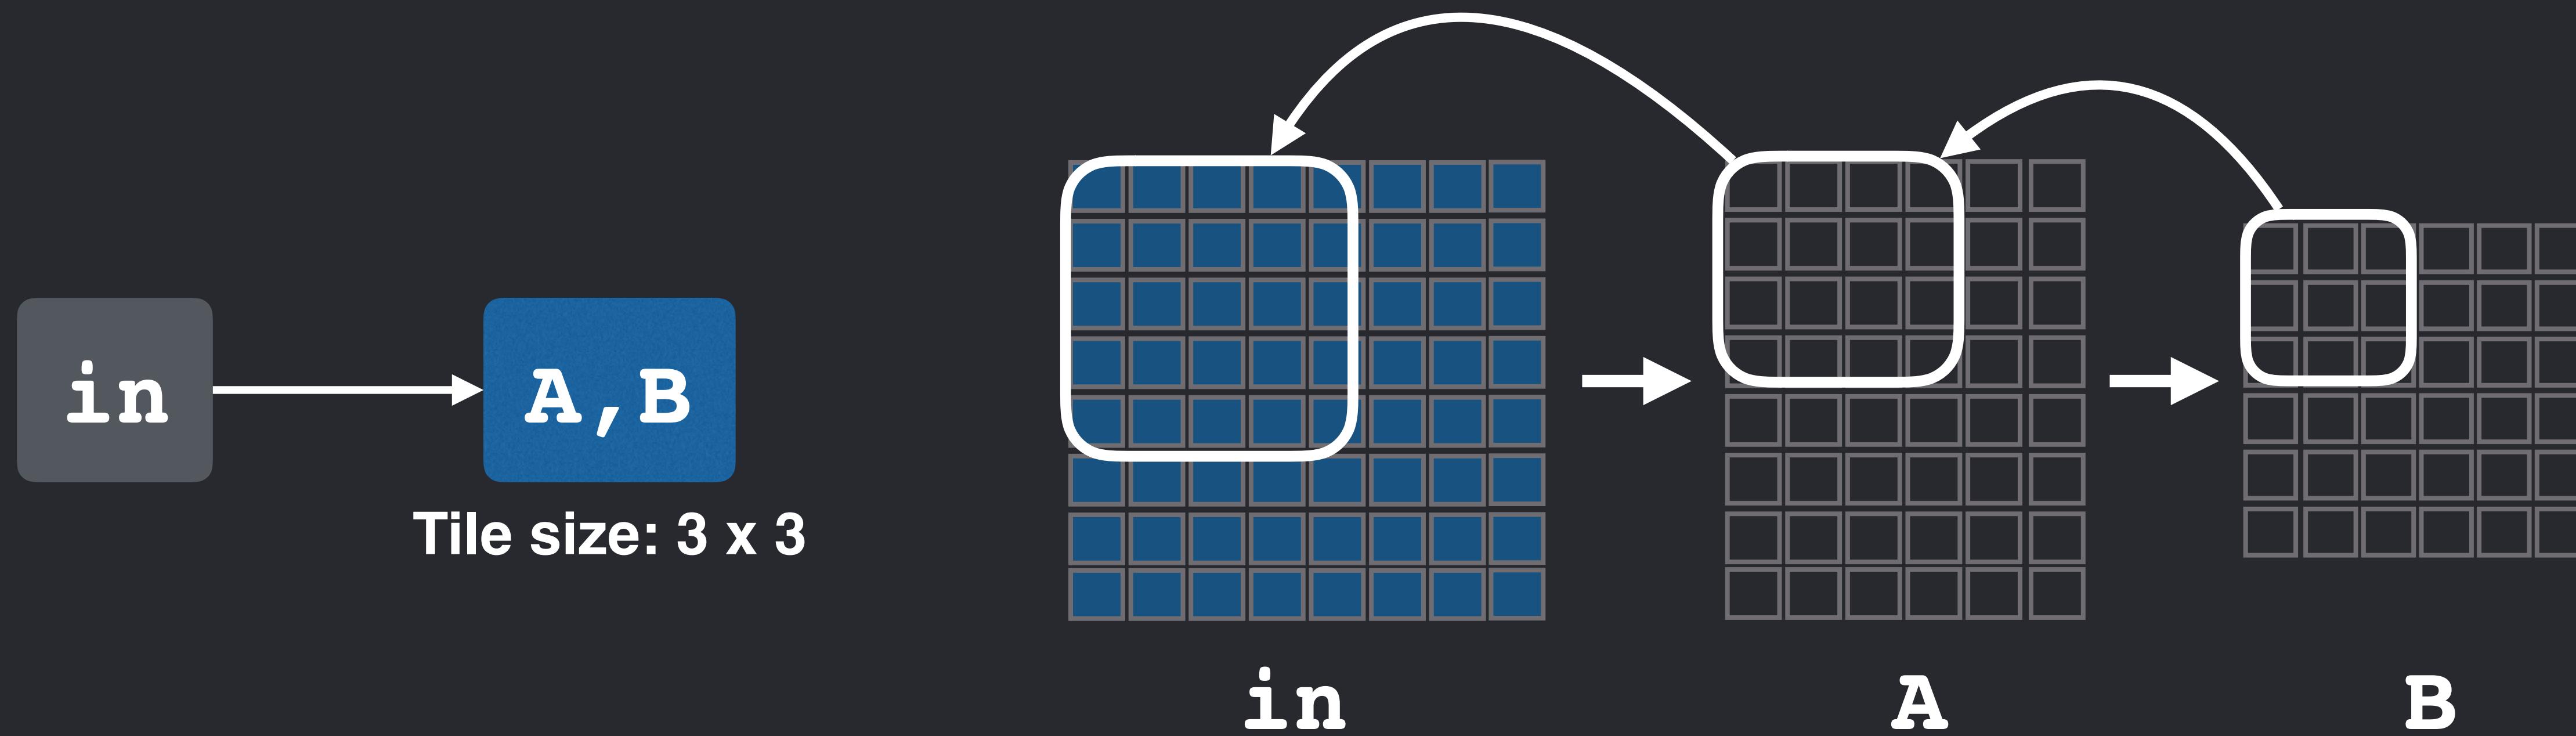
Cost = **Number of arithmetic operations** +
(Number of memory accesses) x (LOAD COST)

Estimating cost using interval analysis



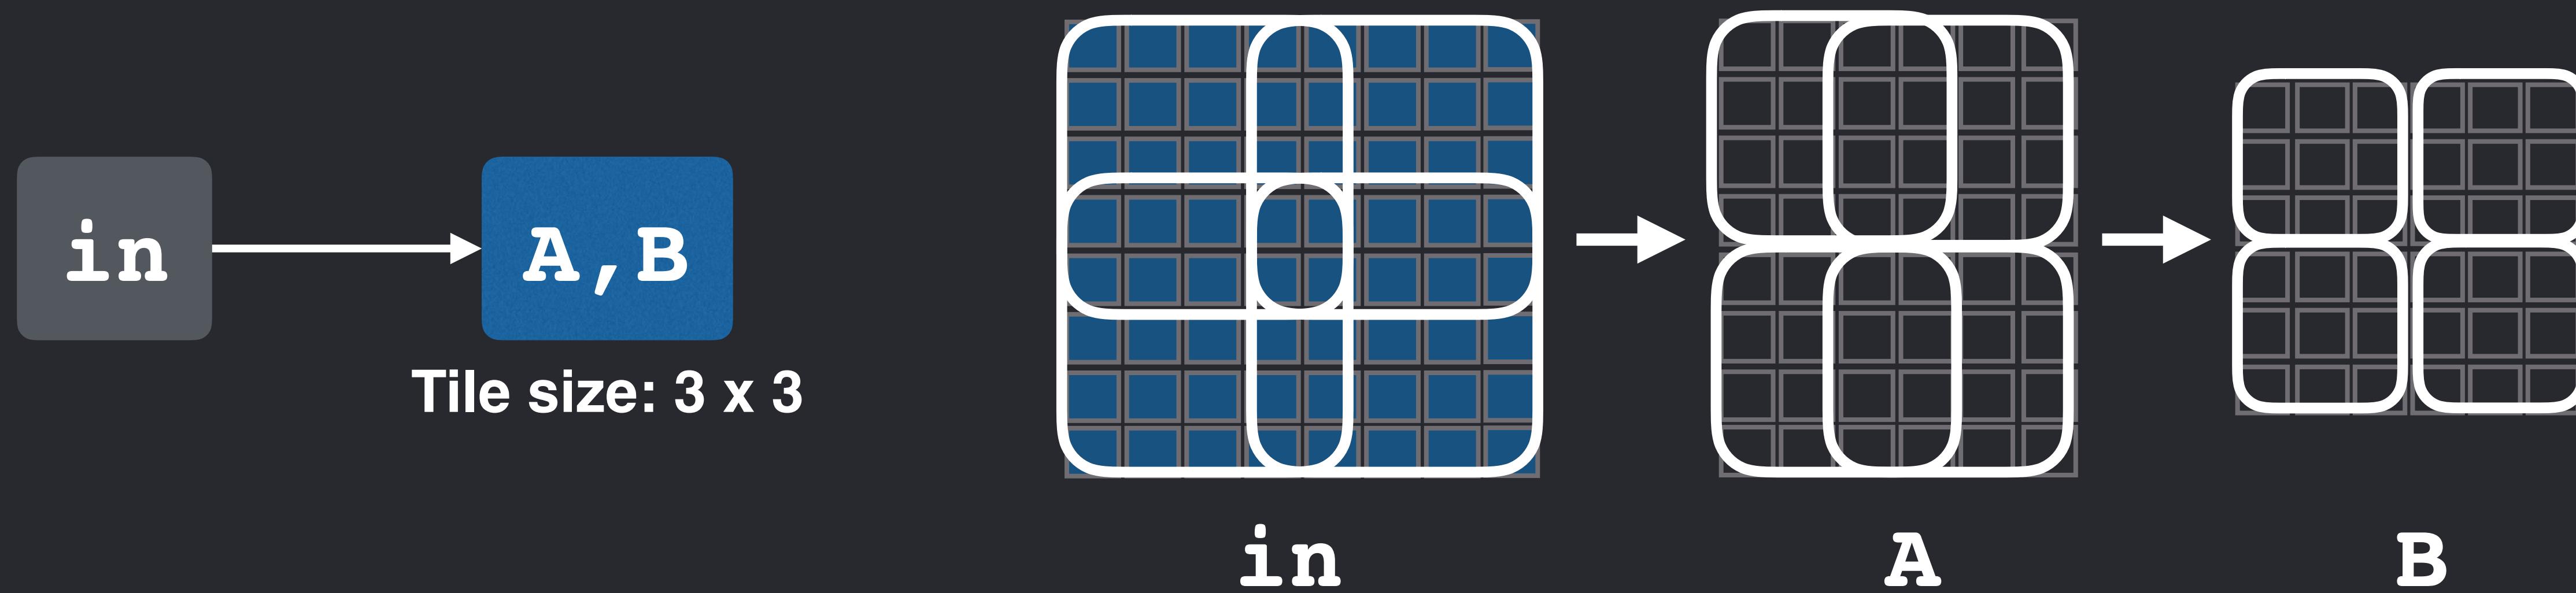
Cost = (Number of arithmetic operations) +
(Number of memory accesses) x (LOAD COST)

Estimating cost using interval analysis



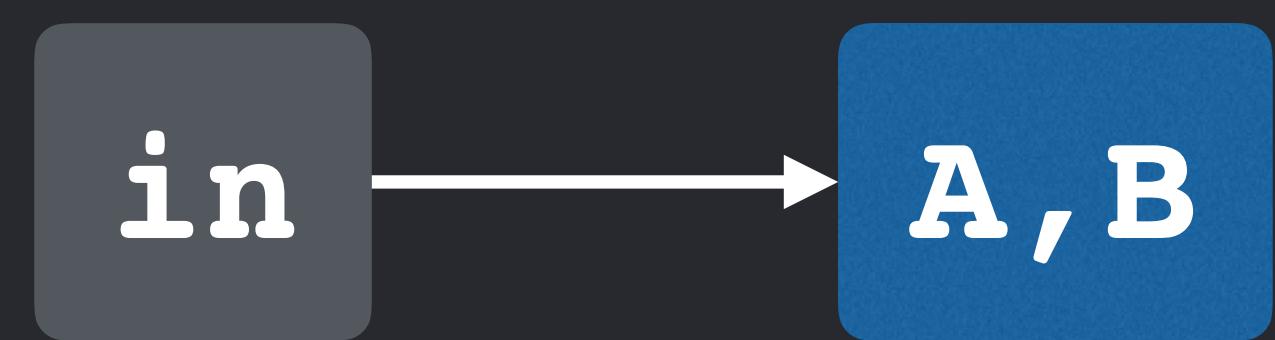
Cost = (Number of arithmetic operations) +
(Number of memory accesses) x (LOAD COST)

Estimating cost using interval analysis

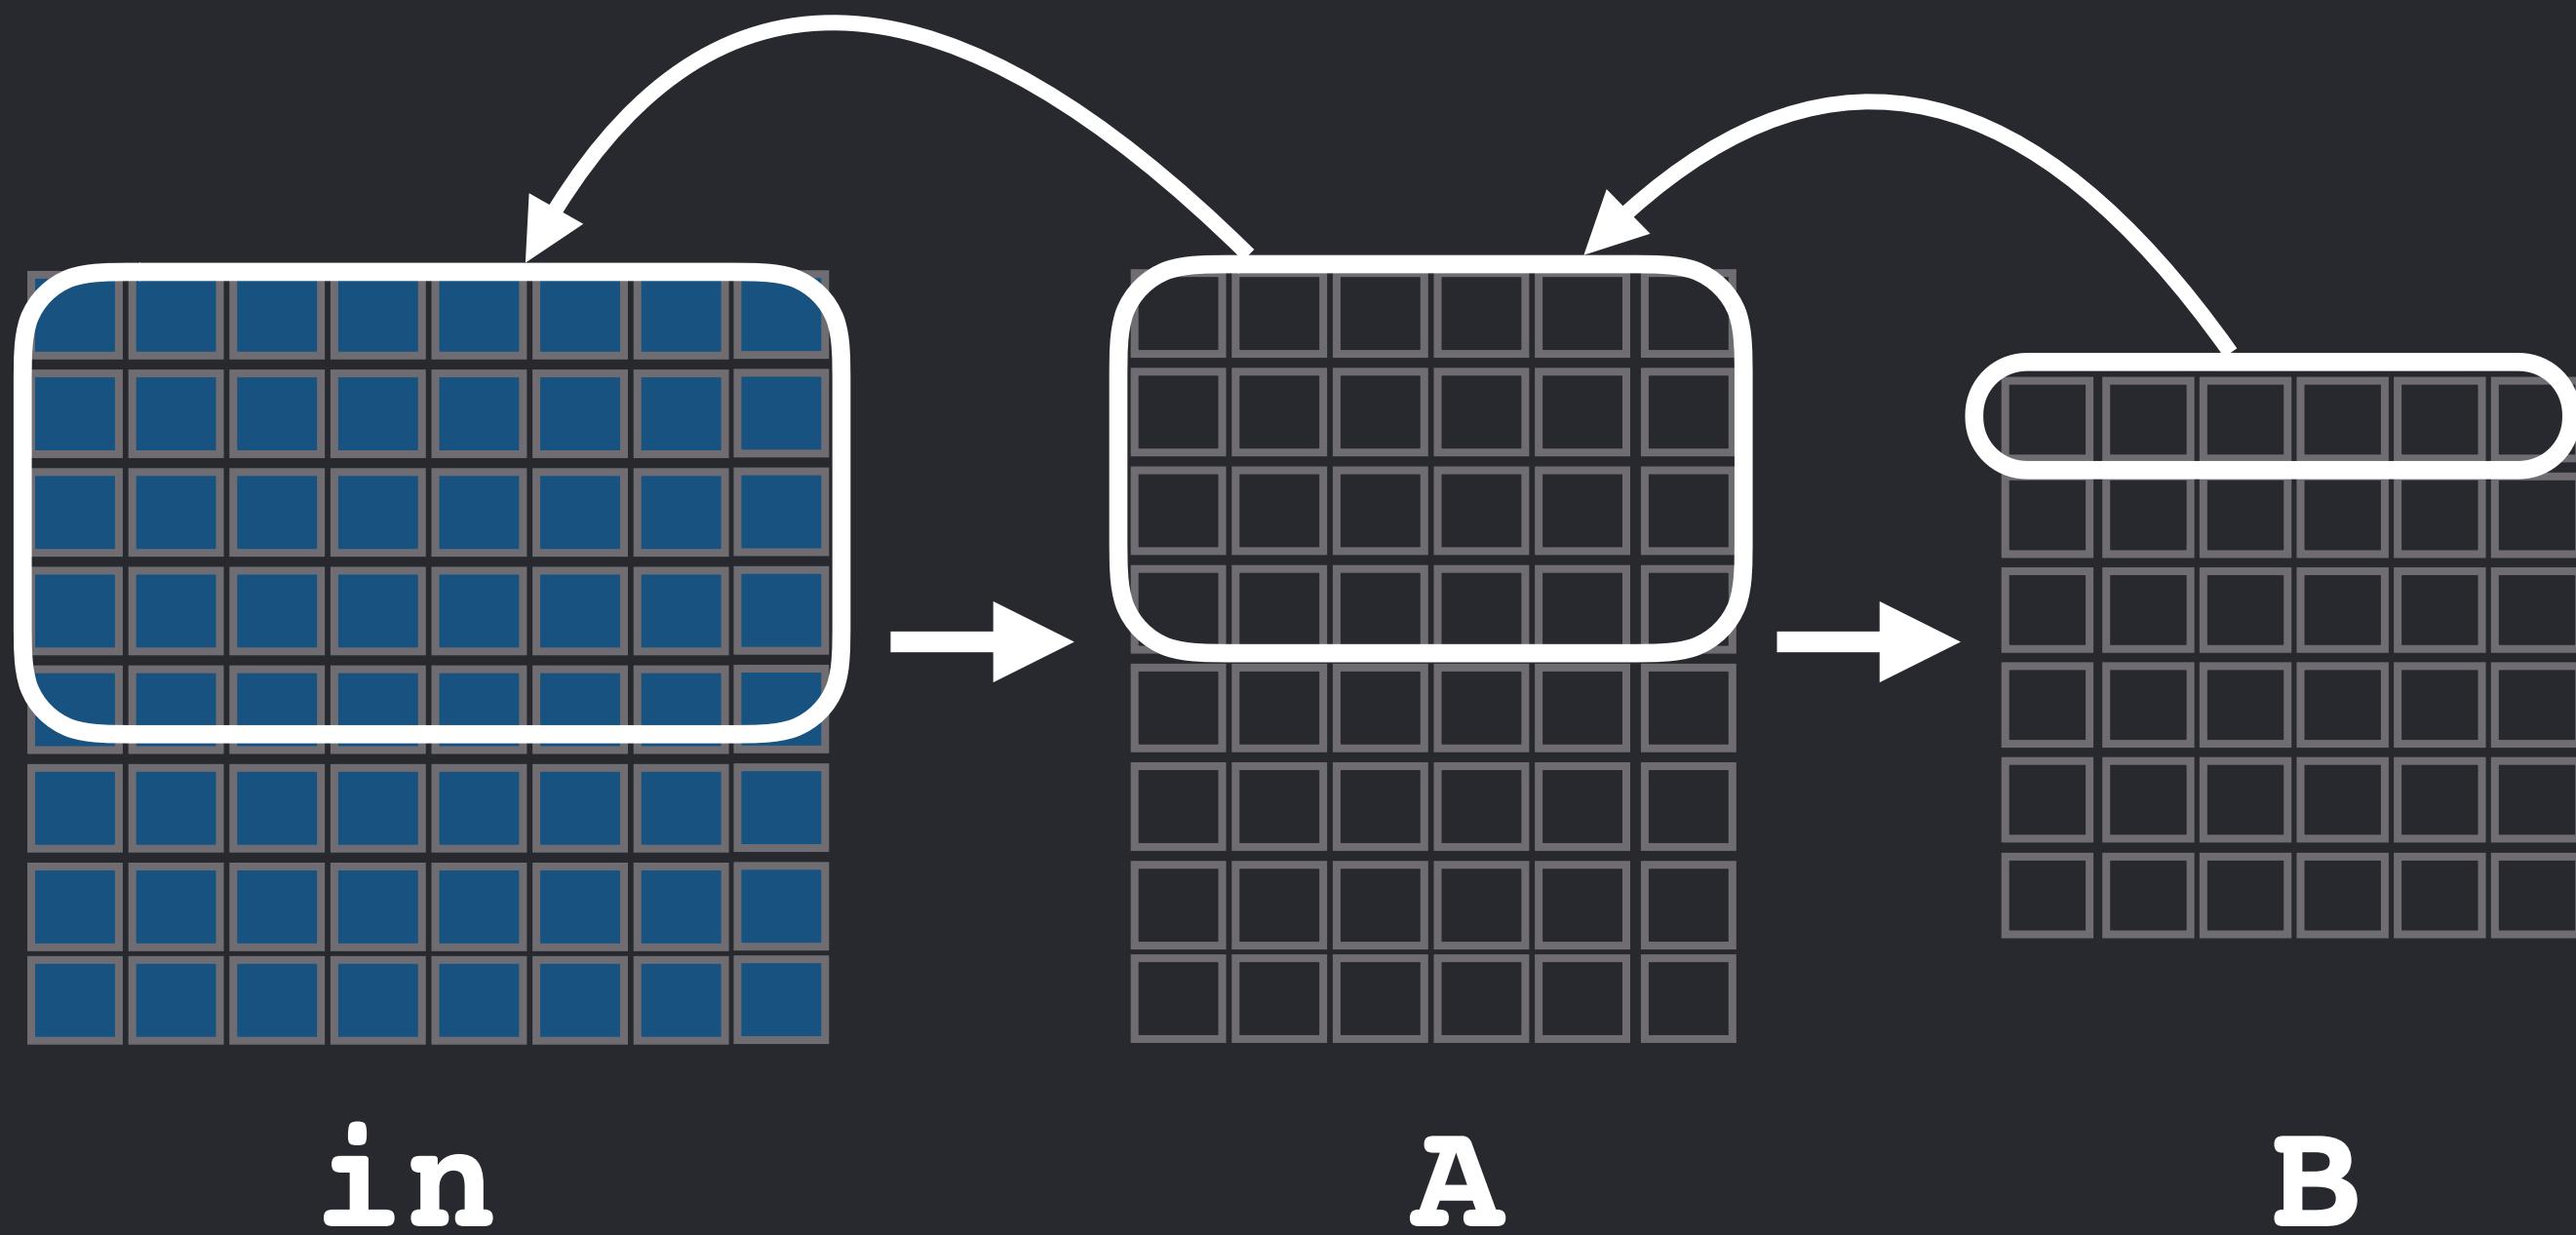


Cost = Number of tiles \times Cost per tile

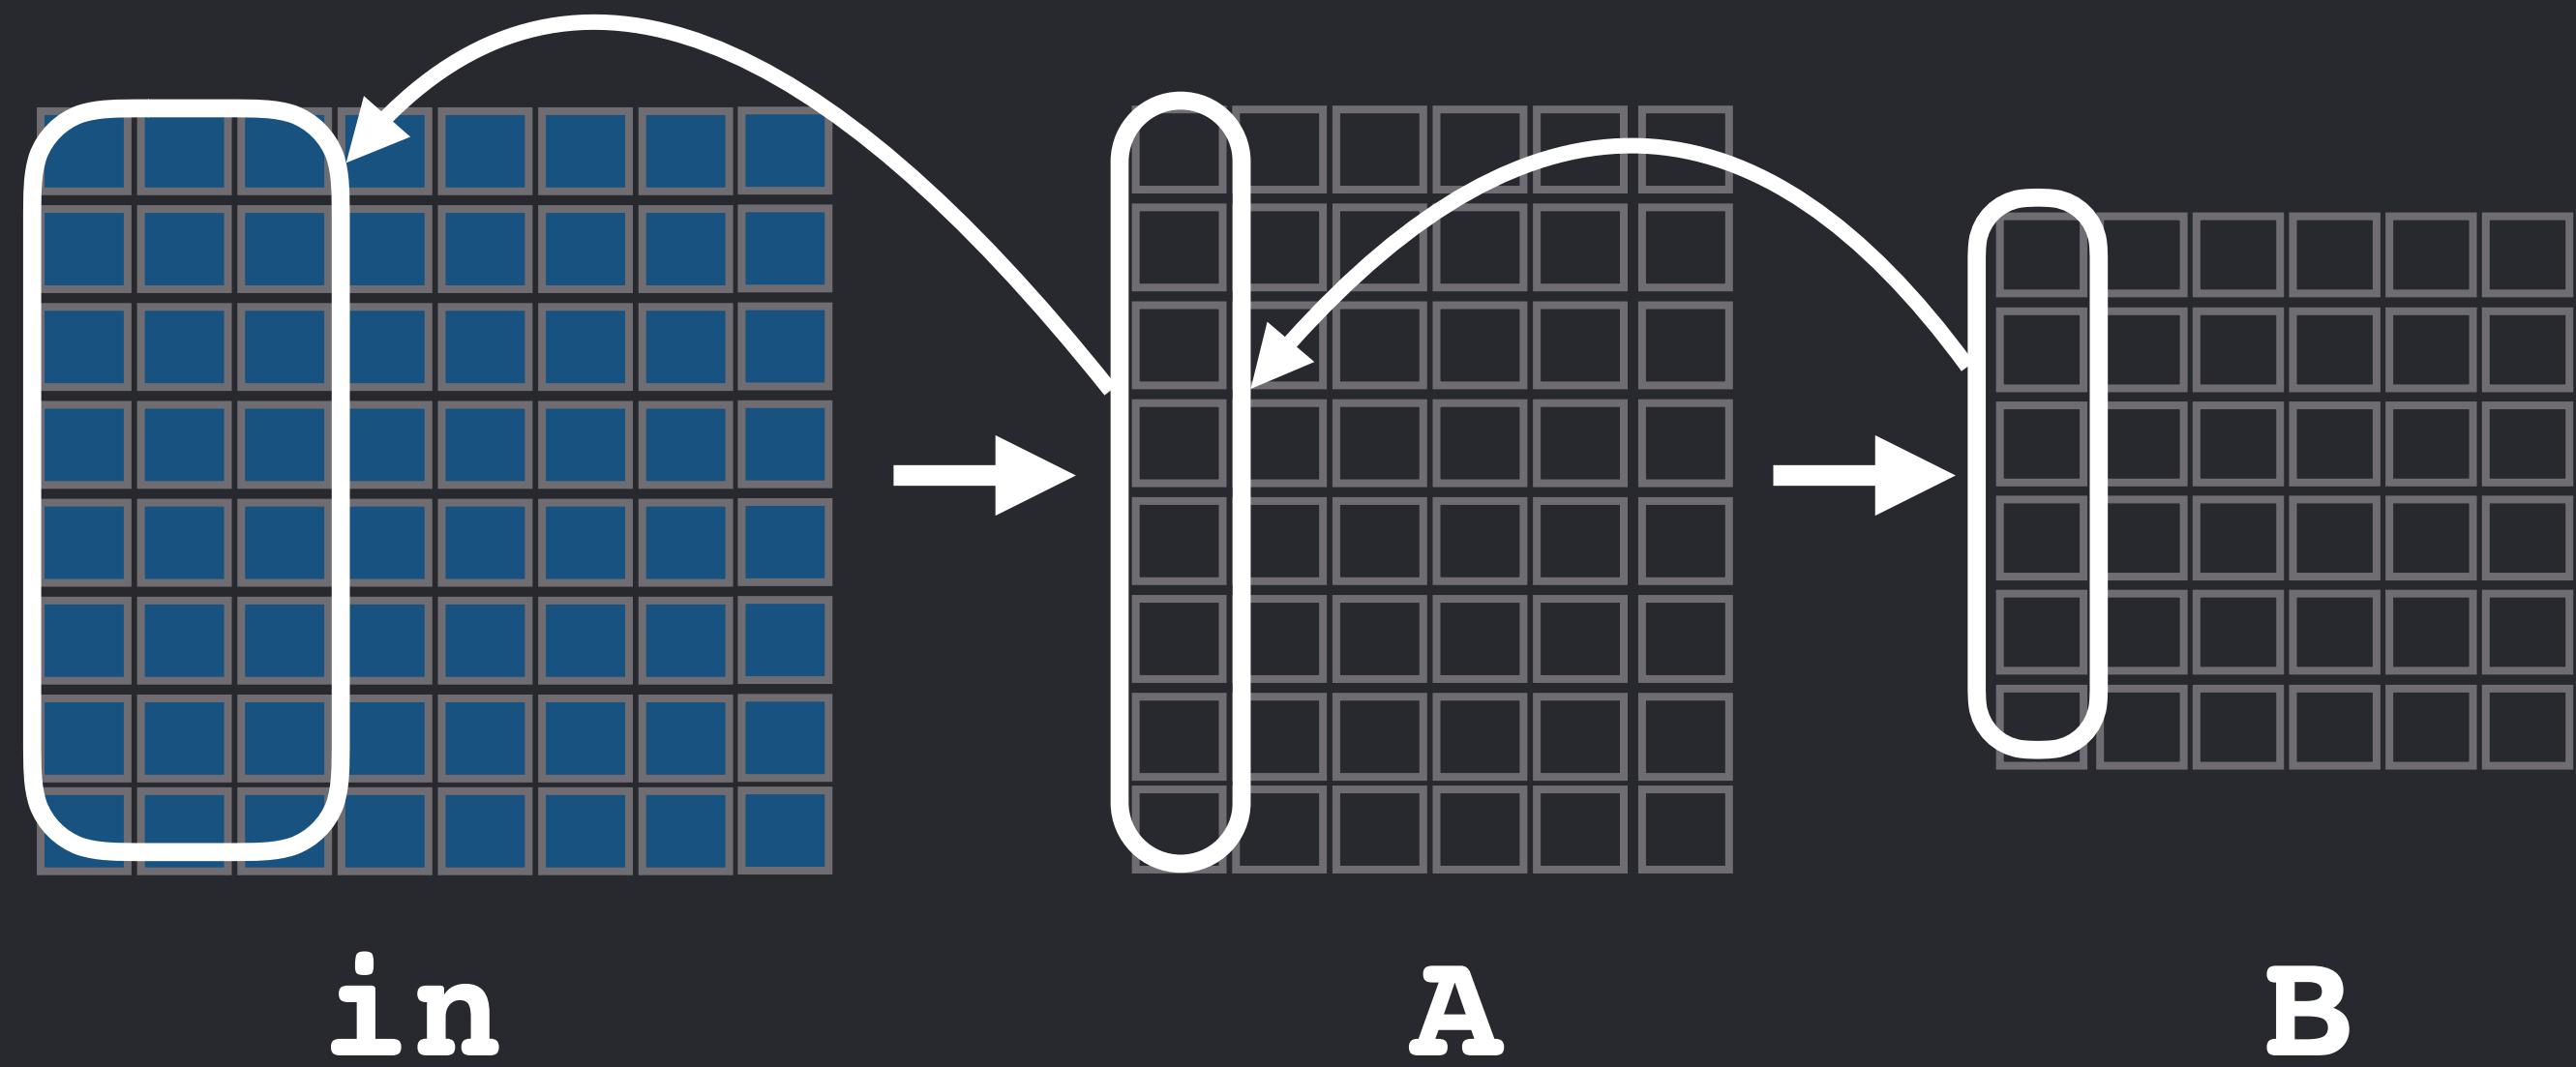
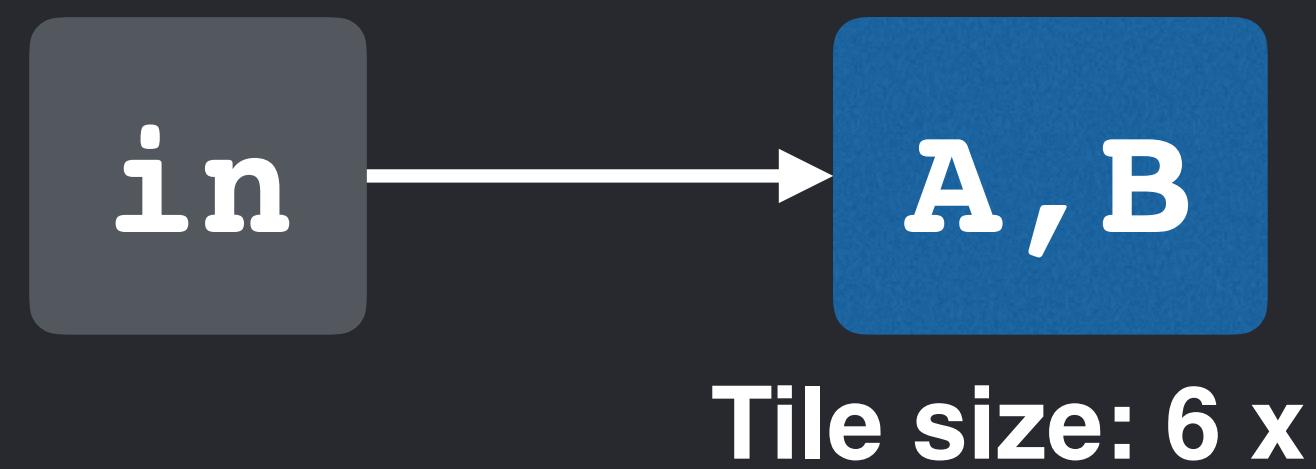
Search for best tile sizes



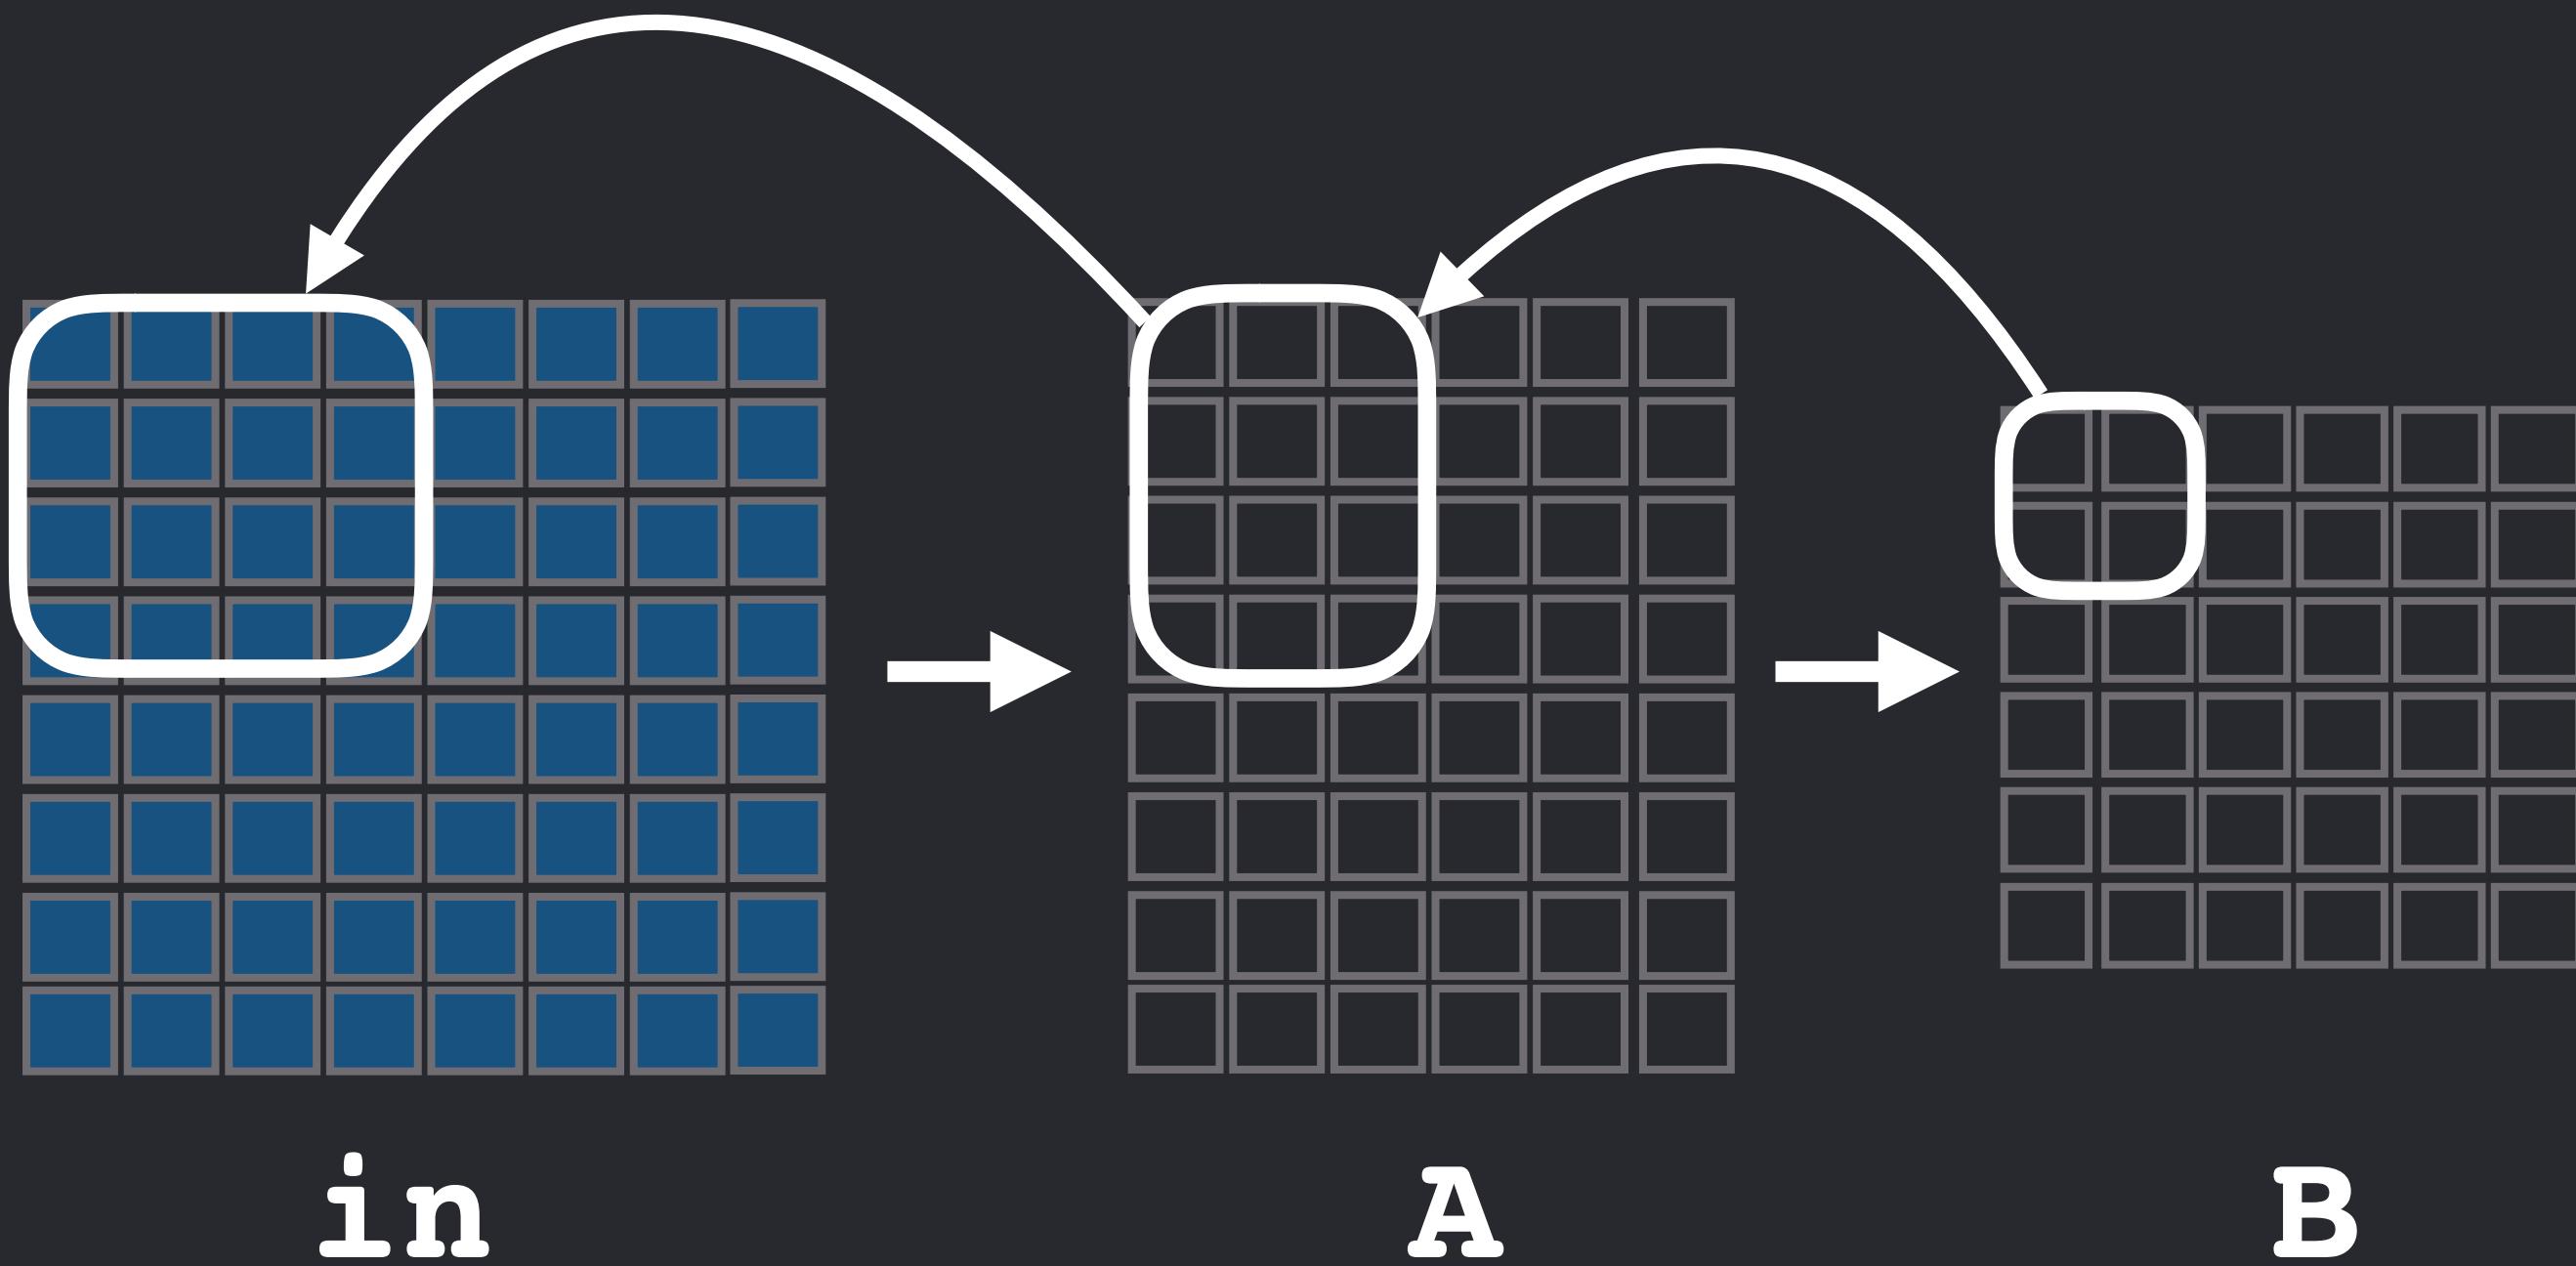
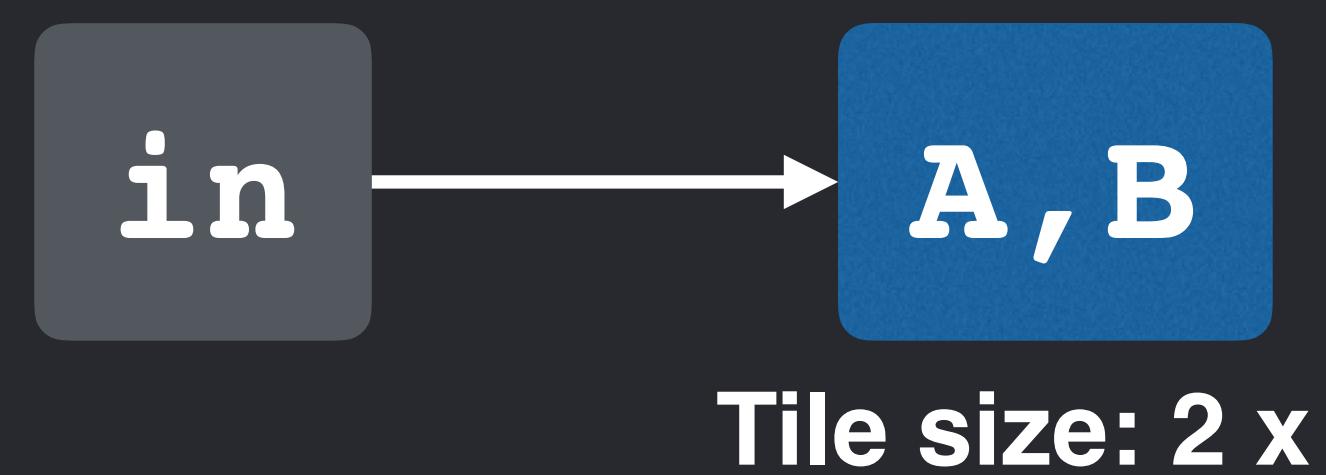
Tile size: 1×6



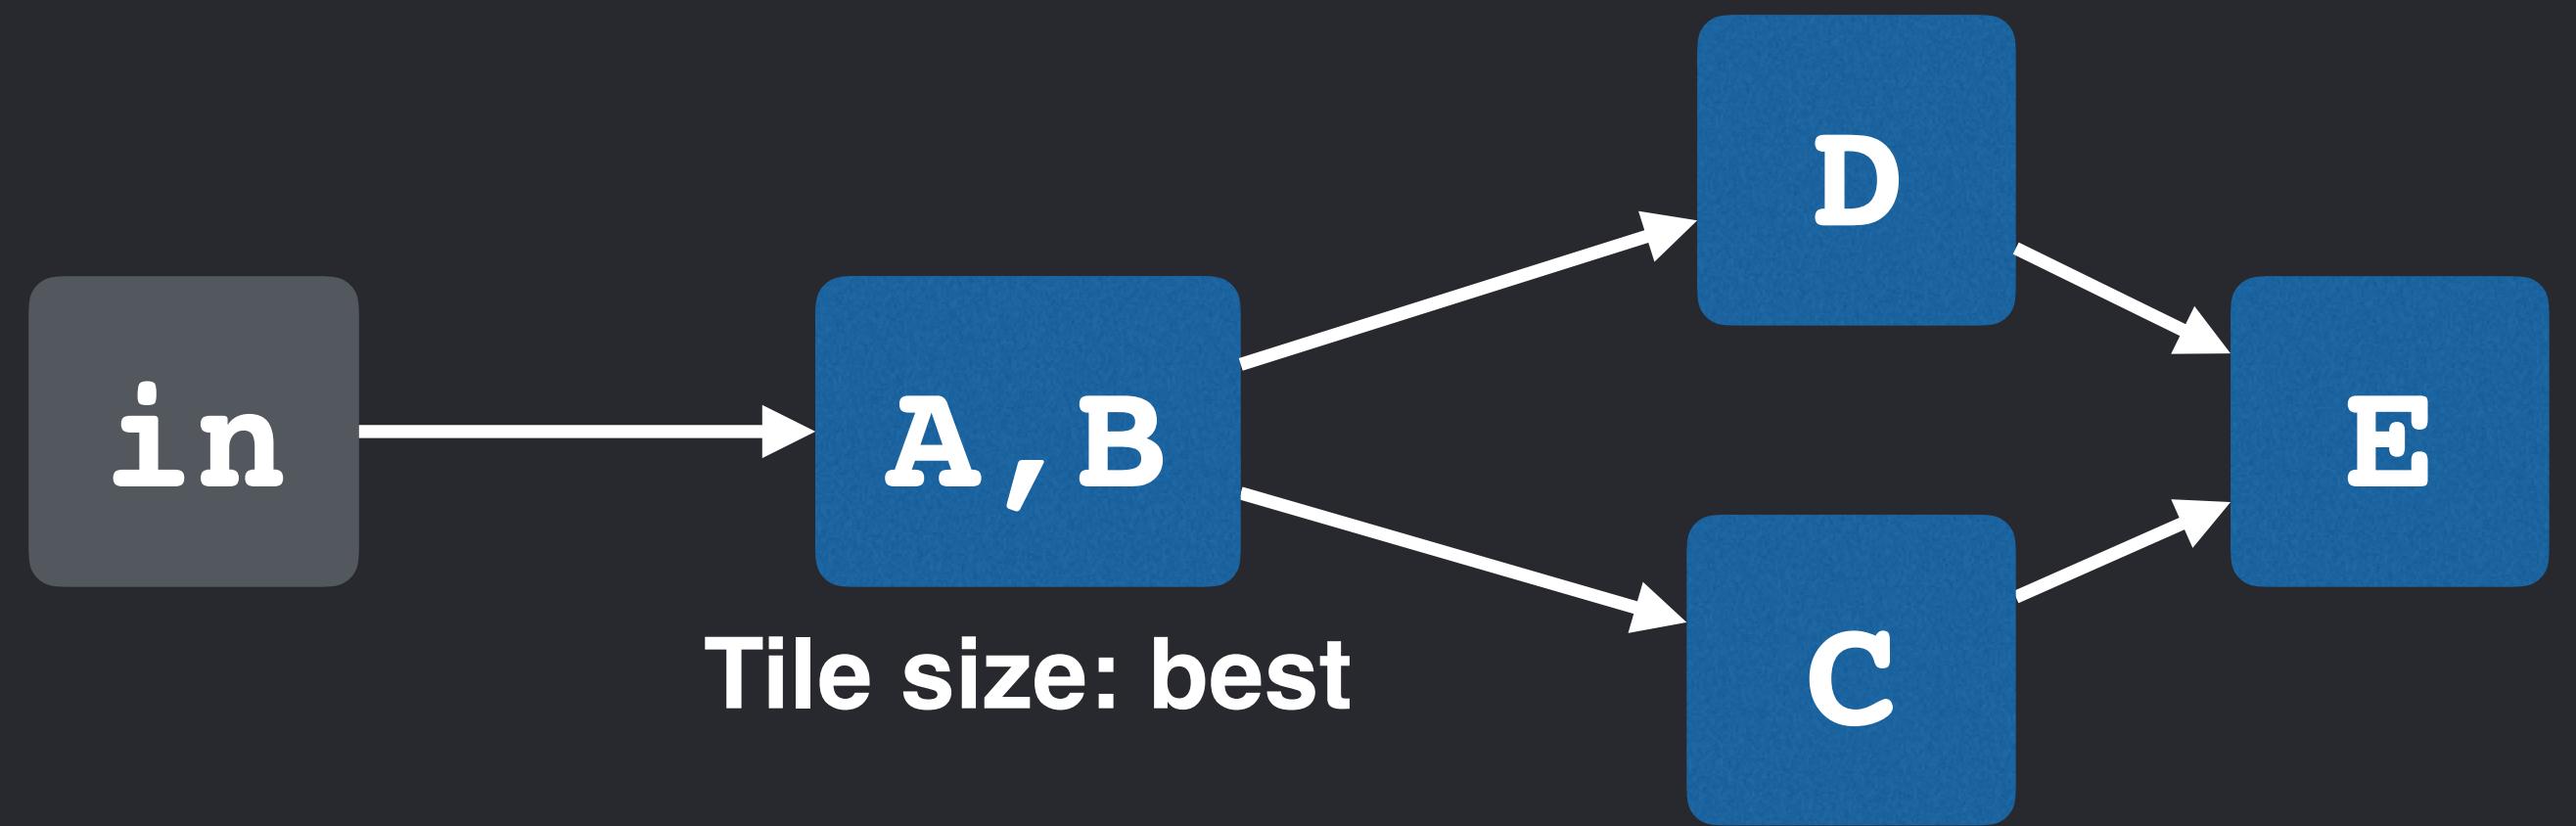
Search for best tile sizes



Search for best tile sizes

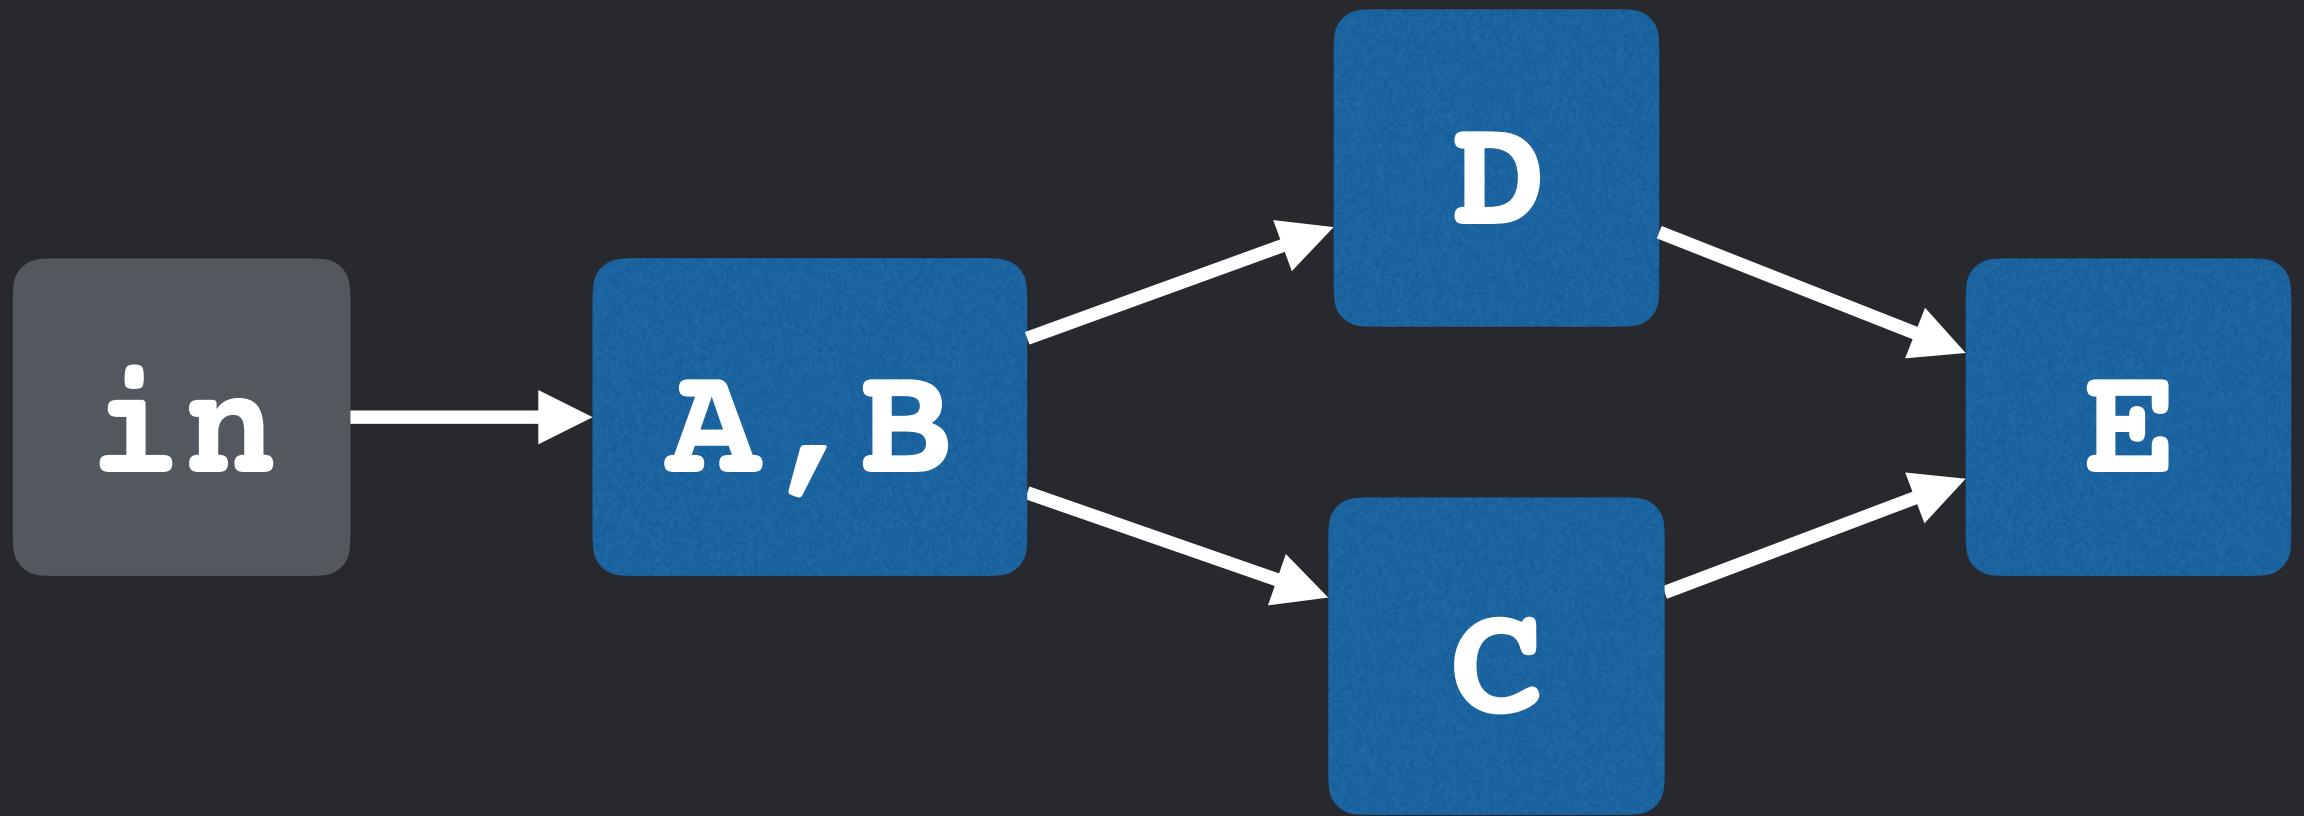
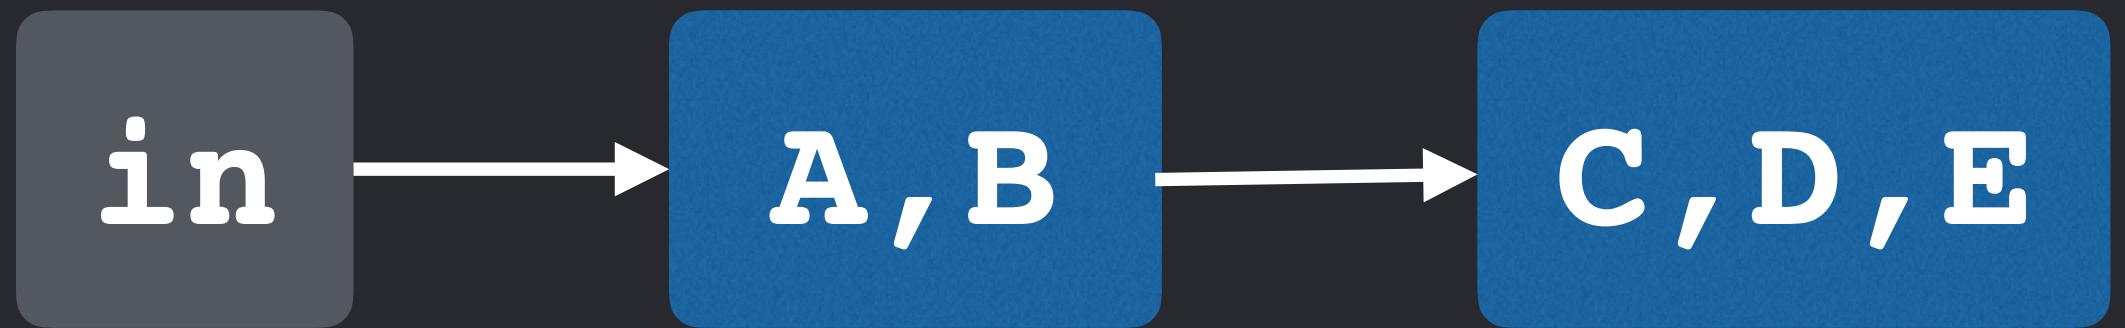
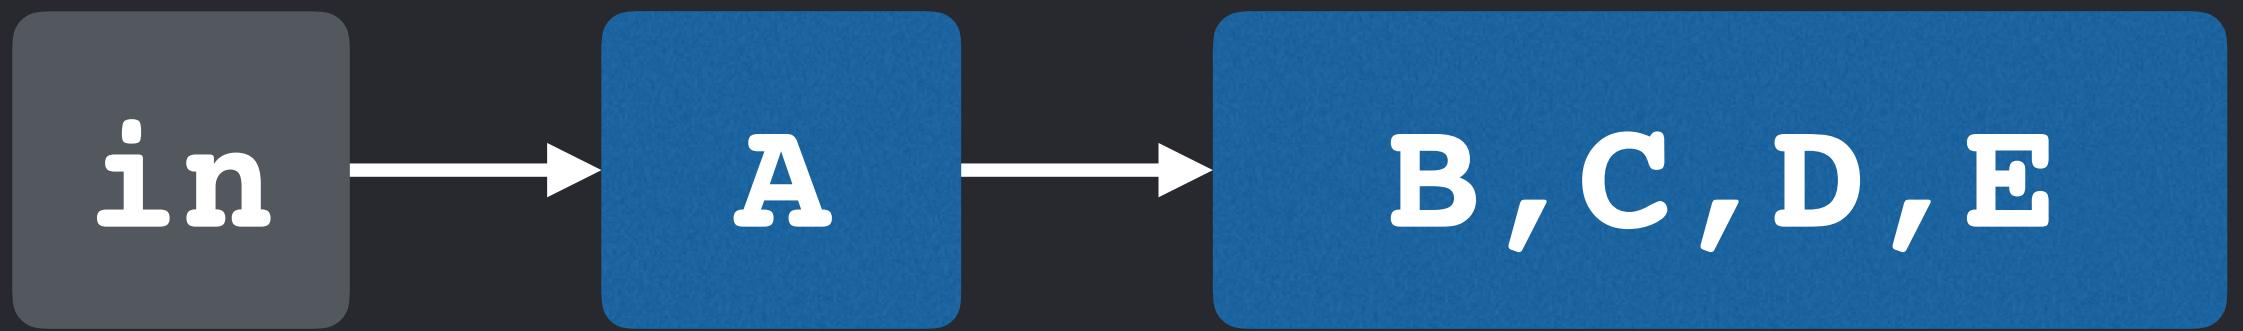
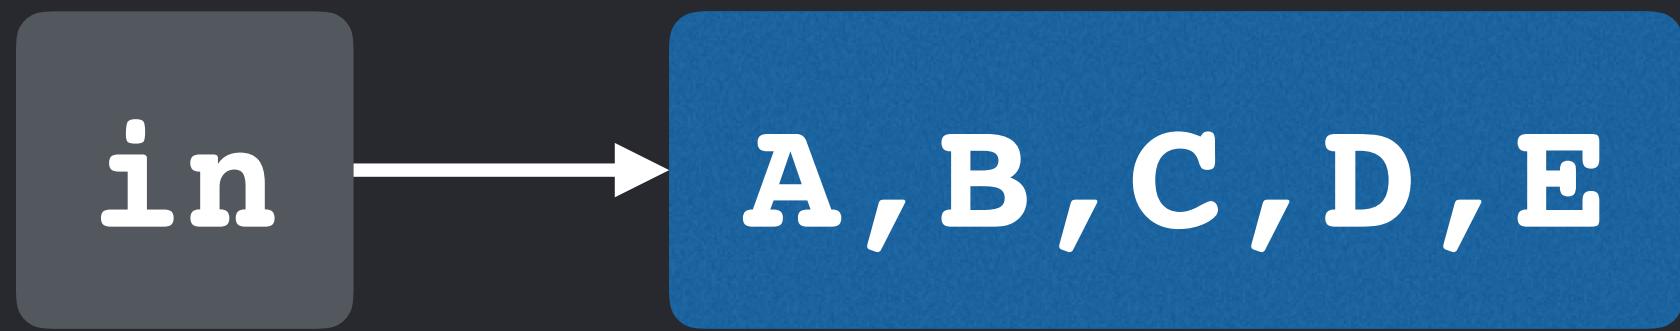


When to group stages?



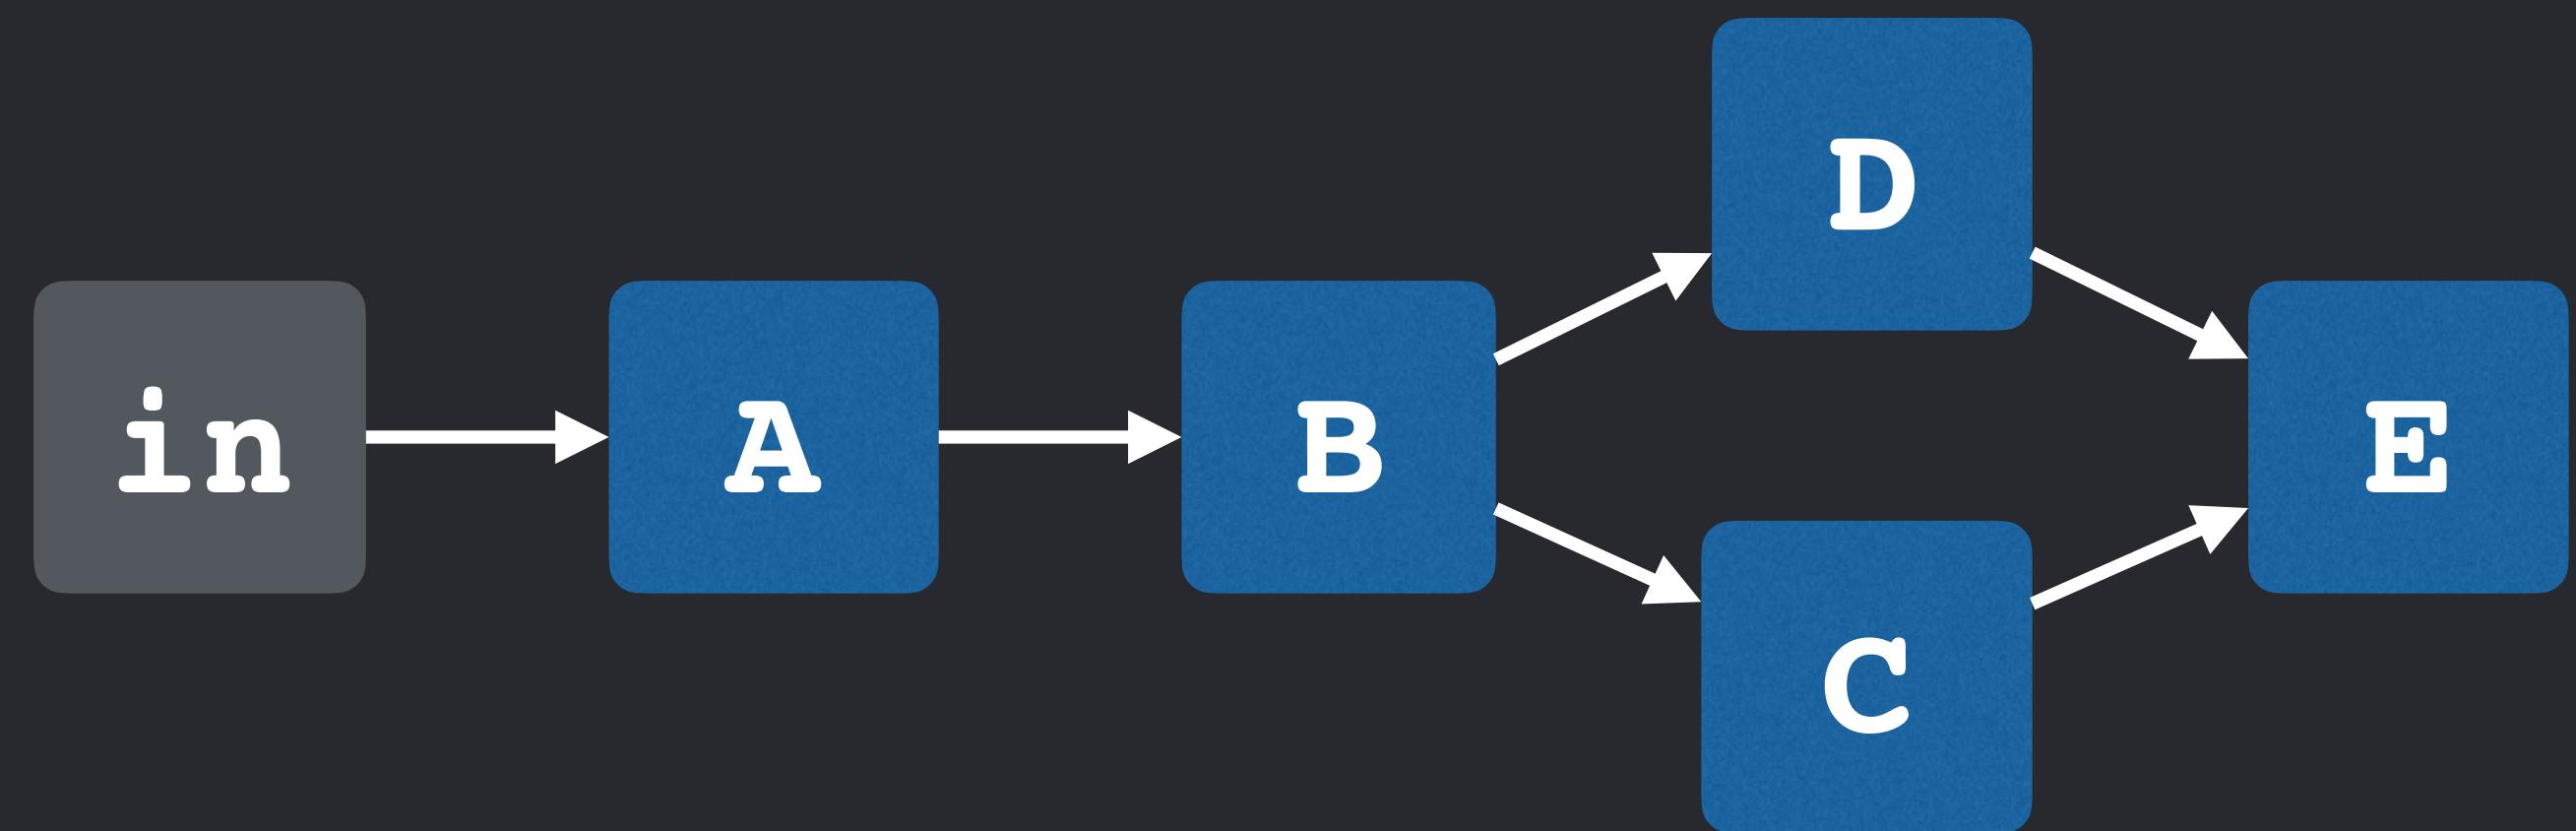
Benefit(**A, B**) = Cost(**A**) + Cost(**B**) - Cost(**A, B**)
Tile size: best

Exhaustive search is infeasible



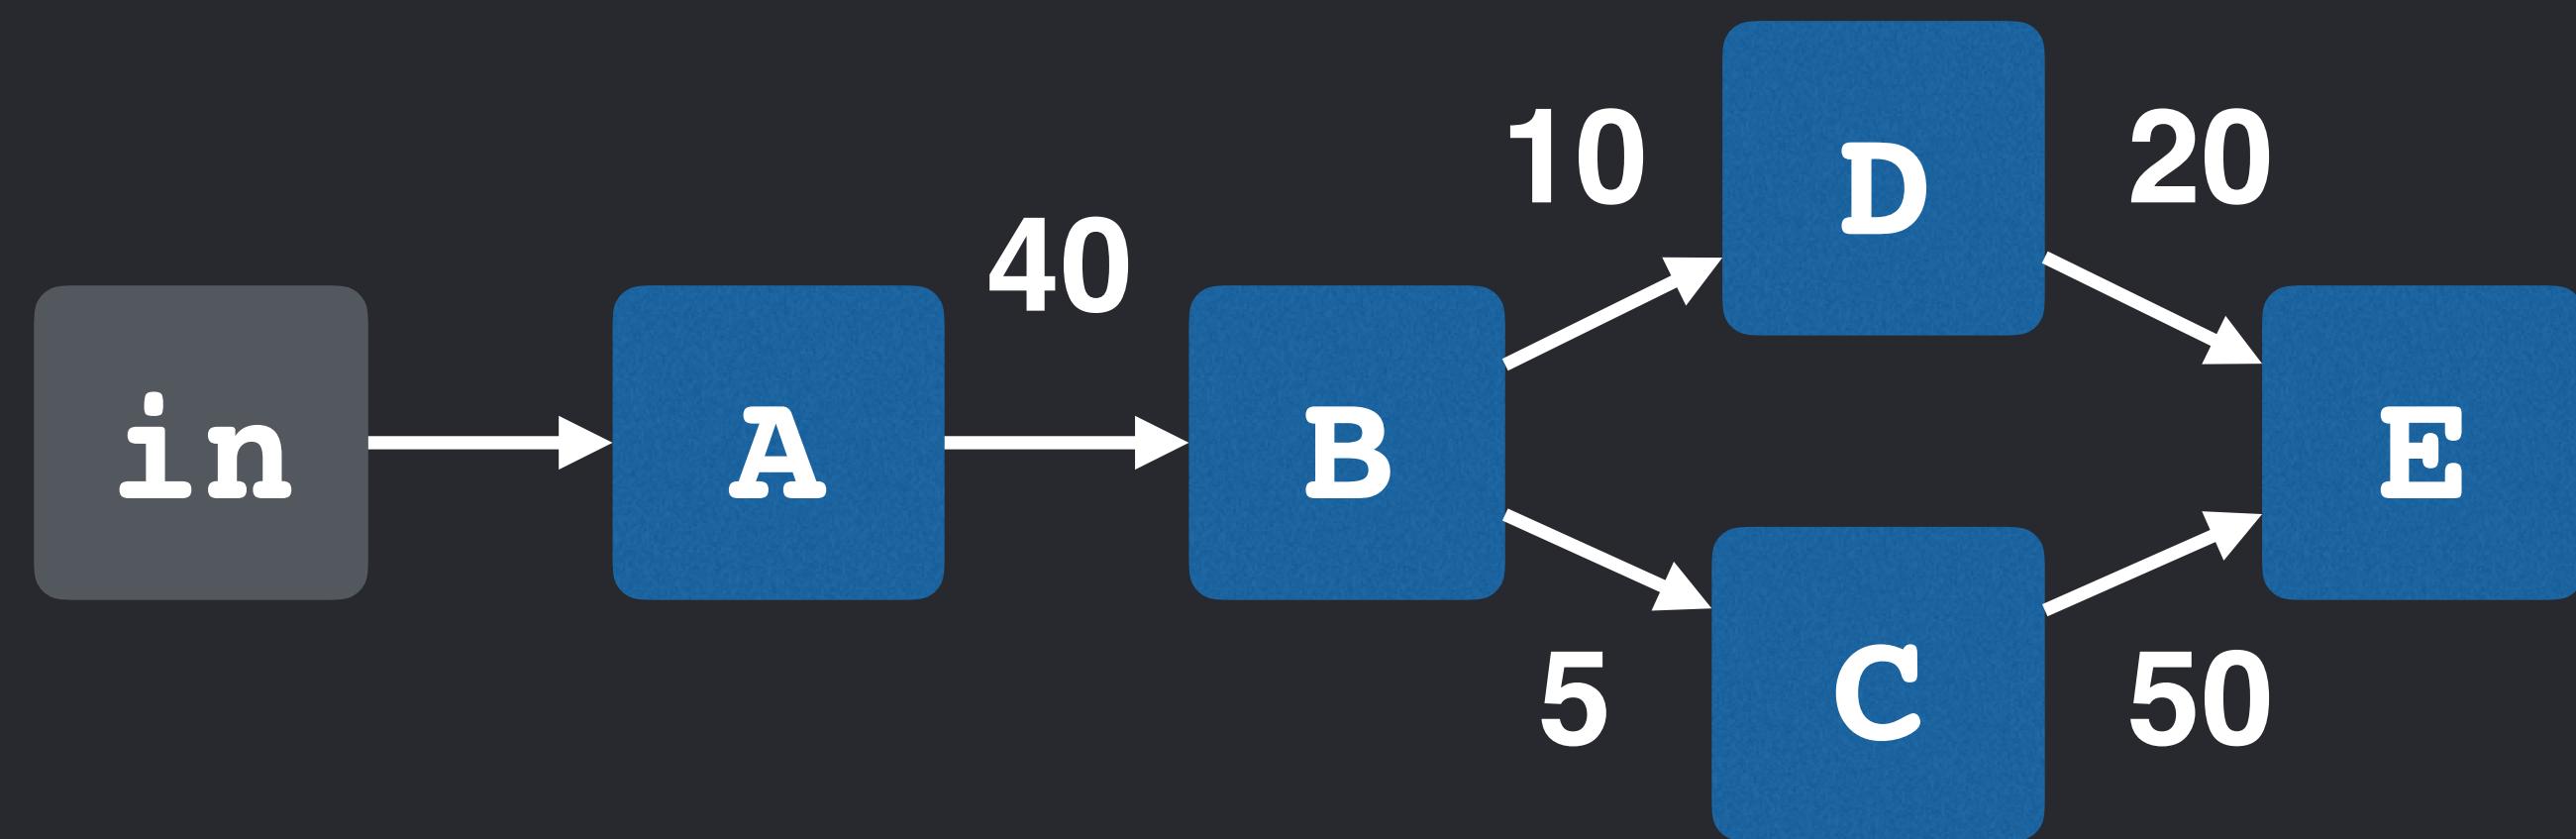
Exponential number of possible groupings

Greedy grouping algorithm



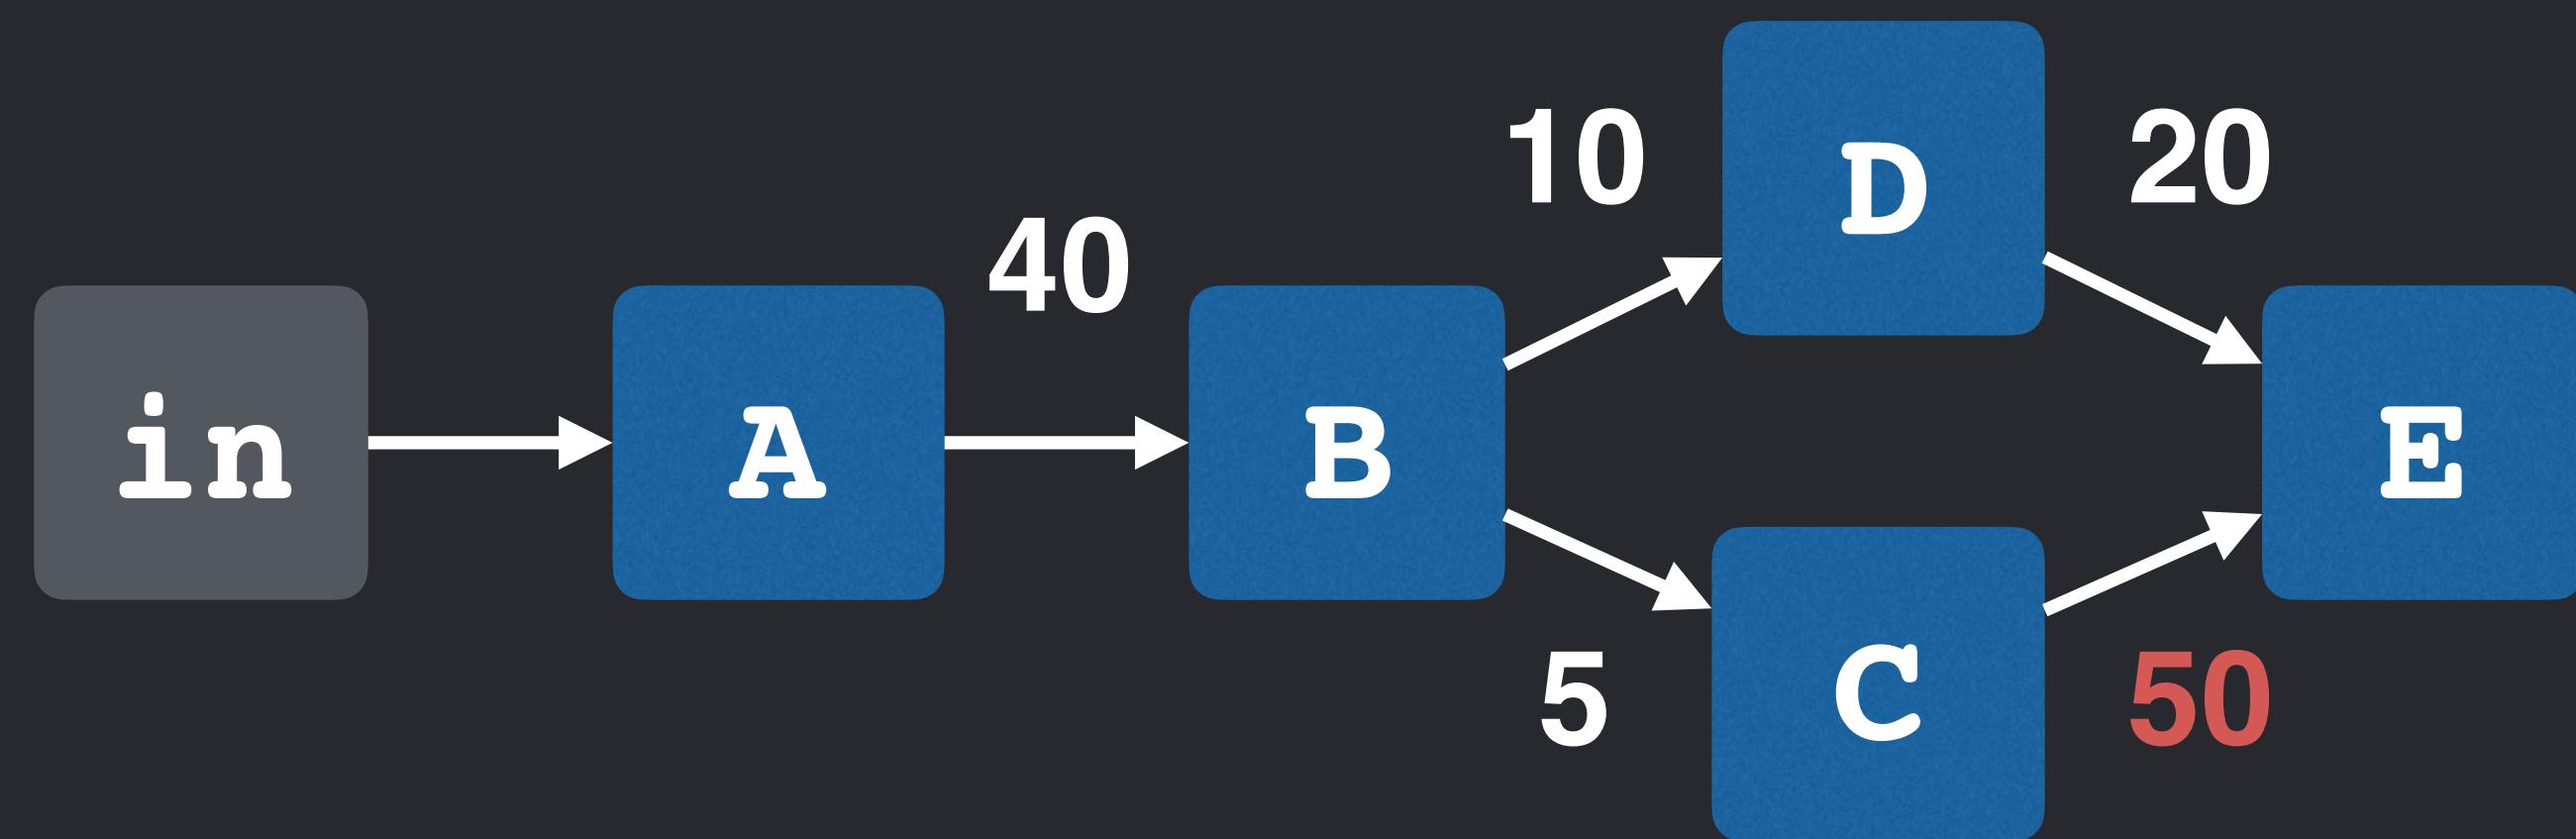
compute all pixels of A, in parallel
compute all pixels of B, in parallel
compute all pixels of C, in parallel
compute all pixels of D, in parallel
compute all pixels of E, in parallel

Greedy grouping algorithm



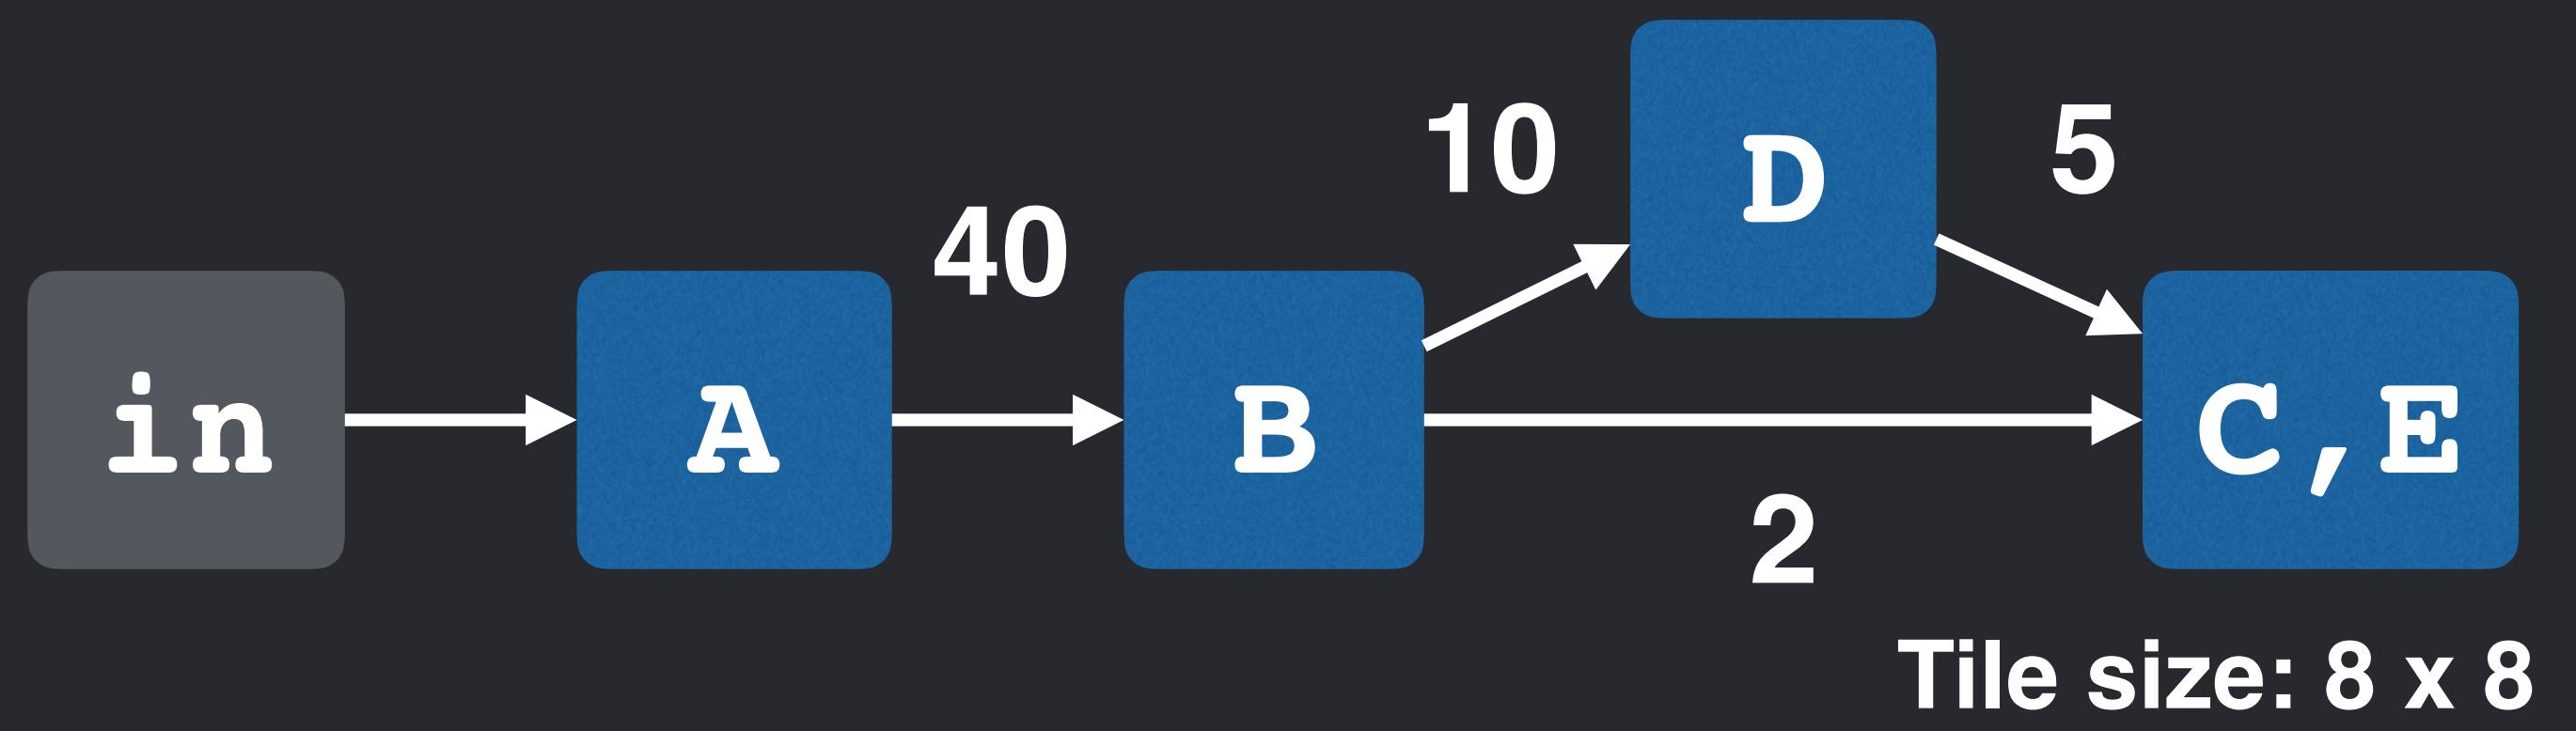
compute all pixels of A, in parallel
compute all pixels of B, in parallel
compute all pixels of C, in parallel
compute all pixels of D, in parallel
compute all pixels of E, in parallel

Greedy grouping algorithm



compute all pixels of A, in parallel
compute all pixels of B, in parallel
compute all pixels of C, in parallel
compute all pixels of D, in parallel
compute all pixels of E, in parallel

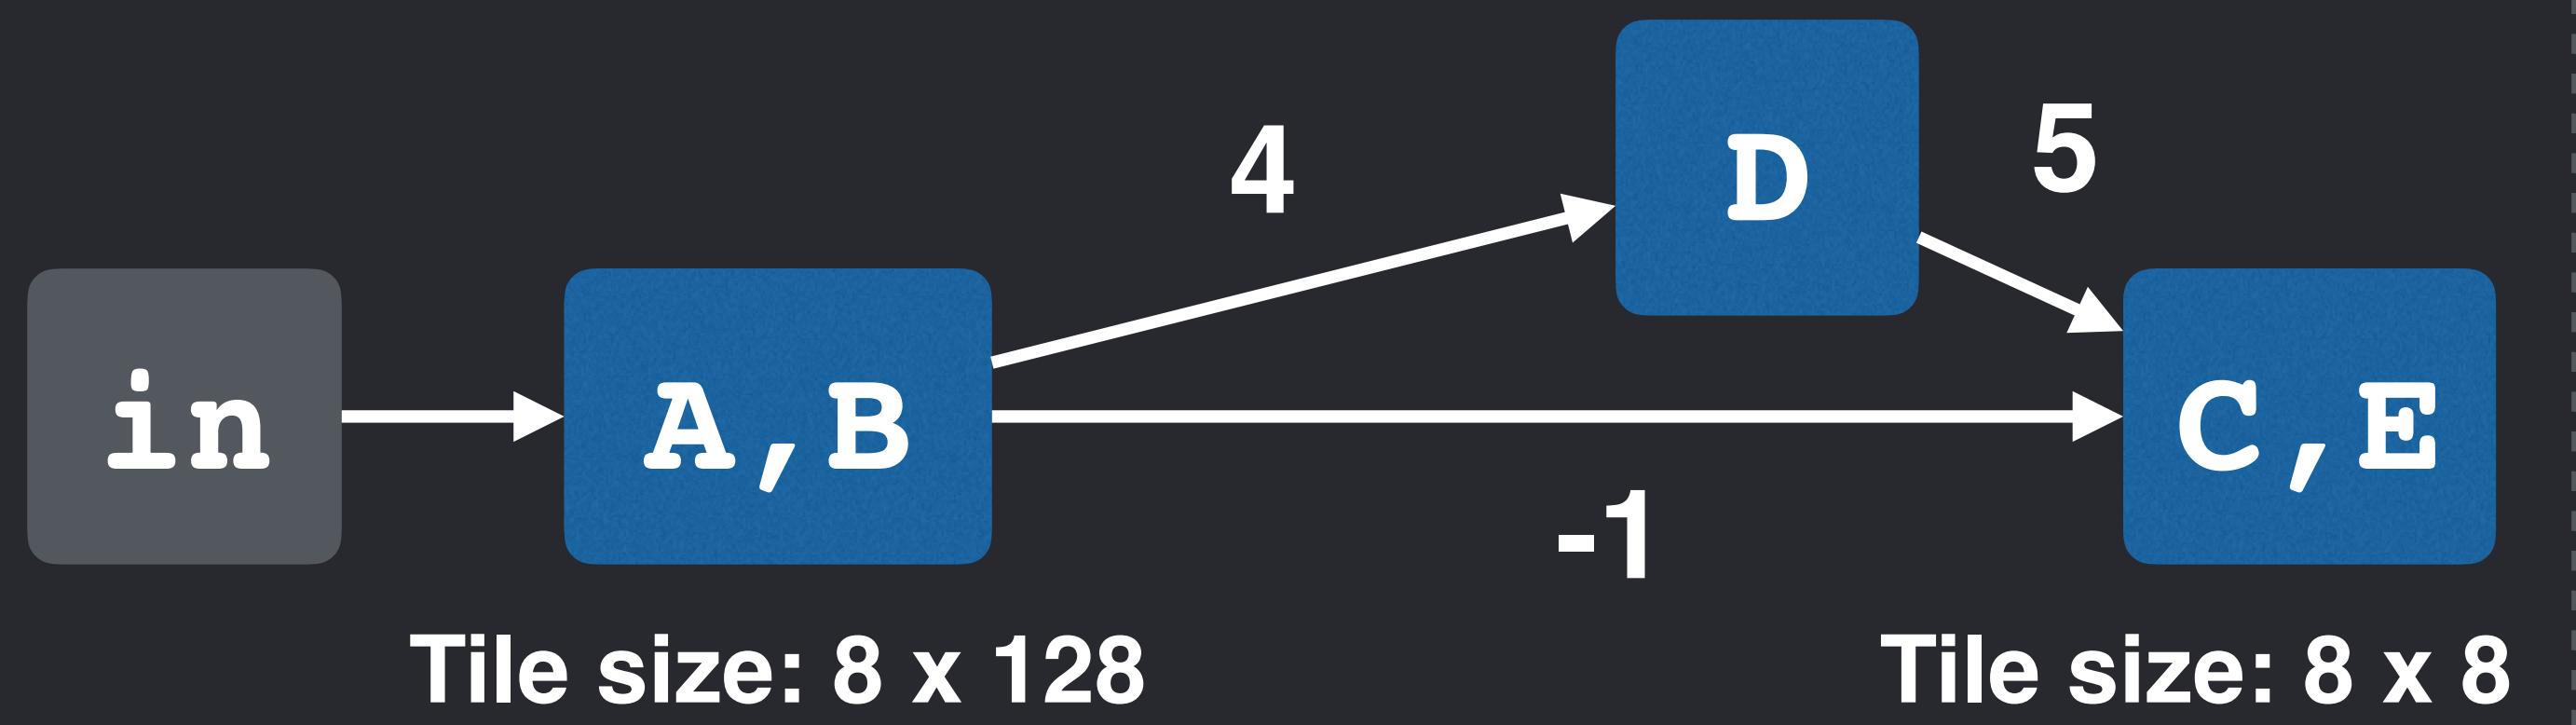
Greedy grouping algorithm



compute all pixels of A, in parallel
compute all pixels of B, in parallel
compute all pixels of D, in parallel

for each 8x8 tile in parallel
compute required pixels of C
compute pixels in tile of E

Greedy grouping algorithm



```
for each 8x128 tile in parallel  
  compute required pixels of A  
  compute pixels in tile of B  
  
compute all pixels of D, in parallel  
  
for each 8x8 tile in parallel  
  compute required pixels of C  
  compute pixels in tile of E
```

Greedy grouping algorithm



for each 8x128 tile in parallel
compute required pixels of A
compute pixels in tile of B

for each 8x8 tile in parallel
compute required pixels of C
compute required pixels of
compute pixels in tile of E

Auto scheduler implementation details

- Multi-core parallelism, vectorization, loop reordering, and unrolling

```
for each 8x128 tile in parallel
    vectorize compute required pixels of A unroll x by 4
    vectorize compute required pixels of B
```

```
    vectorize compute pixels in tile of D
```

```
for each 8x8 tile in parallel
    vectorize compute required pixels of C unroll y by 2
```

```
    vectorize compute pixels in tile of E
```

Evaluation

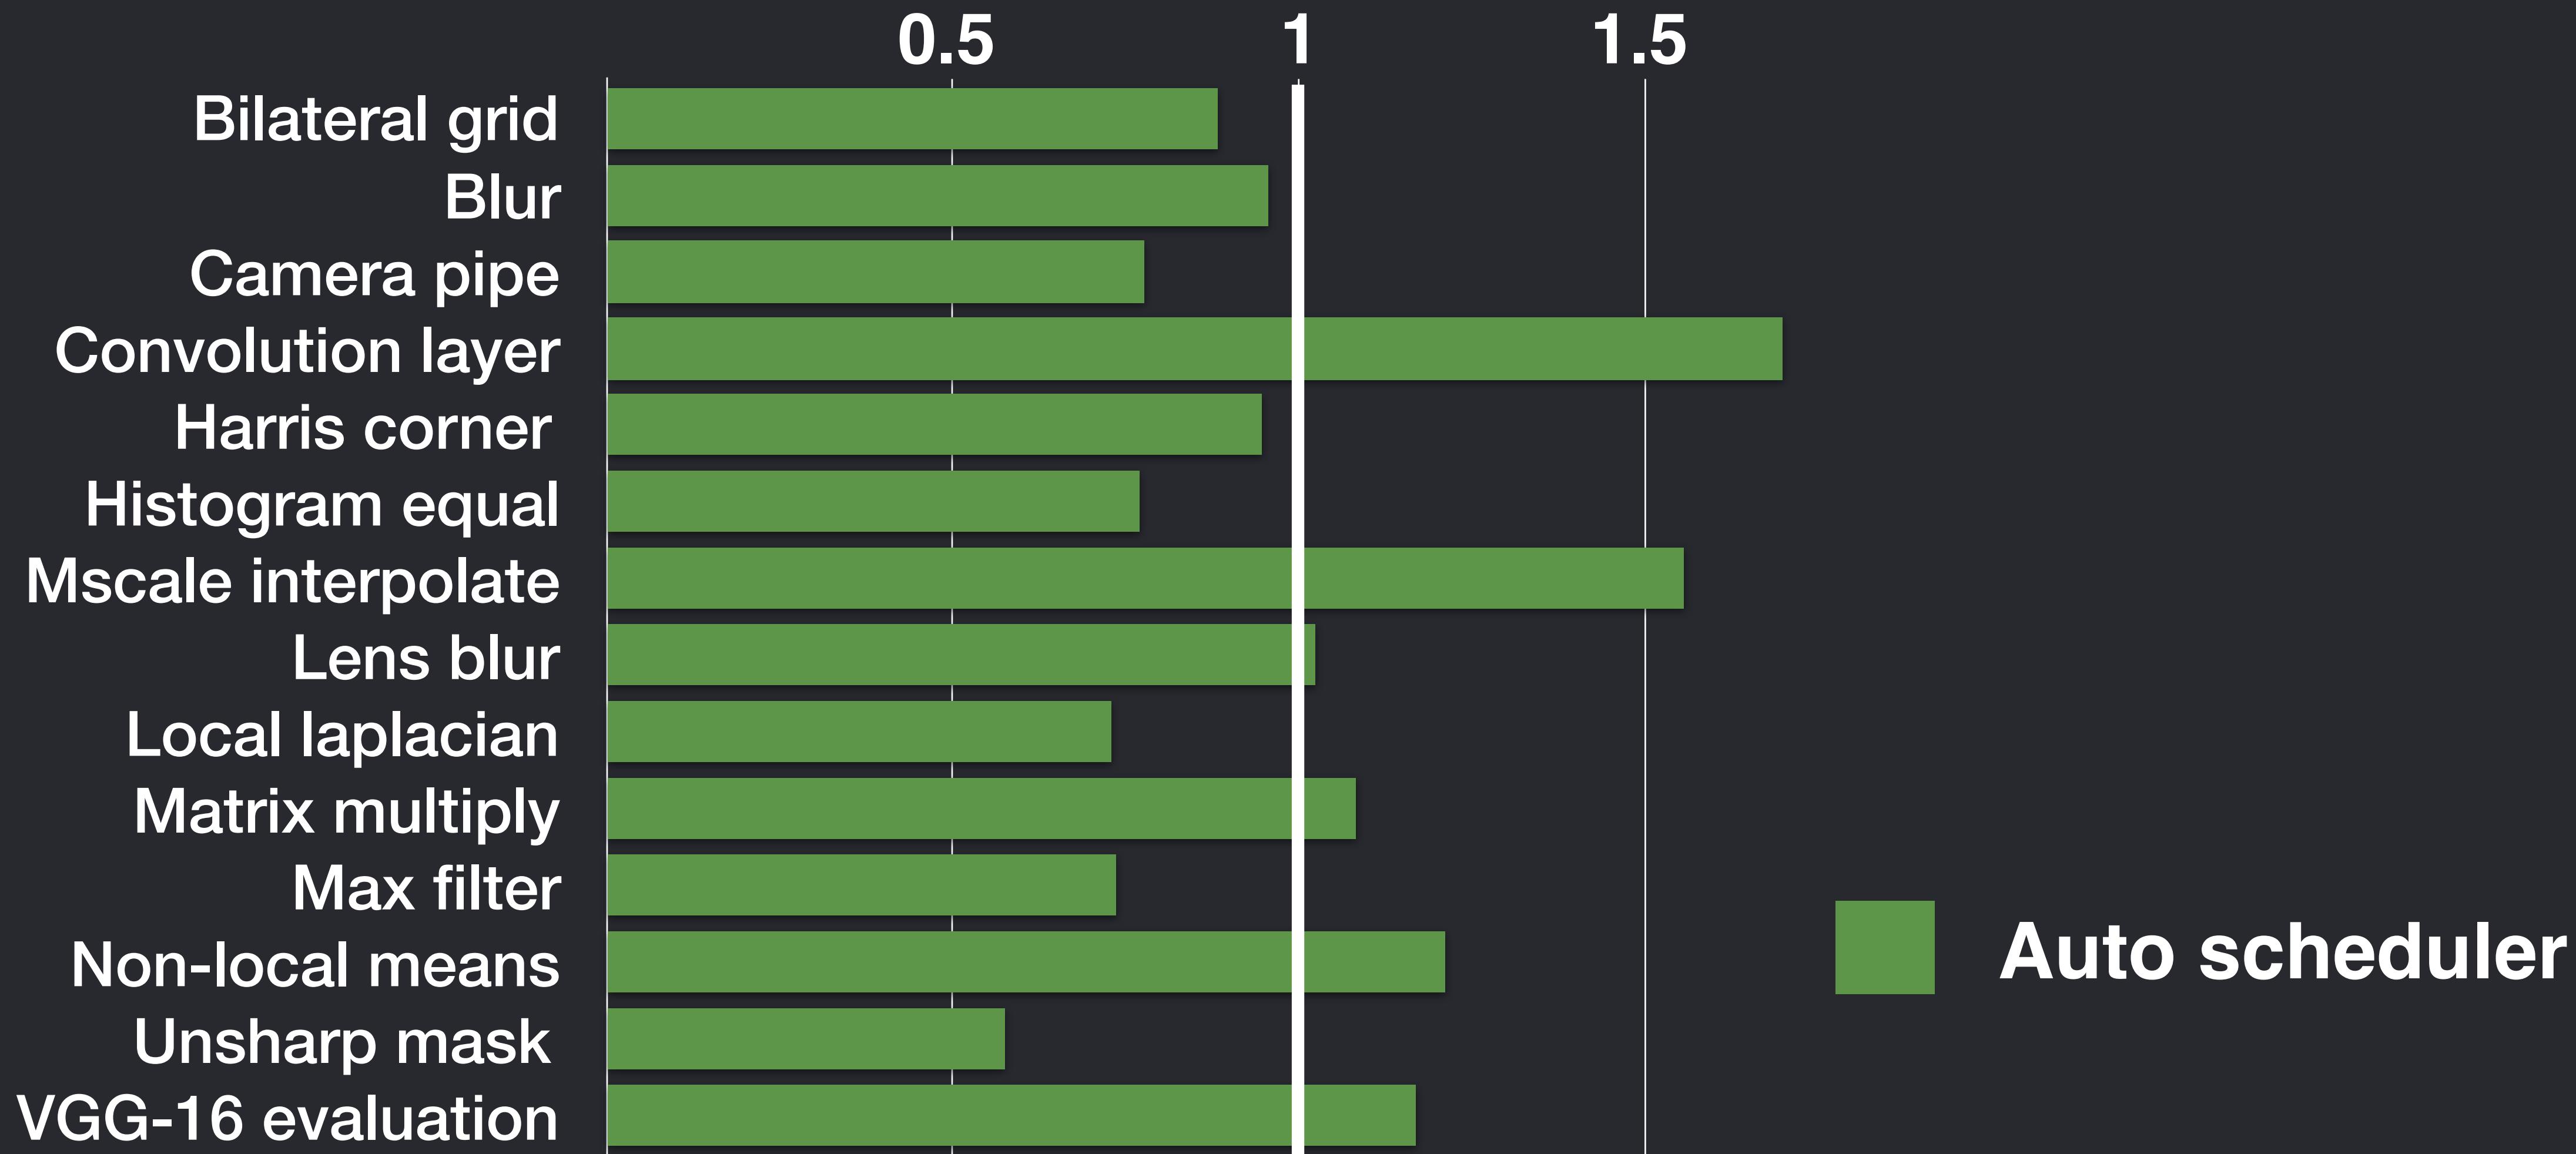
Benchmarks of varying complexity and structure

Benchmark	Stages
Blur	3
Unsharp mask	9
Harris corner detection	13
Camera RAW processing	30
Non-local means denoising	13
Max-brightness filter	9
Multi-scale interpolation	52
Local-laplacian filter	103
Synthetic depth-of-field	74
Bilateral filter	8
Histogram equalization	7
VGG-16 deep network eval	64

Auto scheduler generates schedules in seconds

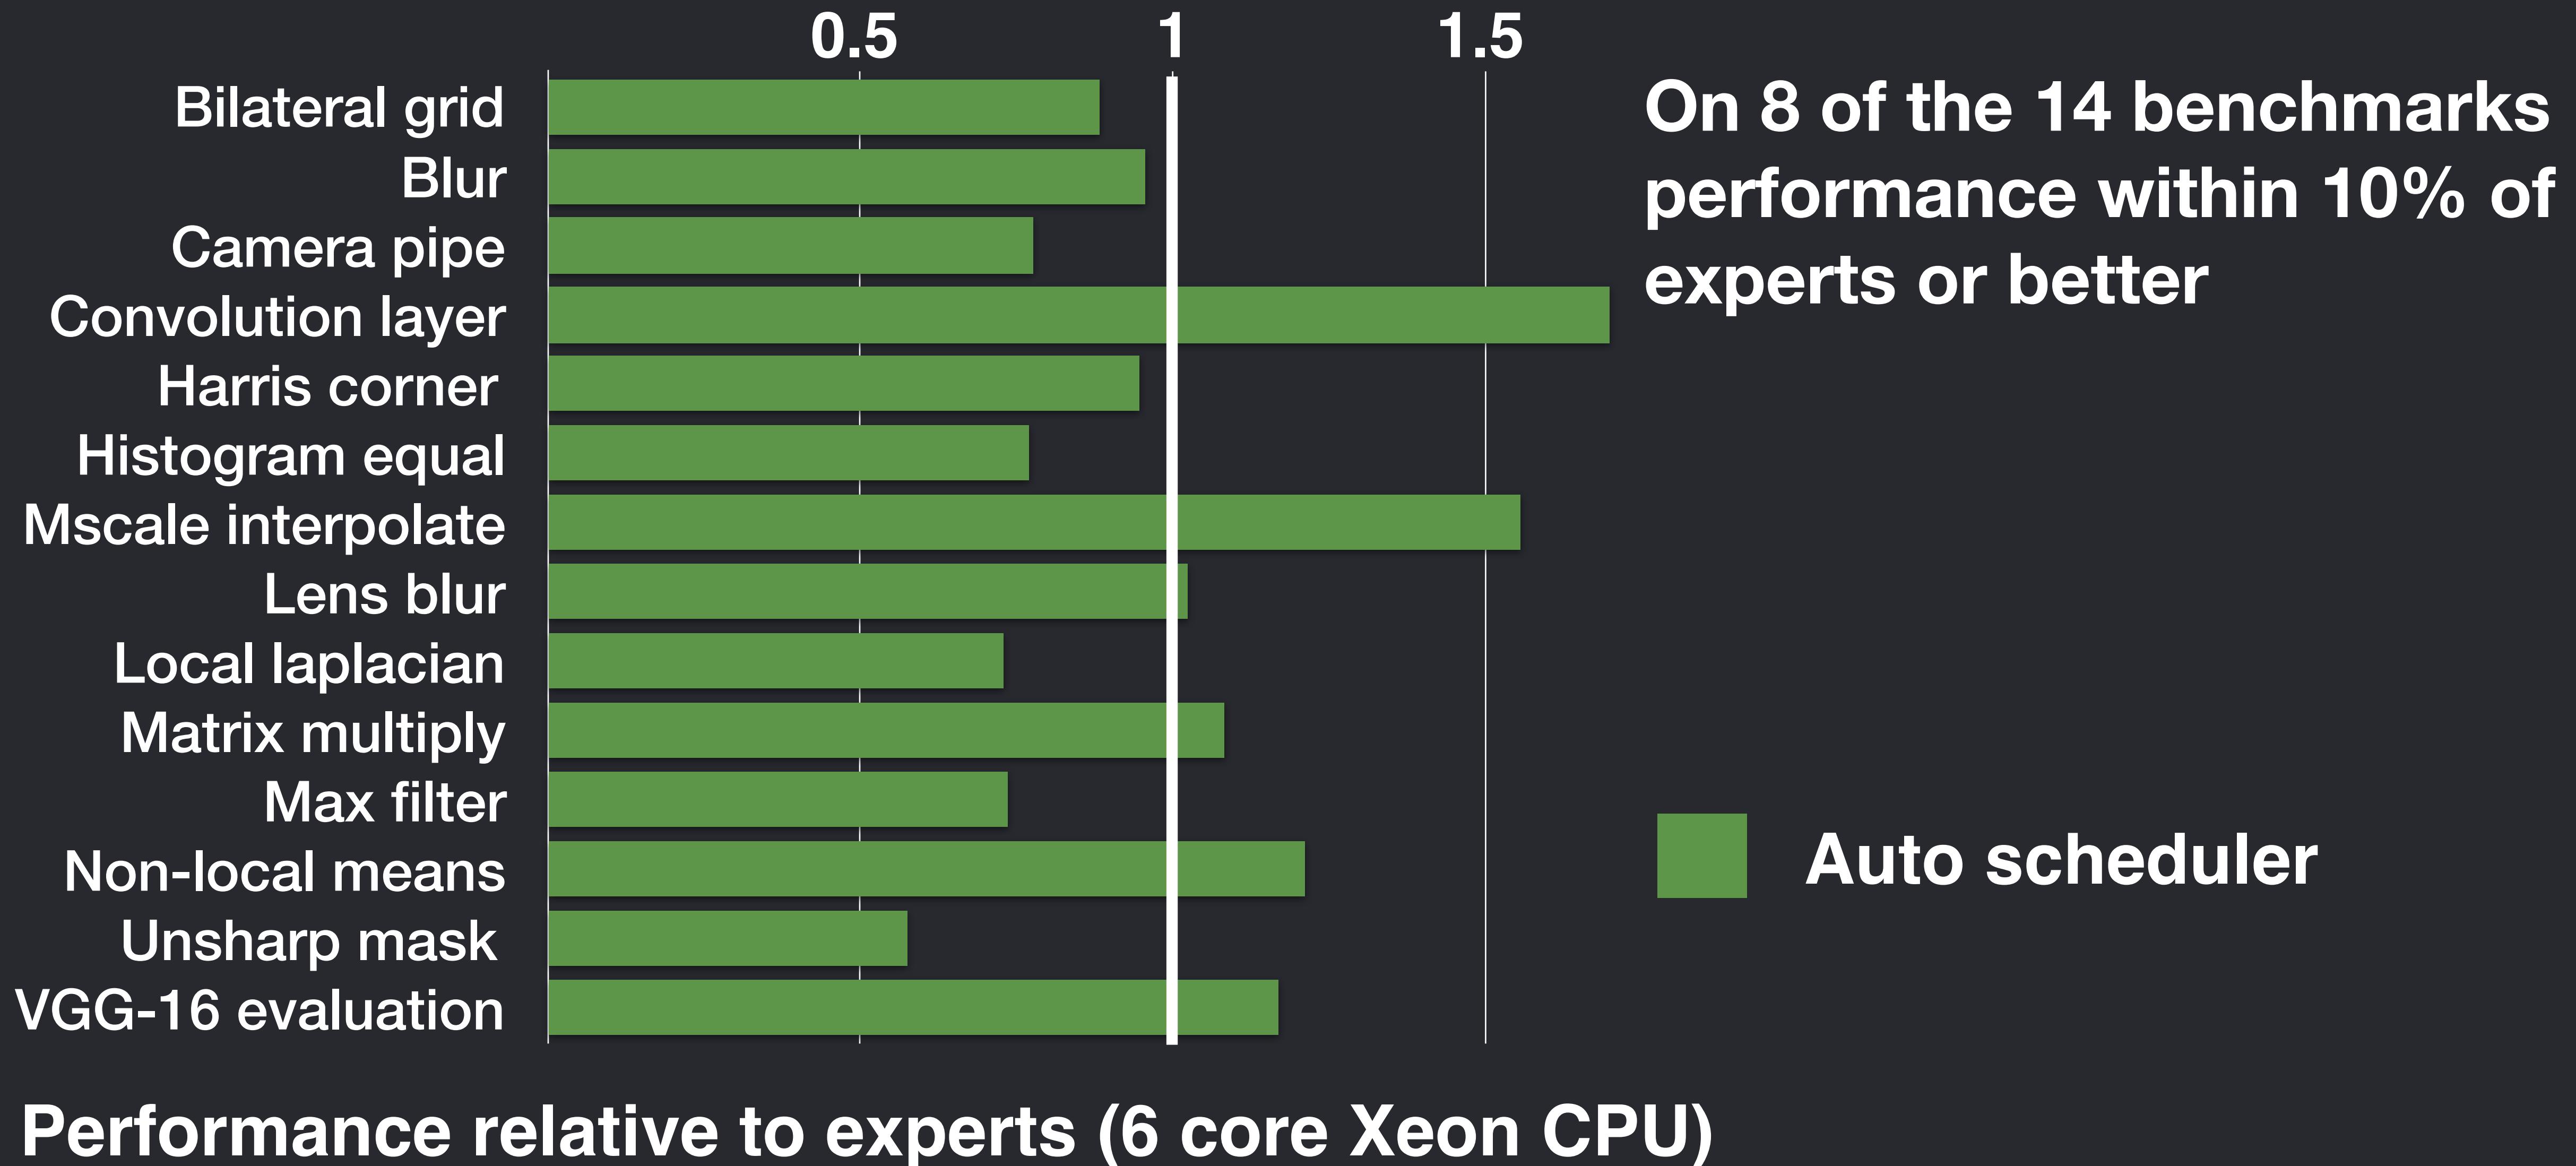
Benchmark	Stages	Compile time (s)
Blur	3	<1
Unsharp mask	9	<1
Harris corner detection	13	<1
Camera RAW processing	30	<1
Non-local means denoising	13	<1
Max-brightness filter	9	<1
Multi-scale interpolation	52	2.6
Local-laplacian filter	103	3.9
Synthetic depth-of-field	74	55
Bilateral filter	8	<1
Histogram equalization	7	<1
VGG-16 deep network eval	64	6.9

Auto scheduler performs comparably to experts

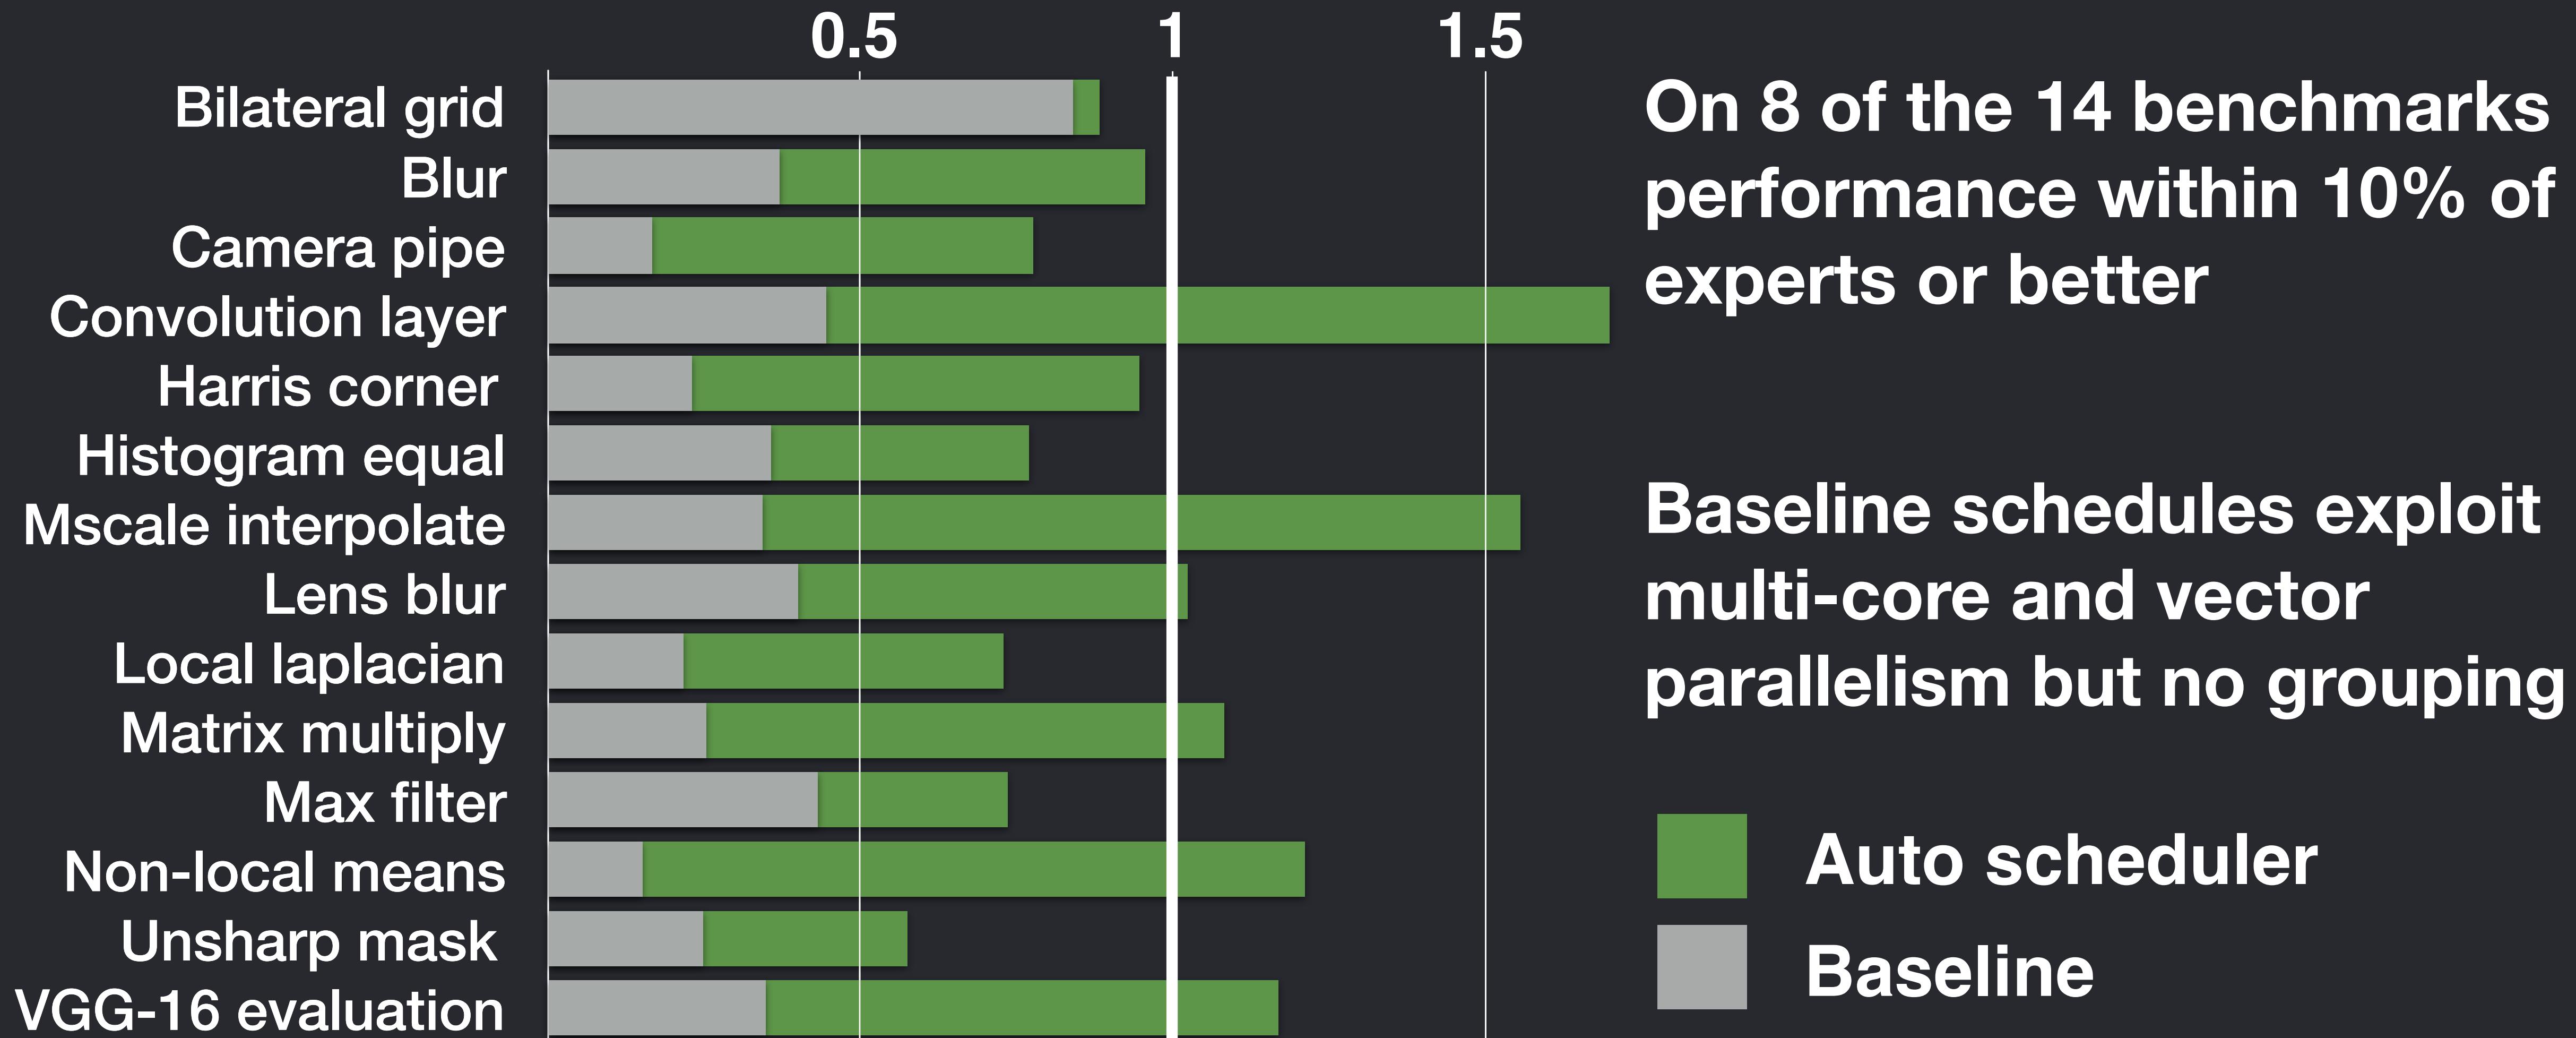


Performance relative to experts (6 core Xeon CPU)

Auto scheduler performs comparably to experts



Auto scheduler performs comparably to experts



Performance relative to experts (6 core Xeon CPU)

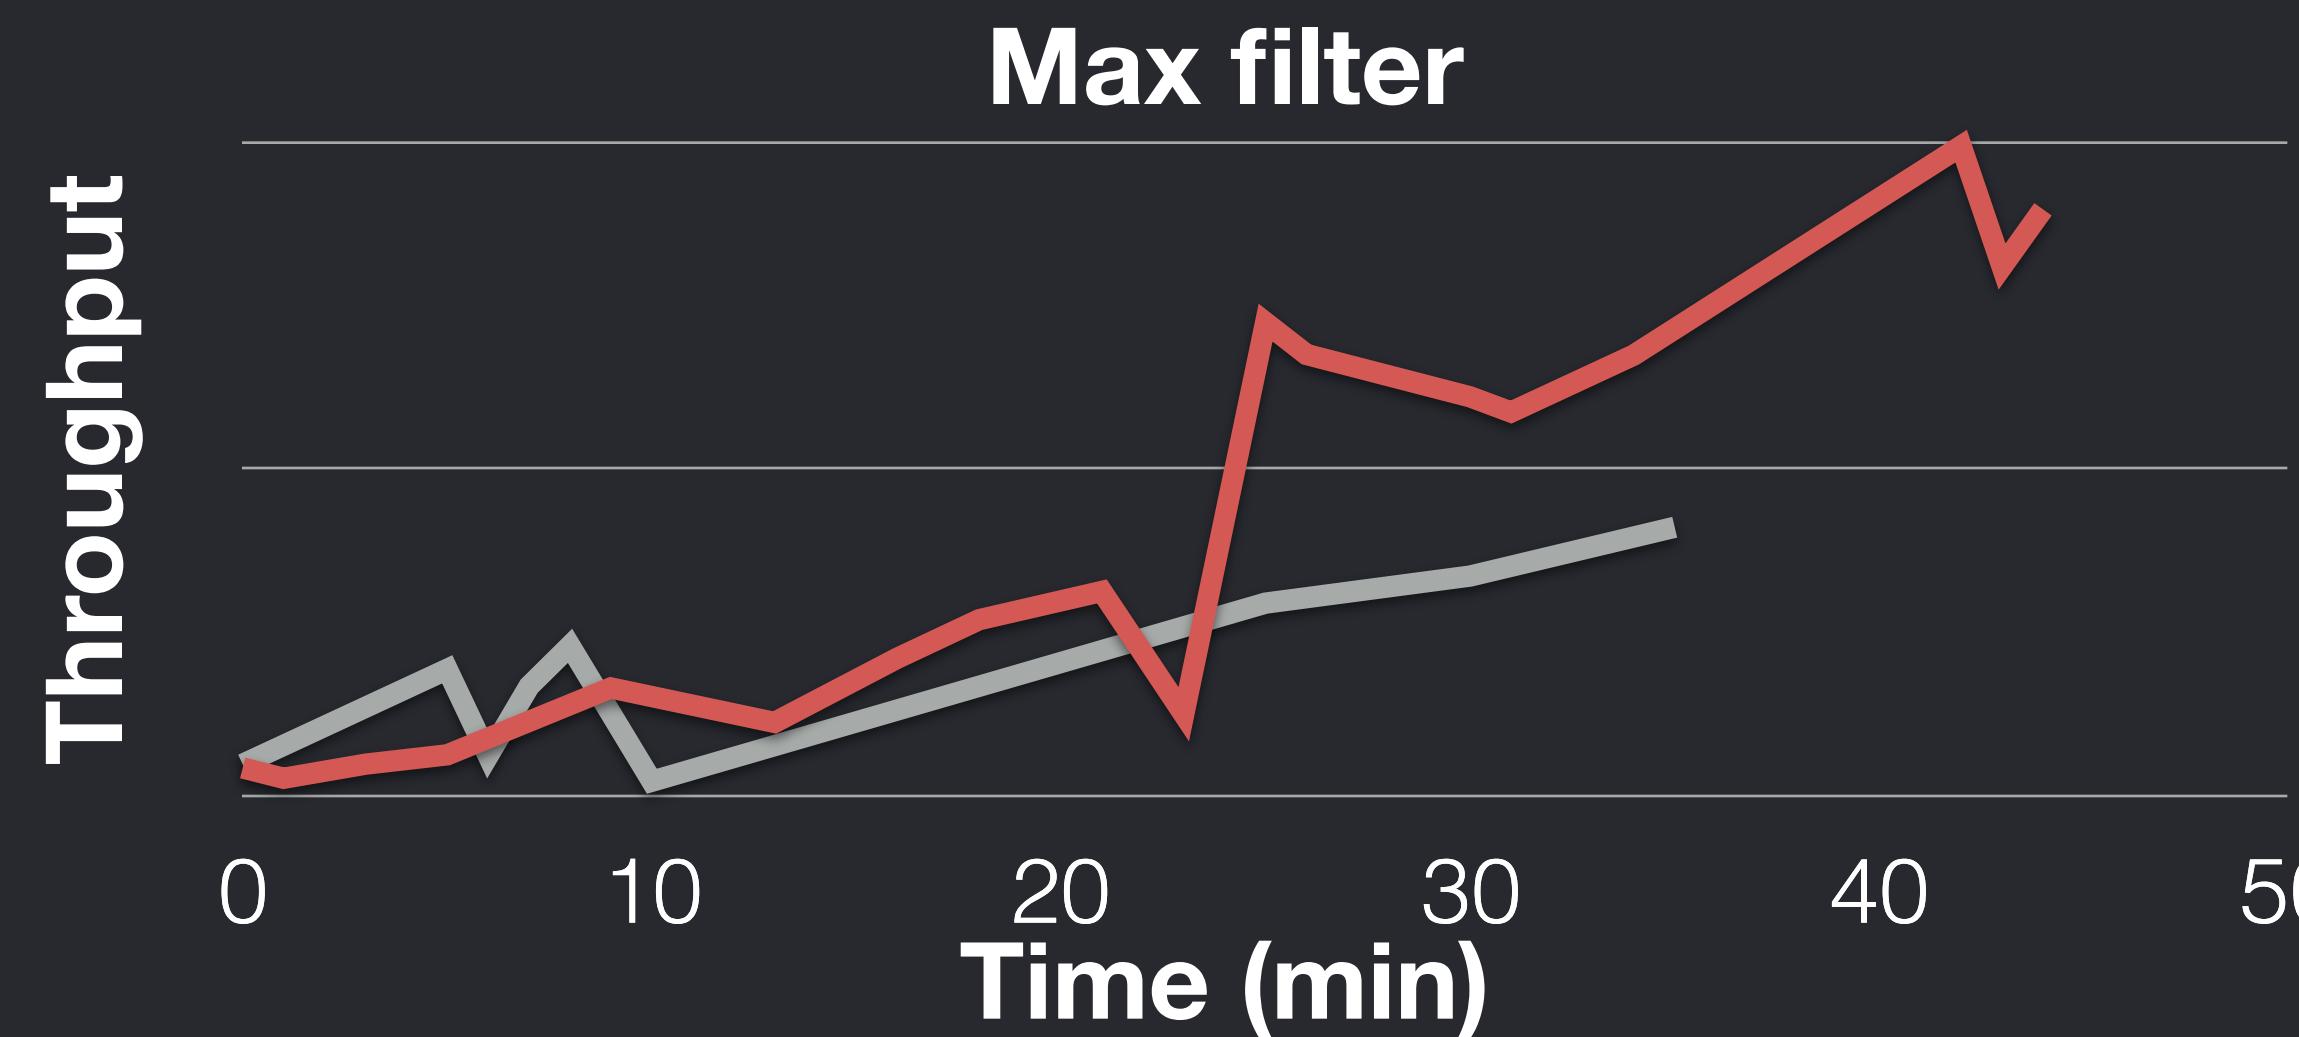
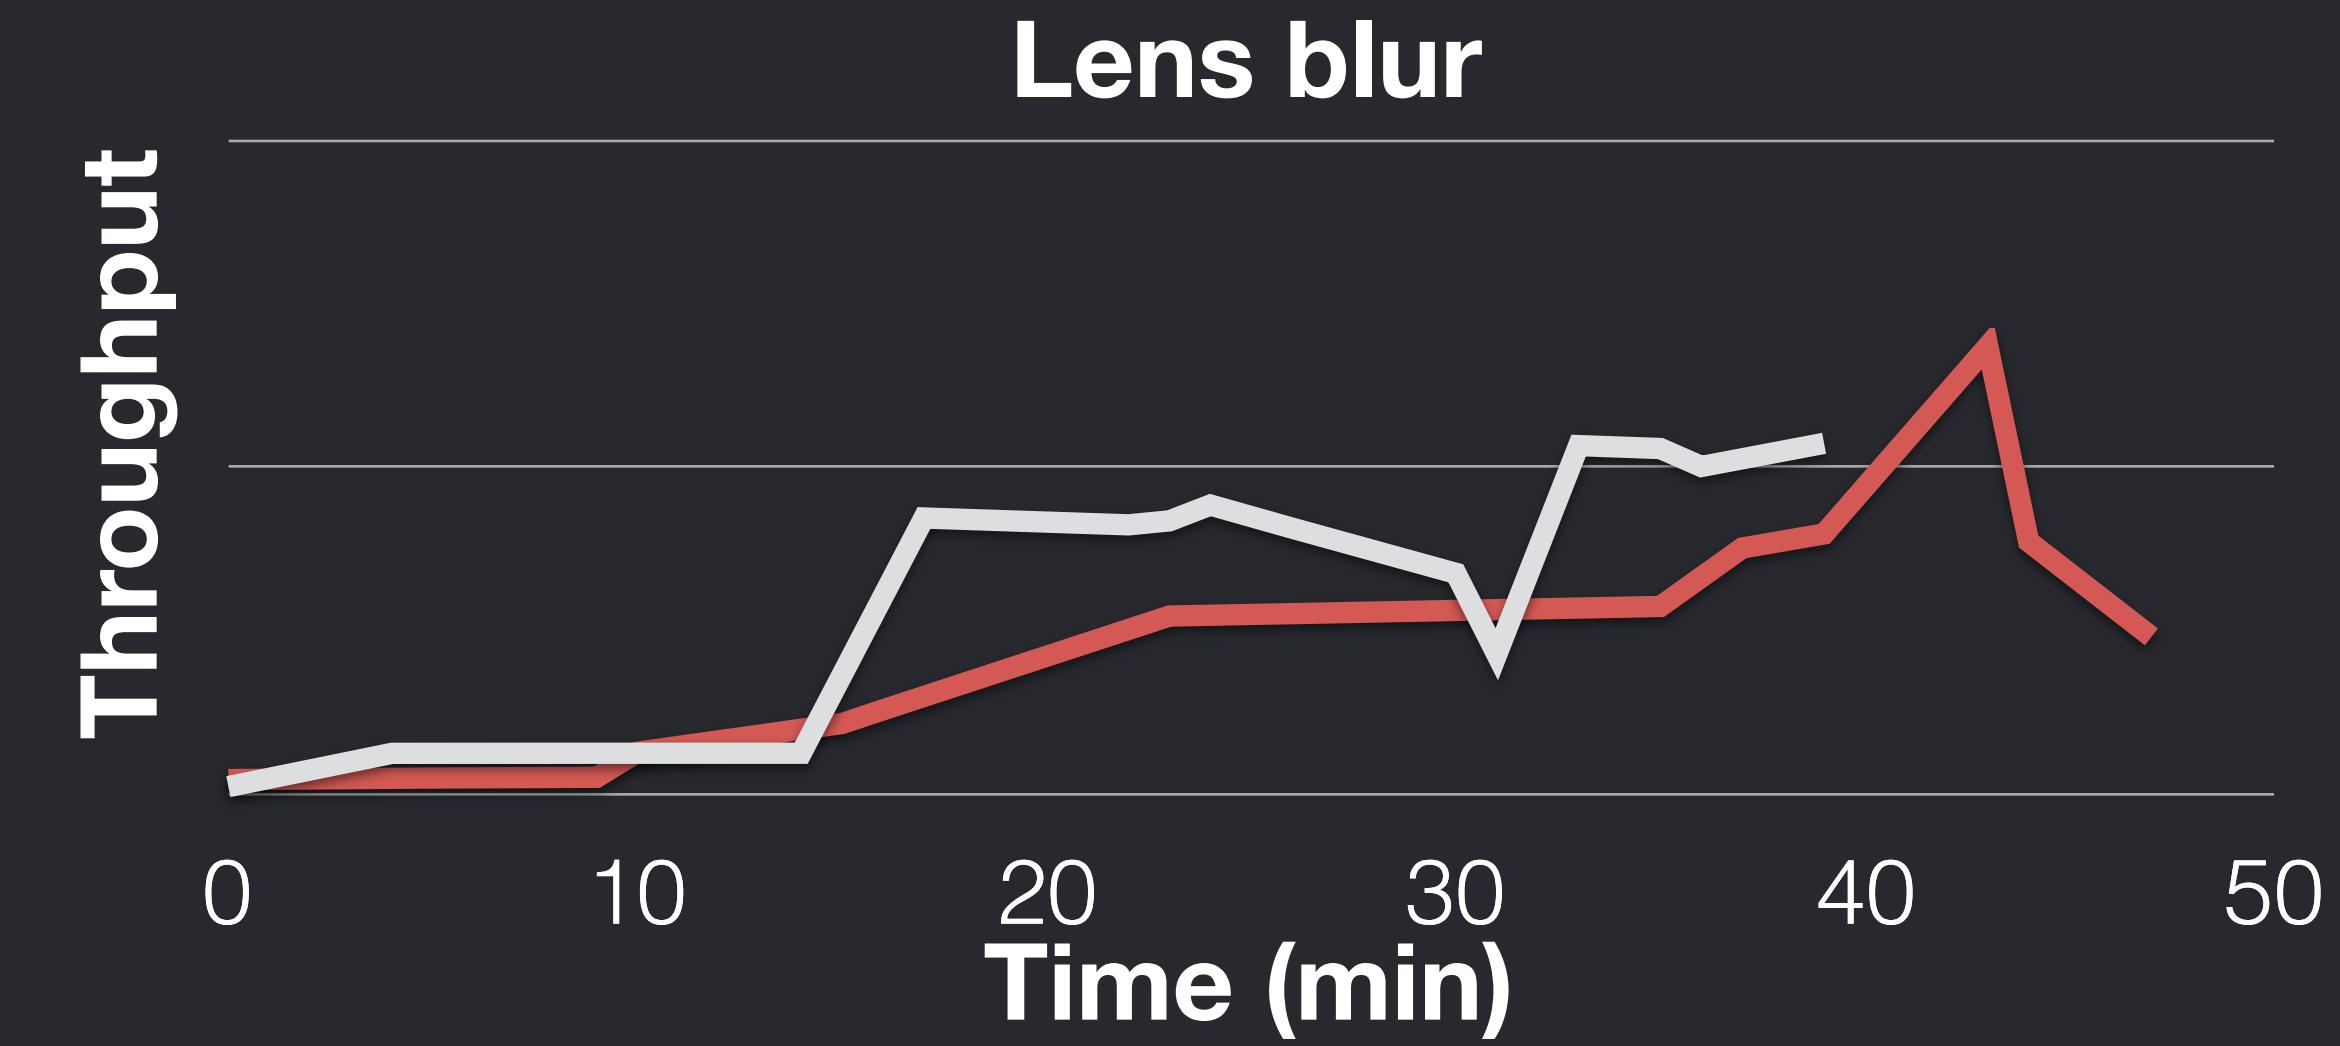
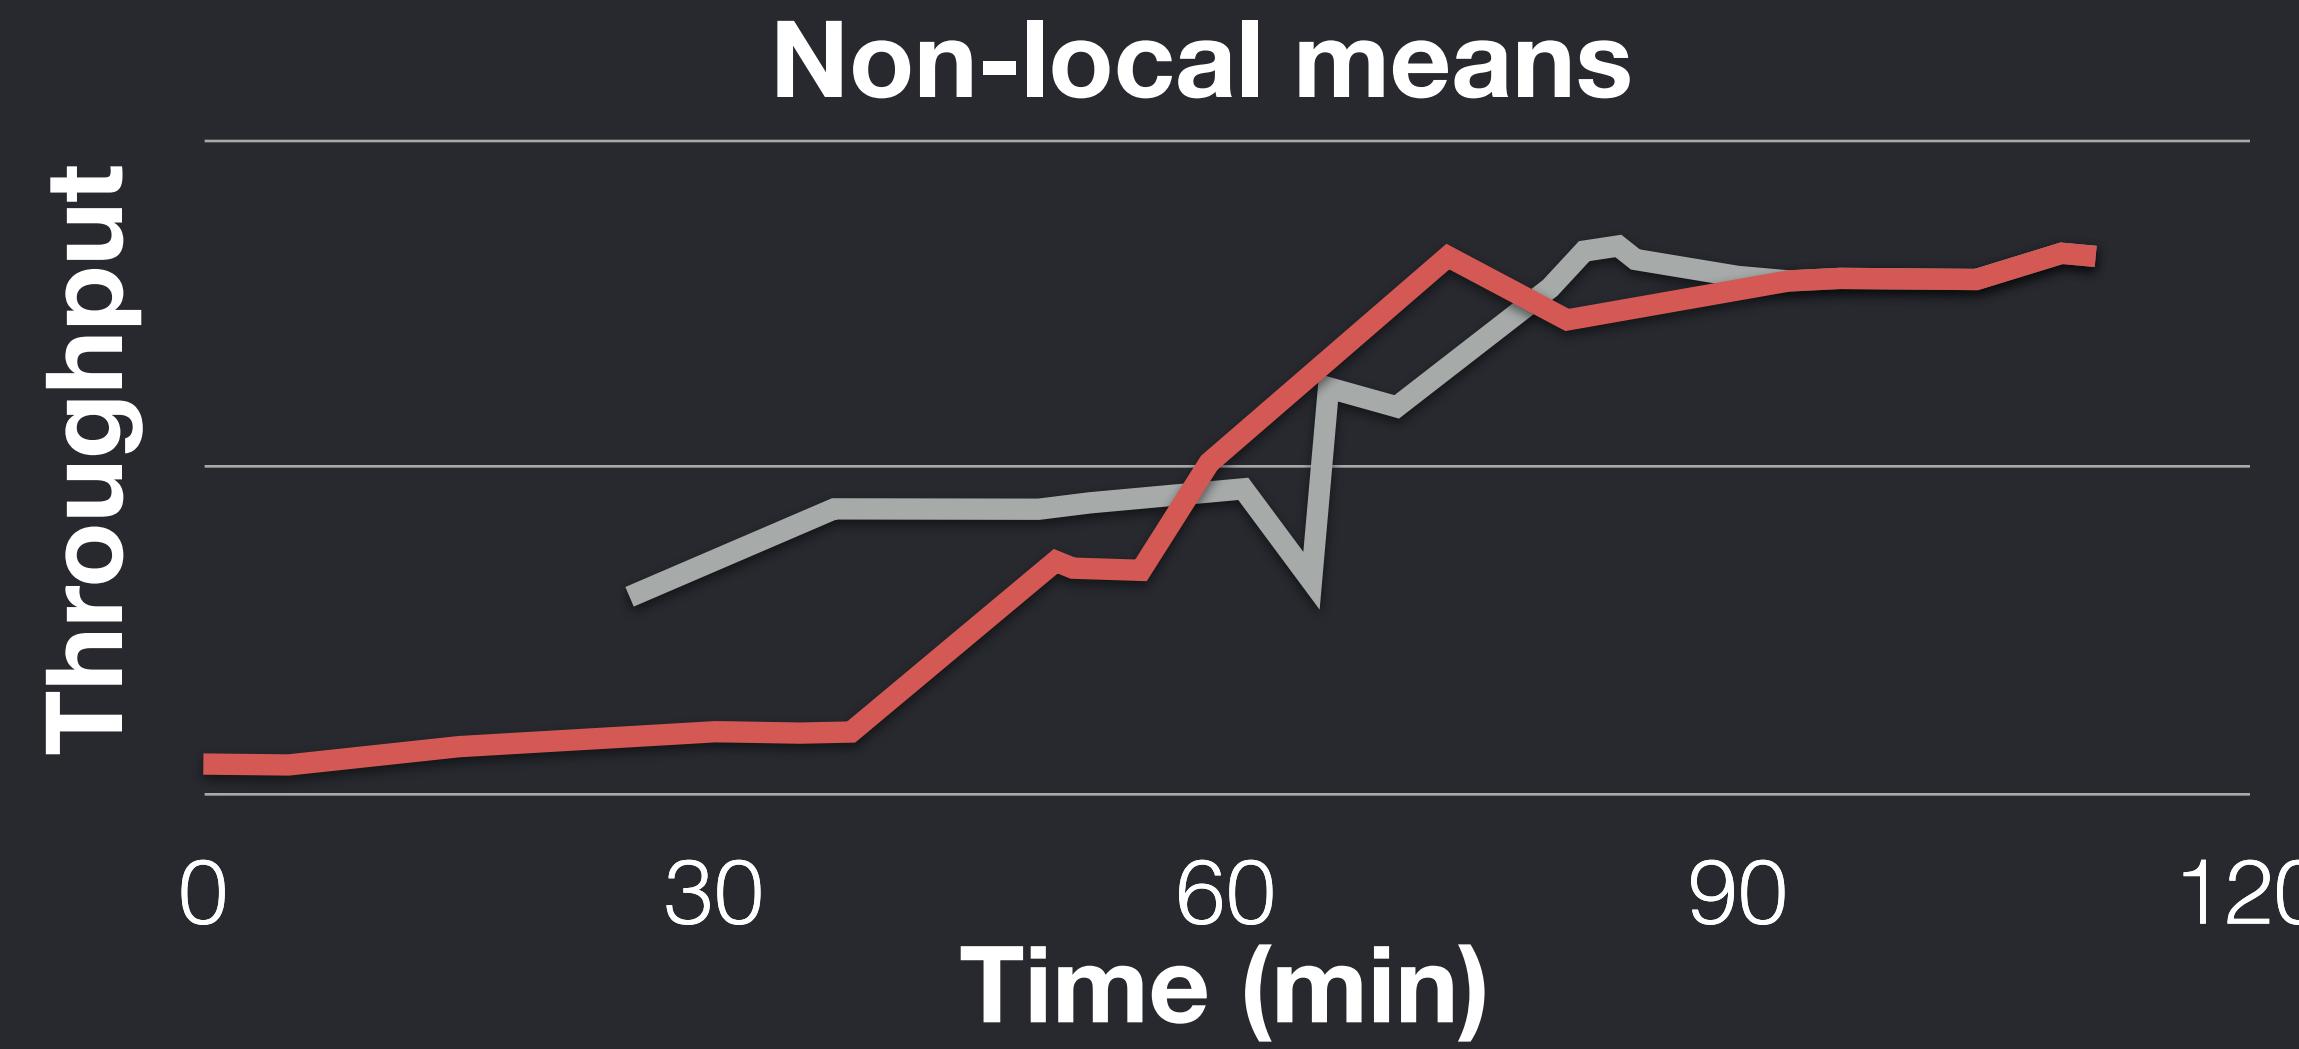
On 8 of the 14 benchmarks
performance within 10% of
experts or better

Baseline schedules exploit
multi-core and vector
parallelism but no grouping

Auto scheduler

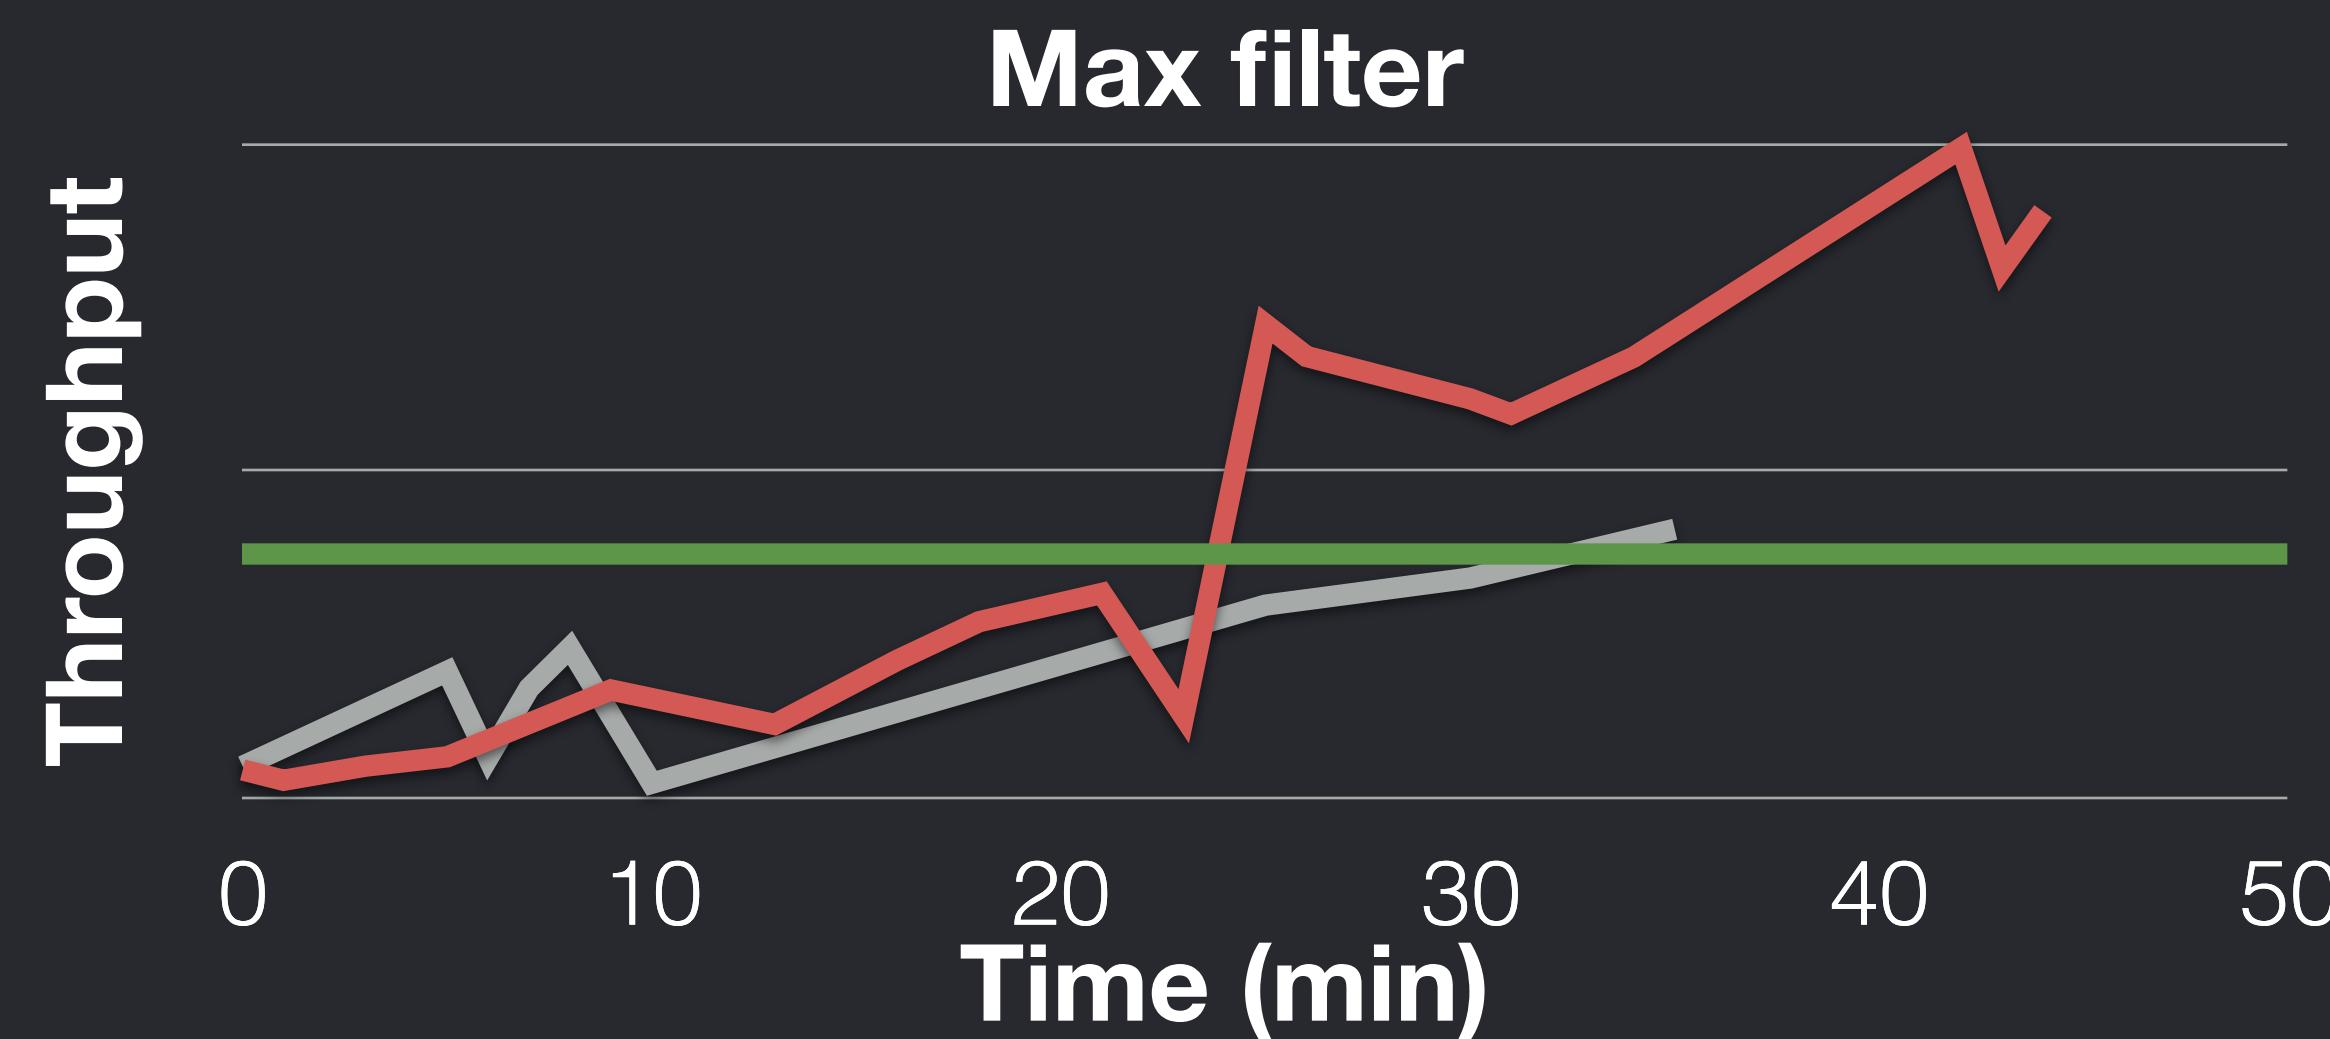
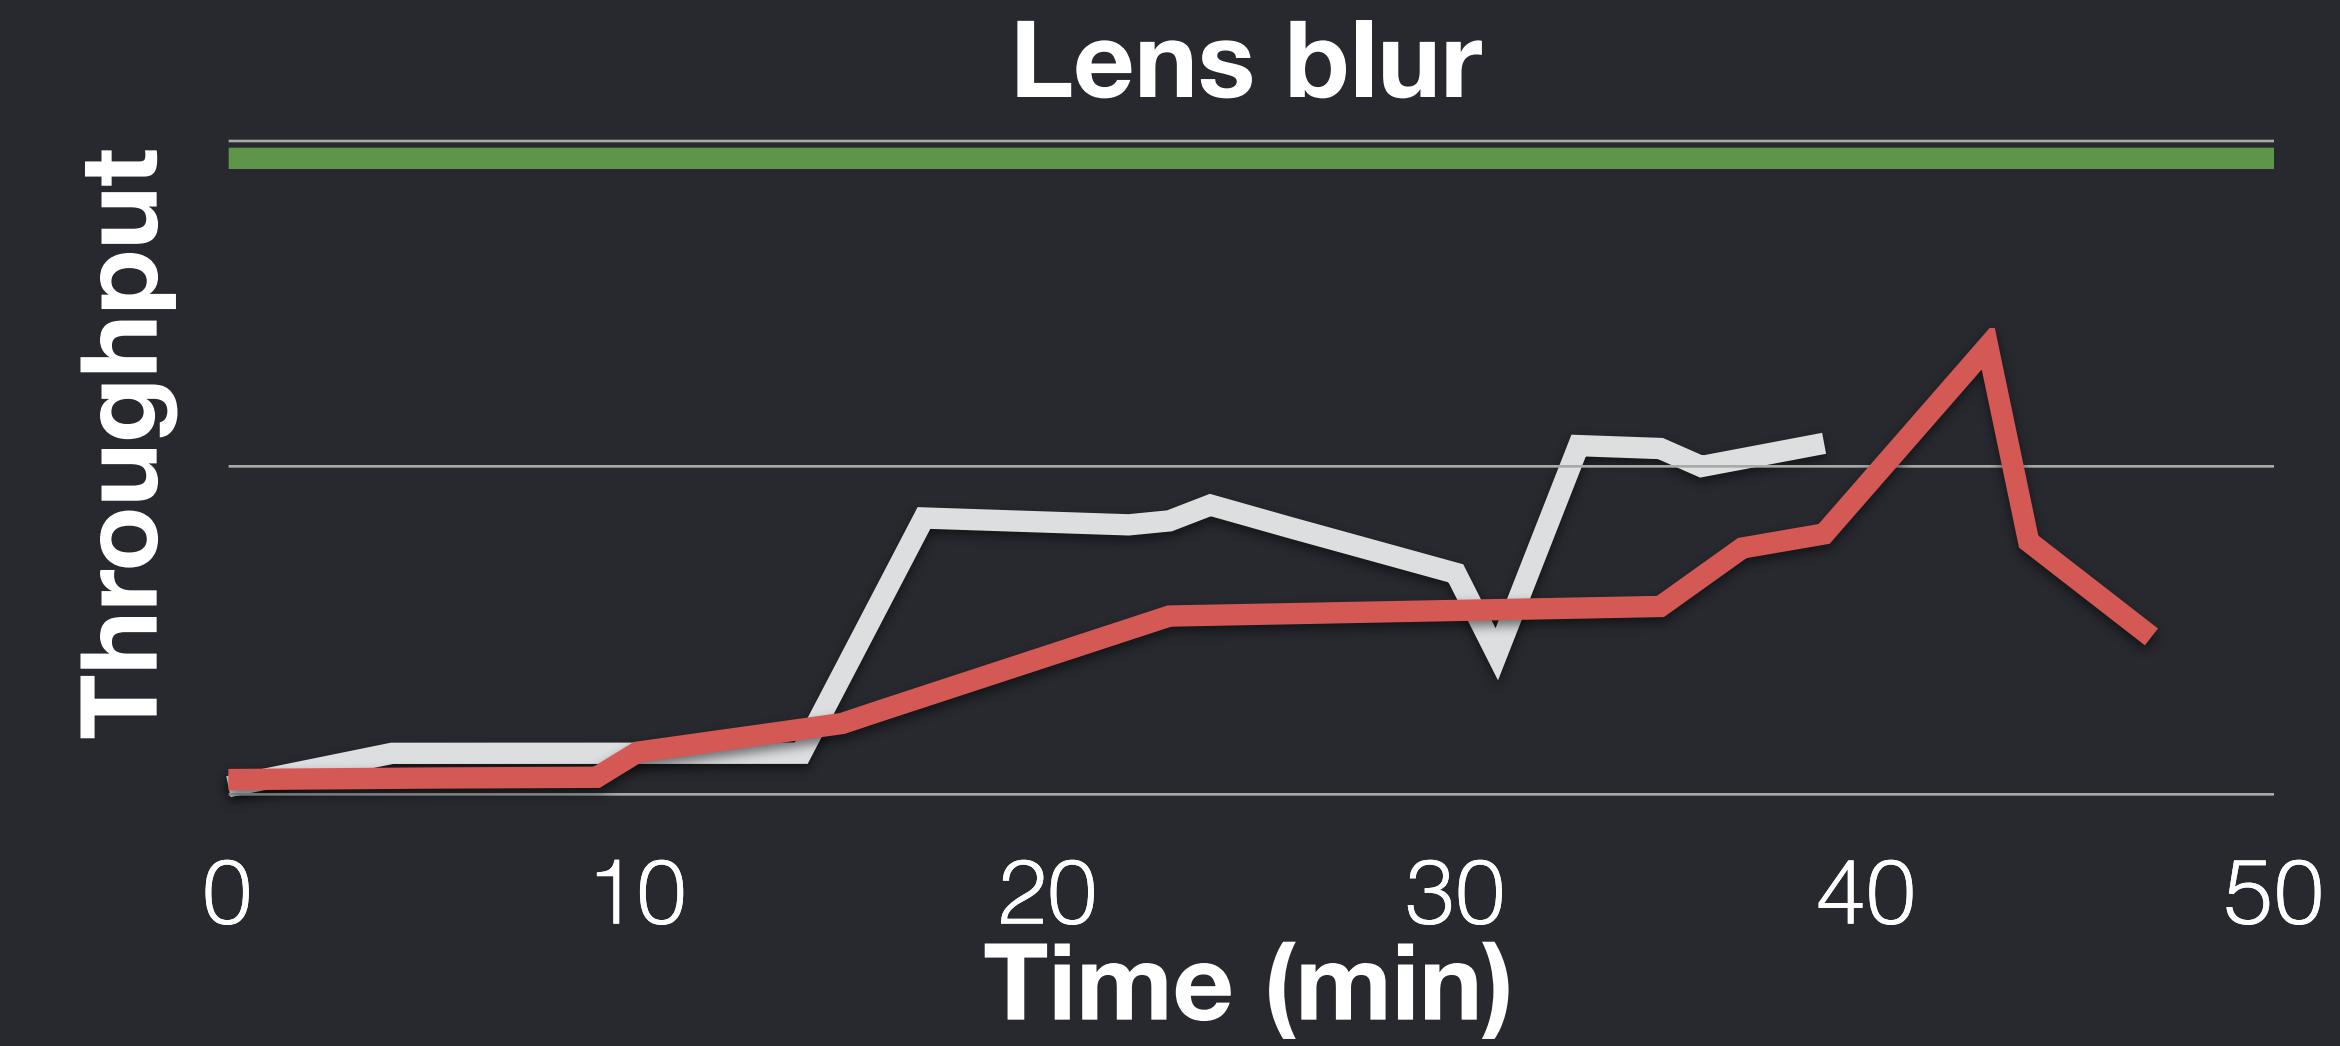
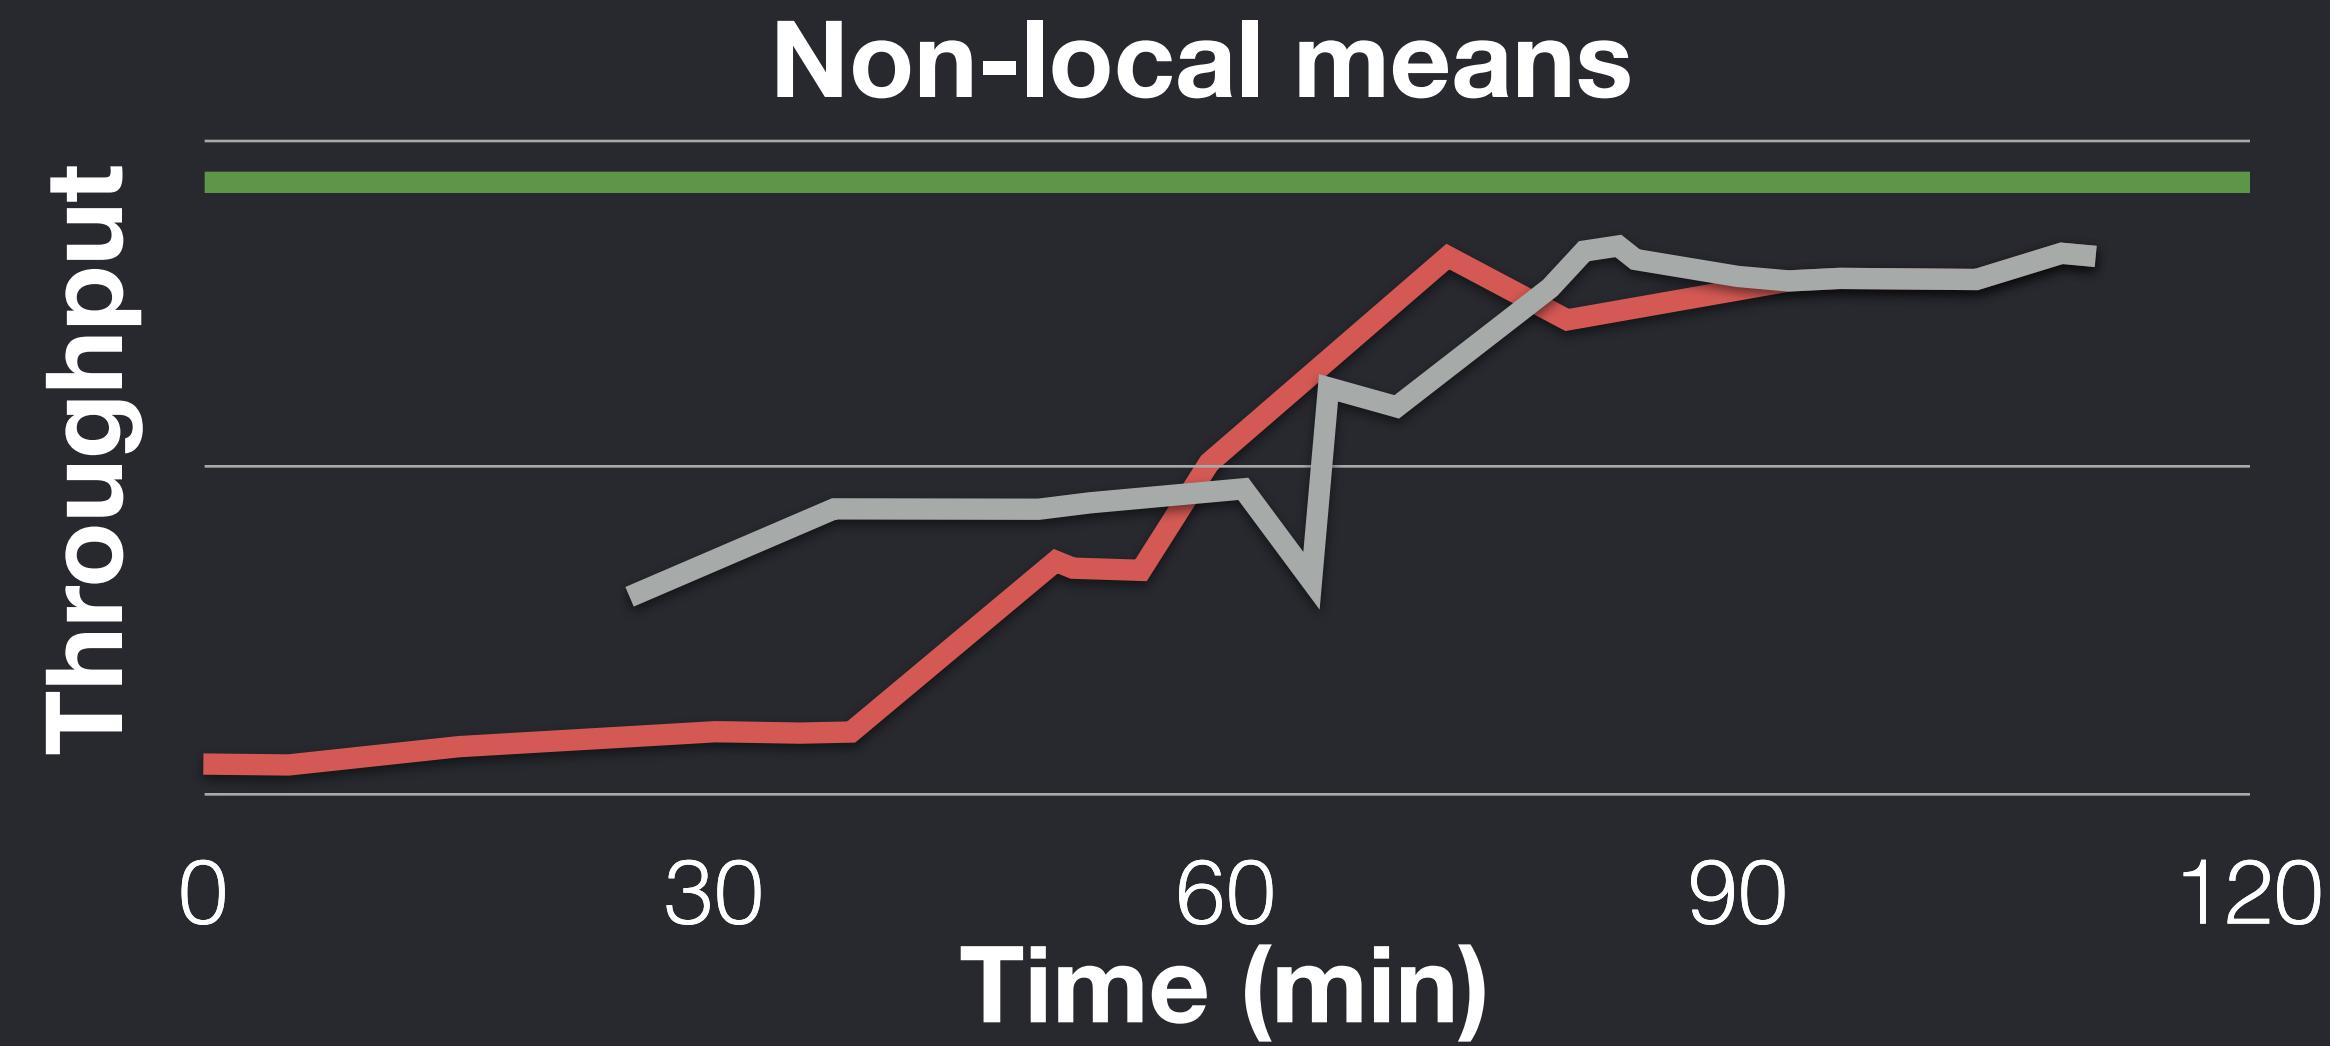
Baseline

Auto scheduler can save time for experts



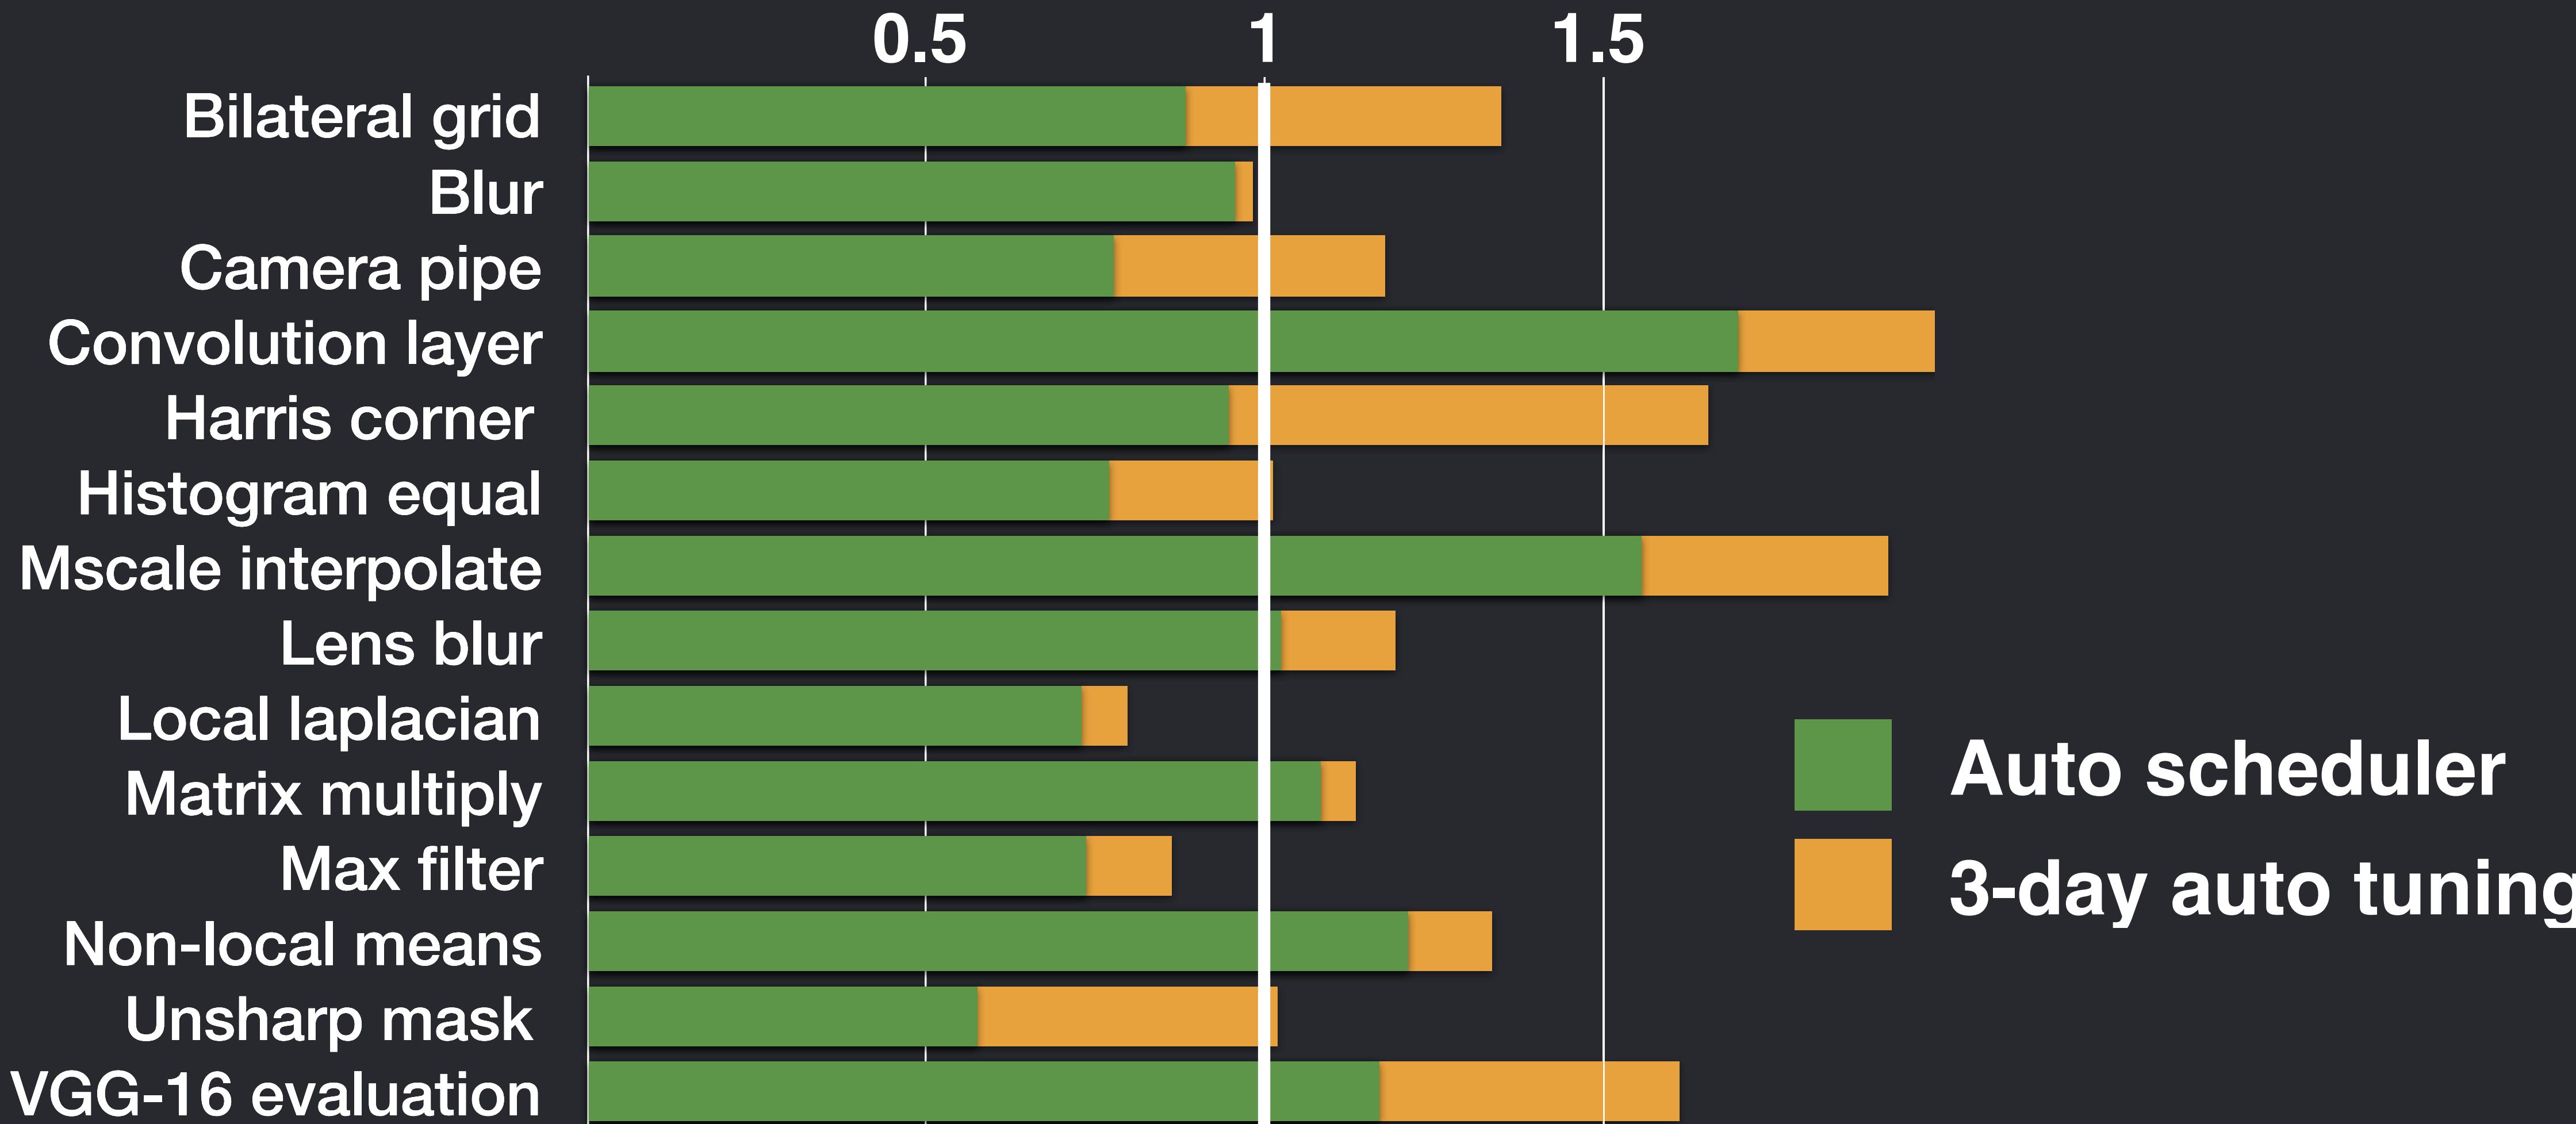
Dillon
Andrew

Auto scheduler can save time for experts



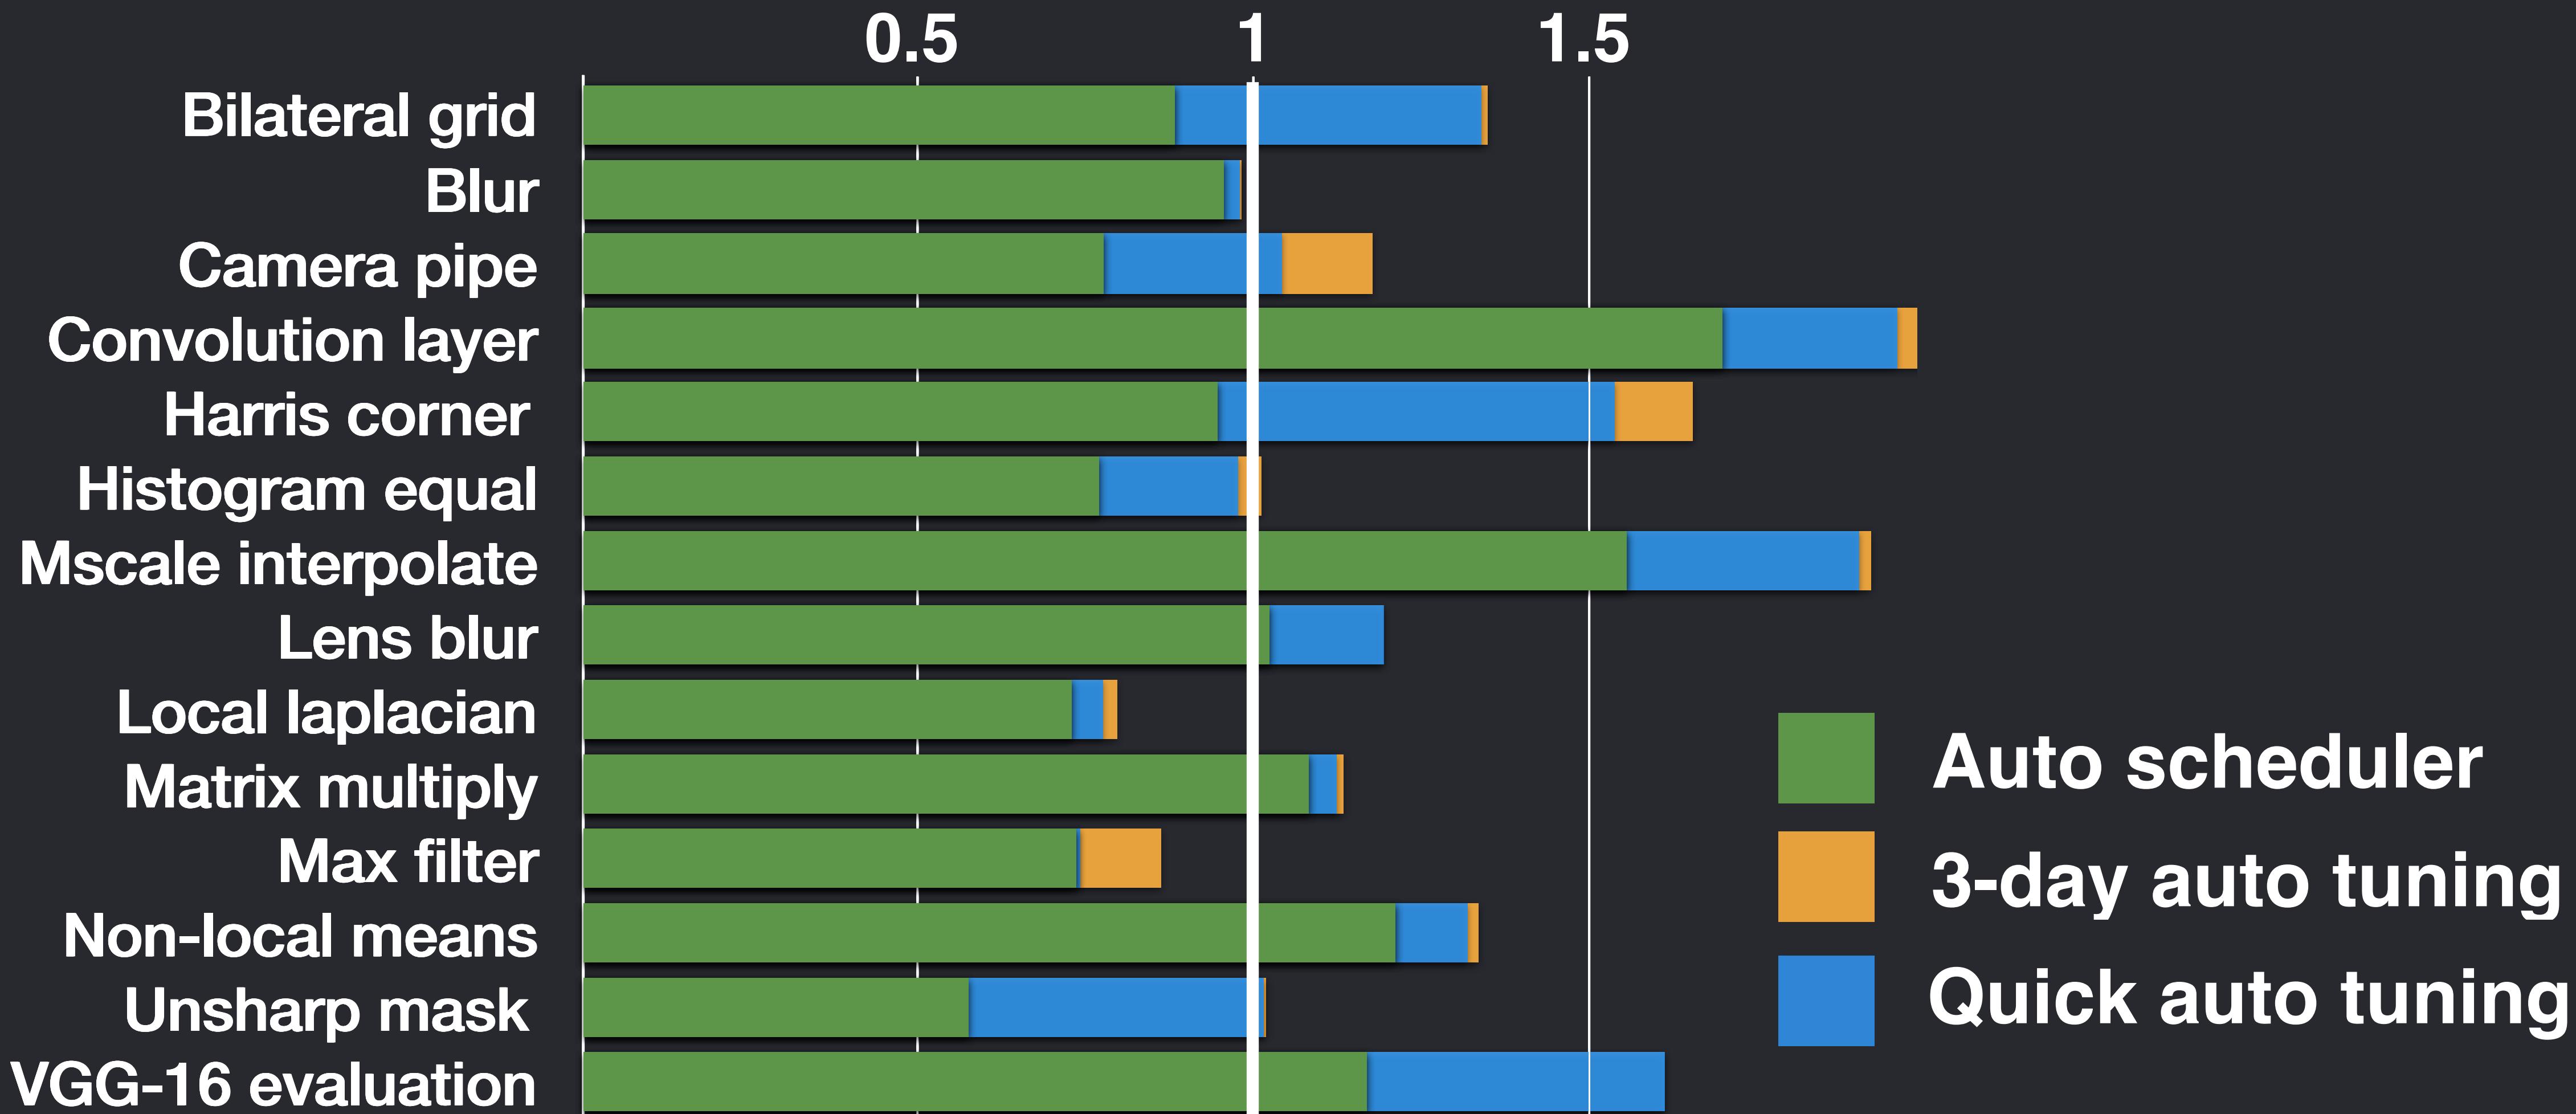
- Auto scheduler
- Dillon
- Andrew

Exploring cost model parameters



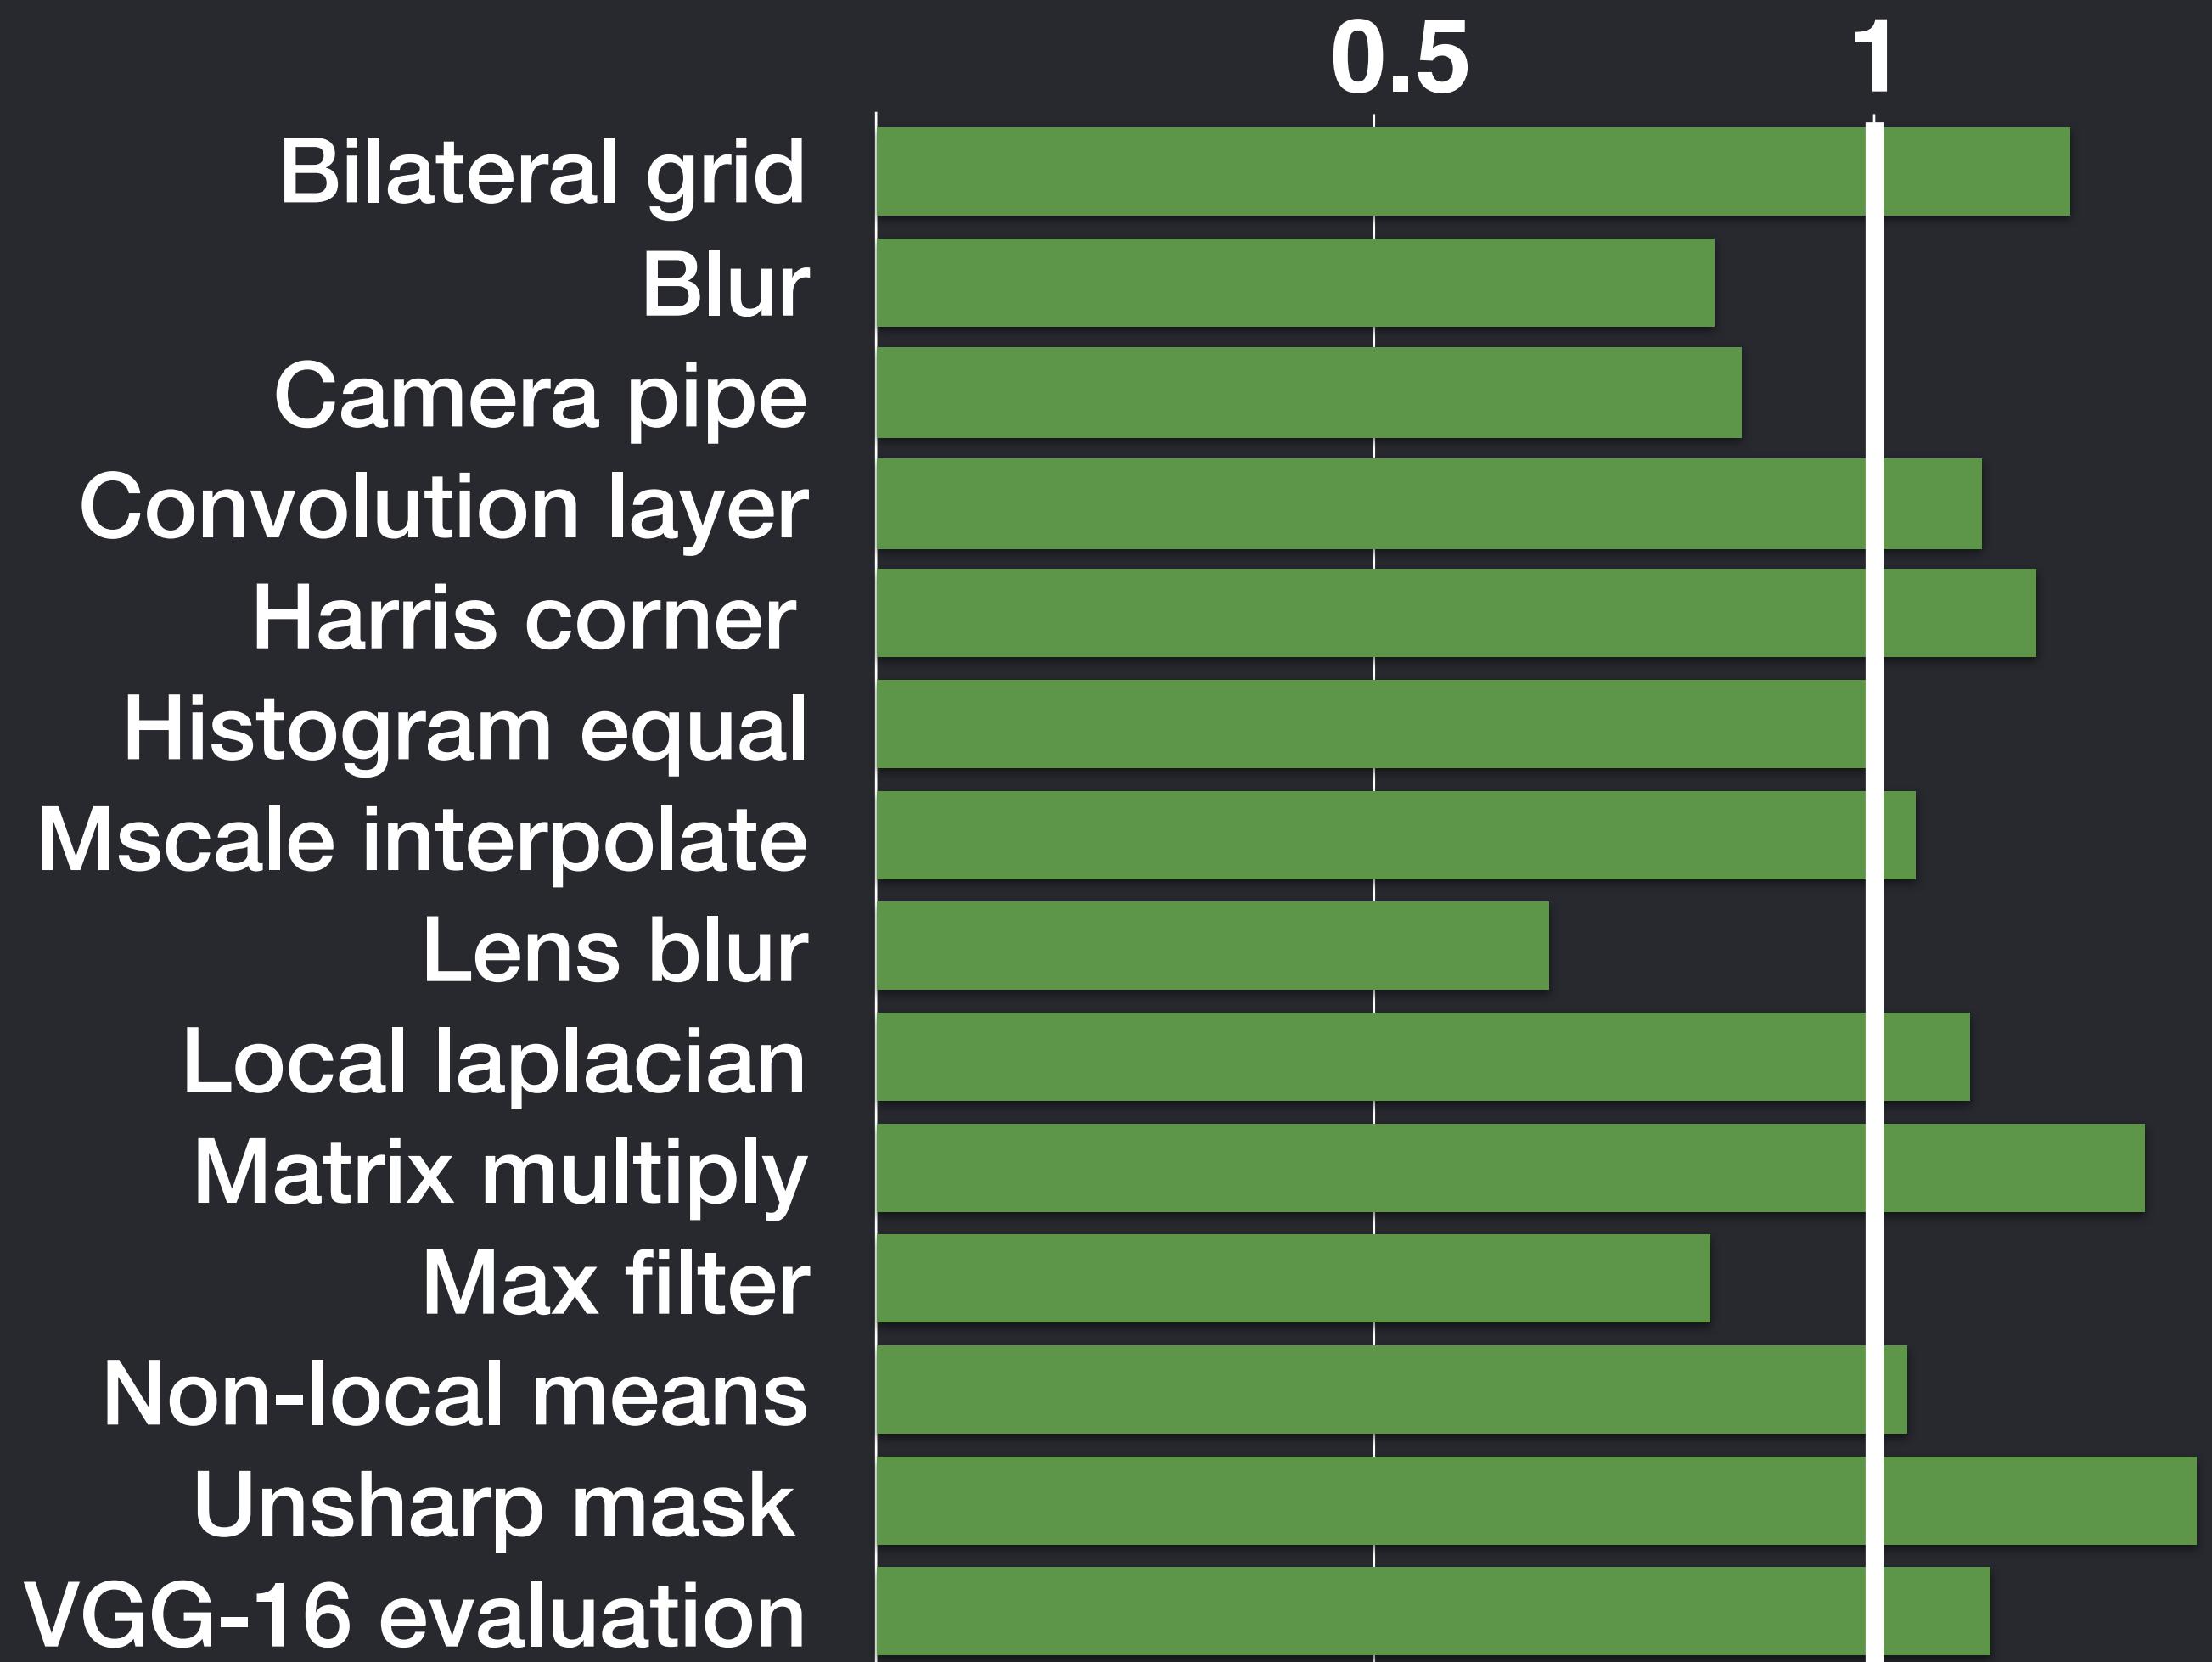
Performance relative to experts (6 core Xeon CPU)

Exploring cost model parameters



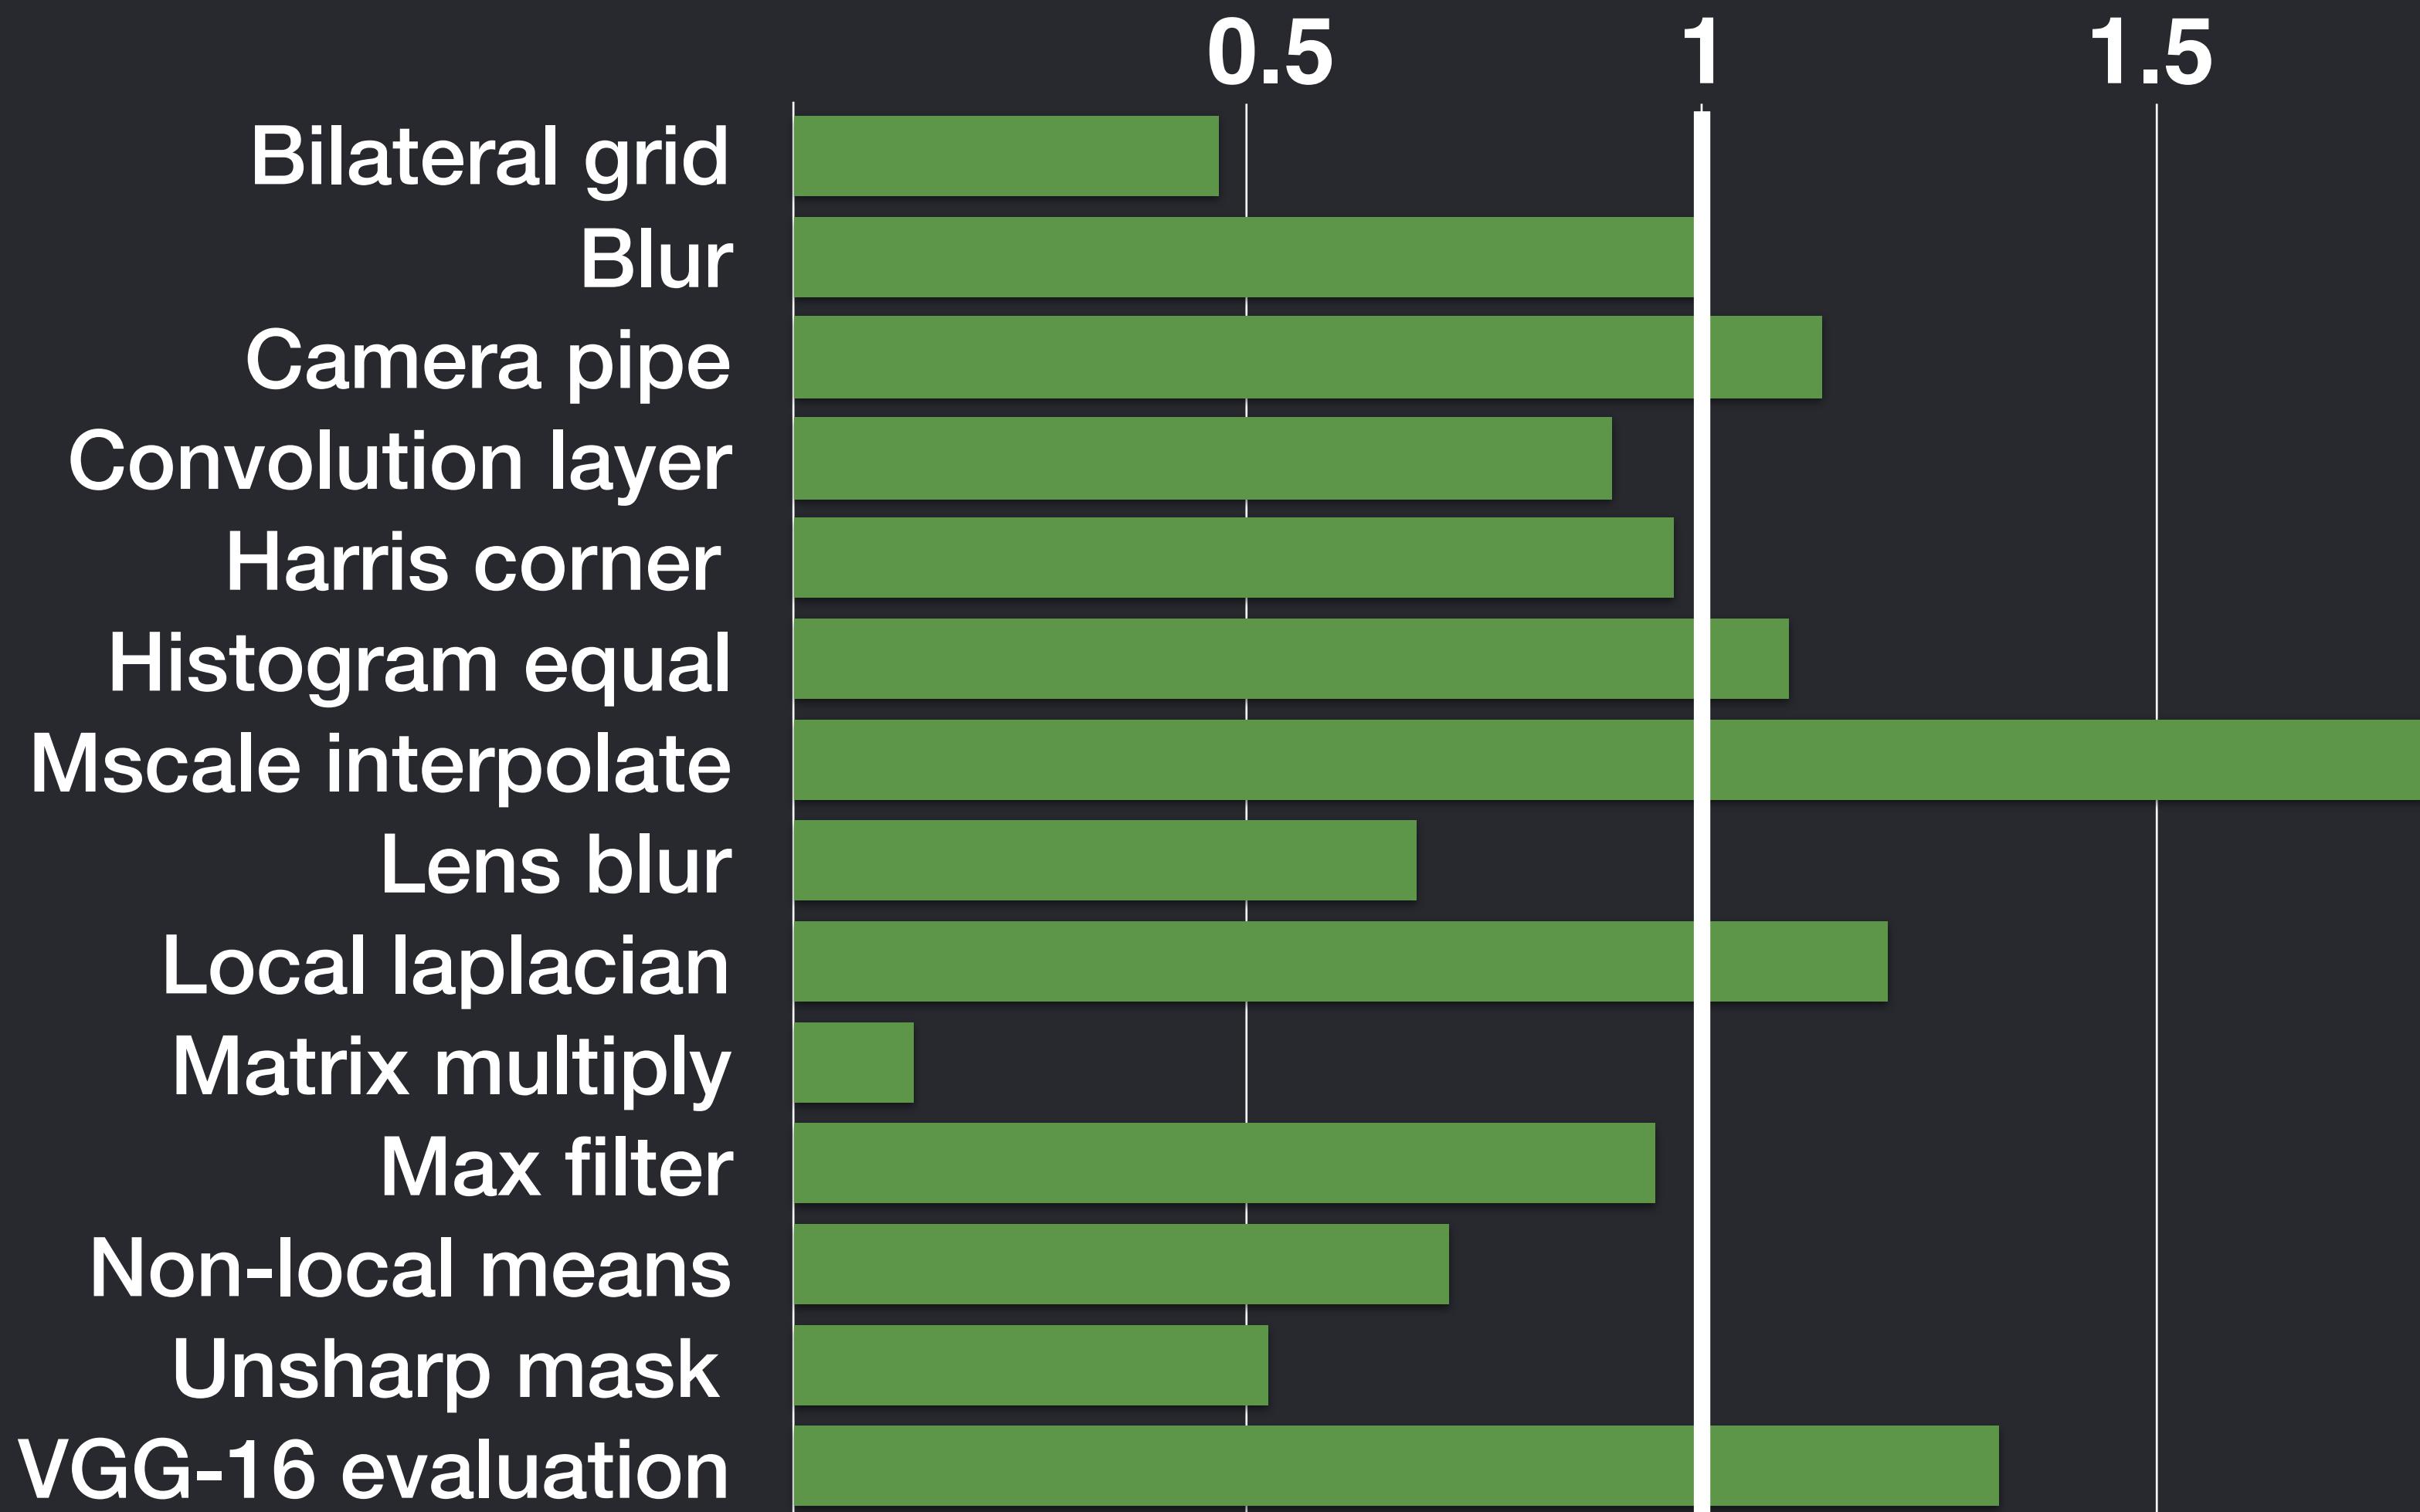
Performance relative to experts (6 core Xeon CPU)

Quad core ARM performance



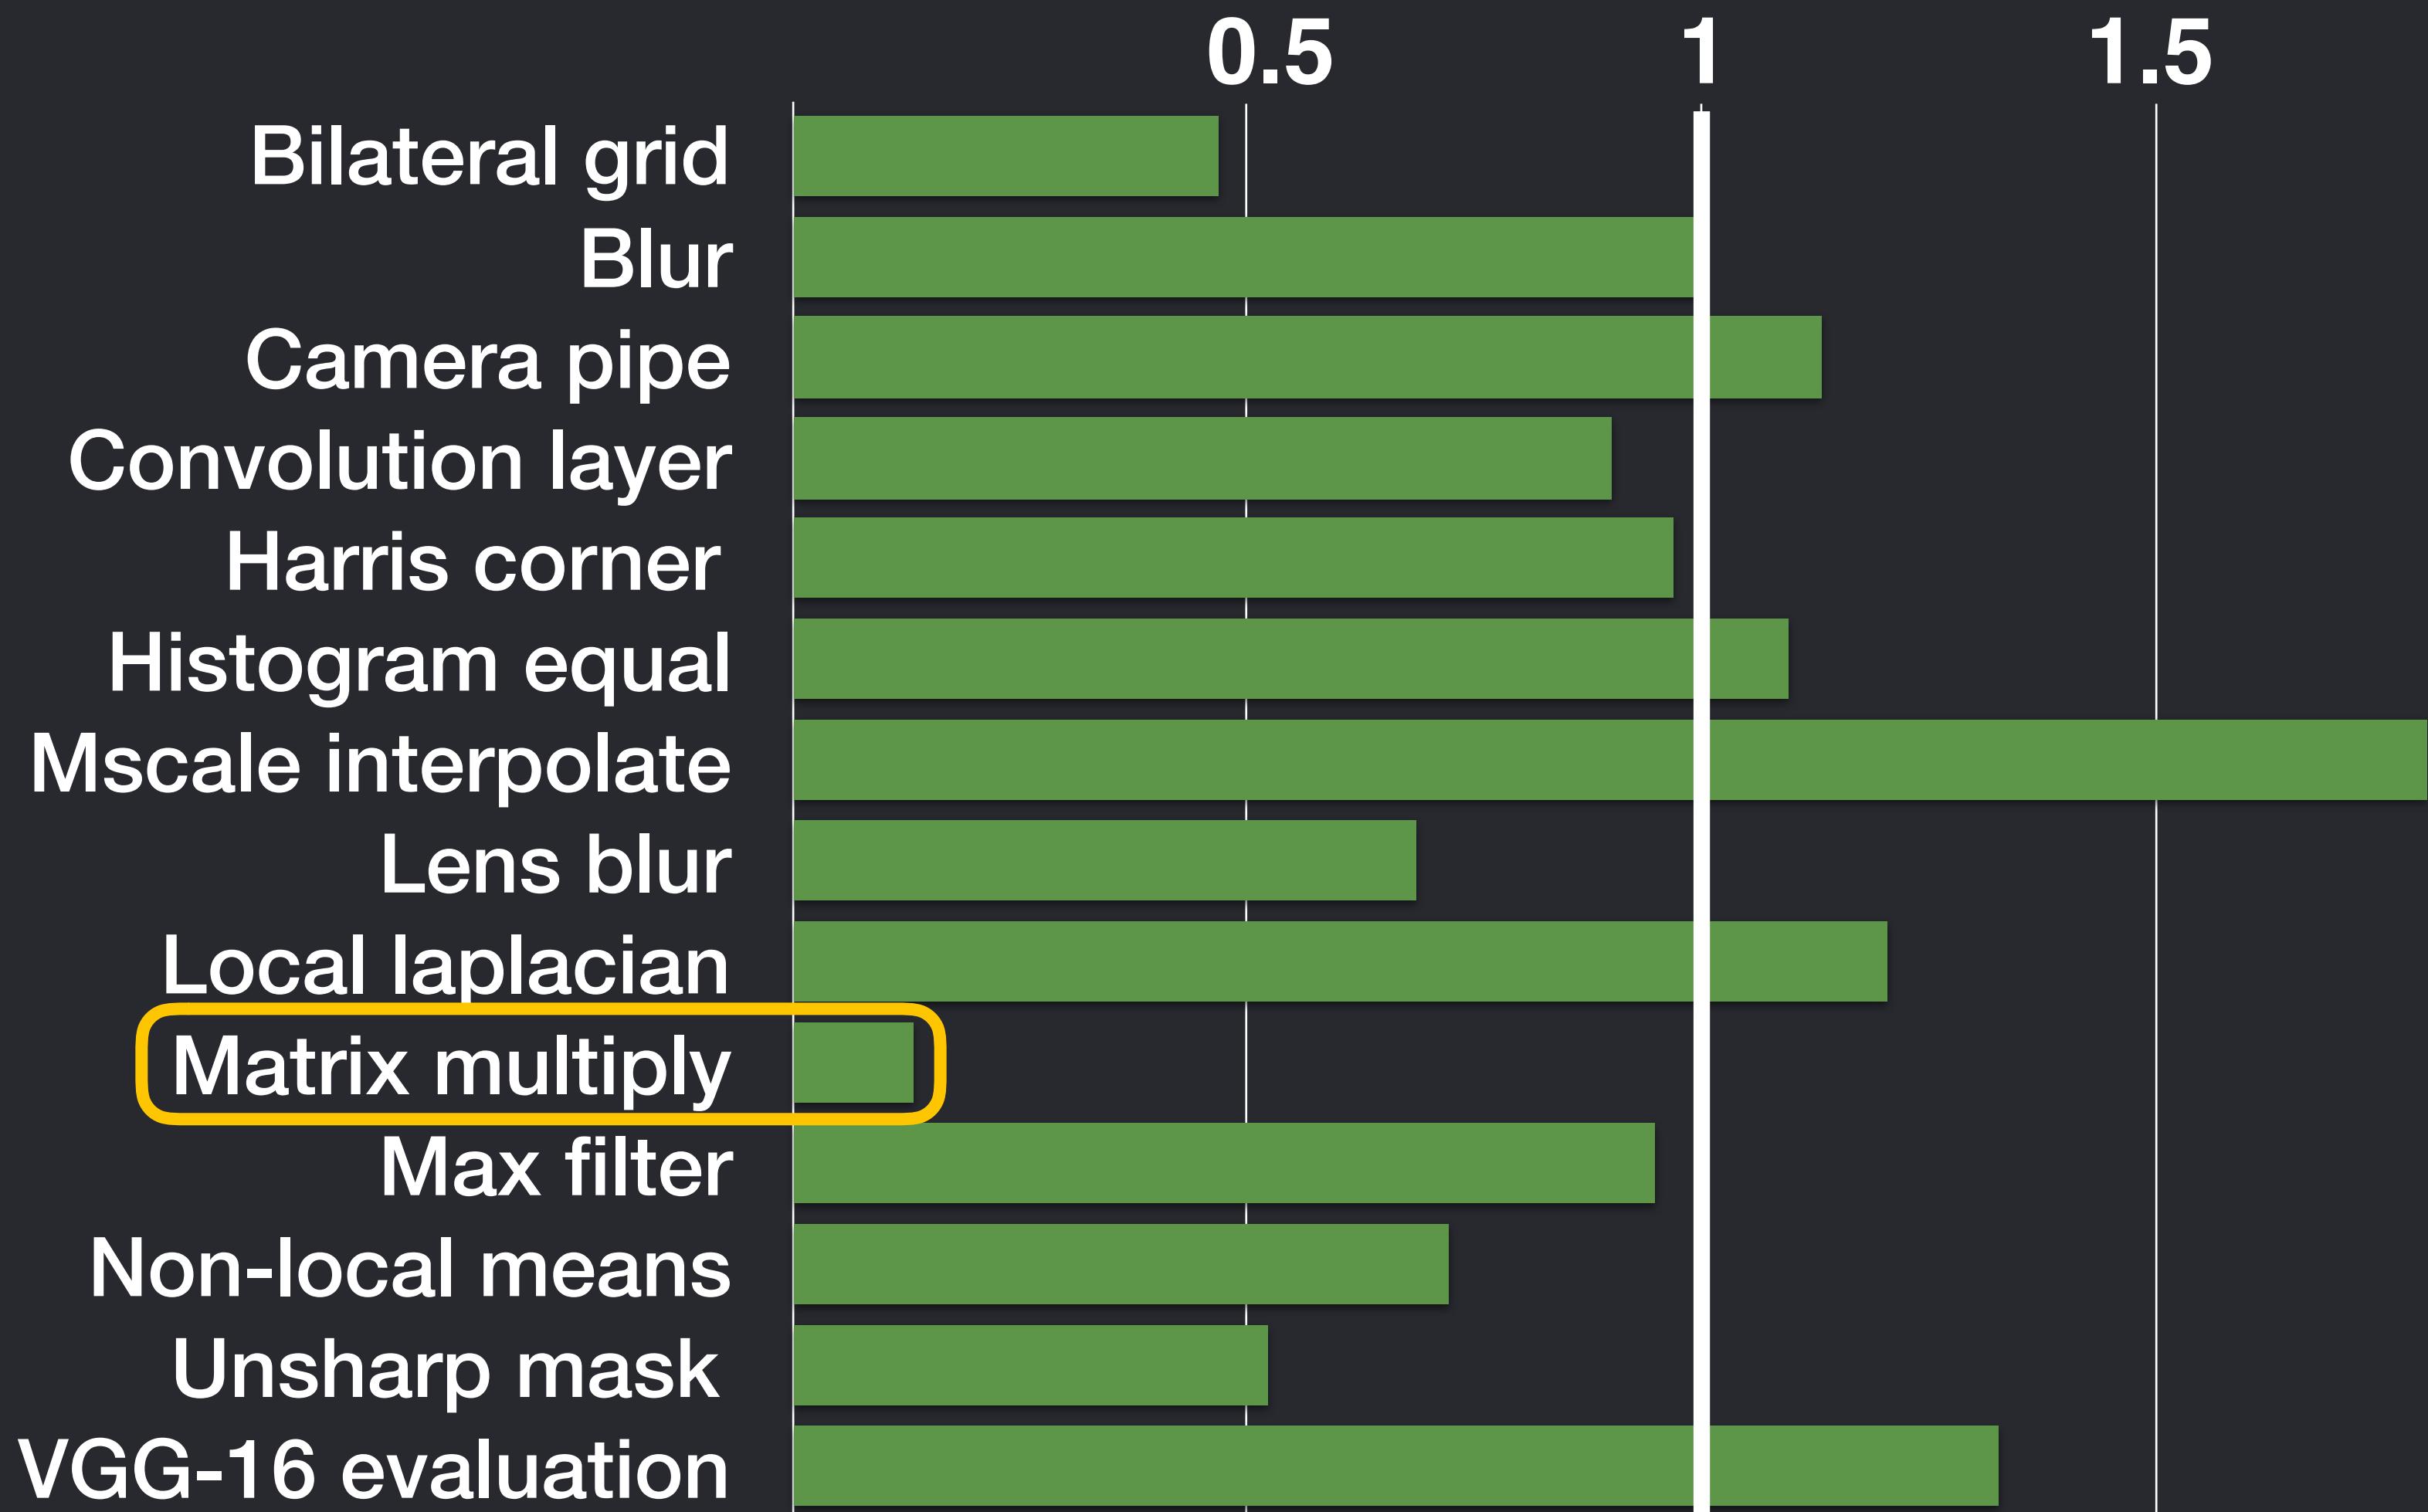
Performance relative to experts (ARM CPU)

K40 GPU performance



Performance relative to experts (K40)

K40 GPU performance



Performance relative to experts (K40)

Prior work

Optimizing Halide via auto-tuning and stochastic search

[Ragan-Kelley 13, Ansel 14]:

- Compilation time: hours to days
- Output up to 5-10x slower than hand-tuned implementations

Darkroom [Hegarty 14]:

- Auto-scheduling assuming applications restricted to fixed-size stencils

PolyMage [Mullapudi 15]: polyhedral-based optimization

- Greedy group-and-tile algorithm was inspired by PolyMage
- Polyhedral approach cannot analyze non-affine and data-dependent computations

Limitations

Restricted space of schedules

- Does not consider sliding windows and multi-level tiling

No human interaction with the auto scheduler

- Enable experts to guide the scheduling process

Summary

Algorithm that generates Halide schedules

- Competitive with experts
- Generated in seconds
- Practical implementation

In the process of being merged into mainline Halide

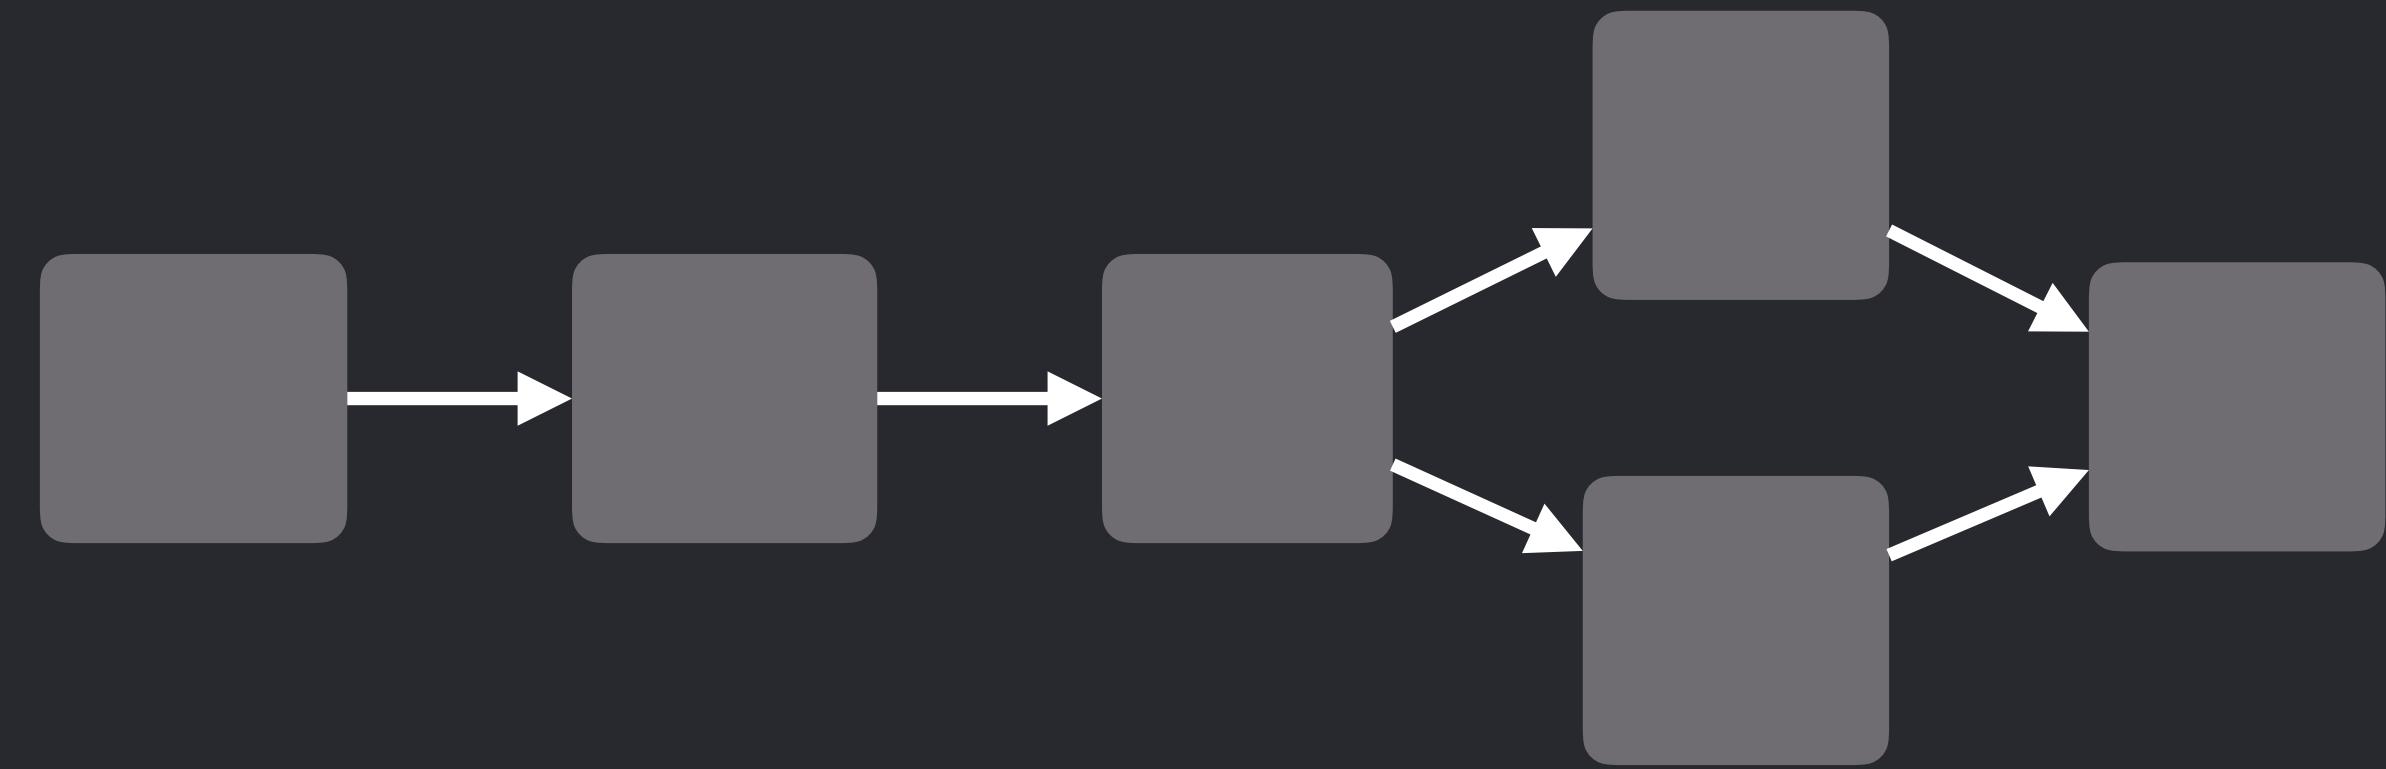
https://github.com/halide/Halide/tree/auto_scheduler

Generalizing the auto scheduler for other DSLs

Tensor Flow

Halide

Opt



Abstract analysis and scheduling techniques into
components that can be used across languages

Thank you

https://github.com/halide/Halide/tree/auto_scheduler