

Lecture 4:

High-performance image processing using Halide

**Visual Computing Systems
Stanford CS348V, Winter 2018**

**Key aspect in the design of any system:
Choosing the “right” representations for the job**

Choosing the “right” representation for the job

- **Good representations are productive to use:**
 - Embody the natural way of thinking about a problem
- **Good representations enable the system to provide the application developer **useful services**:**
 - Validating/providing certain guarantees (correctness, resource bounds, conservation of quantities, type checking)
 - Performance (parallelization, vectorization, use of specialized hardware)
 - Implementations of common, difficult-to-implement functionality (texture mapping and rasterization in 3D graphics, auto-differentiation in ML frameworks)

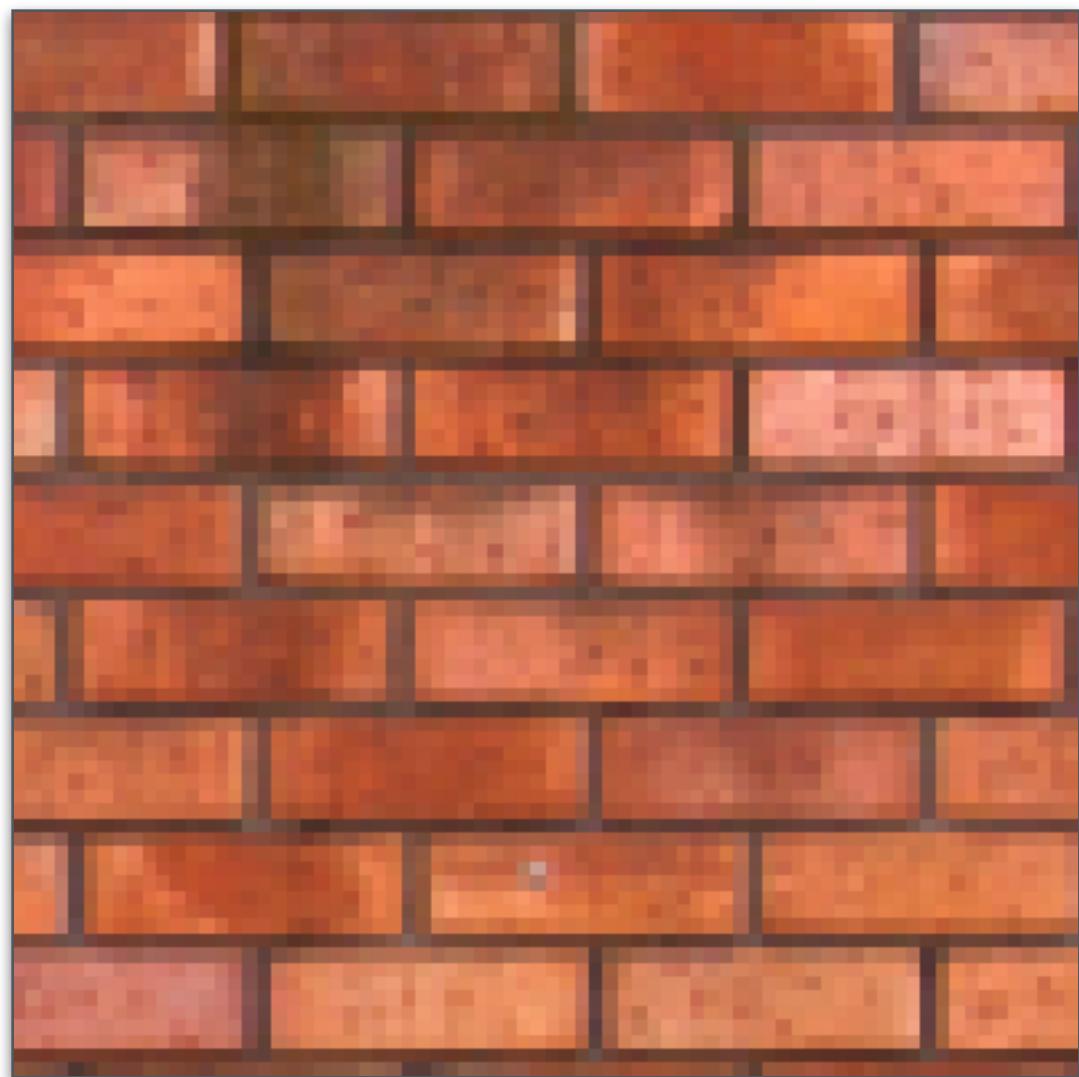
What does this code do?



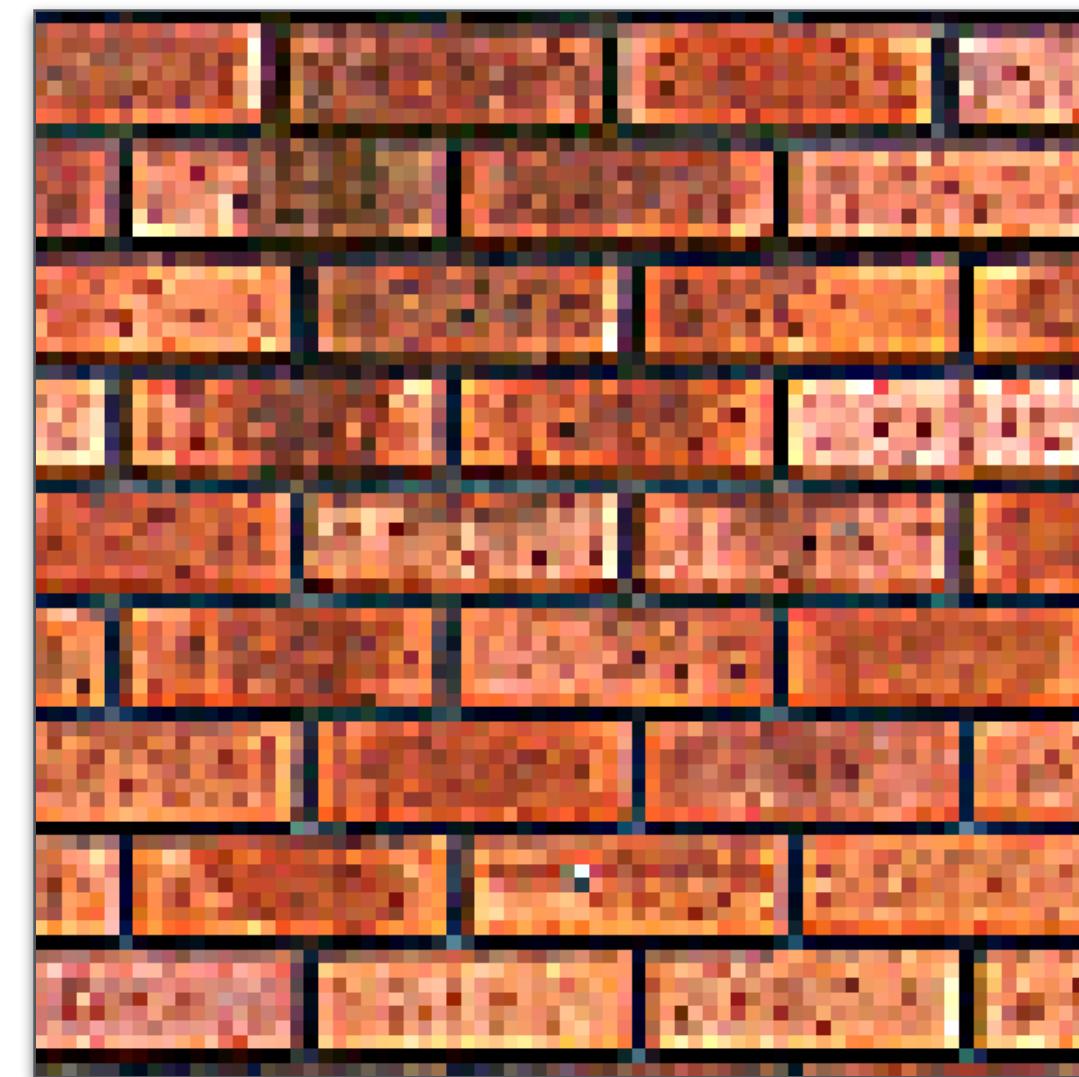
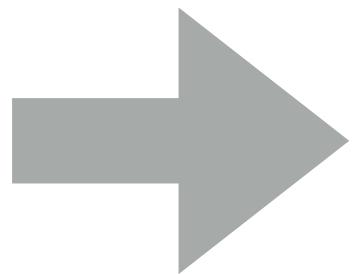
```
void mystery(const Image &in, Image &output) {
    _m128i one_third = _mm_set1_epi16(21846);
#pragma omp parallel for
for (int yTile = 0; yTile < in.height(); yTile += 32) {
    _m128i a, b, c, sum, avg;
    _m128i tmp[(256/8)*(32+2)];
    for (int xTile = 0; xTile < in.width(); xTile += 256) {
        _m128i *tmpPtr = tmp;
        for (int y = -1; y < 32+1; y++) {
            const uint16_t *inPtr = &(in(xTile, yTile+y));
            for (int x = 0; x < 256; x += 8) {
                a = _mm_loadu_si128((m128i*)(inPtr-1));
                b = _mm_loadu_si128((m128i*)(inPtr+1));
                c = _mm_load_si128((m128i*)(inPtr));
                sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                avg = _mm_mulhi_epi16(sum, one_third);
                _mm_store_si128(tmpPtr++, avg);
                inPtr += 8;
            }
        }
        tmpPtr = tmp;
        for (int y = 0; y < 32; y++) {
            _m128i *outPtr = (m128i*)(&(output(xTile, yTile+y)));
            for (int x = 0; x < 256; x += 8) {
                a = _mm_load_si128(tmpPtr+(2*256)/8);
                b = _mm_load_si128(tmpPtr+256/8);
                c = _mm_load_si128(tmpPtr++);
                sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                avg = _mm_mulhi_epi16(sum, one_third);
                _mm_store_si128(outPtr++, avg);
            }
        }
    }
}
}}}}
```

Example task: sharpen an image

$$F = \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



Input



Output

Four different representations of sharpen

Image input;
Image output = sharpen(input);

$$F = \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

Image input;
Image output = convolve(input, F);

Image input;
Image output;
output[i][j]
= F[0][0] * input[i-1][j-1] +
F[0][1] * input[i-1][j] +
F[0][2] * input[i-1][j+1] +
F[1][0] * input[i][j-1] +
F[1][1] * input[i][j] +

1

2

3

4

```
float input[(WIDTH+2) * (HEIGHT+2)];  
float output[WIDTH * HEIGHT];  
  
float weights[] = {0., -1., 0.,  
                   -1., 5, -1.,  
                   0., -1., 0.};  
  
for (int j=0; j<HEIGHT; j++) {  
    for (int i=0; i<WIDTH; i++) {  
        float tmp = 0.f;  
        for (int jj=0; jj<3; jj++)  
            for (int ii=0; ii<3; ii++)  
                tmp += input[(j+jj)*(WIDTH+2) + (i+ii)]  
                      * weights[jj*3 + ii];  
        output[j*WIDTH + i] = tmp;  
    }  
}
```

More image processing tasks from last lecture

Sobel Edge Detection

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} * I$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix} * I$$

$$G = \sqrt{G_x^2 + G_y^2}$$

Local Pixel Clamp

```
float f(image input) {
    float min_value = min( min(input[x-1][y], input[x+1][y]),
                           min(input[x][y-1], input[x][y+1]) );
    float max_value = max( max(input[x-1][y], input[x+1][y]),
                           max(input[x][y-1], input[x][y+1]) );
    output[x][y] = clamp(min_value, max_value, input[x][y]);
    output[x][y] = f(input);
```

3x3 Gaussian blur

$$F = \begin{bmatrix} .075 & .124 & .075 \\ .124 & .204 & .124 \\ .075 & .124 & .075 \end{bmatrix}$$

2x2 downsample (via averaging)

```
output[x][y] = (input[2x][2y] + input[2x+1][2y] +
                  input[2x][2y+1] + input[2x+1][2y+1]) / 4.f;
```

Gamma Correction

```
output[x][y] = pow(input[x][y], 0.5f);
```

LUT-based correction

```
output[x][y] = lookup_table[input[x][y]];
```

Histogram

```
bin[input[x][y]]++;
```

Image processing workload characteristics

- Sequences of operations on images
- Natural to think about algorithms in terms of their local behavior: “pointwise code”: output at pixel (x,y) is function of input pixels in neighborhood around (x,y)
- Common case: access to local “window” of pixels around a point
- But some algorithms require data-dependent data access (e.g., data-dependent access to lookup-tables)
- Multiple rates of computation (upsampling/downsampling)
- Simple inter-pixel communication/reductions (e.g., building a histogram, computing maximum brightness pixel)

Halide language

[Ragan-Kelley 2012]

Simple language embedded in C++ for describing sequences of image processing operations (image processing pipelines)

```
Var x, y;  
Func blurx, blury, out;  
Image<uint8_t> in = load_image("myimage.jpg");
```

Functions map integer coordinates to values
(e.g., colors of corresponding pixels)

```
// perform 3x3 box blur in two-passes (box blur is separable)
```

```
blurx(x,y) = 1/3.f * (in(x-1,y) + in(x,y) + in(x,y));
```

```
blury(x,y) = 1/3.f * (blurx(x,y-1) + blurx(x,y+1) + blurx(x,y+1));
```

```
// brighten blurred result by 25%, then clamp
```

```
out(x,y) = min(blury(x,y) * 1.25f, 255);
```

Value of blurx at coordinate (x,y)
is given by expression accessing
three values of in

```
// execute pipeline on domain of size 800x600
```

```
Image<uint8_t> result = out.realize(800, 600);
```

- Halide function: an infinite (but discrete) set of values
- Halide expression: a side-effect free expression describes how to compute a function's value at a point in its domain in terms of the values of other functions.

Halide language

Update definitions modify function values

Reduction domains provide the ability to iterate

```
Var x;  
Func histogram, modified;  
Image<uint8_t> in = load_image("myimage.jpg");  
  
modified(x,y) = in(x,y) + 10;  
modified(x,3) *= 2; // update definition, modifies 3rd row  
modified(3,y) *= 2; // update definition, modifies 3rd column  
  
// clear all bins of the histogram to 0  
histogram(x) = 0;  
  
// declare "reduction domain" to be size of input image  
RDom r(0, in.width(), 0, in.height());  
  
// update definition on histogram  
// for all points in domain, increment appropriate bin  
histogram(in(r.x, r.y)) += 1;  
  
Image<int> result = histogram.realize(256);
```

Key aspects of Halide's design

- Adopts local “pointwise” view of expressing algorithms
- Language is highly constrained so that iteration over domain points is implicit (no explicit loops in Halide)
 - Halide language is declarative. It does not define order of iteration, or what values in domain are stored!
 - It only defines what operations are needed to compute these values.

```
Var x, y;  
Func blurx, out;  
Image<uint8_t> in = load_image("myimage.jpg");  
  
// perform 3x3 box blur in two-passes (box blur is separable)  
blurx(x,y) = 1/3.f * (in(x-1,y) + in(x,y) + in(x,y+1));  
out(x,y) = 1/3.f * (blurx(x,y-1) + blurx(x,y+1) + blurx(x,y+1));  
  
// execute pipeline on domain of size 800x600  
Image<uint8_t> result = our.realize(800, 600);
```

Efficiently executing Halide programs

Example

Consider writing code for the two-pass 3x3 image blur

```
Var x, y;  
Func blurx, out;  
Image<uint8_t> in = load_image("myimage.jpg");  
  
// perform 3x3 box blur in two-passes (box blur is separable)  
blurx(x,y) = 1/3.f * (in(x-1,y) + in(x,y) + in(x,y));  
out(x,y) = 1/3.f * (blurx(x,y-1) + blurx(x,y+1) + blurx(x,y+1));  
  
// execute pipeline on domain of size 1024x1024  
Image<uint8_t> result = out.realize(1024, 1024);
```

Two-pass 3x3 blur

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/3, 1.0/3, 1.0/3};

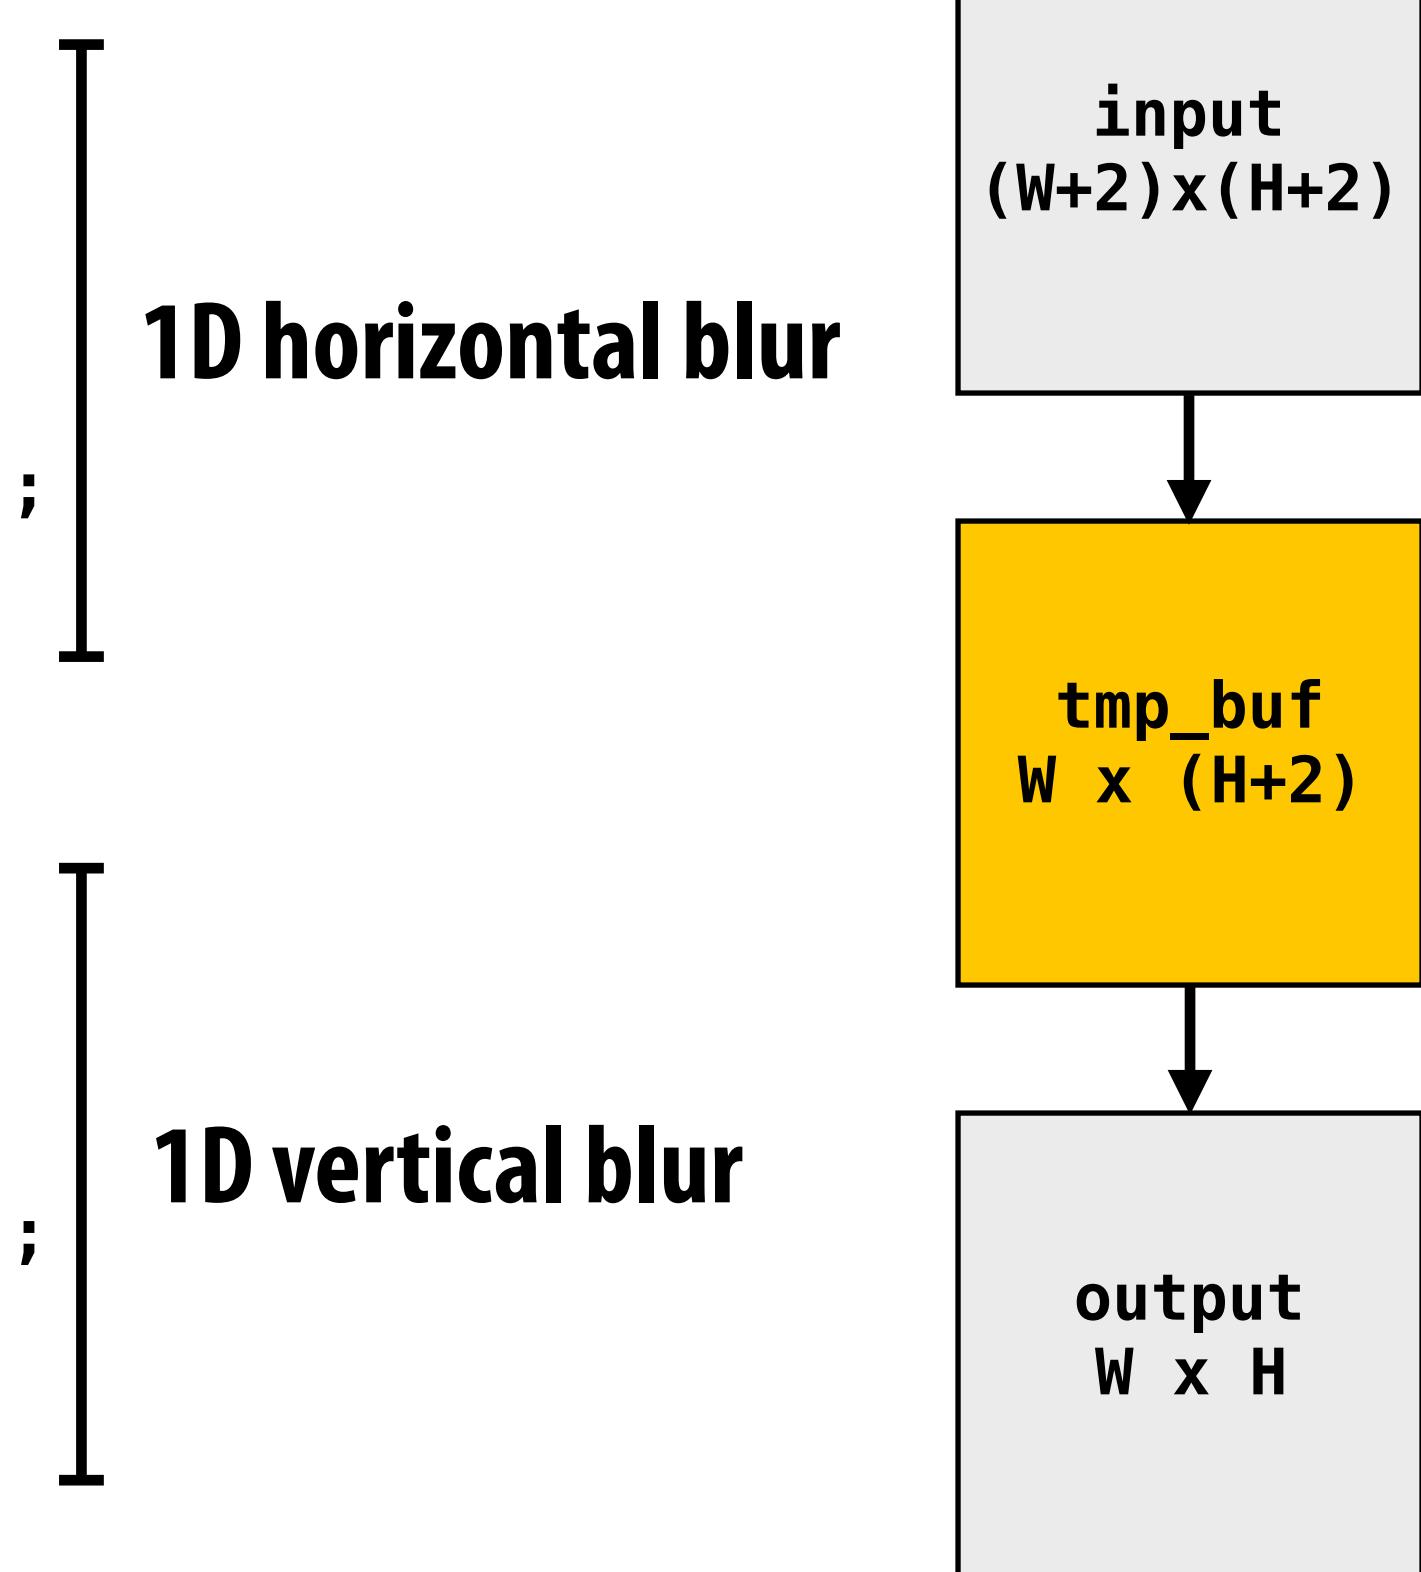
for (int j=0; j<(HEIGHT+2); j++)
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int ii=0; ii<3; ii++)
            tmp += input[j*(WIDTH+2) + i+ii] * weights[ii];
        tmp_buf[j*WIDTH + i] = tmp;
    }

for (int j=0; j<HEIGHT; j++) {
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int jj=0; jj<3; jj++)
            tmp += tmp_buf[(j+jj)*WIDTH + i] * weights[jj];
        output[j*WIDTH + i] = tmp;
    }
}
```

Total work per image = $6 \times \text{WIDTH} \times \text{HEIGHT}$

For NxN filter: $2N \times \text{WIDTH} \times \text{HEIGHT}$

$\text{WIDTH} \times \text{HEIGHT}$ extra storage



Two-pass image blur: locality

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/3, 1.0/3, 1.0/3};

for (int j=0; j<(HEIGHT+2); j++)
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int ii=0; ii<3; ii++)
            tmp += input[j*(WIDTH+2) + i+ii] * weights[ii];
        tmp_buf[j*WIDTH + i] = tmp;
    }

for (int j=0; j<HEIGHT; j++) {
    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int jj=0; jj<3; jj++)
            tmp += tmp_buf[(j+jj)*WIDTH + i] * weights[jj];
        output[j*WIDTH + i] = tmp;
    }
}
```

Intrinsic bandwidth requirements of algorithm:
Application must read each element of input image and must write each element of output image.

Data from **input** reused three times. (immediately reused in next two i-loop iterations after first load, never loaded again.)

- Perfect cache behavior: never load required data more than once
- Perfect use of cache lines (don't load unnecessary data into cache)

Two pass: loads/stores to **tmp_buf** are overhead (this memory traffic is an artifact of the two-pass implementation: it is not intrinsic to computation being performed)

Data from **tmp_buf** reused three times (but three rows of image data are accessed in between)

- Never load required data more than once... if cache has capacity for three rows of image
- Perfect use of cache lines (don't load unnecessary data into cache)

Two-pass image blur, “chunked” (version 1)

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * 3]; ←
float output[WIDTH * HEIGHT];

float weights[] = {1.0/3, 1.0/3, 1.0/3};

for (int j=0; j<HEIGHT; j++) {

    for (int j2=0; j2<3; j2++)
        for (int i=0; i<WIDTH; i++) {
            float tmp = 0.f;
            for (int ii=0; ii<3; ii++)
                tmp += input[(j+j2)*(WIDTH+2) + i+ii] * weights[ii];
            tmp_buf[j2*WIDTH + i] = tmp;
        }

    for (int i=0; i<WIDTH; i++) {
        float tmp = 0.f;
        for (int jj=0; jj<3; jj++)
            tmp += tmp_buf[jj*WIDTH + i] * weights[jj];
        output[j*WIDTH + i] = tmp;
    }
}
```

Only 3 rows of intermediate buffer need to be allocated

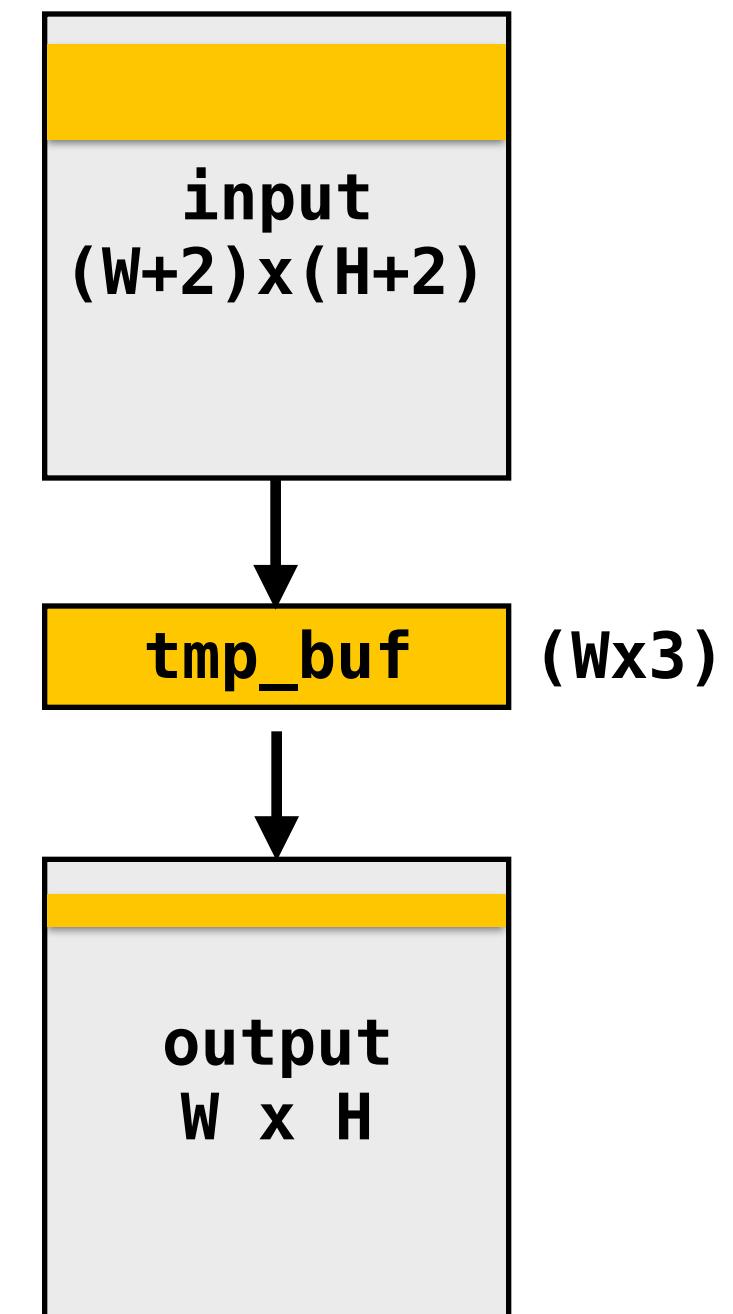
Produce 3 rows of tmp_buf (only what's needed for one row of output)

Combine them together to get one row of output

Total work per row of output:

- step 1: $3 \times 3 \times \text{WIDTH}$ work
- step 2: $3 \times \text{WIDTH}$ work

Total work per image = $12 \times \text{WIDTH} \times \text{HEIGHT}$????



Loads from tmp_buffer are cached (assuming tmp_buffer fits in cache)

Two-pass image blur, “chunked” (version 2)

```

int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * (CHUNK_SIZE+2)]; ←
float output[WIDTH * HEIGHT];

float weights[] = {1.0/3, 1.0/3, 1.0/3};

for (int j=0; j<HEIGHT; j+CHUNK_SIZE) {

    for (int j2=0; j2<CHUNK_SIZE+2; j2++) ←
        for (int i=0; i<WIDTH; i++) {
            float tmp = 0.f;
            for (int ii=0; ii<3; ii++)
                tmp += input[(j+j2)*(WIDTH+2) + i+ii] * weights[ii];
            tmp_buf[j2*WIDTH + i] = tmp;

    for (int j2=0; j2<CHUNK_SIZE; j2++) ← Produce CHUNK_SIZE rows of output
        for (int i=0; i<WIDTH; i++) {
            float tmp = 0.f;
            for (int jj=0; jj<3; jj++)
                tmp += tmp_buf[(j2+jj)*WIDTH + i] * weights[jj];
            output[(j+j2)*WIDTH + i] = tmp;
        }
    }
}

```

Sized so entire buffer fits in cache (capture all producer-consumer locality)

Produce enough rows of tmp_buf to produce a **CHUNK_SIZE number of rows of output**

Produce **CHUNK_SIZE rows of output**

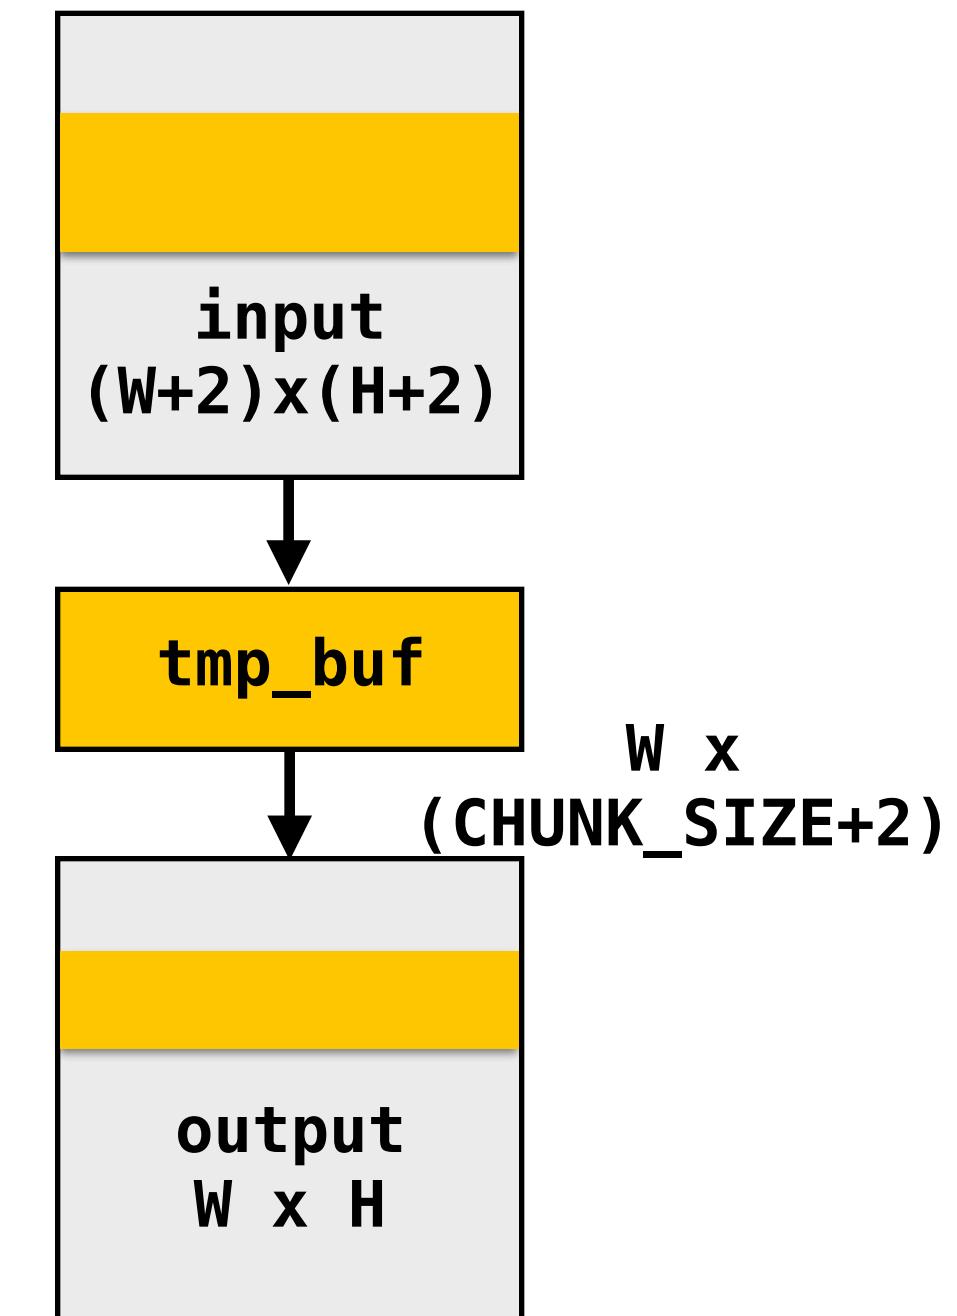
Total work per chuck of output:
(assume **CHUNK_SIZE = 16)**

- Step 1: $18 \times 3 \times \text{WIDTH}$ work
- Step 2: $16 \times 3 \times \text{WIDTH}$ work

Total work per image: $(34/16) \times 3 \times \text{WIDTH} \times \text{HEIGHT}$

→ = $6.4 \times \text{WIDTH} \times \text{HEIGHT}$

Trends to idea $6 \times \text{WIDTH} \times \text{HEIGHT}$ as **CHUNK_SIZE is increased!**



Still not done

- We have not parallelized loops for multi-core execution
- We have not used SIMD instructions to execute loops bodies
- Other basic optimizations: loop unrolling, etc...

Optimized x86 implementation

Good: ~10x faster on a quad-core CPU than my original two-pass code

Bad: specific to SSE (not AVX2), CPU-code only, hard to tell what is going on at all!

```
void fast_blur(const Image &in, Image &blurred) {
    _m128i one_third = _mm_set1_epi16(21846);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        _m128i a, b, c, sum, avg;
        _m128i tmp[(256/8)*(32+2)];
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            _m128i *tmpPtr = tmp;
            for (int y = -1; y < 32+1; y++) {
                const uint16_t *inPtr = &(in(xTile, yTile+y));
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_loadu_si128((__m128i*) (inPtr-1));
                    b = _mm_loadu_si128((__m128i*) (inPtr+1));
                    c = _mm_load_si128((__m128i*) (inPtr));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulhi_epi16(sum, one_third);
                    _mm_store_si128(tmpPtr++, avg);
                    inPtr += 8;
                }
                tmpPtr = tmp;
                for (int y = 0; y < 32; y++) {
                    _m128i *outPtr = (__m128i *) (&(blurred(xTile, yTile+y)));
                    for (int x = 0; x < 256; x += 8) {
                        a = _mm_load_si128(tmpPtr+(2*256)/8);
                        b = _mm_load_si128(tmpPtr+256/8);
                        c = _mm_load_si128(tmpPtr++);
                        sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                        avg = _mm_mulhi_epi16(sum, one_third);
                        _mm_store_si128(outPtr++, avg);
                    }
                }
            }
        }
    }
}
```

Multi-core execution (partition image vertically)

Modified iteration order: 256x32 tiled iteration (to maximize cache hit rate)

use of SIMD vector intrinsics

two passes fused into one: tmp data read from cache

Image processing pipelines feature complex sequences of functions

Benchmark	Number of Halide functions
Two-pass blur	2
Unsharp mask	9
Harris Corner detection	13
Camera RAW processing	30
Non-local means denoising	13
Max-brightness filter	9
Multi-scale interpolation	52
Local-laplacian filter	103
Synthetic depth-of-field	74
Bilateral filter	8
Histogram equalization	7
VGG-16 deep network eval	64

Real-world production applications may features hundreds to thousands of functions!
Google HDR+ pipeline: over 2000 Halide functions.

Key aspect in the design of any system: Choosing the “right” representations for the job

**Now the job is not expressing an image processing
computation, but generating an efficient
implementation of a specific Halide program.**

A second set of representations for “scheduling”

```
Func blurx, out;  
Var x, y, xi, yi;  
Image<uint8_t> in = load_image("myimage.jpg");  
  
// the “algorithm description” (declaration of what to do)  
blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;  
out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;
```

```
// “the schedule” (how to do it)  
out.tile(x, y, xi, yi, 256, 32).vectorize(xi,8).parallel(y);
```

~~blurx.compute_at(x).vectorize(x, 8);~~

Produce elements **blurx** on demand for each tile of output.

Vectorize the x (innermost) loop

When evaluating **out**, use 2D tiling order (loops named by x, y, xi, yi). Use tile size 256 x 32.

Vectorize the xi loop (8-wide)

Use threads to parallelize the y loop

```
// execute pipeline on domain of size 1024x1024  
Image<uint8_t> result = out.realize(1024, 1024);
```

Scheduling primitives allow the programmer to specify a global “sketch” of how to schedule the algorithm onto a parallel machine, but leave the details of emitting the low-level platform-specific code to the Halide compiler

Primitives for iterating over domains

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36

serial y, serial x

1	7	13	19	25	31
2	8	14	20	26	32
3	9	15	21	27	33
4	10	16	22	28	34
5	11	17	23	29	35
6	12	18	24	30	36

serial x, serial y

**Specify both order and how to parallelize
(multi-thread, vectorize via SIMD instr)**

1	2
3	4
5	6
7	8
9	10
11	12

serial y
vectorized x

1	2
1	2
1	2
1	2
1	2

parallel y
vectorized x

1	2	5	6	9	10
3	4	7	8	11	12
13	14	17	18	21	22
15	16	19	20	23	24
25	26	29	30	33	34
27	28	31	32	35	36

split x into $2x_o + x_i$,
split y into $2y_o + y_i$,
serial y_o, x_o, y_i, x_i

2D blocked iteration order

Specifying loop iteration order and parallelism

```
blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;  
out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;
```

Given this schedule for the function “out”...

```
out.tile(x, y, xi, yi, 256, 32).vectorize(xi,8).parallel(y);
```

Halide compiler will generate this parallel, vectorized loop nest for computing elements of out...

```
for y=0 to num_tiles_y:           // parallelize this loop over multiple threads  
  for x=0 to num_tiles_x:  
    for yi=0 to 32:  
      for xi=0 to 256:           // vectorize this loop with SIMD instructions  
        idx_x = x*256+xi;  
        idx_y = y*32+yi  
        out(idx_y, idx_y) = ...
```

Primitives for how to interleave producer/consumer processing

```
blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;  
out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;  
  
out.tile(x, y, xi, yi, 256, 32);
```

blurx.compute_root();

Do not compute blurx within out's loop nest.
Compute all of blurx, then all of out

```
allocate buffer for all of blur(x,y)  
for y=0 to HEIGHT:  
  for x=0 to WIDTH:  
    blurx(x,y) = ...
```

all of blurx is computed here

```
for y=0 to num_tiles_y:  
  for x=0 to num_tiles_x:  
    for yi=0 to 32:  
      for xi=0 to 256:  
        idx_x = x*256+xi;  
        idx_y = y*32+yi  
        out(idx_y, idx_y) = ...
```

values of blurx consumed here

Primitives for how to interleave producer/consumer processing

```
blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y)) / 3.0f;  
out(x, y)   = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1)) / 3.0f;
```

```
out.tile(x, y, xi, yi, 256, 32);
```

```
blurx.compute_at(xi);
```

Compute necessary elements of blurx within
out's xi loop nest

```
for y=0 to num_tiles_y:  
  for x=0 to num_tiles_x:  
    for yi=0 to 32:  
      for xi=0 to 256:  
        idx_x = x*256+xi;  
        idx_y = y*32+yi
```

allocate 3-element buffer for blurx

```
// compute 3 elements of blurx needed for out(idx_x, idx_y) here  
out(idx_y, idx_y) = ...
```

Primitives for how to interleave producer/consumer processing

```
blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y)) / 3.0f;  
out(x, y)   = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1)) / 3.0f;
```

```
out.tile(x, y, xi, yi, 256, 32);
```

```
blurx.compute_at(x);
```

Compute necessary elements of blurx within out's x loop nest (all necessary elements for one tile of out)

```
for y=0 to num_tiles_y:  
  for x=0 to num_tiles_x:
```

```
    allocate 258x34 buffer for tile blurx
```

```
    for yi=0 to 32+2:  
      for xi=0 to 256+2:
```

```
        blur(xi,yi) = // compute blurx from in
```



tile of blurx is computed here

```
    for yi=0 to 32:  
      for xi=0 to 256:  
        idx_x = x*256+xi;  
        idx_y = y*32+yi  
        out(idx_y, idx_y) = ...
```



tile of blurx is consumed here

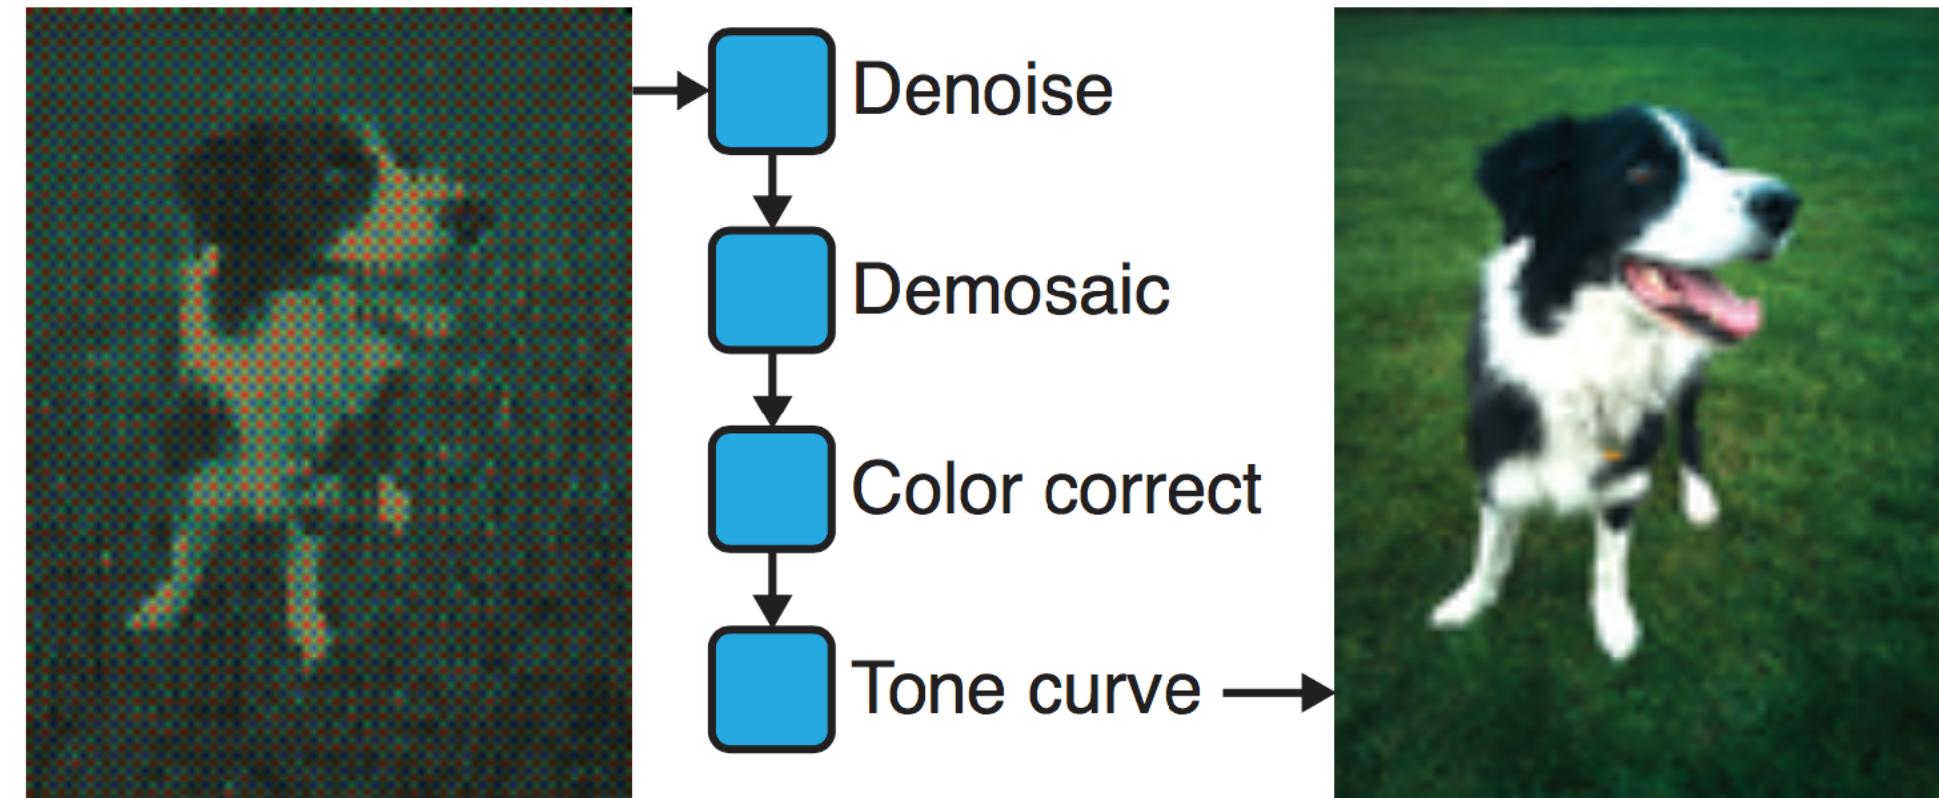
Halide: two domain-specific co-languages

- **Functional primitives for describing image processing operations**
- **Additional primitives for describing schedules**
- **Design principle: separate “algorithm specification” from schedule**
 - **Programmer’s responsibility: provide a high-performance schedule**
 - **Compiler’s responsibility: carry out mechanical process of generating threads, SIMD instructions, managing buffers, etc.**
 - **Result: enable programmer to rapidly exploration of space of schedules (“tile these loops”, “vectorize this loop”, “parallelize this loop across cores”)**
- **Application domain scope:**
 - **All computation on regular N-D coordinate spaces**
 - **Only feed-forward pipelines (includes special support for reductions and fixed recursion depth)**
 - **All dependencies inferable by compiler**

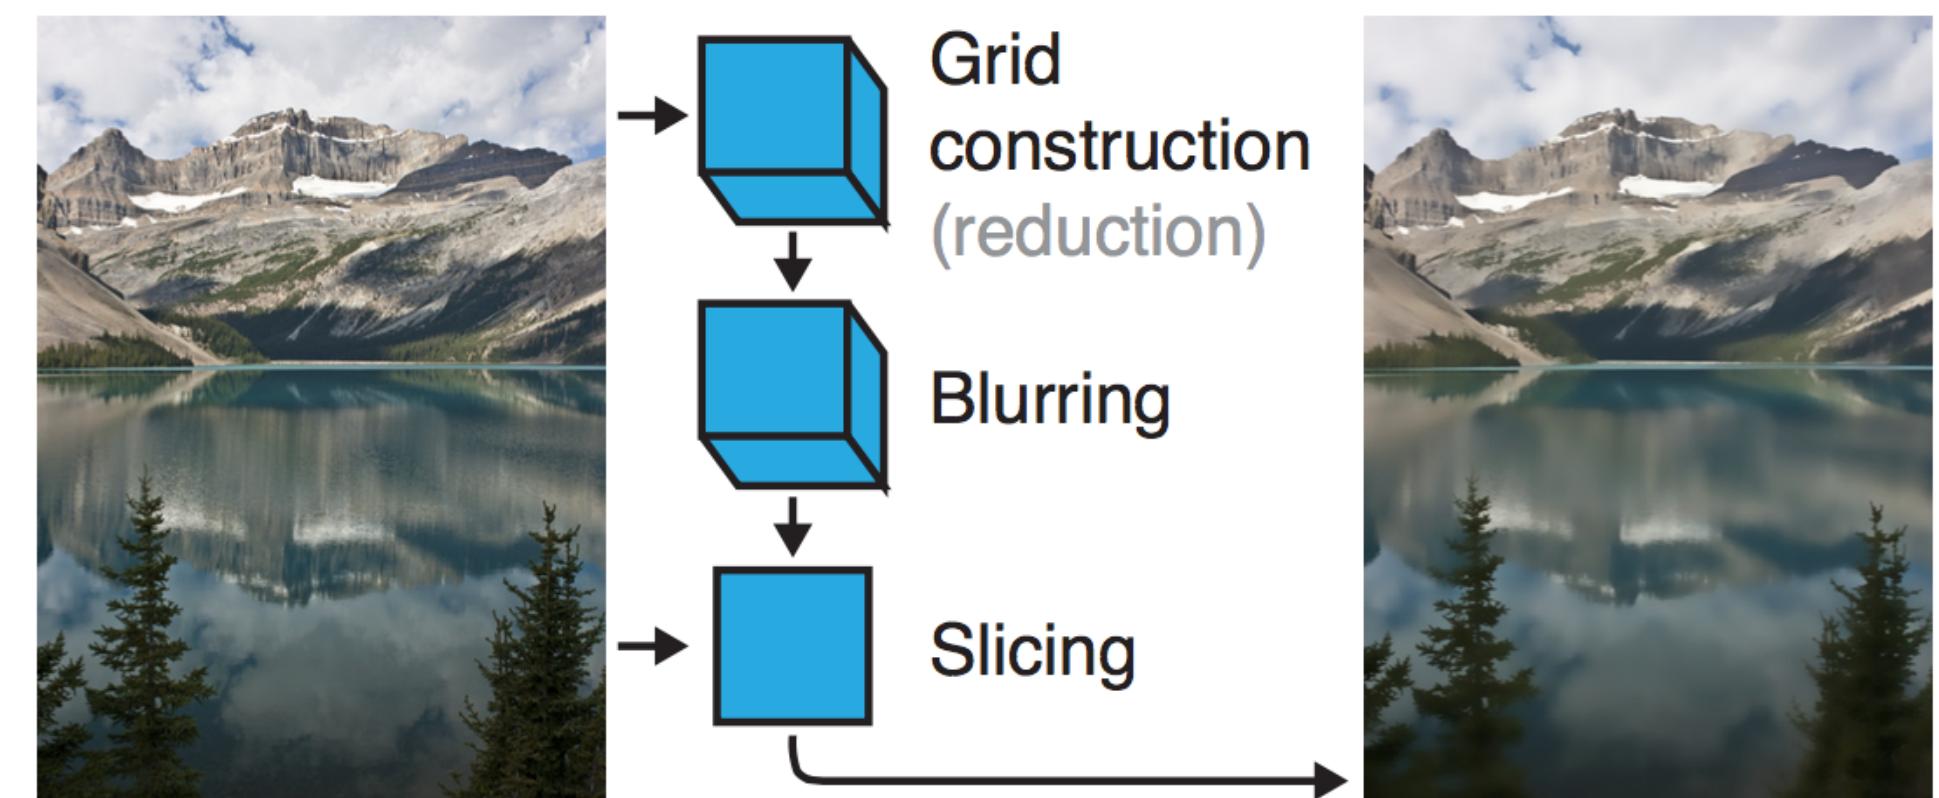
Initial Halide results

[Ragan-Kelley 2012]

- **Camera RAW processing pipeline**
(Convert RAW sensor data to RGB image)
 - **Original: 463 lines of hand-tuned ARM NEON assembly**
 - **Halide: 2.75x less code, 5% faster**

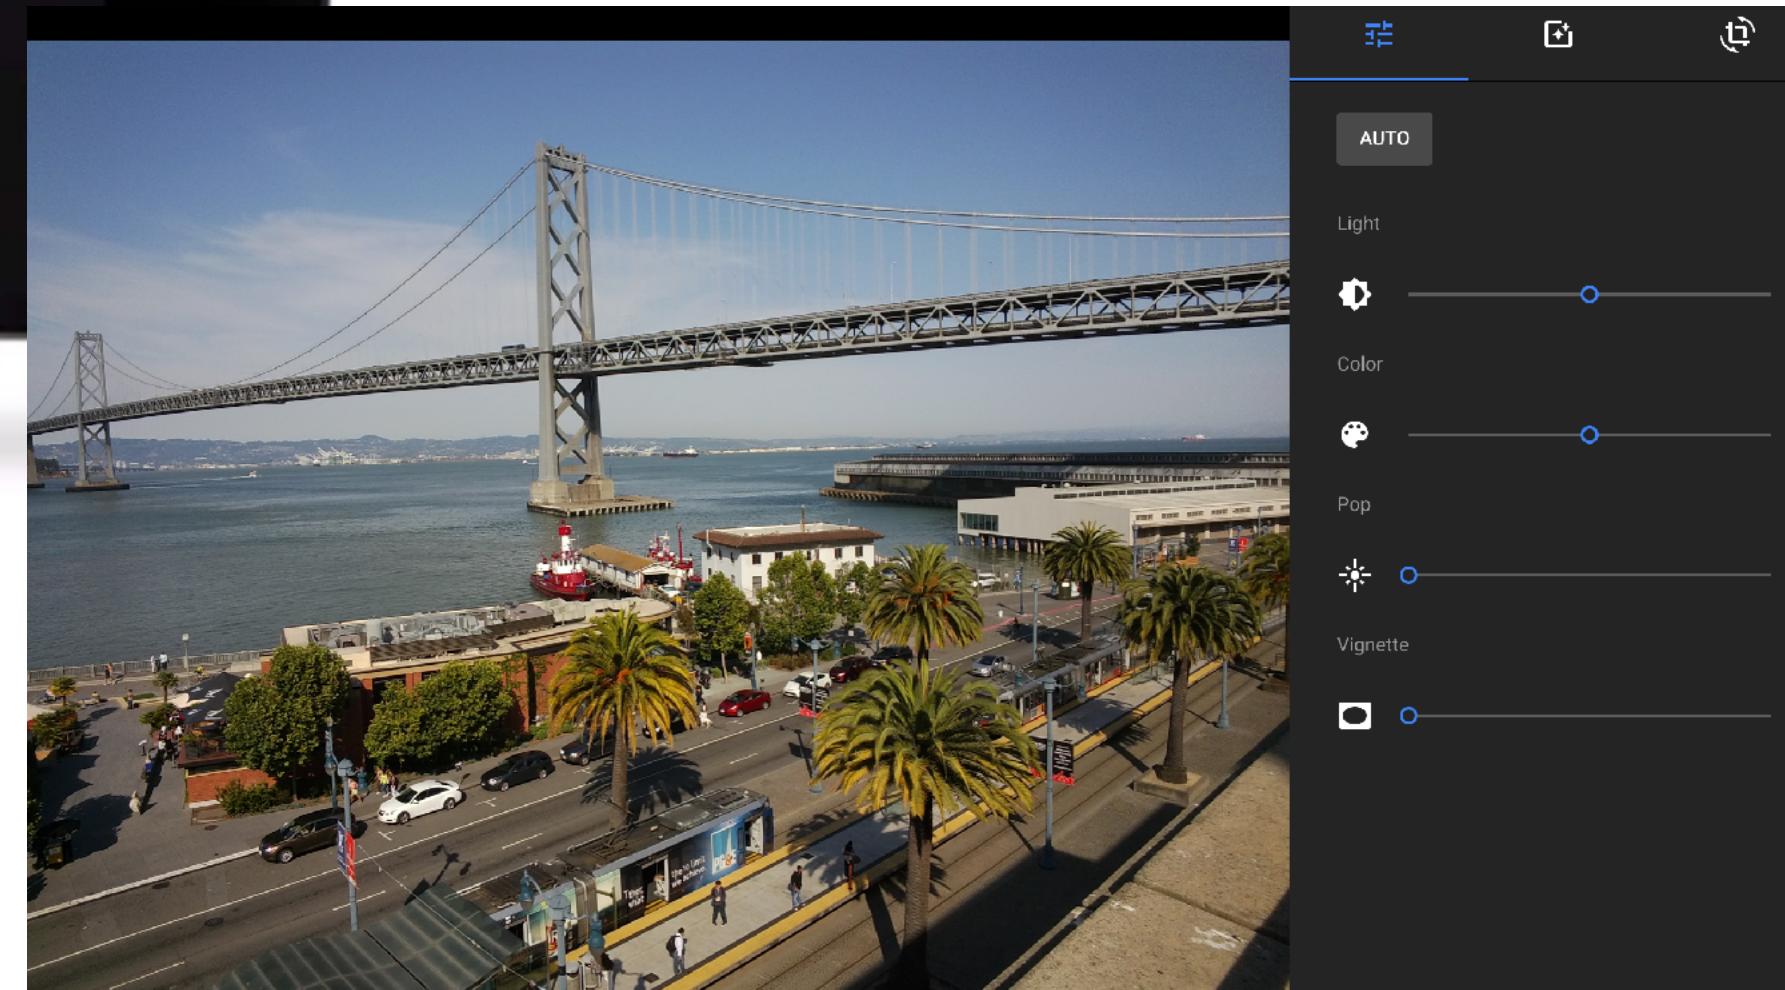


- **Bilateral filter**
(Common image filtering operation used in many applications)
 - **Original 122 lines of C++**
 - **Halide: 34 lines algorithm + 6 lines schedule**
 - **CPU implementation: 5.9x faster**
 - **GPU implementation: 2x faster than hand-written CUDA**



Halide used in practice

- Halide used to implement Google Pixel Photos app
- Halide code used to process images uploaded to Google Photos



Stepping back: what is Halide?

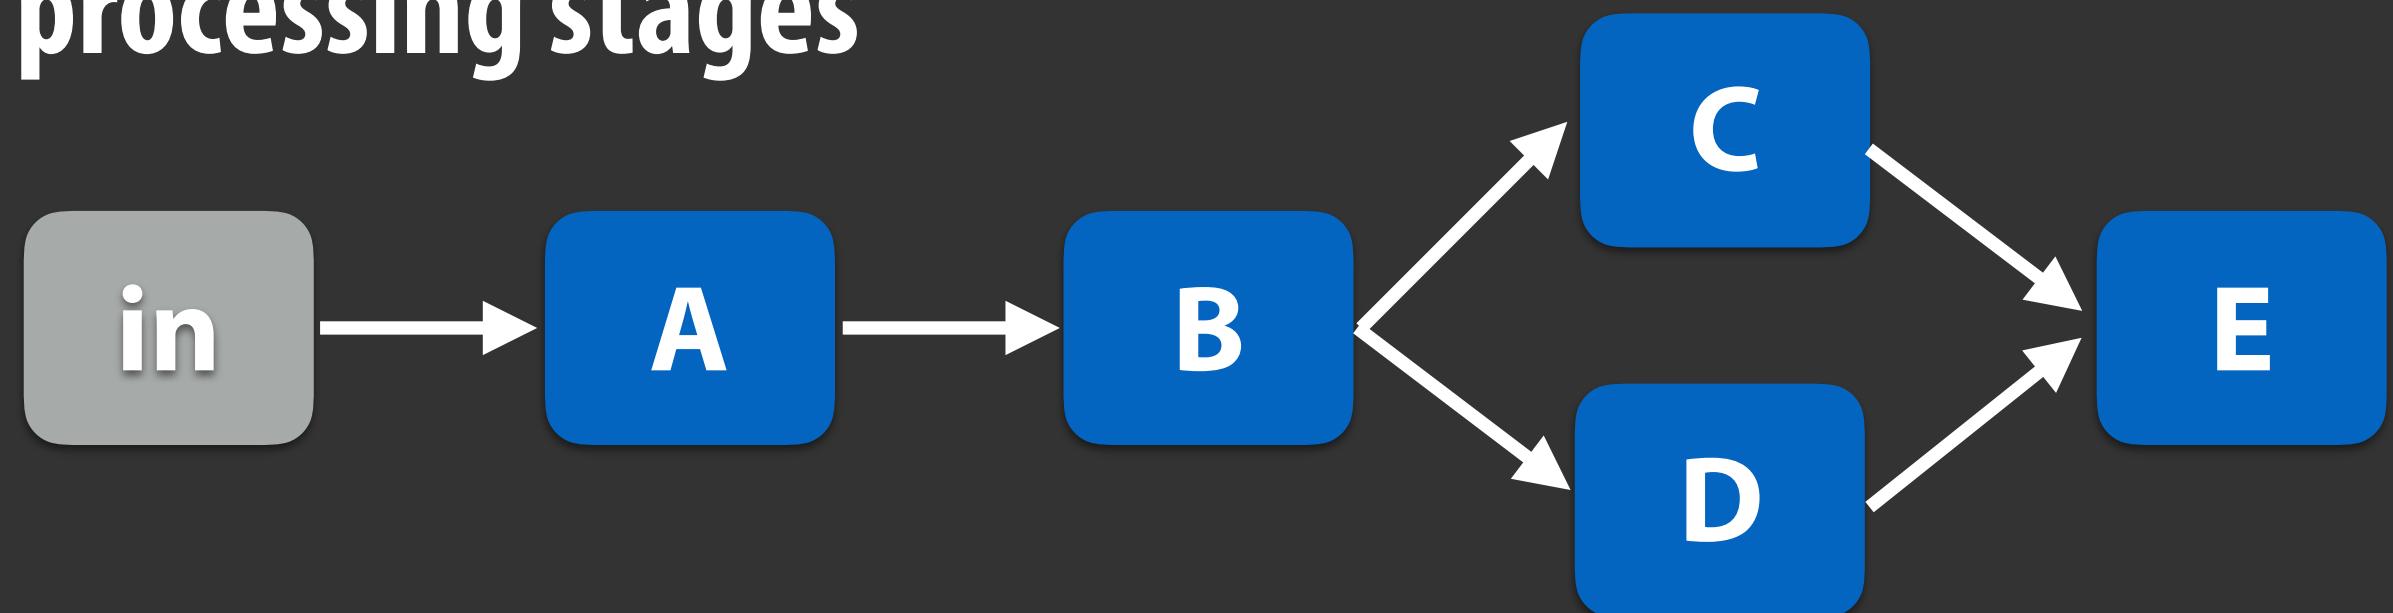
- **Halide is a DSL for helping expert developers optimize image processing code more rapidly**
 - **Halide does not decide how to optimize a program for a novice programmer**
 - **Halide provides primitives for a programmer (that has strong knowledge of code optimization, such as a 15-418 student) to rapidly express what optimizations the system should apply**
 - **Halide compiler carries out the nitty-gritty of mapping that strategy to a machine**

Automatically generating Halide schedules

- **Problem: it turned out that very few programmers have the ability to write good Halide schedules**
 - 80+ programmers at Google write Halide
 - Very small number trusted to write schedules
- **Recent work: compiler analyzes the Halide program to automatically generate efficient schedules for the programmer [Mullapudi 2016]**

Problem definition

Input: DAG of image processing stages



Output: optimized schedule

for each 8x128 tile in parallel
compute required pixels of A
compute pixels in tile of B

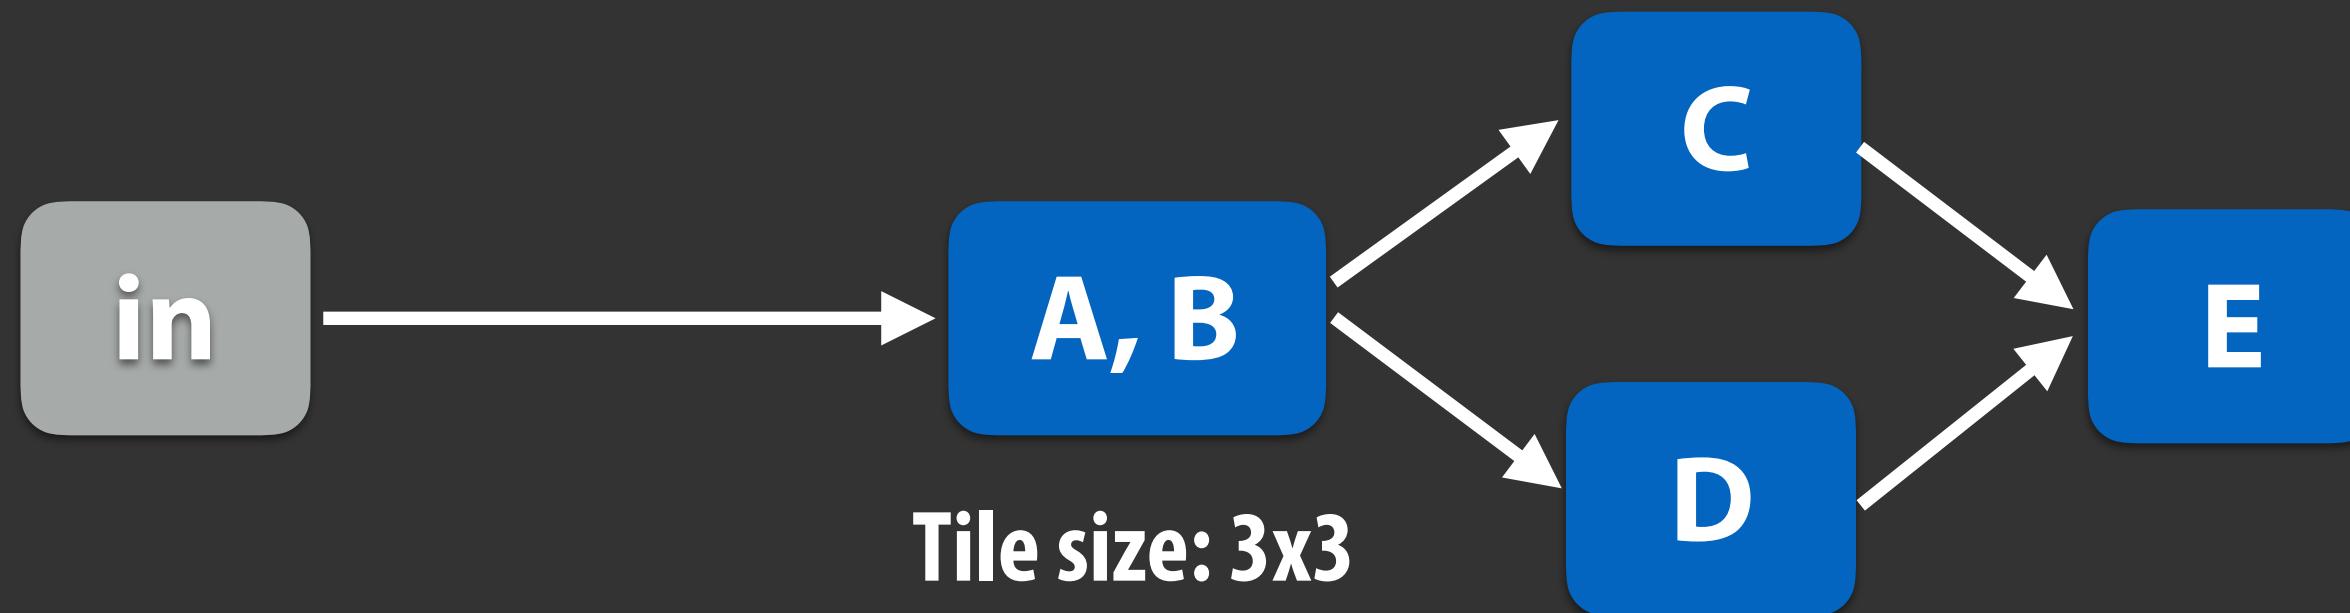
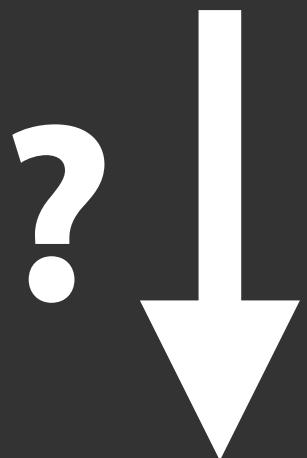
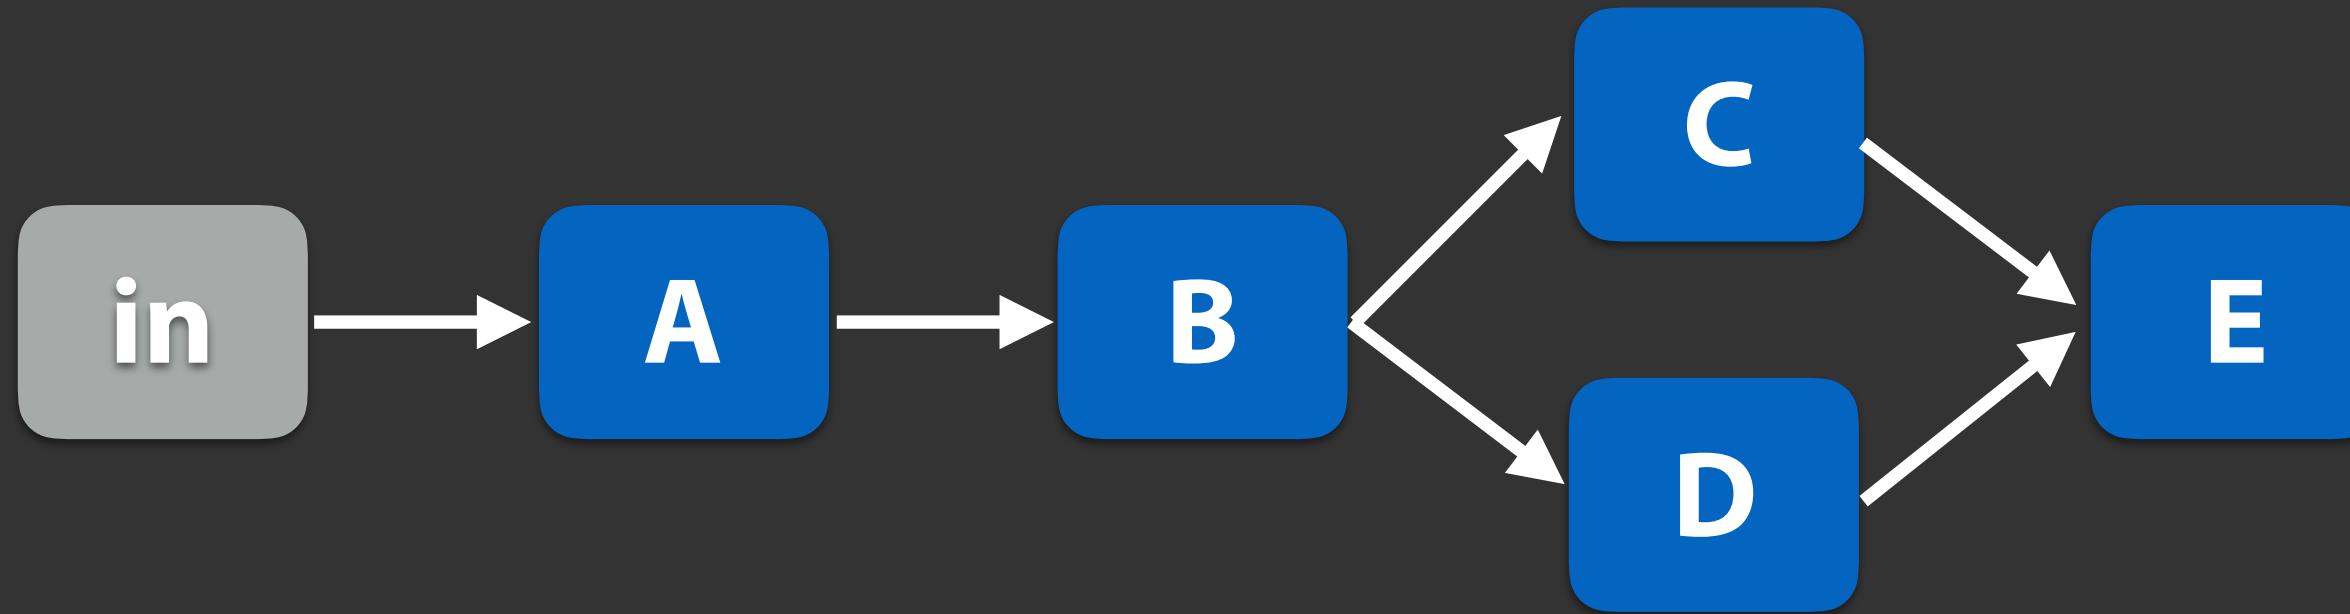
for each 8x8 tile in parallel
compute required pixels of C
compute required pixels of D
compute pixels in tile of E



Tile size: 8 x 128

Tile size: 8 x 8

How to tile a group



The autoscheduler fits programs to a single schedule template

For a single sub-DAG

```
for each tile_y: // multi-core parallel
  for each tile_x: // multi-core parallel
    allocate tmpC; // buffer for C for tile
    allocate tmpD; // buffer for D for tile
```

```
for each y in required region of C
  for each x in required region of C
    tmpC = .... // DRAM accesses to B
```

```
for each y in required region of D
  for each x in required region of D
    tmpD = .... // DRAM accessed to B
```

```
for each y in tile:
  for each x in tile: // optionally SIMD vector parallel
```

```
x' = tile_x * TILE_WIDTH + x;
y' = tile_y * TILE_HEIGHT + y;
E(x',y') = ...
```

Cost model to evaluate a tiling:

- cost of data access is proportional to buffer size
- cost arithmetic op = 1

Estimating costs via interval analysis

$[x-1, x+tx+1],$
 $[y-1, y+ty+1]$

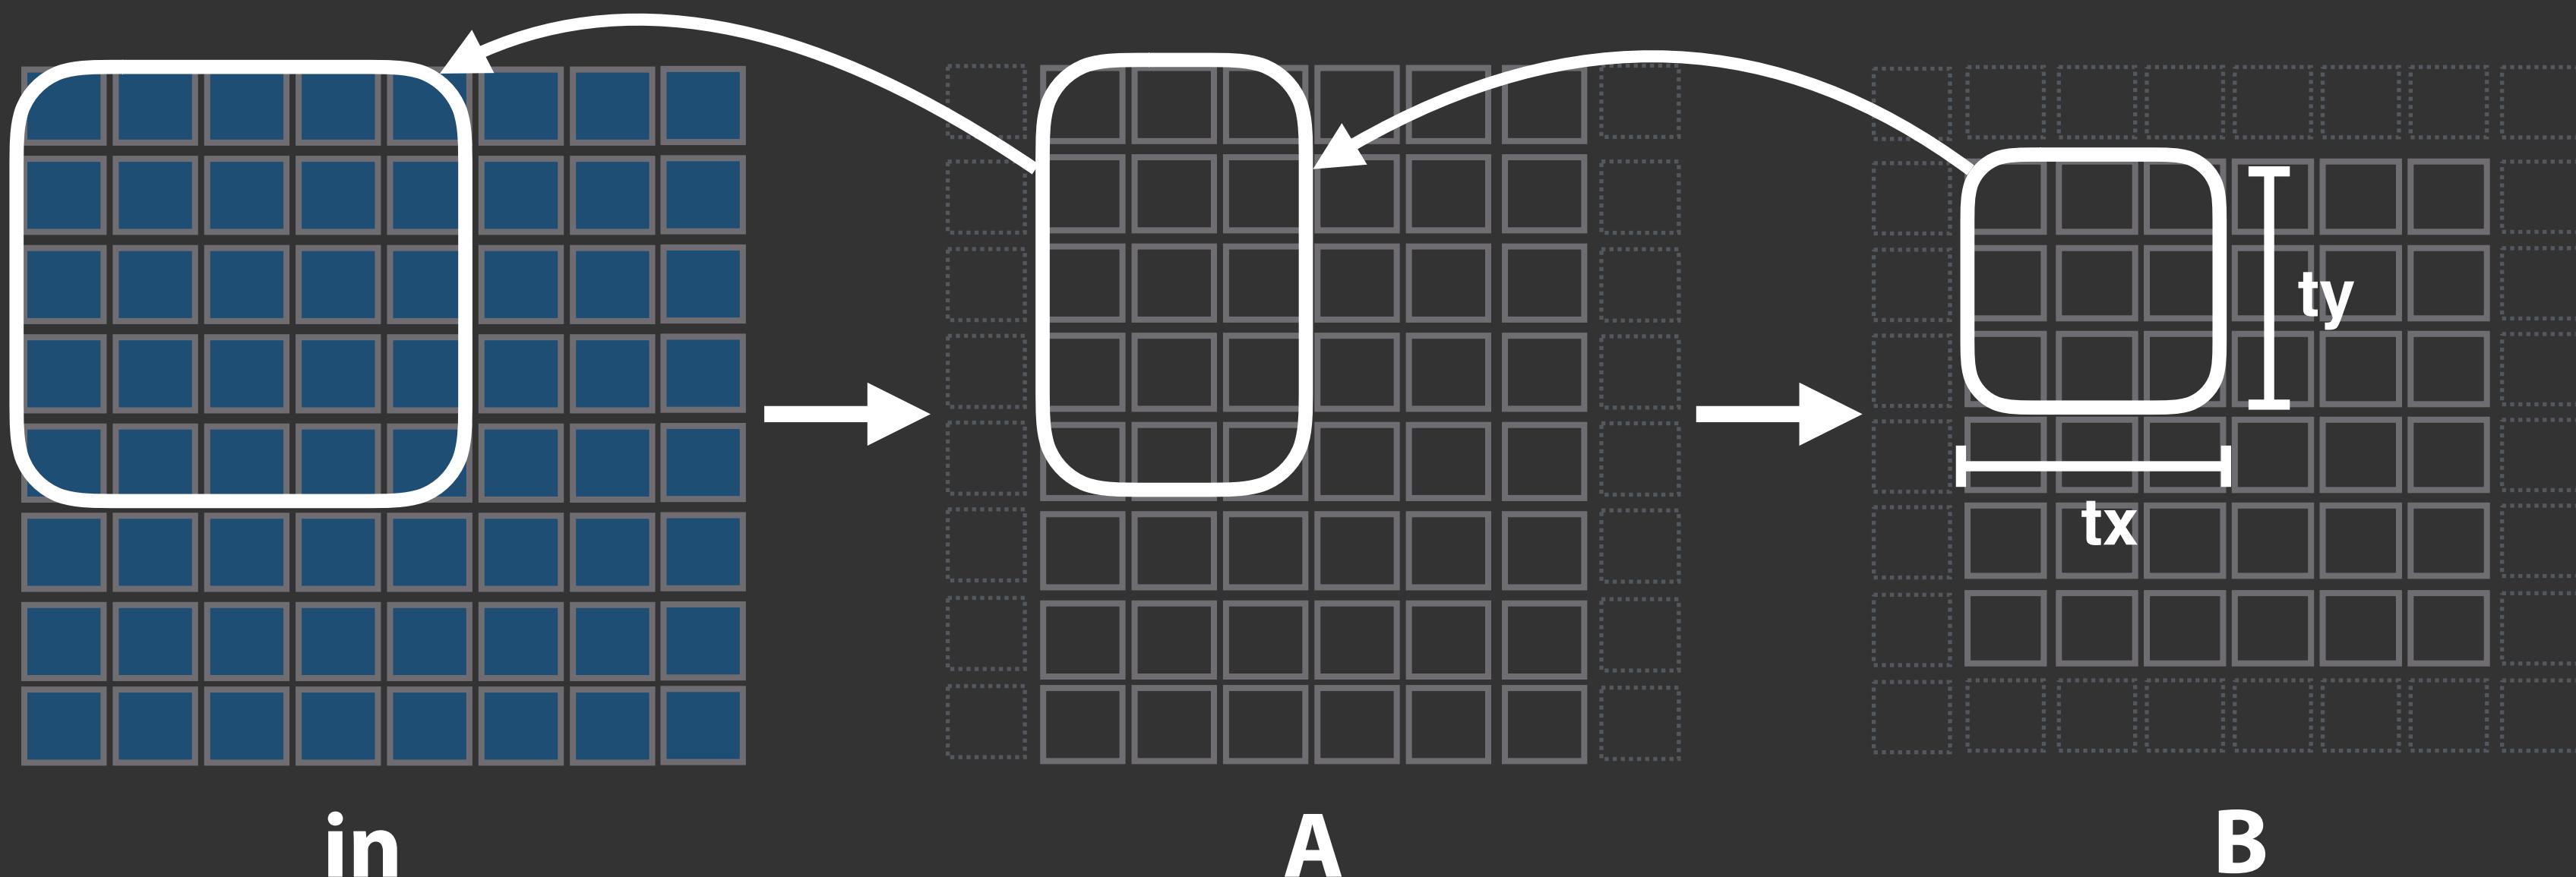
[0,5], [0,5]

$[x, x+tx],$
 $[y-1, y+ty+1]$

[1,4], [0,5]

$[x, x+tx],$
 $[y, y+ty]$

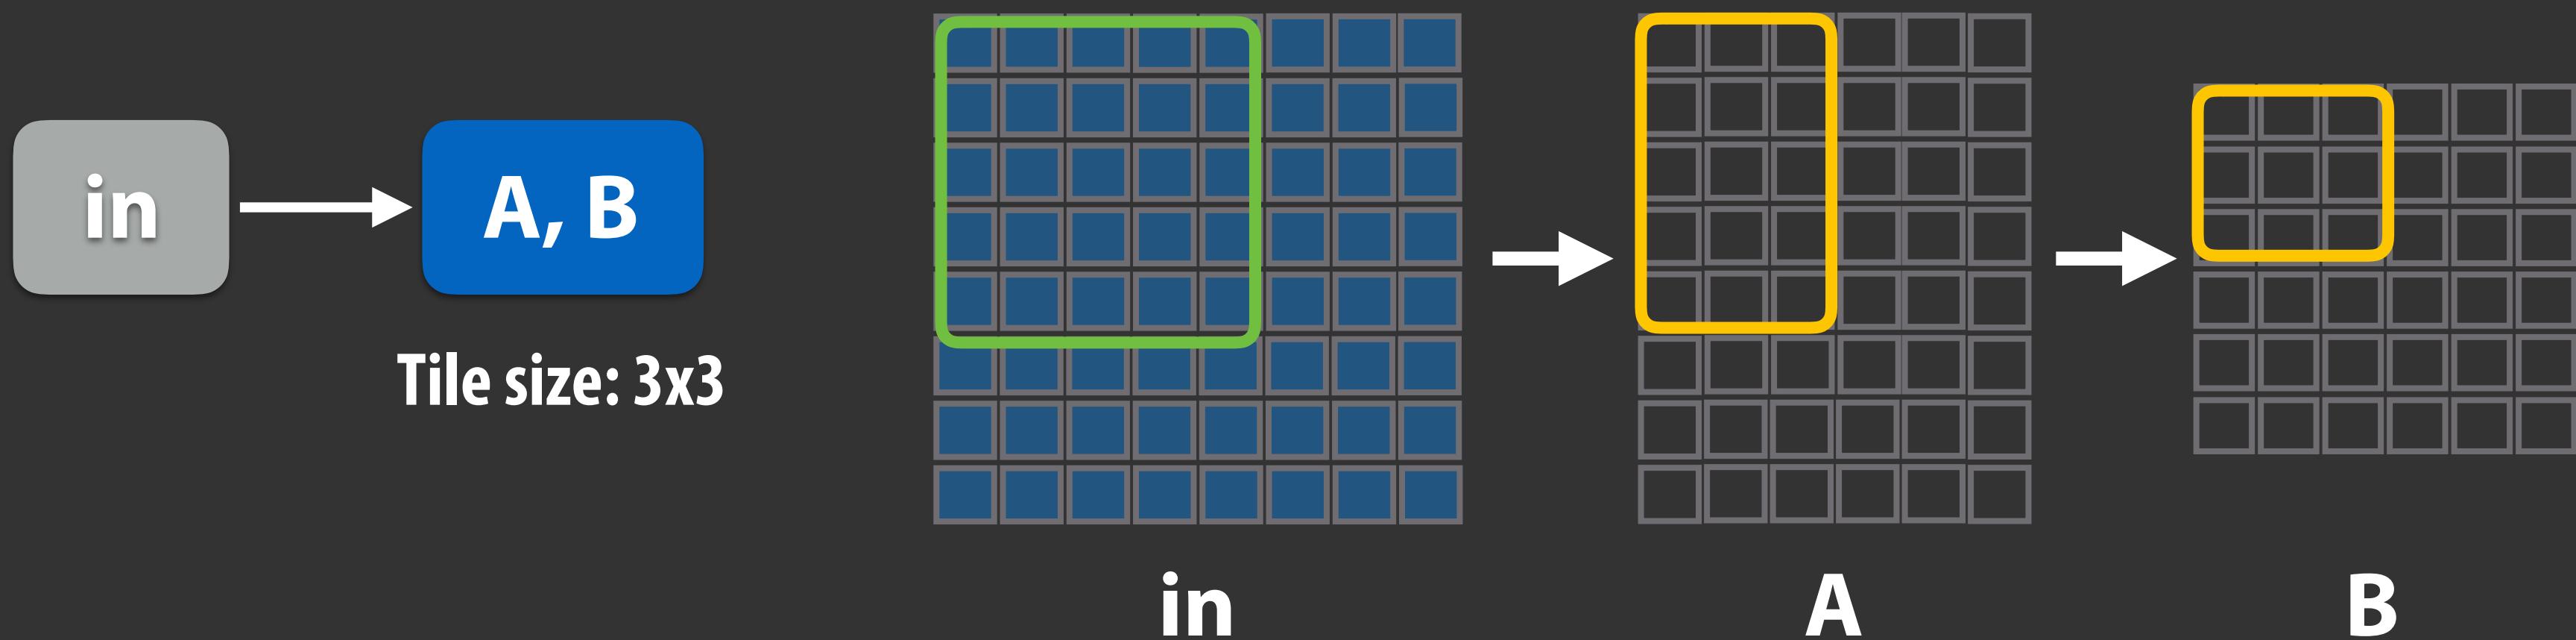
[1,4], [1,4]



$$A(x, y) = \text{in}(x-1, y) + \text{in}(x, y) + \text{in}(x+1, y)$$

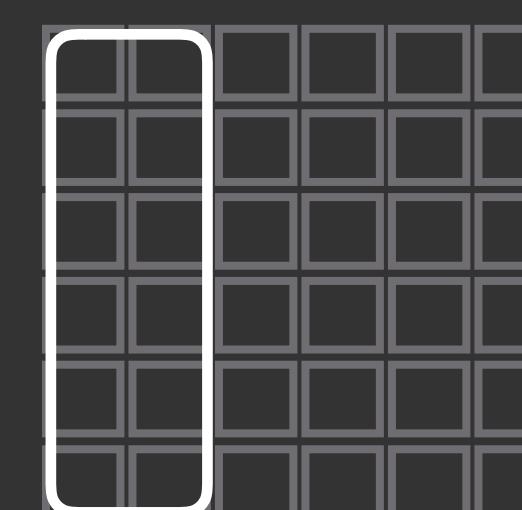
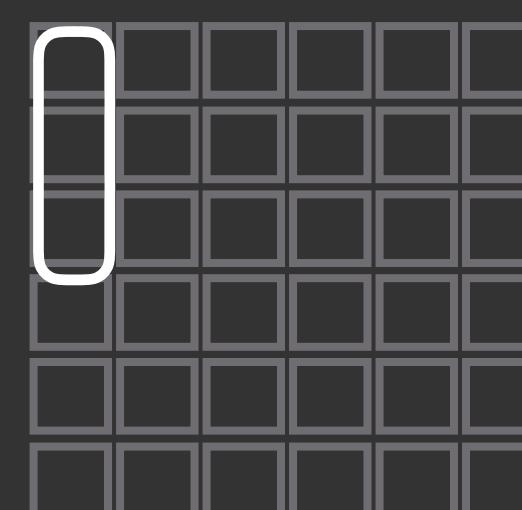
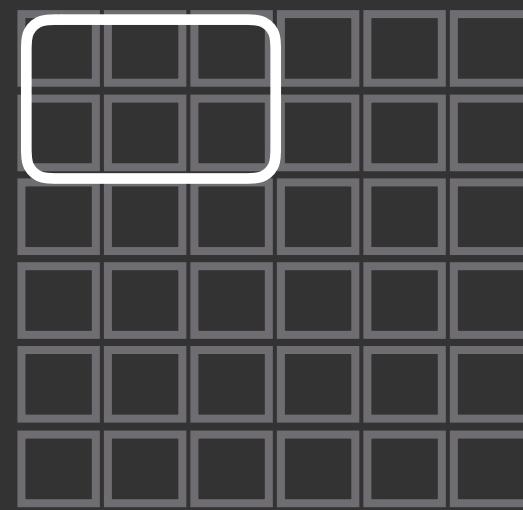
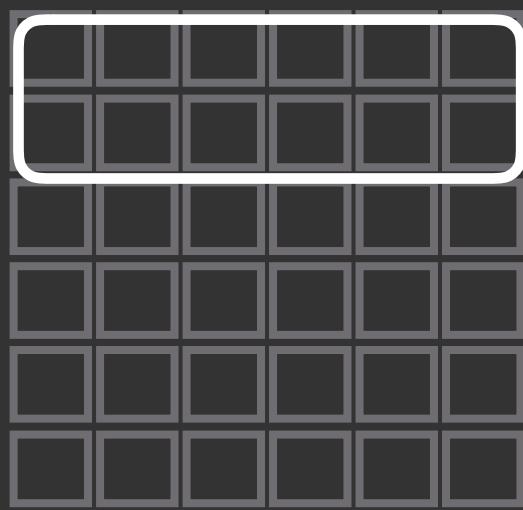
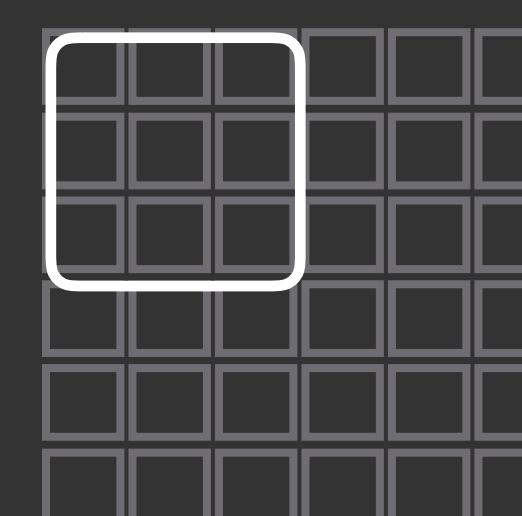
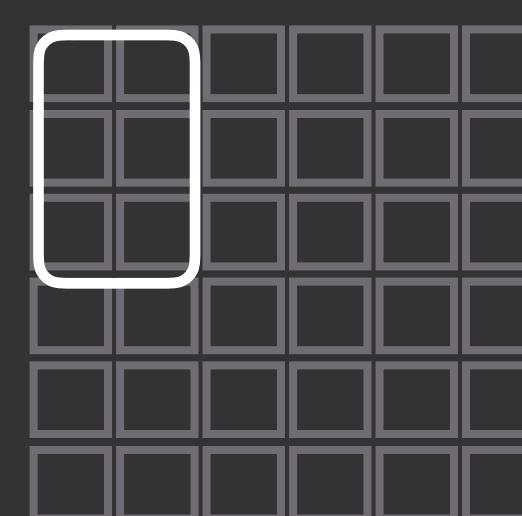
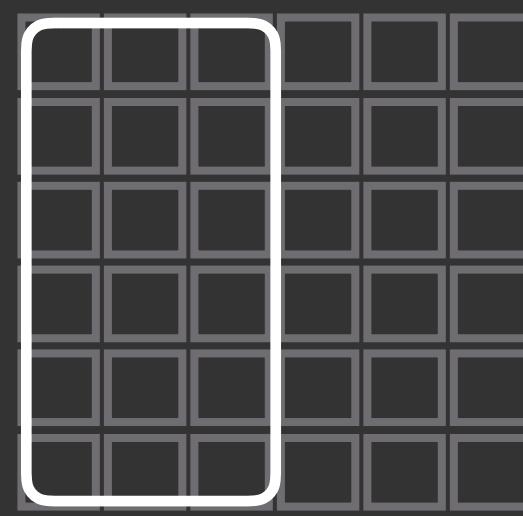
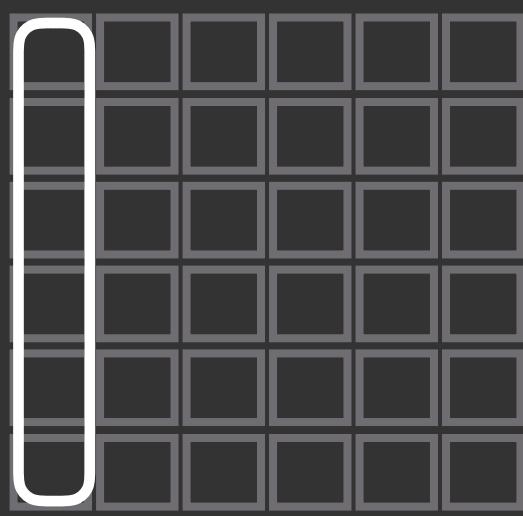
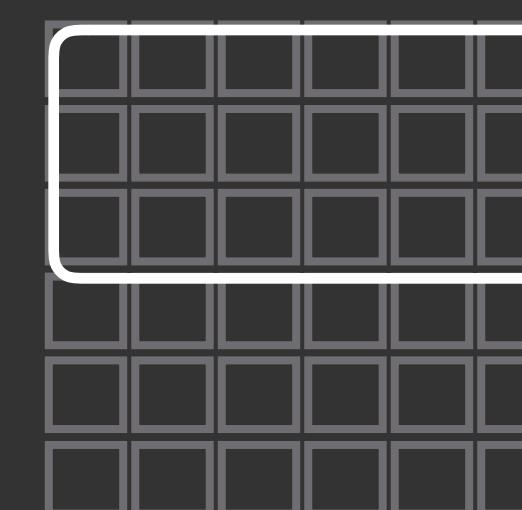
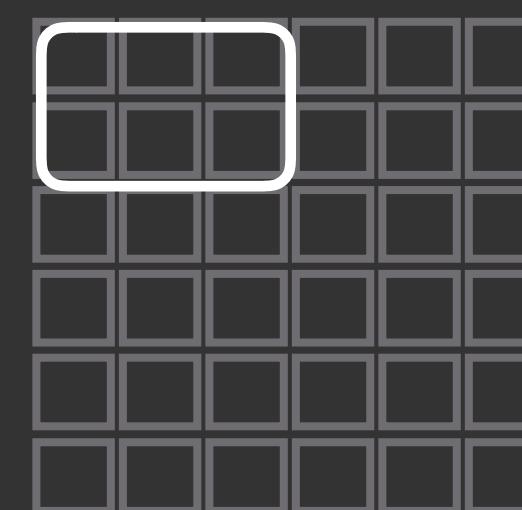
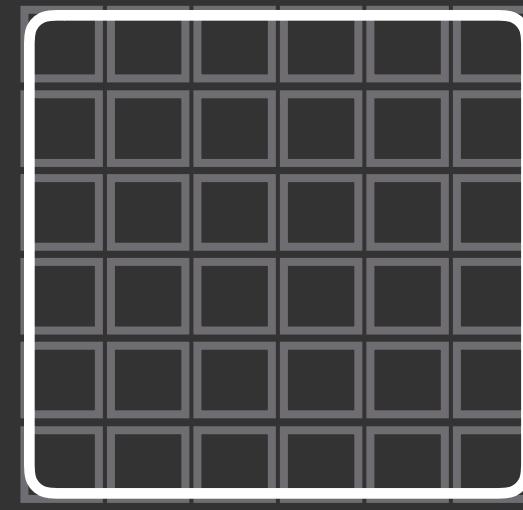
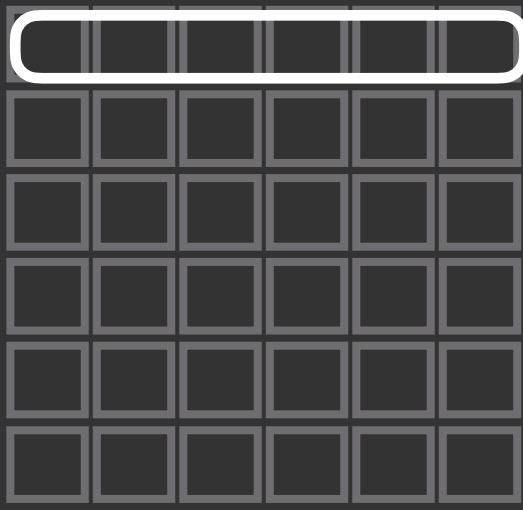
$$B(x, y) = A(x, y-1) + A(x, y) + A(x, y+1)$$

Estimating costs via interval analysis

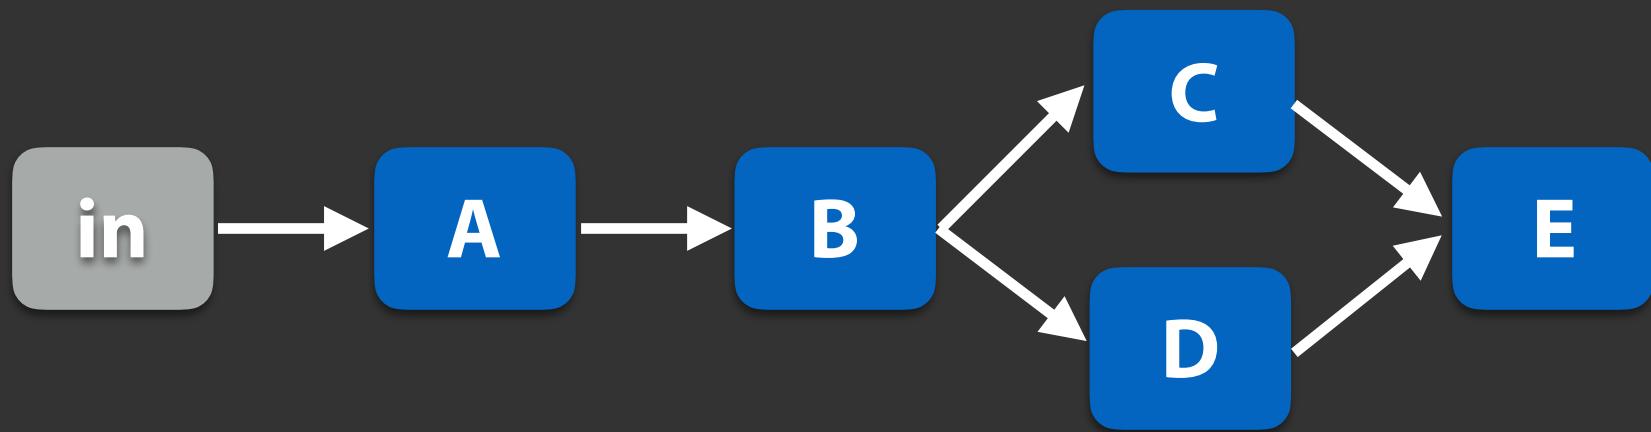


Cost = **Number of arithmetic operations** +
Number of memory accesses x **COST OF ELEMENT LOAD**

Search for most efficient tile size



Finding efficient groupings



Given the ability to compute a good tiling of groups...

Which stages should be grouped together?

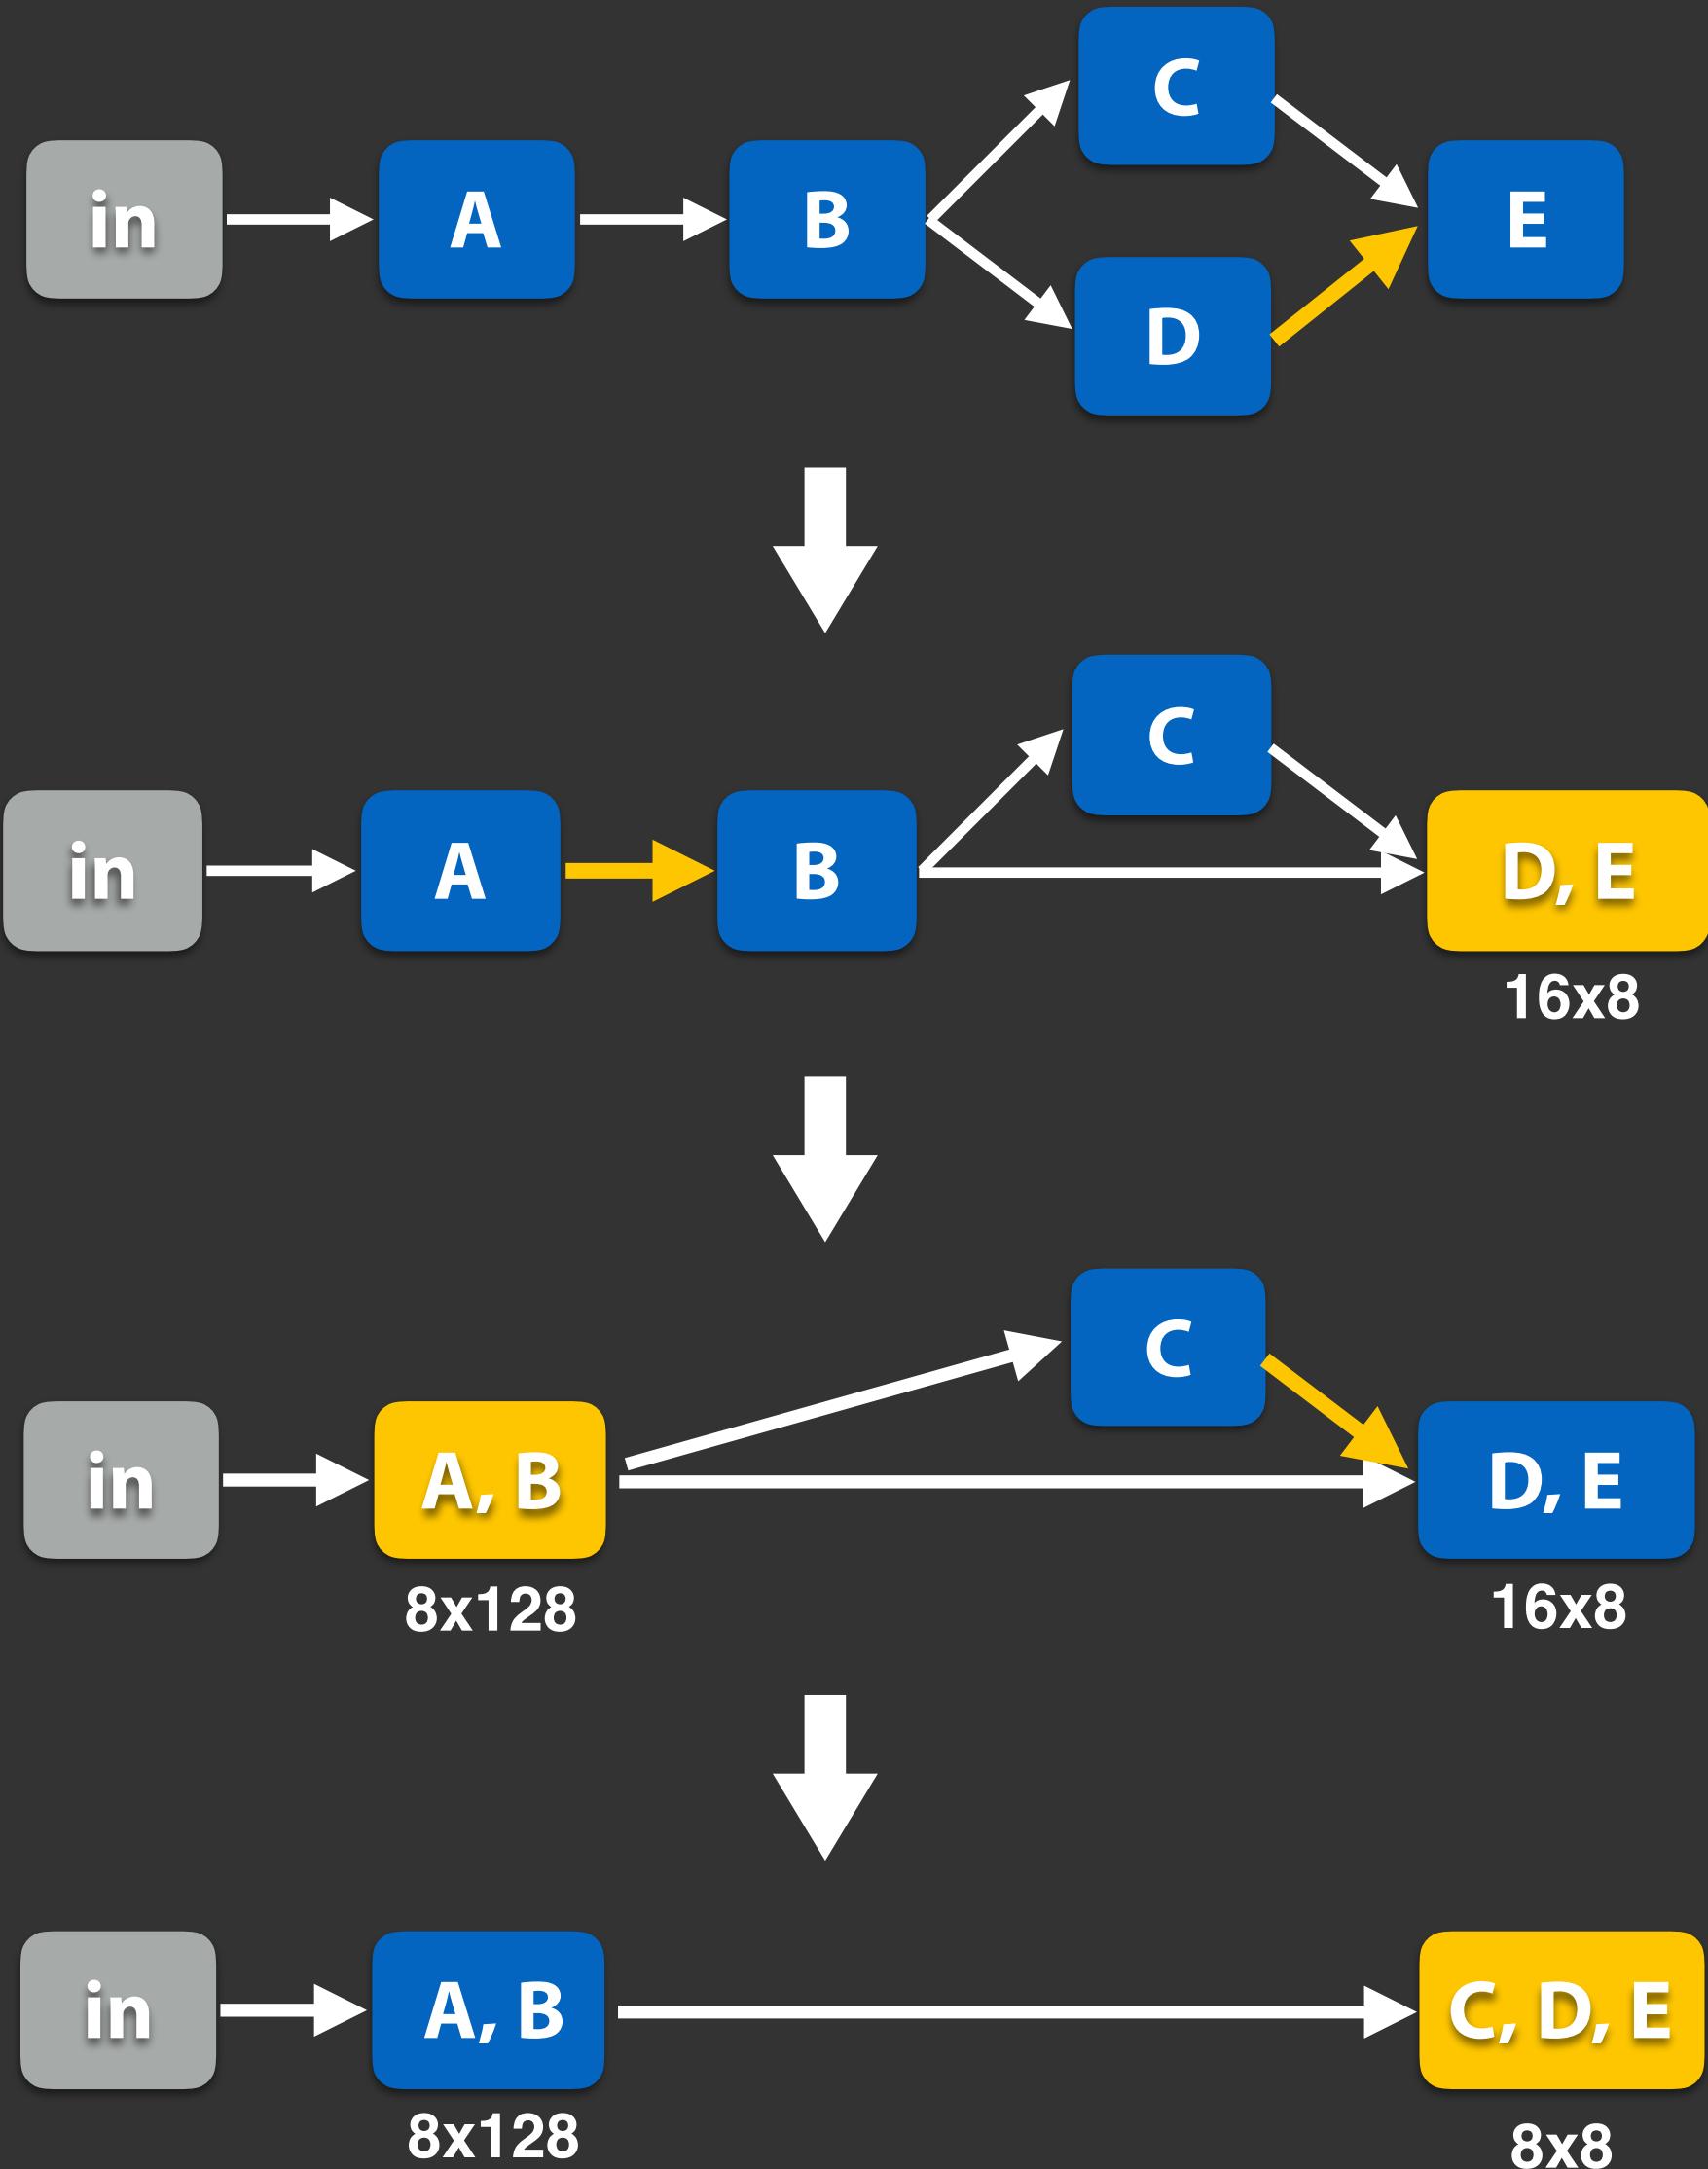
Greedy grouping strategy

**Greedily form groups
(take best merge)**

**But re-evaluate local
optimization to determine
tile size each step.**

for each 8x128 tile
compute required pixels of A
compute required pixels of B
compute pixels in tile of D

for each 8x8 tile
compute required pixels of C
compute pixels in tile of E



Complete schedule

Add standard compiler optimizations: multi-core parallelism, vectorization, loop unrolling, etc.

for each 8x128 tile in parallel

 vectorize compute required pixels of A unroll x by 4

 vectorize compute required pixels of B

 vectorize compute pixels in tile of D

for each 8x8 tile in parallel

 vectorize compute required pixels of C unroll y by 2

 vectorize compute pixels in tile of E

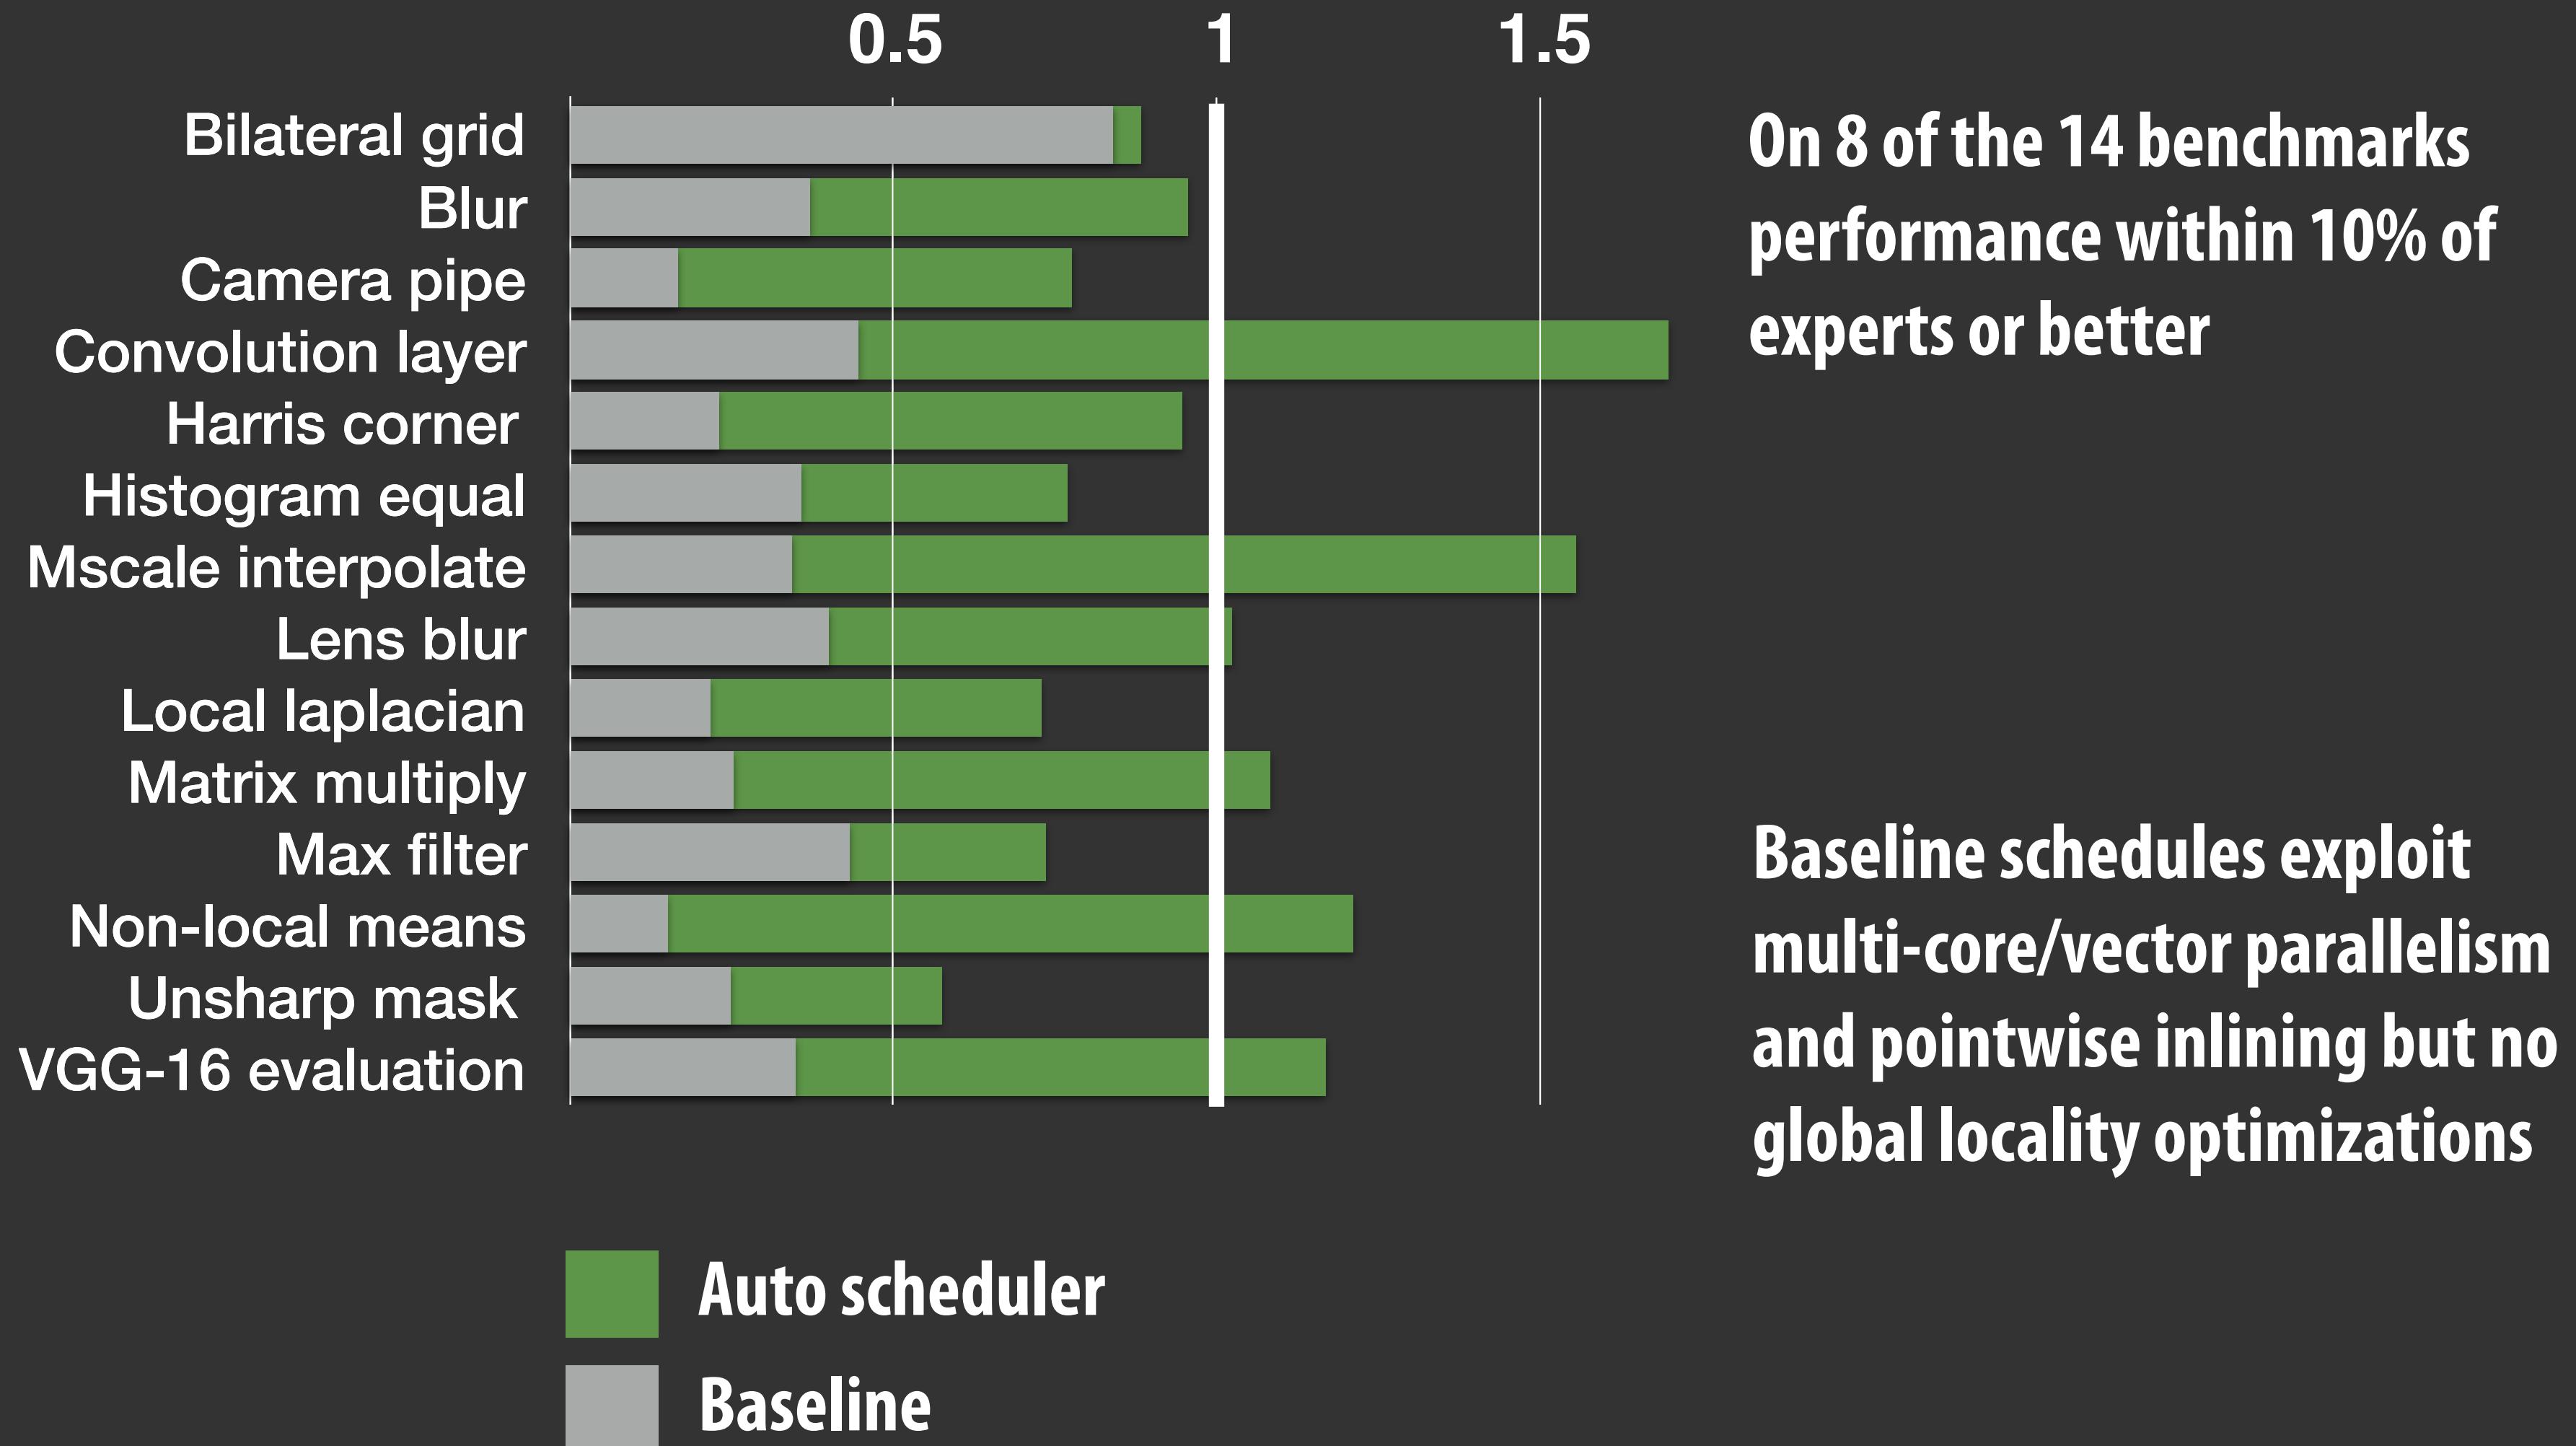
Autoscheduler generates schedules in seconds

Benchmark	Stages	Compile time (s)
Blur	3	<1
Unsharp mask	9	<1
Harris corner detection	13	<1
Camera RAW processing	30	<1
Non-local means denoising	13	<1
Max-brightness filter	9	<1
Multi-scale interpolation	52	2.6
Local-laplacian filter	103	3.9
Synthetic depth-of-field	74	55
Bilateral filter	8	<1
Histogram equalization	7	<1
VGG-16 deep network eval	64	6.9

Autoscheduler performs comparably to experts

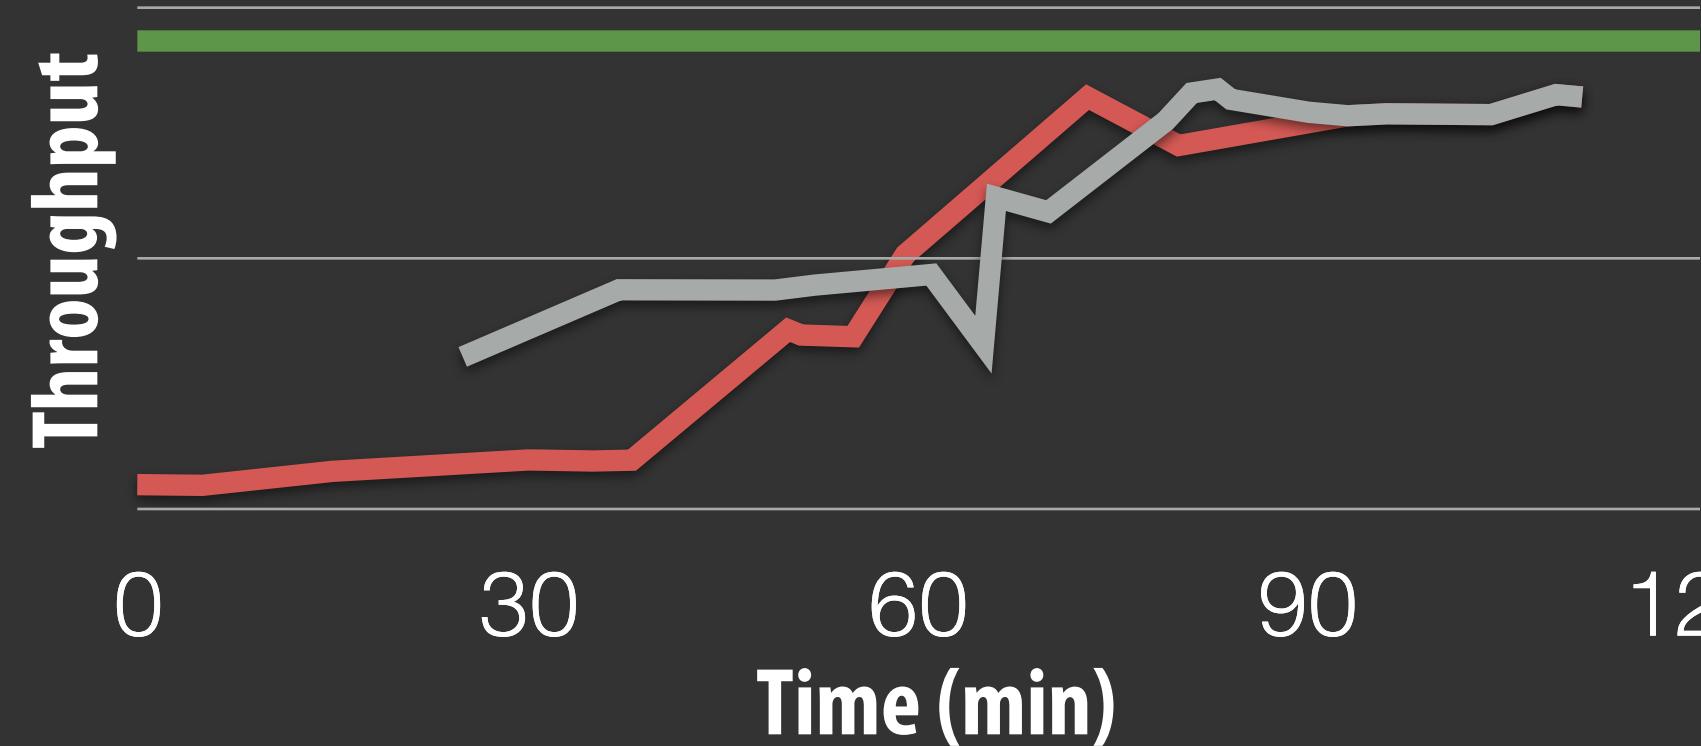
Performance relative to schedules authored by experts

(6 core Xeon CPU)

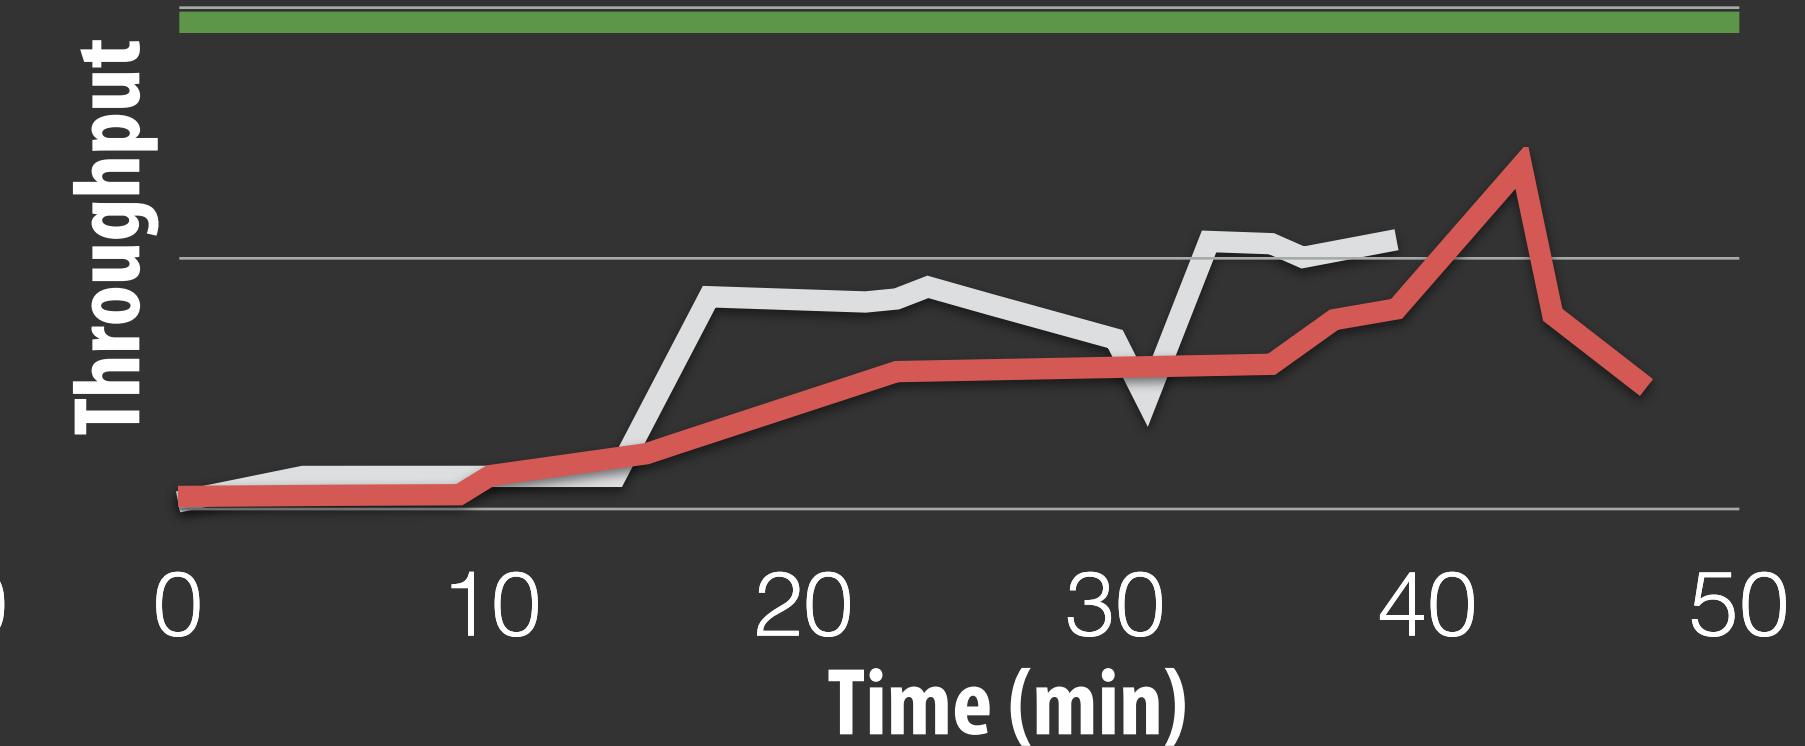


Autoscheduler saves time for experts

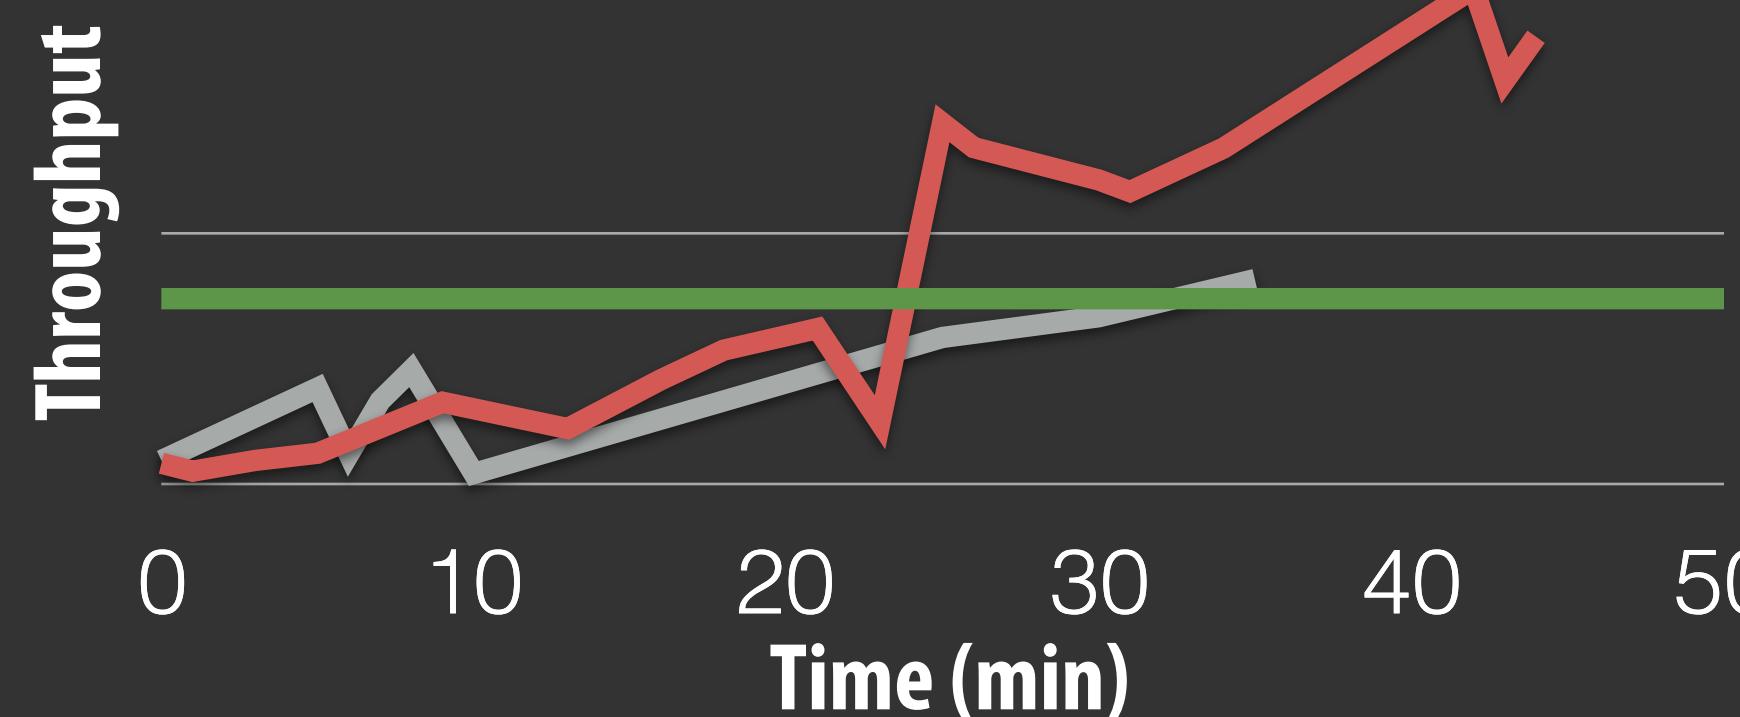
Non-local means denoising



Lens blur



Max filter



Auto scheduler
Dillon
Andrew

Impact / tech transfer

Auto scheduler now maintained by Google's Halide team



- If you write a Halide program, you can now schedule it with `.autoschedule()`
- Works quite well for CPU targets (x86/ARM)
- Works okay for GPU targets
- Will generate better schedulers in <1 sec than I would be able to create manually in a few days
- May not beat top developers at a major company, but gives them a strong head start

Tonight's Halide readings

- **What is the key intellectual idea of the Halide system?**
 - Hint: it is not the declarative language syntax
- **What services does Halide provide its users?**
- **What aspects of the design of Halide allow it to provide those services?**
- **Keep in mind: the key aspect in the design of any system usually is choosing the “right” representations for the job**